

## autogluon\_each\_location

October 19, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*30
presets = 'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = [] # ["hour", "year"]
use_estimated_diff_attr = False
use_is_estimated_attr = False

use_groups = False
n_groups = 8

auto_stack = False
num_stack_levels = 0
num_bag_folds = 8
num_bag_sets = 20

use_tune_data = True
use_test_data = True
tune_and_test_length = 0.5 # 3 months from end
holdout_frac = None
use_bag_holdout = True # Enable this if there is a large gap between score_val_
    ↪ and score_test in stack models.

sample_weight = None # 'sample_weight' # None
weight_evaluation = False #
sample_weight_estimated = 1
sample_weight_may_july = 1

run_analysis = True

shift_predictions_by_average_of_negatives_then_clip = False
```

```
clip_predictions = True
shift_predictions = False
```

```
[2]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift all columns with "1h" in them by 1 hour, so that for index 16:00,
    ↳ we have the values from 17:00
    # but only for the columns with "1h" in the name
    # X_shifted = X.filter(regex="\dh").shift(-1, axis=1)
    # print(f"Number of columns with 1h in name: {X_shifted.columns}")

    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

    X_shifted = X[X.index.minute==0][columns].copy()
    # loop through all rows and check if index + 1 hour is in the index, if so
    ↳ get that value, else nan
    count1 = 0
    count2 = 0
    for i in range(len(X_shifted)):
        if X_shifted.index[i] + pd.Timedelta('1 hour') in X.index:
            count1 += 1
            X_shifted.iloc[i] = X.loc[X_shifted.index[i] + pd.Timedelta('1
    ↳ hour')][columns]
        else:
            count2 += 1
            X_shifted.iloc[i] = np.nan

    print("COUNT1", count1)
    print("COUNT2", count2)

    X_old_unshifted = X[X.index.minute==0][columns]
    # rename X_old_unshifted columns to have _not_shifted at the end
    X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
    ↳ columns]

    # put the shifted columns back into the original dataframe
```

```

X[columns] = X_shifted[columns]

date_calc = None
if "date_calc" in X.columns:
    date_calc = X[X.index.minute == 0]['date_calc']

# resample to hourly
print("index: ", X.index[0])
X = X.resample('H').mean()
print("index AFTER: ", X.index[0])

X[columns] = X_shifted[columns]
#X[X_old_unshifted.columns] = X_old_unshifted

if date_calc is not None:
    X['date_calc'] = date_calc

return X

def fix_X(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    # with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    X = feature_engineering(X)

    return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    if weight_evaluation:
        # add sample weights, which are 1 for observed and 3 for estimated

```

```

X_train_observed["sample_weight"] = 1
X_train_estimated["sample_weight"] = sample_weight_estimated
X_test["sample_weight"] = sample_weight_estimated

y_train['ds'] = pd.to_datetime(y_train['time'])
y_train.drop(columns=['time'], inplace=True)
y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
        X_test["estimated_diff_hours"] = (X_test.index - pd.
to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

        X_train_estimated["estimated_diff_hours"] =
X_train_estimated["estimated_diff_hours"].astype('int64')
        # the filled once will get dropped later anyways, when we drop y nans
        X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
fillna(-50).astype('int64')

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0
        X_train_estimated["is_estimated"] = 1
        X_test["is_estimated"] = 1

    # drop date_calc
    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)

```

```

X_train = pd.concat([X_train_observed, X_train_estimated])

# clip all y values to 0 if negative
y_train["y"] = y_train["y"].clip(lower=0)

X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

# print number of nans in y
print(f"Number of nans in y: {X_train['y'].isna().sum()}")

X_train["location"] = location
X_test["location"] = location

return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↳X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

Processing location A...  
COUNT1 29667  
COUNT2 1  
index: 2019-06-02 22:00:00  
index AFTER: 2019-06-02 22:00:00  
COUNT1 4392  
COUNT2 2  
index: 2022-10-28 22:00:00  
index AFTER: 2022-10-28 22:00:00  
COUNT1 702  
COUNT2 18  
index: 2023-05-01 00:00:00  
index AFTER: 2023-05-01 00:00:00  
Number of nans in y: 0  
Processing location B...  
COUNT1 29232  
COUNT2 1  
index: 2019-01-01 00:00:00  
index AFTER: 2019-01-01 00:00:00  
COUNT1 4392  
COUNT2 2  
index: 2022-10-28 22:00:00  
index AFTER: 2022-10-28 22:00:00  
COUNT1 702  
COUNT2 18  
index: 2023-05-01 00:00:00  
index AFTER: 2023-05-01 00:00:00  
Number of nans in y: 4  
Processing location C...  
COUNT1 29206  
COUNT2 1  
index: 2019-01-01 00:00:00  
index AFTER: 2019-01-01 00:00:00  
COUNT1 4392  
COUNT2 2  
index: 2022-10-28 22:00:00  
index AFTER: 2022-10-28 22:00:00  
COUNT1 702  
COUNT2 18  
index: 2023-05-01 00:00:00  
index AFTER: 2023-05-01 00:00:00  
Number of nans in y: 6059

# 1 Feature engineering

```
[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
               inplace=True)

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

#print(X_train.head())

# If the "sample_weight" column is present and weight_evaluation is True,
# multiply sample_weight with sample_weight_may_july if the ds is between
# 05-01 00:00:00 and 07-03 23:00:00, else add sample_weight as a column to
# X_train
if weight_evaluation:
    if "sample_weight" not in X_train.columns:
        X_train["sample_weight"] = 1

    X_train.loc[((X_train.index.month >= 5) & (X_train.index.month <= 6)) |
               ((X_train.index.month == 7) & (X_train.index.day <= 3)), "sample_weight"] *=
    sample_weight_may_july

print(X_train.iloc[200])
print(X_train[((X_train.index.month >= 5) & (X_train.index.month <= 6)) |
              ((X_train.index.month == 7) & (X_train.index.day <= 3))].head(1))

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    # of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()
```

```

# Calculate the size of each group for this location
group_size = len(loc_df) // n_groups

# Create a new 'group' column for this location
loc_df['group'] = np.repeat(range(n_groups),
↪repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

# Append to list of grouped DataFrames
grouped_dfs.append(loc_df)

# Concatenate all the grouped DataFrames back together
X_train = pd.concat(grouped_dfs)
X_train.sort_index(inplace=True)
print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
↪"dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm",
↪"wind_speed_u_10m:ms", "wind_speed_v_10m:ms", "snow_melt_10min:mm",
↪"rain_water:kgm2", "dew_point_2m:K", "precip_5min:mm", "absolute_humidity_2m:
↪gm3", "air_density_2m:kgm3", "msl_pressure:hPa", "pressure_100m:hPa",
↪"pressure_50m:hPa", "clear_sky_rad:W"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

```

absolute_humidity_2m:gm3          7.825
air_density_2m:kgm3              1.245
ceiling_height_agl:m            2085.774902
clear_sky_energy_1h:J            1685498.875
clear_sky_rad:W                  452.100006
cloud_base_agl:m                2085.774902
dew_or_rime:idx                  0.0
dew_point_2m:K                  280.549988
diffuse_rad:W                   140.800003
diffuse_rad_1h:J                538581.625
direct_rad:W                    102.599998
direct_rad_1h:J                 439453.8125
effective_cloud_cover:p          71.849998
elevation:m                     6.0

```



```

fresh_snow_12h:cm          0.0
fresh_snow_1h:cm           0.0
fresh_snow_24h:cm          0.0
fresh_snow_3h:cm           0.0
fresh_snow_6h:cm           0.0
is_day:idx                 1.0
is_in_shadow:idx           0.0
msl_pressure:hPa           1026.349976
precip_5min:mm             0.0
precip_type_5min:idx       0.0
pressure_100m:hPa          1013.325012
pressure_50m:hPa           1019.450012
prob_rime:p                0.0
rain_water:kgm2            0.0
relative_humidity_1000hPa:p 77.099998
sfc_pressure:hPa           1025.550049
snow_density:kgm3          NaN
snow_depth:cm              0.0
snow_drift:idx             0.0
snow_melt_10min:mm         0.0
snow_water:kgm2            0.0
sun_azimuth:d              93.415253
sun_elevation:d            27.633499
super_cooled_liquid_water:kgm2 0.025
t_1000hPa:K               282.625
total_cloud_cover:p         71.849998
visibility:m               44177.875
wind_speed_10m:ms          2.675
wind_speed_u_10m:ms        -2.3
wind_speed_v_10m:ms        -1.4
wind_speed_w_1000hPa:ms    0.0
y                           2991.12
location                    A
Name: 2019-06-11 06:00:00, dtype: object
      absolute_humidity_2m:kgm3  air_density_2m:kgm3  \
ds
2019-06-02 22:00:00            7.7            1.22825

      ceiling_height_agl:m  clear_sky_energy_1h:J  \
ds
2019-06-02 22:00:00      1728.949951            0.0

      clear_sky_rad:W  cloud_base_agl:m  dew_or_rime:idx  \
ds
2019-06-02 22:00:00      0.0      1728.949951            0.0

      dew_point_2m:K  diffuse_rad:W  diffuse_rad_1h:J  ...  \
ds
...
```

```

2019-06-02 22:00:00      280.299988      0.0      0.0 ...
                                super_cooled_liquid_water:kgm2  t_1000hPa:K  \
ds
2019-06-02 22:00:00      0.0      286.225006
                                total_cloud_cover:p  visibility:m  wind_speed_10m:ms  \
ds
2019-06-02 22:00:00      100.0  40386.476562      3.6
                                wind_speed_u_10m:ms  wind_speed_v_10m:ms  \
ds
2019-06-02 22:00:00      -3.575      -0.5
                                wind_speed_w_1000hPa:ms  y  location
ds
2019-06-02 22:00:00      0.0  0.0      A

[1 rows x 47 columns]

```

```

[5]: # Create a plot of X_train showing its "y" and color it based on the value of
      ↪ the sample_weight column.
      #import matplotlib.pyplot as plt
      #import seaborn as sns
      #sns.scatterplot(data=X_train, x=X_train.index, y="y", hue="sample_weight",
      ↪ palette="deep", size=3)
      #plt.show()

```

```

[6]: def normalize_sample_weights_per_location(df):
      for loc in locations:
          loc_df = df[df["location"] == loc]
          loc_df["sample_weight"] = loc_df["sample_weight"] /
          ↪ loc_df["sample_weight"].sum() * loc_df.shape[0]
          df[df["location"] == loc] = loc_df
      return df

import pandas as pd
import numpy as np

def split_and_shuffle_data(input_data, num_bins, frac1):
    """
    Splits the input_data into num_bins and shuffles them, then divides the
    ↪ bins into two datasets based on the given fraction for the first set.

    Args:
        input_data (pd.DataFrame): The data to be split and shuffled.
        num_bins (int): The number of bins to split the data into.

```

```

    frac1 (float): The fraction of each bin to go into the first output_
↳ dataset.

Returns:
    pd.DataFrame, pd.DataFrame: The two output datasets.
"""
# Validate the input fraction
if frac1 < 0 or frac1 > 1:
    raise ValueError("frac1 must be between 0 and 1.")

if frac1==1:
    return input_data, pd.DataFrame()

# Calculate the fraction for the second output set
frac2 = 1 - frac1

# Calculate bin size
bin_size = len(input_data) // num_bins

# Initialize empty DataFrames for output
output_data1 = pd.DataFrame()
output_data2 = pd.DataFrame()

for i in range(num_bins):
    # Shuffle the data in the current bin
    np.random.seed(i)
    current_bin = input_data.iloc[i * bin_size: (i + 1) * bin_size].
↳ sample(frac=1)

    # Calculate the sizes for each output set
    size1 = int(len(current_bin) * frac1)

    # Split and append to output DataFrames
    output_data1 = pd.concat([output_data1, current_bin.iloc[:size1]])
    output_data2 = pd.concat([output_data2, current_bin.iloc[size1:]])

# Shuffle and split the remaining data
remaining_data = input_data.iloc[num_bins * bin_size:].sample(frac=1)
remaining_size1 = int(len(remaining_data) * frac1)

output_data1 = pd.concat([output_data1, remaining_data.iloc[:
↳ remaining_size1]])
output_data2 = pd.concat([output_data2, remaining_data.iloc[remaining_size1:
↳ ]])

return output_data1, output_data2

```

```

[7]: from autogluon.tabular import TabularDataset, TabularPredictor
      from autogluon.timeseries import TimeSeriesDataFrame
      import numpy as np
      data = TabularDataset('X_train_raw.csv')
      # set group column of train_data be increasing from 0 to 7 based on time, the
      ↪ first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
      data['ds'] = pd.to_datetime(data['ds'])
      data = data.sort_values(by='ds')

      # # print size of the group for each location
      # for loc in locations:
      #     print(f"Location {loc}:")
      #     print(train_data[train_data["location"] == loc].groupby('group').size())

      # get end date of train data and subtract 3 months
      # split_time = pd.to_datetime(train_data["ds"]).max() - pd.
      ↪ Timedelta(hours=tune_and_test_length)
      # 2022-10-28 22:00:00
      split_time = pd.to_datetime("2022-10-28 22:00:00")
      train_set = TabularDataset(data[data["ds"] < split_time])
      test_set = TabularDataset(data[data["ds"] >= split_time])

      # shuffle test_set and only grab tune_and_test_length percent of it, rest goes
      ↪ to train_set
      test_set, new_train_set = split_and_shuffle_data(test_set, 40,
      ↪ tune_and_test_length)

      print("Length of train set before adding test set", len(train_set))
      # add rest to train_set
      train_set = pd.concat([train_set, new_train_set])
      print("Length of train set after adding test set", len(train_set))
      print("Length of test set", len(test_set))

      if use_groups:
          test_set = test_set.drop(columns=['group'])

      tuning_data = None
      if use_tune_data:
          if use_test_data:
              # split test_set in half, use first half for tuning
              tuning_data, test_data = [], []

```

```

        for loc in locations:
            loc_test_set = test_set[test_set["location"] == loc]
            # randomly shuffle the loc_test_set
            loc_tuning_data, loc_test_data = □
            ↪split_and_shuffle_data(loc_test_set, 40, 0.5)
            tuning_data.append(loc_tuning_data)
            test_data.append(loc_test_data)
            tuning_data = pd.concat(tuning_data)
            test_data = pd.concat(test_data)
            print("Shapes of tuning and test", tuning_data.shape[0], test_data.
            ↪shape[0], tuning_data.shape[0] + test_data.shape[0])

        else:
            tuning_data = test_set
            print("Shape of tuning", tuning_data.shape[0])

            # ensure sample weights for your tuning data sum to the number of rows in
            ↪the tuning data.
            if weight_evaluation:
                tuning_data = normalize_sample_weights_per_location(tuning_data)

    else:
        if use_test_data:
            test_data = test_set
            print("Shape of test", test_data.shape[0])

train_data = train_set

# ensure sample weights for your training (or tuning) data sum to the number of
↪rows in the training (or tuning) data.
if weight_evaluation:
    train_data = normalize_sample_weights_per_location(train_data)
    if use_test_data:
        test_data = normalize_sample_weights_per_location(test_data)

train_data = TabularDataset(train_data)
if use_tune_data:
    tuning_data = TabularDataset(tuning_data)
if use_test_data:
    test_data = TabularDataset(test_data)

```

Length of train set before adding test set 82026

Length of train set after adding test set 87486

Length of test set 5459

Shapes of tuning and test 2728 2731 5459

```
[8]: if run_analysis:
      import autogluon.eda.auto as auto
      auto.dataset_overview(train_data=train_data, test_data=test_data,
                             label="y", sample=None)
```

train\_data dataset summary

|                                | count | unique | top                 | freq  | \ |
|--------------------------------|-------|--------|---------------------|-------|---|
| ceiling_height_agl:m           | 72280 | 59980  |                     |       |   |
| clear_sky_energy_1h:J          | 87486 | 46359  |                     |       |   |
| cloud_base_agl:m               | 81454 | 61360  |                     |       |   |
| diffuse_rad:W                  | 87486 | 11092  |                     |       |   |
| diffuse_rad_1h:J               | 87486 | 46319  |                     |       |   |
| direct_rad:W                   | 87486 | 14181  |                     |       |   |
| direct_rad_1h:J                | 87486 | 40118  |                     |       |   |
| ds                             | 87486 | 36794  | 2021-02-05 14:00:00 | 3     |   |
| effective_cloud_cover:p        | 87486 | 5655   |                     |       |   |
| elevation:m                    | 87486 | 3      |                     |       |   |
| fresh_snow_1h:cm               | 87486 | 39     |                     |       |   |
| fresh_snow_3h:cm               | 87486 | 70     |                     |       |   |
| fresh_snow_6h:cm               | 87486 | 96     |                     |       |   |
| is_day:idx                     | 87486 | 5      |                     |       |   |
| is_in_shadow:idx               | 87486 | 5      |                     |       |   |
| location                       | 87486 | 3      | A                   | 31872 |   |
| precip_type_5min:idx           | 87486 | 15     |                     |       |   |
| relative_humidity_1000hPa:p    | 87486 | 3799   |                     |       |   |
| sfc_pressure:hPa               | 87486 | 3795   |                     |       |   |
| snow_depth:cm                  | 87486 | 487    |                     |       |   |
| snow_water:kgm2                | 87486 | 161    |                     |       |   |
| sun_azimuth:d                  | 87486 | 83179  |                     |       |   |
| sun_elevation:d                | 87486 | 72262  |                     |       |   |
| super_cooled_liquid_water:kgm2 | 87486 | 53     |                     |       |   |
| t_1000hPa:K                    | 87486 | 1989   |                     |       |   |
| total_cloud_cover:p            | 87486 | 5556   |                     |       |   |
| visibility:m                   | 87486 | 85949  |                     |       |   |
| wind_speed_10m:ms              | 87486 | 596    |                     |       |   |
| y                              | 87486 | 11321  |                     |       |   |

|                       | first      | last                | mean          | \ |
|-----------------------|------------|---------------------|---------------|---|
| ceiling_height_agl:m  | NaT        | NaT                 | 2861.929806   |   |
| clear_sky_energy_1h:J | NaT        | NaT                 | 530297.395771 |   |
| cloud_base_agl:m      | NaT        | NaT                 | 1740.241802   |   |
| diffuse_rad:W         | NaT        | NaT                 | 40.267497     |   |
| diffuse_rad_1h:J      | NaT        | NaT                 | 145328.6257   |   |
| direct_rad:W          | NaT        | NaT                 | 51.524847     |   |
| direct_rad_1h:J       | NaT        | NaT                 | 185338.05854  |   |
| ds                    | 2019-01-01 | 2023-04-30 22:00:00 |               |   |

|                                |     |     |              |
|--------------------------------|-----|-----|--------------|
| effective_cloud_cover:p        | NaT | NaT | 67.052836    |
| elevation:m                    | NaT | NaT | 11.414718    |
| fresh_snow_1h:cm               | NaT | NaT | 0.008783     |
| fresh_snow_3h:cm               | NaT | NaT | 0.026713     |
| fresh_snow_6h:cm               | NaT | NaT | 0.05322      |
| is_day:idx                     | NaT | NaT | 0.490147     |
| is_in_shadow:idx               | NaT | NaT | 0.556952     |
| location                       | NaT | NaT |              |
| precip_type_5min:idx           | NaT | NaT | 0.084976     |
| relative_humidity_1000hPa:p    | NaT | NaT | 73.779918    |
| sfc_pressure:hPa               | NaT | NaT | 1008.035963  |
| snow_depth:cm                  | NaT | NaT | 0.197574     |
| snow_water:kgm2                | NaT | NaT | 0.090839     |
| sun_azimuth:d                  | NaT | NaT | 179.584247   |
| sun_elevation:d                | NaT | NaT | -0.705998    |
| super_cooled_liquid_water:kgm2 | NaT | NaT | 0.058256     |
| t_1000hPa:K                    | NaT | NaT | 279.675551   |
| total_cloud_cover:p            | NaT | NaT | 73.72398     |
| visibility:m                   | NaT | NaT | 32944.238197 |
| wind_speed_10m:ms              | NaT | NaT | 3.032943     |
| y                              | NaT | NaT | 294.447861   |

|                                | std           | min     | 25%        | 50% \     |
|--------------------------------|---------------|---------|------------|-----------|
| ceiling_height_agl:m           | 2532.377528   | 27.8    | 1082.3125  | 1856.075  |
| clear_sky_energy_1h:J          | 831839.646633 | 0.0     | 0.0        | 9084.9    |
| cloud_base_agl:m               | 1808.519208   | 27.5    | 598.15625  | 1174.775  |
| diffuse_rad:W                  | 61.119566     | 0.0     | 0.0        | 1.35      |
| diffuse_rad_1h:J               | 218036.296903 | 0.0     | 0.0        | 12531.2   |
| direct_rad:W                   | 114.236728    | 0.0     | 0.0        | 0.0       |
| direct_rad_1h:J                | 406379.39471  | 0.0     | 0.0        | 0.0       |
| ds                             |               |         |            |           |
| effective_cloud_cover:p        | 34.132847     | 0.0     | 42.125     | 79.7      |
| elevation:m                    | 7.881545      | 6.0     | 6.0        | 7.0       |
| fresh_snow_1h:cm               | 0.110515      | 0.0     | 0.0        | 0.0       |
| fresh_snow_3h:cm               | 0.277575      | 0.0     | 0.0        | 0.0       |
| fresh_snow_6h:cm               | 0.474579      | 0.0     | 0.0        | 0.0       |
| is_day:idx                     | 0.486133      | 0.0     | 0.0        | 0.5       |
| is_in_shadow:idx               | 0.484138      | 0.0     | 0.0        | 1.0       |
| location                       |               |         |            |           |
| precip_type_5min:idx           | 0.32995       | 0.0     | 0.0        | 0.0       |
| relative_humidity_1000hPa:p    | 14.200631     | 19.575  | 64.4       | 76.15     |
| sfc_pressure:hPa               | 13.038723     | 941.55  | 1000.05    | 1009.0    |
| snow_depth:cm                  | 1.28439       | 0.0     | 0.0        | 0.0       |
| snow_water:kgm2                | 0.240248      | 0.0     | 0.0        | 0.0       |
| sun_azimuth:d                  | 97.419022     | 6.983   | 94.4135    | 180.00663 |
| sun_elevation:d                | 24.006117     | -49.932 | -17.969875 | -0.453875 |
| super_cooled_liquid_water:kgm2 | 0.106959      | 0.0     | 0.0        | 0.0       |
| t_1000hPa:K                    | 6.551665      | 258.025 | 275.1      | 278.975   |

|                     |             |         |            |          |
|---------------------|-------------|---------|------------|----------|
| total_cloud_cover:p | 33.851299   | 0.0     | 53.475     | 92.825   |
| visibility:m        | 17949.64428 | 132.375 | 16564.5125 | 36910.75 |
| wind_speed_10m:ms   | 1.758713    | 0.025   | 1.675      | 2.7      |
| y                   | 774.531815  | -0.0    | 0.0        | 0.0      |

|                                |            |           |                |
|--------------------------------|------------|-----------|----------------|
|                                | 75%        | max       | dtypes \       |
| ceiling_height_agl:m           | 3916.78125 | 12285.775 | float64        |
| clear_sky_energy_1h:J          | 831169.675 | 3006697.2 | float64        |
| cloud_base_agl:m               | 2080.39375 | 11673.725 | float64        |
| diffuse_rad:W                  | 67.15      | 334.75    | float64        |
| diffuse_rad_1h:J               | 243365.275 | 1182265.4 | float64        |
| direct_rad:W                   | 32.225     | 683.4     | float64        |
| direct_rad_1h:J                | 122454.125 | 2445897.0 | float64        |
| ds                             |            |           | datetime64[ns] |
| effective_cloud_cover:p        | 98.5       | 100.0     | float64        |
| elevation:m                    | 24.0       | 24.0      | float64        |
| fresh_snow_1h:cm               | 0.0        | 7.1       | float64        |
| fresh_snow_3h:cm               | 0.0        | 20.6      | float64        |
| fresh_snow_6h:cm               | 0.0        | 34.0      | float64        |
| is_day:idx                     | 1.0        | 1.0       | float64        |
| is_in_shadow:idx               | 1.0        | 1.0       | float64        |
| location                       |            |           | object         |
| precip_type_5min:idx           | 0.0        | 5.0       | float64        |
| relative_humidity_1000hPa:p    | 85.175     | 100.0     | float64        |
| sfc_pressure:hPa               | 1017.1     | 1043.725  | float64        |
| snow_depth:cm                  | 0.0        | 18.2      | float64        |
| snow_water:kgm2                | 0.1        | 5.65      | float64        |
| sun_azimuth:d                  | 264.601138 | 348.48752 | float64        |
| sun_elevation:d                | 16.004499  | 49.94375  | float64        |
| super_cooled_liquid_water:kgm2 | 0.1        | 1.375     | float64        |
| t_1000hPa:K                    | 284.225    | 303.25    | float64        |
| total_cloud_cover:p            | 99.9       | 100.0     | float64        |
| visibility:m                   | 48289.05   | 75326.58  | float64        |
| wind_speed_10m:ms              | 4.05       | 13.275    | float64        |
| y                              | 183.7125   | 5733.42   | float64        |

|                         |               |               |            |
|-------------------------|---------------|---------------|------------|
|                         | missing_count | missing_ratio | raw_type \ |
| ceiling_height_agl:m    | 15206         | 0.173811      | float      |
| clear_sky_energy_1h:J   |               |               | float      |
| cloud_base_agl:m        | 6032          | 0.068948      | float      |
| diffuse_rad:W           |               |               | float      |
| diffuse_rad_1h:J        |               |               | float      |
| direct_rad:W            |               |               | float      |
| direct_rad_1h:J         |               |               | float      |
| ds                      |               |               | datetime   |
| effective_cloud_cover:p |               |               | float      |
| elevation:m             |               |               | float      |
| fresh_snow_1h:cm        |               |               | float      |



|                                |        |
|--------------------------------|--------|
| fresh_snow_3h:cm               | float  |
| fresh_snow_6h:cm               | float  |
| is_day:idx                     | float  |
| is_in_shadow:idx               | float  |
| location                       | object |
| precip_type_5min:idx           | float  |
| relative_humidity_1000hPa:p    | float  |
| sfc_pressure:hPa               | float  |
| snow_depth:cm                  | float  |
| snow_water:kgm2                | float  |
| sun_azimuth:d                  | float  |
| sun_elevation:d                | float  |
| super_cooled_liquid_water:kgm2 | float  |
| t_1000hPa:K                    | float  |
| total_cloud_cover:p            | float  |
| visibility:m                   | float  |
| wind_speed_10m:ms              | float  |
| y                              | float  |

|                                | variable_type | special_types |
|--------------------------------|---------------|---------------|
| ceiling_height_agl:m           | numeric       |               |
| clear_sky_energy_1h:J          | numeric       |               |
| cloud_base_agl:m               | numeric       |               |
| diffuse_rad:W                  | numeric       |               |
| diffuse_rad_1h:J               | numeric       |               |
| direct_rad:W                   | numeric       |               |
| direct_rad_1h:J                | numeric       |               |
| ds                             |               |               |
| effective_cloud_cover:p        | numeric       |               |
| elevation:m                    | category      |               |
| fresh_snow_1h:cm               | numeric       |               |
| fresh_snow_3h:cm               | numeric       |               |
| fresh_snow_6h:cm               | numeric       |               |
| is_day:idx                     | category      |               |
| is_in_shadow:idx               | category      |               |
| location                       | category      |               |
| precip_type_5min:idx           | category      |               |
| relative_humidity_1000hPa:p    | numeric       |               |
| sfc_pressure:hPa               | numeric       |               |
| snow_depth:cm                  | numeric       |               |
| snow_water:kgm2                | numeric       |               |
| sun_azimuth:d                  | numeric       |               |
| sun_elevation:d                | numeric       |               |
| super_cooled_liquid_water:kgm2 | numeric       |               |
| t_1000hPa:K                    | numeric       |               |
| total_cloud_cover:p            | numeric       |               |
| visibility:m                   | numeric       |               |
| wind_speed_10m:ms              | numeric       |               |

y

numeric

## test\_data dataset summary

|                                | count | unique | top                 | freq | \ |
|--------------------------------|-------|--------|---------------------|------|---|
| ceiling_height_agl:m           | 2100  | 2065   |                     |      |   |
| clear_sky_energy_1h:J          | 2731  | 1138   |                     |      |   |
| cloud_base_agl:m               | 2374  | 2332   |                     |      |   |
| diffuse_rad:W                  | 2731  | 983    |                     |      |   |
| diffuse_rad_1h:J               | 2731  | 1138   |                     |      |   |
| direct_rad:W                   | 2731  | 750    |                     |      |   |
| direct_rad_1h:J                | 2731  | 932    |                     |      |   |
| ds                             | 2731  | 2200   | 2023-04-10 19:00:00 | 3    |   |
| effective_cloud_cover:p        | 2731  | 1348   |                     |      |   |
| elevation:m                    | 2731  | 3      |                     |      |   |
| fresh_snow_1h:cm               | 2731  | 19     |                     |      |   |
| fresh_snow_3h:cm               | 2731  | 36     |                     |      |   |
| fresh_snow_6h:cm               | 2731  | 50     |                     |      |   |
| is_day:idx                     | 2731  | 5      |                     |      |   |
| is_in_shadow:idx               | 2731  | 5      |                     |      |   |
| location                       | 2731  | 3      | A                   | 1094 |   |
| precip_type_5min:idx           | 2731  | 10     |                     |      |   |
| relative_humidity_1000hPa:p    | 2731  | 1523   |                     |      |   |
| sfc_pressure:hPa               | 2731  | 1575   |                     |      |   |
| snow_depth:cm                  | 2731  | 61     |                     |      |   |
| snow_water:kgm2                | 2731  | 61     |                     |      |   |
| sun_azimuth:d                  | 2731  | 2716   |                     |      |   |
| sun_elevation:d                | 2731  | 2652   |                     |      |   |
| super_cooled_liquid_water:kgm2 | 2731  | 29     |                     |      |   |
| t_1000hPa:K                    | 2731  | 774    |                     |      |   |
| total_cloud_cover:p            | 2731  | 1126   |                     |      |   |
| visibility:m                   | 2731  | 2729   |                     |      |   |
| wind_speed_10m:ms              | 2731  | 366    |                     |      |   |
| y                              | 2731  | 895    |                     |      |   |

|                         | first               | last                | \ |
|-------------------------|---------------------|---------------------|---|
| ceiling_height_agl:m    | NaT                 | NaT                 |   |
| clear_sky_energy_1h:J   | NaT                 | NaT                 |   |
| cloud_base_agl:m        | NaT                 | NaT                 |   |
| diffuse_rad:W           | NaT                 | NaT                 |   |
| diffuse_rad_1h:J        | NaT                 | NaT                 |   |
| direct_rad:W            | NaT                 | NaT                 |   |
| direct_rad_1h:J         | NaT                 | NaT                 |   |
| ds                      | 2022-10-28 22:00:00 | 2023-04-30 19:00:00 |   |
| effective_cloud_cover:p | NaT                 | NaT                 |   |
| elevation:m             | NaT                 | NaT                 |   |
| fresh_snow_1h:cm        | NaT                 | NaT                 |   |
| fresh_snow_3h:cm        | NaT                 | NaT                 |   |
| fresh_snow_6h:cm        | NaT                 | NaT                 |   |

|                                |     |     |
|--------------------------------|-----|-----|
| is_day:idx                     | NaT | NaT |
| is_in_shadow:idx               | NaT | NaT |
| location                       | NaT | NaT |
| precip_type_5min:idx           | NaT | NaT |
| relative_humidity_1000hPa:p    | NaT | NaT |
| sfc_pressure:hPa               | NaT | NaT |
| snow_depth:cm                  | NaT | NaT |
| snow_water:kgm2                | NaT | NaT |
| sun_azimuth:d                  | NaT | NaT |
| sun_elevation:d                | NaT | NaT |
| super_cooled_liquid_water:kgm2 | NaT | NaT |
| t_1000hPa:K                    | NaT | NaT |
| total_cloud_cover:p            | NaT | NaT |
| visibility:m                   | NaT | NaT |
| wind_speed_10m:ms              | NaT | NaT |
| y                              | NaT | NaT |

|                                | mean          | std           | min \   |
|--------------------------------|---------------|---------------|---------|
| ceiling_height_agl:m           | 3361.682133   | 2562.862274   | 28.0    |
| clear_sky_energy_1h:J          | 285528.142658 | 573252.521662 | 0.0     |
| cloud_base_agl:m               | 1685.111501   | 1833.56975    | 27.5    |
| diffuse_rad:W                  | 26.230813     | 48.360421     | 0.0     |
| diffuse_rad_1h:J               | 94649.301538  | 172723.786046 | 0.0     |
| direct_rad:W                   | 32.798068     | 92.244541     | 0.0     |
| direct_rad_1h:J                | 118054.008202 | 329601.815692 | 0.0     |
| ds                             |               |               |         |
| effective_cloud_cover:p        | 66.798654     | 36.717964     | 0.0     |
| elevation:m                    | 11.193336     | 7.806119      | 6.0     |
| fresh_snow_1h:cm               | 0.023984      | 0.147555      | 0.0     |
| fresh_snow_3h:cm               | 0.069498      | 0.352141      | 0.0     |
| fresh_snow_6h:cm               | 0.13885       | 0.57722       | 0.0     |
| is_day:idx                     | 0.378341      | 0.472171      | 0.0     |
| is_in_shadow:idx               | 0.677133      | 0.452911      | 0.0     |
| location                       |               |               |         |
| precip_type_5min:idx           | 0.076254      | 0.344931      | 0.0     |
| relative_humidity_1000hPa:p    | 71.631472     | 14.652551     | 21.7    |
| sfc_pressure:hPa               | 1009.397775   | 14.318453     | 971.15  |
| snow_depth:cm                  | 0.119965      | 0.56196       | 0.0     |
| snow_water:kgm2                | 0.080511      | 0.19473       | 0.0     |
| sun_azimuth:d                  | 180.975998    | 94.222121     | 14.914  |
| sun_elevation:d                | -8.945472     | 22.095926     | -49.887 |
| super_cooled_liquid_water:kgm2 | 0.035088      | 0.082444      | 0.0     |
| t_1000hPa:K                    | 275.52987     | 4.271781      | 259.975 |
| total_cloud_cover:p            | 72.32056      | 37.085445     | 0.0     |
| visibility:m                   | 34504.948017  | 17242.154257  | 270.3   |
| wind_speed_10m:ms              | 3.109676      | 1.782531      | 0.125   |
| y                              | 179.379421    | 641.546947    | 0.0     |


|                                | 25%        | 50%       | 75%        | max \     |
|--------------------------------|------------|-----------|------------|-----------|
| ceiling_height_agl:m           | 1245.55625 | 2784.275  | 4919.18125 | 12294.9   |
| clear_sky_energy_1h:J          | 0.0        | 0.0       | 220909.2   | 2551917.2 |
| cloud_base_agl:m               | 516.7625   | 1000.8    | 2066.9875  | 10813.7   |
| diffuse_rad:W                  | 0.0        | 0.0       | 32.0       | 280.5     |
| diffuse_rad_1h:J               | 0.0        | 0.0       | 116208.0   | 986147.0  |
| direct_rad:W                   | 0.0        | 0.0       | 4.7625     | 511.7     |
| direct_rad_1h:J                | 0.0        | 0.0       | 24326.65   | 1844204.9 |
| ds                             |            |           |            |           |
| effective_cloud_cover:p        | 36.3125    | 83.05     | 99.825     | 100.0     |
| elevation:m                    | 6.0        | 7.0       | 24.0       | 24.0      |
| fresh_snow_1h:cm               | 0.0        | 0.0       | 0.0        | 2.3       |
| fresh_snow_3h:cm               | 0.0        | 0.0       | 0.0        | 4.8       |
| fresh_snow_6h:cm               | 0.0        | 0.0       | 0.0        | 6.3       |
| is_day:idx                     | 0.0        | 0.0       | 1.0        | 1.0       |
| is_in_shadow:idx               | 0.0        | 1.0       | 1.0        | 1.0       |
| location                       |            |           |            |           |
| precip_type_5min:idx           | 0.0        | 0.0       | 0.0        | 3.0       |
| relative_humidity_1000hPa:p    | 61.775     | 73.55     | 82.9625    | 99.775    |
| sfc_pressure:hPa               | 999.1      | 1009.5    | 1020.275   | 1040.6    |
| snow_depth:cm                  | 0.0        | 0.0       | 0.0        | 4.9       |
| snow_water:kgm2                | 0.0        | 0.0       | 0.1        | 2.15      |
| sun_azimuth:d                  | 101.047372 | 179.89075 | 260.65412  | 347.81226 |
| sun_elevation:d                | -26.72725  | -8.178    | 6.393625   | 41.09175  |
| super_cooled_liquid_water:kgm2 | 0.0        | 0.0       | 0.0        | 0.75      |
| t_1000hPa:K                    | 272.6625   | 275.45    | 278.5      | 285.825   |
| total_cloud_cover:p            | 44.5875    | 96.775    | 100.0      | 100.0     |
| visibility:m                   | 20902.55   | 36141.85  | 48867.949  | 73937.67  |
| wind_speed_10m:ms              | 1.6        | 2.8       | 4.2375     | 9.9       |
| y                              | 0.0        | 0.0       | 40.397706  | 5043.72   |

|                         | dtypes         | missing_count | missing_ratio \ |
|-------------------------|----------------|---------------|-----------------|
| ceiling_height_agl:m    | float64        | 631           | 0.231051        |
| clear_sky_energy_1h:J   | float64        |               |                 |
| cloud_base_agl:m        | float64        | 357           | 0.130721        |
| diffuse_rad:W           | float64        |               |                 |
| diffuse_rad_1h:J        | float64        |               |                 |
| direct_rad:W            | float64        |               |                 |
| direct_rad_1h:J         | float64        |               |                 |
| ds                      | datetime64[ns] |               |                 |
| effective_cloud_cover:p | float64        |               |                 |
| elevation:m             | float64        |               |                 |
| fresh_snow_1h:cm        | float64        |               |                 |
| fresh_snow_3h:cm        | float64        |               |                 |
| fresh_snow_6h:cm        | float64        |               |                 |
| is_day:idx              | float64        |               |                 |
| is_in_shadow:idx        | float64        |               |                 |
| location                | object         |               |                 |

|                                |         |
|--------------------------------|---------|
| precip_type_5min:idx           | float64 |
| relative_humidity_1000hPa:p    | float64 |
| sfc_pressure:hPa               | float64 |
| snow_depth:cm                  | float64 |
| snow_water:kgm2                | float64 |
| sun_azimuth:d                  | float64 |
| sun_elevation:d                | float64 |
| super_cooled_liquid_water:kgm2 | float64 |
| t_1000hPa:K                    | float64 |
| total_cloud_cover:p            | float64 |
| visibility:m                   | float64 |
| wind_speed_10m:ms              | float64 |
| y                              | float64 |

|                                | raw_type | variable_type | special_types |
|--------------------------------|----------|---------------|---------------|
| ceiling_height_agl:m           | float    | numeric       |               |
| clear_sky_energy_1h:J          | float    | numeric       |               |
| cloud_base_agl:m               | float    | numeric       |               |
| diffuse_rad:W                  | float    | numeric       |               |
| diffuse_rad_1h:J               | float    | numeric       |               |
| direct_rad:W                   | float    | numeric       |               |
| direct_rad_1h:J                | float    | numeric       |               |
| ds                             | datetime |               |               |
| effective_cloud_cover:p        | float    | numeric       |               |
| elevation:m                    | float    | category      |               |
| fresh_snow_1h:cm               | float    | category      |               |
| fresh_snow_3h:cm               | float    | numeric       |               |
| fresh_snow_6h:cm               | float    | numeric       |               |
| is_day:idx                     | float    | category      |               |
| is_in_shadow:idx               | float    | category      |               |
| location                       | object   | category      |               |
| precip_type_5min:idx           | float    | category      |               |
| relative_humidity_1000hPa:p    | float    | numeric       |               |
| sfc_pressure:hPa               | float    | numeric       |               |
| snow_depth:cm                  | float    | numeric       |               |
| snow_water:kgm2                | float    | numeric       |               |
| sun_azimuth:d                  | float    | numeric       |               |
| sun_elevation:d                | float    | numeric       |               |
| super_cooled_liquid_water:kgm2 | float    | numeric       |               |
| t_1000hPa:K                    | float    | numeric       |               |
| total_cloud_cover:p            | float    | numeric       |               |
| visibility:m                   | float    | numeric       |               |
| wind_speed_10m:ms              | float    | numeric       |               |
| y                              | float    | numeric       |               |

### 1.0.1 Feature Distance




The following feature groups are considered as near-duplicates:

Distance threshold:  $\leq 0.01$ . Consider keeping only some of the columns within each group:

- elevation:m, location - distance 0.00
- diffuse\_rad:W, diffuse\_rad\_1h:J - distance 0.00

Feature interaction between diffuse\_rad:W/diffuse\_rad\_1h:J




```
[9]: if run_analysis:
      auto.target_analysis(train_data=train_data, label="y", sample=None)
```

## 1.1 Target variable analysis

|   | count | mean       | std        | min  | 25% | 50% | 75%      | max     | dtypes  | \ |
|---|-------|------------|------------|------|-----|-----|----------|---------|---------|---|
| y | 87486 | 294.447861 | 774.531815 | -0.0 | 0.0 | 0.0 | 183.7125 | 5733.42 | float64 |   |

|   | unique | missing_count | missing_ratio | raw_type | special_types |
|---|--------|---------------|---------------|----------|---------------|
| y | 11321  |               |               | float    |               |




### 1.1.1 Distribution fits for target variable

- none of the [attempted](#) distribution fits satisfy specified minimum p-value threshold: 0.01


### 1.1.2 Target variable correlations

train\_data - spearman correlation matrix; focus: absolute correlation for  $y \geq 0.5$




Feature interaction between `clear_sky_energy_1h:J/y` in `train_data`







Feature interaction between diffuse\_rad:W/y in train\_data




Feature interaction between diffuse\_rad\_1h:J/y in train\_data




Feature interaction between `direct_rad_1h:J/y` in `train_data`




Feature interaction between `direct_rad:W/y` in `train_data`




Feature interaction between sun\_elevation:d/y in train\_data



Feature interaction between is\_day:idx/y in train\_data



Feature interaction between `is_in_shadow:idx`/`y` in `train_data`



## 2 Starting

```
[10]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    ↪filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

```
Last submission number: 99
Now creating submission number: 100
New filename: submission_100
```

```
[11]: predictors = [None, None, None]
```

```
[ ]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪train and tune data and test data
    if weight_evaluation:
        print("Train data sample weight sum:",
            ↪train_data[train_data["location"] == loc]["sample_weight"].sum())
        print("Train data number of rows:", train_data[train_data["location"]
            ↪== loc].shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
            ↪tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:",
            ↪tuning_data[tuning_data["location"] == loc].shape[0])
    if use_test_data:
        print("Test data sample weight sum:",
            ↪test_data[test_data["location"] == loc]["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"]
            ↪== loc].shape[0])
    predictor = TabularPredictor(
        label=label,
```

```

        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        # sample_weight=sample_weight,
        # weight_evaluation=weight_evaluation,
        # groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc].
↳drop(columns=["ds"]),
        time_limit=time_limit,
        # presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
↳somethin, will be overwritten anyways
        num_bag_sets=num_bag_sets,
        tuning_data=tuning_data[tuning_data["location"] == loc].
↳reset_index(drop=True).drop(columns=["ds"]) if use_tune_data else None,
        use_bag_holdout=use_bag_holdout,
        # holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
↳drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

```

Beginning AutoGluon training ... Time limit = 1800s
AutoGluon will save models to "AutogluonModels/submission_100_A/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 150.33 GB / 315.93 GB (47.6%)
Train Data Rows: 31872
Train Data Columns: 27
Tuning Data Rows: 1093
Tuning Data Columns: 27
Label Column: y

```

```

Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (5733.42, 0.0, 649.68162,
1178.37671)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 131984.69 MB
    Train Data (Original) Memory Usage: 8.77 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 25 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'cloud_base_agl:m', 'diffuse_rad:W',
'diffuse_rad_1h:J', ...]
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 25 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'cloud_base_agl:m', 'diffuse_rad:W',
'diffuse_rad_1h:J', ...]
    0.1s = Fit runtime
    25 features in original data used to generate 25 features in processed
data.
    Train Data (Processed) Memory Usage: 6.59 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.14s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()

```

```

use_bag_holdout=True, will use tuning_data as holdout (will not be used for
early stopping).
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Training model for location A...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1799.86s of the 1799.85s of remaining time.

```

-140.7607      = Validation score    (-mean_absolute_error)
0.03s         = Training    runtime
0.38s         = Validation runtime

```

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 1799.34s of the 1799.34s of remaining time.

```

-140.9568      = Validation score    (-mean_absolute_error)
0.03s         = Training    runtime
0.38s         = Validation runtime

```

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 1798.88s of the 1798.88s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-100.7752      = Validation score    (-mean_absolute_error)
32.14s        = Training    runtime
19.88s        = Validation runtime

```

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 1757.97s of the 1757.97s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-100.7435      = Validation score    (-mean_absolute_error)

```



```

    20.84s = Training runtime
    3.9s   = Validation runtime
Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1733.76s of
the 1733.76s of remaining time.
    -109.2719 = Validation score (-mean_absolute_error)
    6.65s    = Training runtime
    1.15s    = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1724.67s of the
1724.66s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -106.8026 = Validation score (-mean_absolute_error)
    187.78s   = Training runtime
    0.1s      = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1535.71s of the
1535.71s of remaining time.
    -114.7323 = Validation score (-mean_absolute_error)
    1.45s     = Training runtime
    1.15s     = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1531.77s of
the 1531.77s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -119.2733 = Validation score (-mean_absolute_error)
    39.85s    = Training runtime
    0.48s     = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 1490.17s of the
1490.17s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -103.765 = Validation score (-mean_absolute_error)
    5.68s    = Training runtime
    0.31s    = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1482.49s of
the 1482.49s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -99.7353 = Validation score (-mean_absolute_error)
    129.1s   = Training runtime
    0.29s    = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1351.79s of the
1351.79s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -98.4663 = Validation score (-mean_absolute_error)
    92.66s   = Training runtime
    29.64s   = Validation runtime
Repeating k-fold bagging: 2/20

```

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 1247.62s of the 1247.62s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with  
ParallelLocalFoldFittingStrategy  
-100.3554 = Validation score (-mean\_absolute\_error)  
61.35s = Training runtime  
40.15s = Validation runtime

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 1211.12s of the 1211.12s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with  
ParallelLocalFoldFittingStrategy  
-99.8771 = Validation score (-mean\_absolute\_error)  
46.32s = Training runtime  
7.68s = Validation runtime

Fitting model: CatBoost\_BAG\_L1 ... Training model for up to 1181.51s of the 1181.51s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```
[ ]: import matplotlib.pyplot as plt

leaderboards = [None, None, None]
def leaderboard_for_location(i, loc):
    if use_test_data:
        lb = predictors[i].leaderboard(test_data[test_data["location"] == loc])
        lb["location"] = loc
        plt.scatter(test_data[test_data["location"] == loc]["y"].index,
↳test_data[test_data["location"] == loc]["y"])
        if use_tune_data:
            plt.scatter(tuning_data[tuning_data["location"] == loc]["y"].index,
↳tuning_data[tuning_data["location"] == loc]["y"])
        plt.show()

    return lb
    else:
        return pd.DataFrame()

leaderboards[0] = leaderboard_for_location(0, loc)
```

```
[ ]: loc = "B"
predictors[1] = fit_predictor_for_location(loc)
leaderboards[1] = leaderboard_for_location(1, loc)
```

```
[ ]: loc = "C"
predictors[2] = fit_predictor_for_location(loc)
leaderboards[2] = leaderboard_for_location(2, loc)
```

```
[ ]: # save leaderboards to csv
pd.concat(leaderboards).to_csv(f"leaderboards/{new_filename}.csv")
```

### 3 Submit

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt

future_test_data = TabularDataset('X_test_raw.csv')
future_test_data["ds"] = pd.to_datetime(future_test_data["ds"])
#test_data
```

```
[ ]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
future_test_data_merged = pd.merge(future_test_data, test_ids, how="inner",
    ↪right_on=["time", "location"], left_on=["ds", "location"])

#test_data_merged
```

```
[ ]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in future_test_data.groupby('location'):
    i = location_map[loc]
    subset = future_test_data_merged[future_test_data_merged["location"] == loc]
    ↪loc].reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)

# get past predictions
train_data.loc[train_data["location"] == loc, "prediction"] = loc
    ↪predictors[i].predict(train_data[train_data["location"] == loc])
    if use_tune_data:
        tuning_data.loc[tuning_data["location"] == loc, "prediction"] = loc
    ↪predictors[i].predict(tuning_data[tuning_data["location"] == loc])
    if use_test_data:
        test_data.loc[test_data["location"] == loc, "prediction"] = loc
    ↪predictors[i].predict(test_data[test_data["location"] == loc])
```

```
[ ]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data[train_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="train data")
    if use_tune_data:
        tuning_data[tuning_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="tune data")
    if use_test_data:
        test_data[test_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="test data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    #train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='prediction', ax=ax, label="past predictions")
    train_data[train_data["location"]==loc].plot(x='ds', y='prediction', ax=ax,
    ↪label="past predictions train")
    if use_tune_data:
        tuning_data[tuning_data["location"]==loc].plot(x='ds', y='prediction',
    ↪ax=ax, label="past predictions tune")
    if use_test_data:
        test_data[test_data["location"]==loc].plot(x='ds', y='prediction',
    ↪ax=ax, label="past predictions test")

    # title
    ax.set_title(f"Predictions for location {loc}")
```

```
[ ]: temp_predictions = [prediction.copy() for prediction in predictions]
if clip_predictions:
    # clip predictions smaller than 0 to 0
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        pred.loc[pred["prediction"] < 0, "prediction"] = 0
        print("Smallest prediction after clipping:", pred["prediction"].min())

# Instead of clipping, shift all prediction values up by the largest negative
↪number.
# This way, the smallest prediction will be 0.
elif shift_predictions:
    for pred in temp_predictions:
```

```

        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        pred["prediction"] = pred["prediction"] - pred["prediction"].min()
        print("Smallest prediction after clipping:", pred["prediction"].min())

elif shift_predictions_by_average_of_negatives_then_clip:
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        mean_negative = pred[pred["prediction"] < 0]["prediction"].mean()
        # if not nan
        if mean_negative == mean_negative:
            pred["prediction"] = pred["prediction"] - mean_negative

    pred.loc[pred["prediction"] < 0, "prediction"] = 0
    print("Smallest prediction after clipping:", pred["prediction"].min())

# concatenate predictions
submissions_df = pd.concat(temp_predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df

```

```

[ ]: # Save the submission DataFrame to submissions folder, create new name based on
      ↪ last submission, format is submission_<last_submission_number + 1>.csv

```

```

# Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
    ↪ index=False)
print("jallia")

```

```

[ ]: train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])
# feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
    ↪ time_limit=60*10)

```

```

[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]

```

```

predictors[0].feature_importance(feature_stage="original", data=observed,
    ↪time_limit=60*10)

```

```

[ ]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)

```

```

[ ]: # save this notebook to submissions folder
import subprocess
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
    ↪ipynb"])

```

```

[ ]: # display(Javascript("IPython.notebook.save_checkpoint();"))
# time.sleep(3)

# subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↪"autogluon_each_location.ipynb"])

```

```

[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
            ↪stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
            ↪strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
    ↪'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', 'hello if hello is
    ↪not None else 'Henrik eller Jørgen'])

```

```

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b',branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m',commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push',
↳ 'origin',branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream',
↳ 'origin',branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])

```

[ ]:

[ ]:

[ ]:

[ ]: