

## autogluon\_each\_location

October 7, 2023

```
[11]: # config
run_analysis = False

[12]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings("ignore")

def fix_datetime(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    ↪with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    # Drop rows where the minute part of the time is not 0
    X = X[X.index.minute == 0]
    return X

def convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_datetime(X_train_observed, "X_train_observed")
    X_train_estimated = fix_datetime(X_train_estimated, "X_train_estimated")
    X_test = fix_datetime(X_test, "X_test")

    # add sample weights, which are 1 for observed and 3 for estimated
    X_train_observed["sample_weight"] = 1
    X_train_estimated["sample_weight"] = 3
    X_test["sample_weight"] = 3

    X_train_observed["estimated_diff_hours"] = 0
```

```

X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index - pd.
↳to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
X_test["estimated_diff_hours"] = (X_test.index - pd.
↳to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

X_train_estimated["estimated_diff_hours"] =
↳X_train_estimated["estimated_diff_hours"].astype('int64')
# the filled once will get dropped later anyways, when we drop y nans
X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].fillna(-50).
↳astype('int64')

X_train_estimated.drop(columns=['date_calc'], inplace=True)
X_test.drop(columns=['date_calc'], inplace=True)

y_train['ds'] = pd.to_datetime(y_train['time'])
y_train.drop(columns=['time'], inplace=True)
y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
↳location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
↳convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # fill missng sample_weight with 3
    #X_train["sample_weight"] = X_train["sample_weight"].fillna(0)

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

    # print number of nans in sample_weight

```

```

    print(f"Number of nans in sample_weight: {X_train['sample_weight'].isna().
↪sum()}")
    # print number of nans in y
    print(f"Number of nans in y: {X_train['y'].isna().sum()}")

    X_train["location"] = location
    X_test["location"] = location

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↪X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

```

Processing location A...
Number of nans in sample_weight: 0
Number of nans in y: 0
Processing location B...
Number of nans in sample_weight: 0
Number of nans in y: 4

```

Processing location C...  
Number of nans in sample\_weight: 0  
Number of nans in y: 6059

## 1 Feature engineering

```
[13]: # temporary
X_train["hour"] = X_train.index.hour
X_train["weekday"] = X_train.index.weekday
# weekday or is_weekend
X_train["is_weekend"] = X_train["weekday"].apply(lambda x: 1 if x >= 5 else 0)

# drop weekday
#X_train.drop(columns=["weekday"], inplace=True)
X_train["month"] = X_train.index.month
X_train["year"] = X_train.index.year

X_test["hour"] = X_test.index.hour
X_test["weekday"] = X_test.index.weekday

# weekday or is_weekend
X_test["is_weekend"] = X_test["weekday"].apply(lambda x: 1 if x >= 5 else 0)

# drop weekday
#X_test.drop(columns=["weekday"], inplace=True)
X_test["month"] = X_test.index.month
X_test["year"] = X_test.index.year

to_drop = ["snow_drift:idx", "snow_density:kgm3"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.dropna(subset=['y'], inplace=True)
X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)
```

```
[14]: import autogluon.eda.auto as auto

if run_analysis:
    auto.dataset_overview(train_data=X_train, test_data=X_test, label="y",
        ↪sample=None)
```

```
[15]: if run_analysis:
    auto.target_analysis(train_data=X_train, label="y")
```

## 2 Starting

```
[16]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    ↪filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

Last submission number: 78  
Now creating submission number: 79  
New filename: submission\_79

```
[17]: from autogluon.tabular import TabularDataset, TabularPredictor
train_data = TabularDataset('X_train_raw.csv')
train_data.drop(columns=['ds'], inplace=True)

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60
presets = 'best_quality'

sample_weight = 'sample_weight' #None
weight_evaluation = True #False
```

Loaded data from: X\_train\_raw.csv | Columns = 53 / 53 | Rows = 92951 -> 92951

```
[18]: predictors = [None, None, None]
```

```
[19]: loc = "A"
print(f"Training model for location {loc}...")
predictor = TabularPredictor(label=label, eval_metric=metric,
    ↪path=f"AutogluonModels/{new_filename}_{loc}", sample_weight=sample_weight,
    ↪weight_evaluation=weight_evaluation).fit(train_data[train_data["location"]
    ↪== loc], time_limit=time_limit, presets=presets)
predictors[0] = predictor
```

Presets specified: ['best\_quality']

```

Stack configuration (auto_stack=True): num_stack_levels=1, num_bag_folds=8,
num_bag_sets=20
Values in column 'sample_weight' used as sample weights instead of predictive
features. Evaluation will report weighted metrics, so ensure same column exists
in test data.
Beginning AutoGluon training ... Time limit = 60s
AutoGluon will save models to "AutogluonModels/submission_79_A/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Darwin
Platform Machine: arm64
Platform Version: Darwin Kernel Version 22.3.0: Mon Jan 30 20:38:37 PST 2023;
root:xnu-8792.81.3~2/RELEASE_ARM64_T6000
Disk Space Avail: 50.11 GB / 494.38 GB (10.1%)
Train Data Rows: 34061
Train Data Columns: 51
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (5733.42, 0.0, 631.01116,
1166.20607)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 3515.74 MB
    Train Data (Original) Memory Usage: 15.33 MB (0.4% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 4 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.

```

```

        These features do not need to be present at inference time.
Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', []) : 6 | ['estimated_diff_hours', 'hour', 'weekday',
'is_weekend', 'month', ...]
Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', []) : 5 | ['estimated_diff_hours', 'hour',
'weekday', 'month', 'year']
    ('int', ['bool']) : 4 | ['is_day:idx', 'is_in_shadow:idx',
'wind_speed_w_1000hPa:ms', 'is_weekend']
0.1s = Fit runtime
48 features in original data used to generate 48 features in processed
data.

Train Data (Processed) Memory Usage: 12.13 MB (0.3% of available memory)
Data preprocessing and feature engineering runtime = 0.13s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'

This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor()
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
AutoGluon will fit 2 stack levels (L1 to L2) ...
Fitting 11 L1 models ...

```

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 39.9s of the 59.87s of remaining time.

Training model for location A...

Not enough time to generate out-of-fold predictions for model. Estimated time required was 389.26s compared to 51.85s of available time.

Time limit exceeded... Skipping KNeighborsUnif\_BAG\_L1.

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 34.14s of the 54.11s of remaining time.

Not enough time to generate out-of-fold predictions for model. Estimated time required was 410.69s compared to 44.36s of available time.

Time limit exceeded... Skipping KNeighborsDist\_BAG\_L1.

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 28.07s of the 48.04s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

2023-10-07 12:41:52,159 ERROR services.py:1169 -- Failed to start the dashboard , return code 1

2023-10-07 12:41:52,160 ERROR services.py:1194 -- Error should be written to 'dashboard.log' or 'dashboard.err'. We are printing the last 20 lines for you. See '<https://docs.ray.io/en/master/ray-observability/ray-logging.html#logging-directory-structure>' to find where the log file is.

2023-10-07 12:41:52,161 ERROR services.py:1204 -- Couldn't read dashboard.log file. Error: [Errno 2] No such file or directory: '/tmp/ray/session\_2023-10-07\_12-41-46\_146066\_8526/logs/dashboard.log'. It means the dashboard is broken even before it initializes the logger (mostly dependency issues). Reading the dashboard.err file which contains stdout/stderr.

2023-10-07 12:41:52,166 ERROR services.py:1238 --

The last 20 lines of

/tmp/ray/session\_2023-10-07\_12-41-46\_146066\_8526/logs/dashboard.err (it contains the error message from the dashboard):

File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/ray/\_\_init\_\_.py", line 63, in \_configure\_system

import grpc # noqa: F401

File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/grpc/\_\_init\_\_.py", line 22, in <module>

from grpc import \_compression

File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/grpc/\_compression.py", line 15, in <module>

from grpc.\_cython import cygrpc

ImportError: dlopen(/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/grpc/\_cython/cygrpc.cpython-310-darwin.so, 0x0002): symbol not found in flat namespace '\_CFRelease'

During handling of the above exception, another exception occurred:

Traceback (most recent call last):

File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-



```

packages/ray/dashboard/dashboard.py", line 10, in <module>
    import ray._private.ray_constants as ray_constants
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/__init__.py", line 101, in <module>
    _configure_system()
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/__init__.py", line 65, in _configure_system
    raise ImportError(
ImportError: Failed to import grpc on Apple Silicon. On Apple Silicon machines,
try `pip uninstall grpcio; conda install grpcio`. Check out
https://docs.ray.io/en/master/ray-overview/installation.html#m1-mac-apple-silicon-support for more details.

```

```

[2023-10-07 12:41:52,506 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:41:53,507 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:41:54,508 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:41:55,510 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:41:56,520 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:41:57,526 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:41:58,535 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:41:59,537 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:42:00,543 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?
[2023-10-07 12:42:01,550 W 8526 6030562] global_state_accessor.cc:389: Some
processes that the driver needs to connect to have not registered with GCS, so
retrying. Have you run 'ray start' on this node?

```

```

Warning: Exception caused LightGBMXT_BAG_L1 to fail during training...
Skipping this model.

```

```

    b'GCS has started but no raylets have registered yet.'

```

```

Detailed Traceback:

```

```

Traceback (most recent call last):

```

```

File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/autogluon/core/trainer/abstract_trainer.py", line 1733, in
_train_and_save
    model = self._train_single(X, y, model, X_val, y_val,
total_resources=total_resources, **model_fit_kwargs)
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/autogluon/core/trainer/abstract_trainer.py", line 1684, in
_train_single
    model = model.fit(X=X, y=y, X_val=X_val, y_val=y_val,
total_resources=total_resources, **model_fit_kwargs)
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/autogluon/core/models/abstract/abstract_model.py", line 829, in fit
    out = self._fit(**kwargs)
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/autogluon/core/models/ensemble/stacker_ensemble_model.py", line 169, in
_fit
    return super()._fit(X=X, y=y, time_limit=time_limit, **kwargs)
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/autogluon/core/models/ensemble/bagged_ensemble_model.py", line 266, in
_fit
    self._fit_folds(
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/autogluon/core/models/ensemble/bagged_ensemble_model.py", line 592, in
_fit_folds
    fold_fitting_strategy.after_all_folds_scheduled()
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/autogluon/core/models/ensemble/fold_fitting_strategy.py", line 508, in
after_all_folds_scheduled
    self.ray.init(**ray_init_args)
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/_private/client_mode_hook.py", line 105, in wrapper
    return func(*args, **kwargs)
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/_private/worker.py", line 1451, in init
    _global_node = ray._private.node.Node(
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/_private/node.py", line 303, in __init__
    node_info = ray._private.services.get_node_to_connect_for_driver(
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/_private/services.py", line 444, in get_node_to_connect_for_driver
    return global_state.get_node_to_connect_for_driver(node_ip_address)
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/_private/state.py", line 752, in get_node_to_connect_for_driver
    node_info_str = self.global_state_accessor.get_node_to_connect_for_driver(
File "python/ray/includes/global_state_accessor.pxi", line 156, in
ray._raylet.GlobalStateAccessor.get_node_to_connect_for_driver
RuntimeError: b'GCS has started but no raylets have registered yet.'
Fitting model: LightGBM_BAG_L1 ... Training model for up to 10.15s of the 30.11s

```

of remaining time.

```
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
2023-10-07 12:42:03,792 ERROR services.py:1169 -- Failed to start the dashboard
, return code 1
2023-10-07 12:42:03,793 ERROR services.py:1194 -- Error should be written to
'dashboard.log' or 'dashboard.err'. We are printing the last 20 lines for you.
See 'https://docs.ray.io/en/master/ray-observability/ray-logging.html#logging-
directory-structure' to find where the log file is.
2023-10-07 12:42:03,793 ERROR services.py:1204 -- Couldn't read dashboard.log
file. Error: [Errno 2] No such file or directory:
'/tmp/ray/session_2023-10-07_12-42-02_626935_8526/logs/dashboard.log'. It means
the dashboard is broken even before it initializes the logger (mostly dependency
issues). Reading the dashboard.err file which contains stdout/stderr.
2023-10-07 12:42:03,799 ERROR services.py:1238 --
The last 20 lines of
/tmp/ray/session_2023-10-07_12-42-02_626935_8526/logs/dashboard.err (it contains
the error message from the dashboard):
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/__init__.py", line 63, in _configure_system
    import grpc # noqa: F401
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/grpc/__init__.py", line 22, in <module>
    from grpc import _compression
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/grpc/_compression.py", line 15, in <module>
    from grpc._cython import cygrpc
ImportError: dlopen(/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/grpc/_cython/cygrpc.cpython-310-darwin.so, 0x0002): symbol not found in
flat namespace '_CFRelease'
```

During handling of the above exception, another exception occurred:

```
Traceback (most recent call last):
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/dashboard/dashboard.py", line 10, in <module>
    import ray._private.ray_constants as ray_constants
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/__init__.py", line 101, in <module>
    _configure_system()
File "/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-
packages/ray/__init__.py", line 65, in _configure_system
    raise ImportError(
ImportError: Failed to import grpc on Apple Silicon. On Apple Silicon machines,
try `pip uninstall grpcio; conda install grpcio`. Check out
https://docs.ray.io/en/master/ray-overview/installation.html#m1-mac-apple-
silicon-support for more details.
```

The Kernel crashed while executing code in the the current cell or a previous cell. Please review the code in the cell(s) to identify a possible cause of the failure. Click

```
[ ]: loc = "B"
print(f"Training model for location {loc}...")
predictor = TabularPredictor(label=label, eval_metric=metric,
    path=f"AutogluonModels/{new_filename}_{loc}", sample_weight=sample_weight,
    weight_evaluation=weight_evaluation).fit(train_data[train_data["location"]
    == loc], time_limit=time_limit, presets=presets)
predictors[1] = predictor
```

Warning: path already exists! This predictor may overwrite an existing predictor! path="AutogluonModels/submission\_79\_jorge\_B"

Presets specified: ['best\_quality']

Stack configuration (auto\_stack=True): num\_stack\_levels=1, num\_bag\_folds=8, num\_bag\_sets=20

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 60s

AutoGluon will save models to "AutogluonModels/submission\_79\_jorge\_B/"

AutoGluon Version: 0.8.1

Python Version: 3.10.12

Operating System: Darwin

Platform Machine: arm64

Platform Version: Darwin Kernel Version 22.1.0: Sun Oct 9 20:15:09 PDT 2022; root:xnu-8792.41.9~2/RELEASE\_ARM64\_T6000

Disk Space Avail: 15.97 GB / 494.38 GB (3.2%)

Train Data Rows: 32819

Train Data Columns: 51

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (1152.3, -0.0, 96.89334, 194.00409)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 4394.88 MB

Train Data (Original) Memory Usage: 14.77 MB (0.3% of available memory)

Inferring data type of each feature based on column values. Set feature\_metadata\_in to manually specify special dtypes of the features.

```

Stage 1 Generators:
    Fitting AsTypeFeatureGenerator...
    Note: Converting 4 features to boolean dtype as they
only contain 2 unique values.
Stage 2 Generators:
    Fitting FillNaFeatureGenerator...
Stage 3 Generators:
    Fitting IdentityFeatureGenerator...
Stage 4 Generators:
    Fitting DropUniqueFeatureGenerator...
Stage 5 Generators:
    Fitting DropDuplicatesFeatureGenerator...
Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
    This is typically a feature which has the same value for all
rows.
    These features do not need to be present at inference time.
Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', [])   : 6 | ['estimated_diff_hours', 'hour', 'weekday',
'is_weekend', 'month', ...]
Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', [])   : 5 | ['estimated_diff_hours', 'hour',
'weekday', 'month', 'year']
    ('int', ['bool']) : 4 | ['is_day:idx', 'is_in_shadow:idx',
'wind_speed_w_1000hPa:ms', 'is_weekend']
0.1s = Fit runtime
48 features in original data used to generate 48 features in processed
data.
Train Data (Processed) Memory Usage: 11.68 MB (0.3% of available memory)
Data preprocessing and feature engineering runtime = 0.16s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge',
    'CAT': {},

```

```

    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 39.88s of the 59.84s of remaining time.

Training model for location B...

Not enough time to generate out-of-fold predictions for model. Estimated time required was 160.16s compared to 51.82s of available time.

Time limit exceeded... Skipping KNeighborsUnif\_BAG\_L1.

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 37.4s of the 57.35s of remaining time.

Not enough time to generate out-of-fold predictions for model. Estimated time required was 133.14s compared to 48.59s of available time.

Time limit exceeded... Skipping KNeighborsDist\_BAG\_L1.

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 35.32s of the 55.28s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-23.1587 = Validation score (-mean\_absolute\_error)

24.39s = Training runtime

65.15s = Validation runtime

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 0.11s of the 20.06s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-116.8768 = Validation score (-mean\_absolute\_error)

1.21s = Training runtime

0.02s = Validation runtime

Completed 1/20 k-fold bagging repeats ...

Fitting model: WeightedEnsemble\_L2 ... Training model for up to 59.84s of the 16.45s of remaining time.

-23.1587 = Validation score (-mean\_absolute\_error)

0.12s = Training runtime

```

0.0s      = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 16.32s of the
16.3s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.9002      = Validation score    (-mean_absolute_error)
    3.87s        = Training    runtime
    1.0s         = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 10.01s of the 10.0s
of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.3768      = Validation score    (-mean_absolute_error)
    1.57s        = Training    runtime
    0.18s        = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ... Training model for up to 6.63s of the
6.62s of remaining time.
    -20.3273      = Validation score    (-mean_absolute_error)
    25.22s       = Training    runtime
    0.86s        = Validation runtime
Completed 1/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L3 ... Training model for up to 59.84s of the
-19.69s of remaining time.
    -20.3273      = Validation score    (-mean_absolute_error)
    0.2s         = Training    runtime
    0.0s         = Validation runtime
AutoGluon training complete, total runtime = 79.91s ... Best model:
"WeightedEnsemble_L3"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_79_jorge_B/")

```

```

[ ]: loc = "C"
print(f"Training model for location {loc}...")
predictor = TabularPredictor(label=label, eval_metric=metric,
    ↪path=f"AutogluonModels/{new_filename}_{loc}", sample_weight=sample_weight,
    ↪weight_evaluation=weight_evaluation).fit(train_data[train_data["location"]
    ↪== loc], time_limit=time_limit, presets=presets)
predictors[2] = predictor

```

```

Warning: path already exists! This predictor may overwrite an existing
predictor! path="AutogluonModels/submission_79_jorge_C"
Presets specified: ['best_quality']
Stack configuration (auto_stack=True): num_stack_levels=1, num_bag_folds=8,
num_bag_sets=20
Values in column 'sample_weight' used as sample weights instead of predictive
features. Evaluation will report weighted metrics, so ensure same column exists
in test data.

```

```

Beginning AutoGluon training ... Time limit = 60s
AutoGluon will save models to "AutogluonModels/submission_79_jorge_C/"
AutoGluon Version: 0.8.1
Python Version: 3.10.12
Operating System: Darwin
Platform Machine: arm64
Platform Version: Darwin Kernel Version 22.1.0: Sun Oct 9 20:15:09 PDT 2022;
root:xnu-8792.41.9~2/RELEASE_ARM64_T6000
Disk Space Avail: 15.55 GB / 494.38 GB (3.1%)
Train Data Rows: 26071
Train Data Columns: 51
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and label-values can't be converted to int).
    Label info (max, min, mean, stddev): (999.6, -0.0, 77.70004, 165.87752)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 4399.47 MB
    Train Data (Original) Memory Usage: 11.73 MB (0.3% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 3 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 6 | ['estimated_diff_hours', 'hour', 'weekday',

```



```

'is_weekend', 'month', ...]
    Types of features in processed data (raw dtype, special dtypes):
        ('float', [])      : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', [])        : 5 | ['estimated_diff_hours', 'hour',
'weekday', 'month', 'year']
        ('int', ['bool'])  : 3 | ['is_day:idx', 'is_in_shadow:idx',
'is_weekend']
    0.1s = Fit runtime
    48 features in original data used to generate 48 features in processed
data.

    Train Data (Processed) Memory Usage: 9.46 MB (0.2% of available memory)
Data preprocessing and feature engineering runtime = 0.14s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'

    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.

    To change this, specify the eval_metric parameter of Predictor()
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}},
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}},
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
AutoGluon will fit 2 stack levels (L1 to L2) ...
Fitting 11 L1 models ...
Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 39.9s of the
59.86s of remaining time.

Training model for location C...

    Not enough time to generate out-of-fold predictions for model. Estimated
time required was 105.14s compared to 51.84s of available time.

```

Time limit exceeded... Skipping KNeighborsUnif\_BAG\_L1.  
Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 37.83s of the 57.79s of remaining time.  
Not enough time to generate out-of-fold predictions for model. Estimated time required was 57.03s compared to 49.15s of available time.  
Time limit exceeded... Skipping KNeighborsDist\_BAG\_L1.  
Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 36.69s of the 56.65s of remaining time.  
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy  
-15.6035 = Validation score (-mean\_absolute\_error)  
19.08s = Training runtime  
46.76s = Validation runtime  
Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 8.83s of the 28.79s of remaining time.  
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy  
-17.042 = Validation score (-mean\_absolute\_error)  
8.57s = Training runtime  
14.81s = Validation runtime  
Completed 1/20 k-fold bagging repeats ...  
Fitting model: WeightedEnsemble\_L2 ... Training model for up to 59.86s of the 16.36s of remaining time.  
-15.5646 = Validation score (-mean\_absolute\_error)  
0.1s = Training runtime  
0.0s = Validation runtime  
Fitting 9 L2 models ...  
Fitting model: LightGBMXT\_BAG\_L2 ... Training model for up to 16.26s of the 16.25s of remaining time.  
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy  
-16.1537 = Validation score (-mean\_absolute\_error)  
1.9s = Training runtime  
0.42s = Validation runtime  
Fitting model: LightGBM\_BAG\_L2 ... Training model for up to 12.14s of the 12.13s of remaining time.  
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy  
-15.8514 = Validation score (-mean\_absolute\_error)  
1.43s = Training runtime  
0.13s = Validation runtime  
Fitting model: RandomForestMSE\_BAG\_L2 ... Training model for up to 8.85s of the 8.84s of remaining time.  
-15.5697 = Validation score (-mean\_absolute\_error)  
17.87s = Training runtime  
0.55s = Validation runtime  
Completed 1/20 k-fold bagging repeats ...  
Fitting model: WeightedEnsemble\_L3 ... Training model for up to 59.86s of the

```

-9.78s of remaining time.
    -15.4504      = Validation score    (-mean_absolute_error)
    0.14s        = Training    runtime
    0.0s         = Validation runtime
AutoGluon training complete, total runtime = 69.94s ... Best model:
"WeightedEnsemble_L3"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_79_jorge_C/")

```

### 3 Submit

```

[ ]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data

```

```

Loaded data from: X_train_raw.csv | Columns = 53 / 53 | Rows = 92951 -> 92951
Loaded data from: X_test_raw.csv | Columns = 52 / 52 | Rows = 2160 -> 2160

```

```

[ ]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
↪ "location"], left_on=["ds", "location"])

#test_data_merged

```

```

Loaded data from: test.csv | Columns = 4 / 4 | Rows = 2160 -> 2160

```

```

[ ]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in test_data.groupby('location'):
    i = location_map[loc]
    subset = test_data_merged[test_data_merged["location"] == loc].
↪reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred


```



```
predictions.append(subset)
```

```
[ ]: if run_analysis:
    # plot predictions for location A, in addition to train data for A
    for loc, idx in location_map.items():
        fig, ax = plt.subplots(figsize=(20, 10))
        # plot train data
        train_data_with_dates[train_data_with_dates["location"]==loc].
        plot(x='ds', y='y', ax=ax, label="train data")

        # plot predictions
        predictions[idx].plot(x='ds', y='prediction', ax=ax,
        label="predictions")

        # title
        ax.set_title(f"Predictions for location {loc}")
```





```
[ ]: # concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df
```

```
[ ]:      id  prediction
0      0    1.353503
1      1    1.377337
2      2    1.535772
3      3   50.274967
```

```

4          4  298.462311
..      ...      ...
715  2155    91.790527
716  2156    58.327251
717  2157    25.278416
718  2158     3.892292
719  2159     2.036365

```

[2160 rows x 2 columns]

```

[ ]: # Save the submission DataFrame to submissions folder, create new name based on
      ↳ last submission, format is submission_<last_submission_number + 1>.csv

      # Save the submission
      print(f"Saving submission to submissions/{new_filename}.csv")
      submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
      ↳ index=False)

```

Saving submission to submissions/submission\_79\_jorge.csv

```

[ ]: # save this notebook to submissions folder
      import subprocess
      import os
      subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
      ↳ join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
      ↳ ipynb"])

```

```

[NbConvertApp] Converting notebook autogluon_each_location.ipynb to pdf
[NbConvertApp] Support files will be in notebook_pdfs/submission_79_jorge_files/
[NbConvertApp] Making directory
./notebook_pdfs/submission_79_jorge_files/notebook_pdfs
[NbConvertApp] Writing 152138 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 1471956 bytes to notebook_pdfs/submission_79_jorge.pdf

```

```

[ ]: CompletedProcess(args=['jupyter', 'nbconvert', '--to', 'pdf', '--output',
'notebook_pdfs/submission_79_jorge.pdf', 'autogluon_each_location.ipynb'],
returncode=0)

```

```

[ ]: # feature importance
      location="A"
      split_time = pd.Timestamp("2022-10-28 22:00:00")
      estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]

```

```

estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
↳time_limit=60*10)

```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'sample\_weight', 'location']

Computing feature importance via permutation shuffling for 48 features using 4394 rows with 10 shuffle sets... Time limit: 600s...

3094.63s = Expected runtime (309.46s per shuffle set)

419.07s = Actual runtime (Completed 2 of 10 shuffle sets) (Early stopping due to lack of time...)

```

[ ]:

```

	importance	stddev	p_value	n	p99_high \
direct_rad:W	179.527752	0.089041	0.000112	2	183.535682
clear_sky_rad:W	102.034121	0.763313	0.001684	2	136.392434
diffuse_rad:W	76.319510	0.242082	0.000714	2	87.216140
sun_elevation:d	48.306937	1.434587	0.006683	2	112.880735
hour	37.283222	0.841950	0.005082	2	75.181160
sun_azimuth:d	34.812837	0.901175	0.005826	2	75.376611
direct_rad_1h:J	28.456567	0.184929	0.001463	2	36.780617
cloud_base_agl:m	23.171833	0.354225	0.003441	2	39.116231
clear_sky_energy_1h:J	22.080754	0.614640	0.006264	2	49.747021
total_cloud_cover:p	21.249354	0.411635	0.004360	2	39.777903
effective_cloud_cover:p	17.259895	0.014190	0.000185	2	17.898613
month	16.052628	0.649876	0.009110	2	45.304928
ceiling_height_agl:m	15.125792	0.676116	0.010058	2	45.559223
diffuse_rad_1h:J	14.148167	0.011894	0.000189	2	14.683522
relative_humidity_1000hPa:p	13.232798	0.297330	0.005057	2	26.616249
is_day:idx	12.805033	0.807725	0.014188	2	49.162422
wind_speed_u_10m:ms	12.050327	0.464144	0.008667	2	32.942436
weekday	11.374554	0.931762	0.018417	2	53.315108
is_in_shadow:idx	10.136370	0.018201	0.000404	2	10.955648
msl_pressure:hPa	9.303426	0.654040	0.015810	2	38.743145
t_1000hPa:K	8.931381	0.362485	0.009132	2	25.247605
visibility:m	8.547172	0.518337	0.013641	2	31.878598
is_weekend	8.137093	0.320401	0.008860	2	22.559020
wind_speed_10m:ms	7.407126	0.125153	0.003803	2	13.040537
sfc_pressure:hPa	7.345058	0.528305	0.016175	2	31.125179
pressure_100m:hPa	6.942622	0.545800	0.017677	2	31.510253
pressure_50m:hPa	6.667661	0.192459	0.006496	2	15.330645
wind_speed_v_10m:ms	6.375395	0.245266	0.008657	2	17.415344
fresh_snow_24h:cm	5.625596	0.543399	0.021708	2	30.085121
dew_point_2m:K	4.835024	0.656765	0.030480	2	34.397422
estimated_diff_hours	4.296988	0.021430	0.001123	2	5.261618
snow_water:kgm2	4.248560	0.146461	0.007758	2	10.841074
precip_type_5min:idx	2.011457	0.735394	0.080526	2	35.113079
air_density_2m:kgm3	1.468686	0.179319	0.027413	2	9.540235

fresh_snow_12h:cm	1.399083	0.050244	0.008081	2	3.660664
absolute_humidity_2m:gm3	1.260919	0.006155	0.001099	2	1.537983
super_cooled_liquid_water:kgm2	0.855587	0.032606	0.008576	2	2.323244
snow_depth:cm	0.446138	0.040745	0.020528	2	2.280167
precip_5min:mm	0.422909	0.164053	0.085216	2	7.807265
fresh_snow_6h:cm	0.291019	0.010971	0.008483	2	0.784855
dew_or_rime:idx	0.104030	0.001518	0.003285	2	0.172374
prob_rime:p	0.038235	0.014263	0.082092	2	0.680243
snow_melt_10min:mm	0.016844	0.102404	0.427249	2	4.626258
fresh_snow_1h:cm	0.012910	0.006818	0.113768	2	0.319820
fresh_snow_3h:cm	0.002094	0.028007	0.466464	2	1.262740
wind_speed_w_1000hPa:ms	0.000000	0.000000	0.500000	2	0.000000
rain_water:kgm2	-0.032241	0.039702	0.771960	2	1.754814
year	-0.048251	0.020571	0.906799	2	0.877698

	p99_low
direct_rad:W	175.519822
clear_sky_rad:W	67.675808
diffuse_rad:W	65.422880
sun_elevation:d	-16.266861
hour	-0.614717
sun_azimuth:d	-5.750937
direct_rad_1h:J	20.132516
cloud_base_agl:m	7.227436
clear_sky_energy_1h:J	-5.585513
total_cloud_cover:p	2.720806
effective_cloud_cover:p	16.621176
month	-13.199671
ceiling_height_agl:m	-15.307639
diffuse_rad_1h:J	13.612812
relative_humidity_1000hPa:p	-0.150654
is_day:idx	-23.552357
wind_speed_u_10m:ms	-8.841782
weekday	-30.566001
is_in_shadow:idx	9.317091
msl_pressure:hPa	-20.136293
t_1000hPa:K	-7.384843
visibility:m	-14.784253
is_weekend	-6.284834
wind_speed_10m:ms	1.773715
sfc_pressure:hPa	-16.435063
pressure_100m:hPa	-17.625010
pressure_50m:hPa	-1.995323
wind_speed_v_10m:ms	-4.664554
fresh_snow_24h:cm	-18.833928
dew_point_2m:K	-24.727373
estimated_diff_hours	3.332359



```

snow_water:kgm2                -2.343955
precip_type_5min:idx           -31.090164
air_density_2m:kgm3            -6.602862
fresh_snow_12h:cm              -0.862498
absolute_humidity_2m:gm3       0.983856
super_cooled_liquid_water:kgm2 -0.612070
snow_depth:cm                  -1.387891
precip_5min:mm                 -6.961447
fresh_snow_6h:cm               -0.202818
dew_or_rime:idx                0.035686
prob_rime:p                    -0.603772
snow_melt_10min:mm            -4.592570
fresh_snow_1h:cm               -0.294000
fresh_snow_3h:cm               -1.258551
wind_speed_w_1000hPa:ms        0.000000
rain_water:kgm2                -1.819296
year                           -0.974200

```

```

[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictor.feature_importance(feature_stage="original", data=observed,
    ↪time_limit=60*10)

```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'sample\_weight', 'location']

Computing feature importance via permutation shuffling for 48 features using 5000 rows with 10 shuffle sets... Time limit: 600s...

4293.33s = Expected runtime (429.33s per shuffle set)

359.81s = Actual runtime (Completed 1 of 10 shuffle sets) (Early stopping due to lack of time...)

```

[ ]:
importance  stddev  p_value  n  p99_high  \
clear_sky_rad:W      34.800616    NaN    NaN  1      NaN
sun_elevation:d      21.053896    NaN    NaN  1      NaN
clear_sky_energy_1h:J 17.959037    NaN    NaN  1      NaN
direct_rad:W         15.484854    NaN    NaN  1      NaN
diffuse_rad:W         8.571135    NaN    NaN  1      NaN
direct_rad_1h:J       3.283643    NaN    NaN  1      NaN
relative_humidity_1000hPa:p 1.839748    NaN    NaN  1      NaN
sun_azimuth:d         1.549566    NaN    NaN  1      NaN
hour                0.997634    NaN    NaN  1      NaN
wind_speed_v_10m:ms   0.907447    NaN    NaN  1      NaN
msl_pressure:hPa      0.676269    NaN    NaN  1      NaN
snow_water:kgm2       0.627690    NaN    NaN  1      NaN
is_day:idx            0.620870    NaN    NaN  1      NaN
precip_type_5min:idx  0.312266    NaN    NaN  1      NaN
wind_speed_10m:ms     0.258658    NaN    NaN  1      NaN

```

pressure_100m:hPa	0.258171	NaN	NaN	1	NaN
ceiling_height_agl:m	0.255112	NaN	NaN	1	NaN
sfc_pressure:hPa	0.224268	NaN	NaN	1	NaN
pressure_50m:hPa	0.223026	NaN	NaN	1	NaN
precip_5min:mm	0.207318	NaN	NaN	1	NaN
snow_depth:cm	0.079170	NaN	NaN	1	NaN
air_density_2m:kgm3	0.078937	NaN	NaN	1	NaN
fresh_snow_12h:cm	0.077565	NaN	NaN	1	NaN
effective_cloud_cover:p	0.044103	NaN	NaN	1	NaN
snow_melt_10min:mm	0.018210	NaN	NaN	1	NaN
fresh_snow_6h:cm	0.014053	NaN	NaN	1	NaN
fresh_snow_3h:cm	0.004762	NaN	NaN	1	NaN
estimated_diff_hours	0.000000	NaN	NaN	1	NaN
wind_speed_w_1000hPa:ms	-0.000056	NaN	NaN	1	NaN
prob_rime:p	-0.000317	NaN	NaN	1	NaN
fresh_snow_1h:cm	-0.001555	NaN	NaN	1	NaN
dew_or_rime:idx	-0.003530	NaN	NaN	1	NaN
fresh_snow_24h:cm	-0.015575	NaN	NaN	1	NaN
rain_water:kgm2	-0.020746	NaN	NaN	1	NaN
year	-0.021817	NaN	NaN	1	NaN
super_cooled_liquid_water:kgm2	-0.096266	NaN	NaN	1	NaN
cloud_base_agl:m	-0.115793	NaN	NaN	1	NaN
t_1000hPa:K	-0.135829	NaN	NaN	1	NaN
wind_speed_u_10m:ms	-0.223229	NaN	NaN	1	NaN
visibility:m	-0.255399	NaN	NaN	1	NaN
is_weekend	-0.301423	NaN	NaN	1	NaN
weekday	-0.339798	NaN	NaN	1	NaN
total_cloud_cover:p	-0.380572	NaN	NaN	1	NaN
diffuse_rad_1h:J	-0.514786	NaN	NaN	1	NaN
absolute_humidity_2m:gm3	-0.560996	NaN	NaN	1	NaN
month	-0.563007	NaN	NaN	1	NaN
is_in_shadow:idx	-0.585143	NaN	NaN	1	NaN
dew_point_2m:K	-1.026145	NaN	NaN	1	NaN

#### p99\_low

clear_sky_rad:W	NaN
sun_elevation:d	NaN
clear_sky_energy_1h:J	NaN
direct_rad:W	NaN
diffuse_rad:W	NaN
direct_rad_1h:J	NaN
relative_humidity_1000hPa:p	NaN
sun_azimuth:d	NaN
hour	NaN
wind_speed_v_10m:ms	NaN
msl_pressure:hPa	NaN
snow_water:kgm2	NaN

is_day:idx	NaN
precip_type_5min:idx	NaN
wind_speed_10m:ms	NaN
pressure_100m:hPa	NaN
ceiling_height_agl:m	NaN
sfc_pressure:hPa	NaN
pressure_50m:hPa	NaN
precip_5min:mm	NaN
snow_depth:cm	NaN
air_density_2m:kgm3	NaN
fresh_snow_12h:cm	NaN
effective_cloud_cover:p	NaN
snow_melt_10min:mm	NaN
fresh_snow_6h:cm	NaN
fresh_snow_3h:cm	NaN
estimated_diff_hours	NaN
wind_speed_w_1000hPa:ms	NaN
prob_rime:p	NaN
fresh_snow_1h:cm	NaN
dew_or_rime:idx	NaN
fresh_snow_24h:cm	NaN
rain_water:kgm2	NaN
year	NaN
super_cooled_liquid_water:kgm2	NaN
cloud_base_agl:m	NaN
t_1000hPa:K	NaN
wind_speed_u_10m:ms	NaN
visibility:m	NaN
is_weekend	NaN
weekday	NaN
total_cloud_cover:p	NaN
diffuse_rad_1h:J	NaN
absolute_humidity_2m:gm3	NaN
month	NaN
is_in_shadow:idx	NaN
dew_point_2m:K	NaN

```
[ ]: subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↪"autogluon_each_location.ipynb"])
```

```
[NbConvertApp] Converting notebook autogluon_each_location.ipynb to pdf
[NbConvertApp] Support files will be in
notebook_pdfs/submission_79_jorge_with_feature_importance_files/
[NbConvertApp] Making directory
./notebook_pdfs/submission_79_jorge_with_feature_importance_files/notebook_pdfs
[NbConvertApp] Writing 152954 bytes to notebook.tex
[NbConvertApp] Building PDF
```

```
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 1471953 bytes to
notebook_pdfs/submission_79_jorge_with_feature_importance.pdf
```

```
[ ]: CompletedProcess(args=['jupyter', 'nbconvert', '--to', 'pdf', '--output',
'notebook_pdfs/submission_79_jorge_with_feature_importance.pdf',
'autogluon_each_location.ipynb'], returncode=0)
```

```
[ ]: import subprocess

def execute_git_command(directory, command):
    """Execute a Git command in the specified directory."""
    try:
        result = subprocess.check_output(['git', '-C', directory] + command,
        ↪stderr=subprocess.STDOUT)
        return result.decode('utf-8').strip(), True
    except subprocess.CalledProcessError as e:
        print(f"Git command failed with message: {e.output.decode('utf-8')}.
        ↪strip()}"
        return e.output.decode('utf-8').strip(), False

git_repo_path = "."

execute_git_command(git_repo_path, ['config', 'user.email', 'henrikskog01@gmail.
        ↪com'])
execute_git_command(git_repo_path, ['config', 'user.name', hello if hello is
        ↪not None else 'Henrik eller Jørgen'])

branch_name = new_filename

# add datetime to branch name
branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

commit_msg = "run result"

execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# Navigate to your repo and commit changes
execute_git_command(git_repo_path, ['add', '.'])
execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# Push to remote
```

```

output, success = execute_git_command(git_repo_path, ['push',
↳'origin',branch_name])

# If the push fails, try setting an upstream branch and push again
if not success and 'upstream' in output:
    print("Attempting to set upstream and push again...")
    execute_git_command(git_repo_path, ['push', '--set-upstream',
↳'origin',branch_name])
    execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

execute_git_command(git_repo_path, ['checkout', 'main'])

```

```

[ ]: ('Switched to branch \'main\'
Your branch is behind \'origin/main\' by 2
commits, and can be fast-forwarded.
(use "git pull" to update your local
branch)',
True)

```