autogluon_each_location

October 9, 2023

```python
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*10
presets = 'best_quality'

do_drop_ds = True

use_groups = False
n_groups = 8

auto_stack = True
num_stack_levels = 1
num_bag_folds = 0
if auto_stack:
    num_stack_levels = None
    num_bag_folds = None

use_tune_data = False
use_test_data = True
tune_and_test_length = 24*30*3 # 3 months from end, this changes the␣
 ↪evaluations for only test
holdout_frac = None
use_bag_holdout = False # Enable this if there is a large gap between score_val␣
 ↪and score_test in stack models.

sample_weight = 'sample_weight' #None
weight_evaluation = True #False
sample_weight_estimated = 1 # this changes evaluations for test and tune WTF,␣
 ↪cant find a fix

run_analysis = False
```

```python
[2]: import pandas as pd
import numpy as np
```

```python
import warnings
warnings.filterwarnings("ignore")

def fix_datetime(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    ↪with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    # Drop rows where the minute part of the time is not 0
    X = X[X.index.minute == 0].copy()
    return X



def convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_datetime(X_train_observed, "X_train_observed")
    X_train_estimated = fix_datetime(X_train_estimated, "X_train_estimated")
    X_test = fix_datetime(X_test, "X_test")

    # add sample weights, which are 1 for observed and 3 for estimated
    X_train_observed["sample_weight"] = 1
    X_train_estimated["sample_weight"] = sample_weight_estimated
    X_test["sample_weight"] = sample_weight_estimated

    X_train_observed["estimated_diff_hours"] = 0
    X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index - pd.
    ↪to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
    X_test["estimated_diff_hours"] = (X_test.index - pd.
    ↪to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

    X_train_estimated["estimated_diff_hours"] =␣
    ↪X_train_estimated["estimated_diff_hours"].astype('int64')
    # the filled once will get dropped later anyways, when we drop y nans
    X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].fillna(-50).
    ↪astype('int64')


    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train['ds'] = pd.to_datetime(y_train['time'])
    y_train.drop(columns=['time'], inplace=True)
```

```python
        y_train.sort_values(by='ds', inplace=True)
        y_train.set_index('ds', inplace=True)

        return X_train_observed, X_train_estimated, X_test, y_train




def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
↪location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
↪convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # fill missng sample_weight with 3
    #X_train["sample_weight"] = X_train["sample_weight"].fillna(0)


    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↪right_index=True)

    # print number of nans in sample_weight
    print(f"Number of nans in sample_weight: {X_train['sample_weight'].isna().
↪sum()}")
    # print number of nans in y
    print(f"Number of nans in y: {X_train['y'].isna().sum()}")


    X_train["location"] = location
    X_test["location"] = location

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
```

```python
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed= pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,␣
 ↪X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)
```

```
Processing location A…
Number of nans in sample_weight: 0
Number of nans in y: 0
Processing location B…
Number of nans in sample_weight: 0
Number of nans in y: 4
Processing location C…
Number of nans in sample_weight: 0
Number of nans in y: 6059
```

1 Feature enginering

```python
import numpy as np
import pandas as pd

X_train.dropna(subset=['y'], inplace=True)

if not do_drop_ds:
    # add hour datetime feature
    X_train["hour"] = X_train.index.hour
    X_test["hour"] = X_test.index.hour

#print(X_train.head())
```

```python
if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique()  # Assuming 'location' is the name↵
 ↪of the column representing locations

    grouped_dfs = []  # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()

        # Calculate the size of each group for this location
        group_size = len(loc_df) // n_groups

        # Create a new 'group' column for this location
        loc_df['group'] = np.repeat(range(n_groups),↵
 ↪repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

        # Append to list of grouped DataFrames
        grouped_dfs.append(loc_df)

    # Concatenate all the grouped DataFrames back together
    X_train = pd.concat(grouped_dfs)
    X_train.sort_index(inplace=True)
    print(X_train["group"].head())




to_drop = ["snow_drift:idx", "snow_density:kgm3"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)
```

```python
[4]: from autogluon.tabular import TabularDataset, TabularPredictor
     from autogluon.timeseries import TimeSeriesDataFrame
     import numpy as np
     train_data = TabularDataset('X_train_raw.csv')
```

```python
# set group column of train_data be increasing from 0 to 7 based on time, the
 ↪first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
train_data['ds'] = pd.to_datetime(train_data['ds'])
train_data = train_data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())


# get end date of train data and subtract 3 months
split_time = pd.to_datetime(train_data["ds"]).max() - pd.
 ↪Timedelta(hours=tune_and_test_length)
train_set = TabularDataset(train_data[train_data["ds"] < split_time])
test_set = TabularDataset(train_data[train_data["ds"] >= split_time])
if use_groups:
    test_set = test_set.drop(columns=['group'])

if do_drop_ds:
    train_set = train_set.drop(columns=['ds'])
    test_set = test_set.drop(columns=['ds'])
    train_data = train_data.drop(columns=['ds'])


def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
 ↪loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

tuning_data = None
if use_tune_data:
    train_data = train_set
    if use_test_data:
        # split test_set in half, use first half for tuning
        tuning_data, test_data = [], []
        for loc in locations:
            loc_test_set = test_set[test_set["location"] == loc]
            loc_tuning_data = loc_test_set.iloc[:len(loc_test_set)//2]
            loc_test_data = loc_test_set.iloc[len(loc_test_set)//2:]
            tuning_data.append(loc_tuning_data)
            test_data.append(loc_test_data)
        tuning_data = pd.concat(tuning_data)
        test_data = pd.concat(test_data)
```

```
        print("Shapes of tuning and test", tuning_data.shape[0], test_data.
 ↪shape[0], tuning_data.shape[0] + test_data.shape[0])

    else:
        tuning_data = test_set
        print("Shape of tuning", tuning_data.shape[0])

    # ensure sample weights for your tuning data sum to the number of rows in
 ↪the tuning data.
    tuning_data = normalize_sample_weights_per_location(tuning_data)


else:
    if use_test_data:
        train_data = train_set
        test_data = test_set
        print("Shape of test", test_data.shape[0])

# ensure sample weights for your training (or tuning) data sum to the number of
 ↪rows in the training (or tuning) data.
train_data = normalize_sample_weights_per_location(train_data)
if use_test_data:
    test_data = normalize_sample_weights_per_location(test_data)
```

Shape of test 5791

```
[5]: if run_analysis:
        import autogluon.eda.auto as auto
        auto.dataset_overview(train_data=train_data, test_data=test_data,
     ↪label="y", sample=None)
```

```
[6]: if run_analysis:
        auto.target_analysis(train_data=train_data, label="y")
```

2 Starting

```
[7]: import os


     # Get the last submission number
     last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
      ↪filename in os.listdir('submissions') if "submission" in filename]))
     print("Last submission number:", last_submission_number)
     print("Now creating submission number:", last_submission_number + 1)

     # Create the new filename
     new_filename = f'submission_{last_submission_number + 1}'
```

```python
hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

```
Last submission number: 79
Now creating submission number: 80
New filename: submission_80
```

[8]:
```python
predictors = [None, None, None]
```

[9]:
```python
def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪train and tune data and test data
    print("Train data sample weight sum:", train_data[train_data["location"] ==
    ↪loc]["sample_weight"].sum())
    print("Train data number of rows:", train_data[train_data["location"] ==
    ↪loc].shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
    ↪tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:", tuning_data[tuning_data["location"]
    ↪== loc].shape[0])
    if use_test_data:
        print("Test data sample weight sum:", test_data[test_data["location"]
    ↪== loc]["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"] ==
    ↪loc].shape[0])
    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        sample_weight=sample_weight,
        weight_evaluation=weight_evaluation,
        groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc],
        time_limit=time_limit,
        #presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2,# just put
    ↪somethin, will be overwritten anyways
        tuning_data=tuning_data[tuning_data["location"] == loc] if
    ↪use_tune_data else None,
```

```
            use_bag_holdout=use_bag_holdout,
            holdout_frac=holdout_frac,
        )

        # evaluate on test data
        if use_test_data:
            # drop sample_weight column
            t = test_data[test_data["location"] == loc]#.
    ↪drop(columns=["sample_weight"])
            perf = predictor.evaluate(t)
            print("Evaluation on test data:")
            print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)
```

Values in column 'sample_weight' used as sample weights instead of predictive
features. Evaluation will report weighted metrics, so ensure same column exists
in test data.
Beginning AutoGluon training … Time limit = 600s
AutoGluon will save models to "AutogluonModels/submission_80_A/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.191-1 (2023-08-16)
Disk Space Avail: 311.54 GB / 315.93 GB (98.6%)
Train Data Rows: 31900
Train Data Columns: 46
Label Column: y
Preprocessing data …
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).

Training model for location A…
Train data sample weight sum: 31900
Train data number of rows: 31900
Test data sample weight sum: 2161
Test data number of rows: 2161

        Label info (max, min, mean, stddev): (5733.42, 0.0, 633.132, 1165.64686)
        If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data …
Fitting AutoMLPipelineFeatureGenerator…
        Available Memory: 132448.77 MB

9

```
Train Data (Original)  Memory Usage: 13.08 MB (0.0% of available memory)
        Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
        Stage 1 Generators:
                Fitting AsTypeFeatureGenerator…
                        Note: Converting 3 features to boolean dtype as they
only contain 2 unique values.
        Stage 2 Generators:
                Fitting FillNaFeatureGenerator…
        Stage 3 Generators:
                Fitting IdentityFeatureGenerator…
        Stage 4 Generators:
                Fitting DropUniqueFeatureGenerator…
        Stage 5 Generators:
                Fitting DropDuplicatesFeatureGenerator…
        Useless Original Features (Count: 2): ['elevation:m', 'location']
                These features carry no predictive signal and should be manually
investigated.
                This is typically a feature which has the same value for all
rows.
                These features do not need to be present at inference time.
        Types of features in original data (raw dtype, special dtypes):
                ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', …]
                ('int', [])   :  1 | ['estimated_diff_hours']
        Types of features in processed data (raw dtype, special dtypes):
                ('float', [])     : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', …]
                ('int', [])       :  1 | ['estimated_diff_hours']
                ('int', ['bool']) :  3 | ['is_day:idx', 'is_in_shadow:idx',
'wind_speed_w_1000hPa:ms']
        0.2s = Fit runtime
        43 features in original data used to generate 43 features in processed
data.
        Train Data (Processed) Memory Usage: 10.3 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.19s …
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
        This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
        To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with
holdout_frac=0.07836990595611286, Train Rows: 29400, Val Rows: 2500
User-specified model hyperparameters to be fit:
{
        'NN_TORCH': {},
```

```
        'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {},
'GBMLarge'],
        'CAT': {},
        'XGB': {},
        'FASTAI': {},
        'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
        'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
        'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Fitting 11 L1 models …
Fitting model: KNeighborsUnif … Training model for up to 599.81s of the 599.8s
of remaining time.
        -285.3795        = Validation score    (-mean_absolute_error)
        0.04s    = Training   runtime
        0.07s    = Validation runtime
Fitting model: KNeighborsDist … Training model for up to 599.69s of the
599.68s of remaining time.
        -288.0059        = Validation score    (-mean_absolute_error)
        0.04s    = Training   runtime
        0.04s    = Validation runtime
Fitting model: LightGBMXT … Training model for up to 599.6s of the 599.6s of
remaining time.

[1000]  valid_set's l1: 178.37
[2000]  valid_set's l1: 174.975
[3000]  valid_set's l1: 173.514
[4000]  valid_set's l1: 172.456
[5000]  valid_set's l1: 172.017
[6000]  valid_set's l1: 171.584
[7000]  valid_set's l1: 171.168
[8000]  valid_set's l1: 170.818
[9000]  valid_set's l1: 170.597
[10000] valid_set's l1: 170.345

        -170.3454        = Validation score    (-mean_absolute_error)
        13.42s   = Training   runtime
        0.16s    = Validation runtime
Fitting model: LightGBM … Training model for up to 585.75s of the 585.74s of
remaining time.
```

```
[1000]  valid_set's l1: 182.44
[2000]  valid_set's l1: 180.615
[3000]  valid_set's l1: 180.212

        -180.107          = Validation score    (-mean_absolute_error)
        4.7s      = Training    runtime
        0.04s     = Validation runtime
Fitting model: RandomForestMSE … Training model for up to 580.94s of the
580.94s of remaining time.
        -187.2246         = Validation score    (-mean_absolute_error)
        7.49s     = Training    runtime
        0.1s      = Validation runtime
Fitting model: CatBoost … Training model for up to 572.87s of the 572.87s of
remaining time.
        -181.3073         = Validation score    (-mean_absolute_error)
        111.24s  = Training    runtime
        0.01s     = Validation runtime
Fitting model: ExtraTreesMSE … Training model for up to 461.59s of the 461.58s
of remaining time.
        -186.6021         = Validation score    (-mean_absolute_error)
        1.65s     = Training    runtime
        0.1s      = Validation runtime
Fitting model: NeuralNetFastAI … Training model for up to 459.35s of the
459.35s of remaining time.
        -192.4111         = Validation score    (-mean_absolute_error)
        27.11s   = Training    runtime
        0.04s     = Validation runtime
Fitting model: XGBoost … Training model for up to 432.16s of the 432.16s of
remaining time.
        -186.0713         = Validation score    (-mean_absolute_error)
        2.37s     = Training    runtime
        0.01s     = Validation runtime
Fitting model: NeuralNetTorch … Training model for up to 429.76s of the
429.75s of remaining time.
        -176.1157         = Validation score    (-mean_absolute_error)
        51.95s   = Training    runtime
        0.04s     = Validation runtime
Fitting model: LightGBMLarge … Training model for up to 377.76s of the 377.76s
of remaining time.

[1000]  valid_set's l1: 172.273
[2000]  valid_set's l1: 170.86
[3000]  valid_set's l1: 170.486
[4000]  valid_set's l1: 170.39
[5000]  valid_set's l1: 170.319
[6000]  valid_set's l1: 170.298
[7000]  valid_set's l1: 170.29
[8000]  valid_set's l1: 170.284
[9000]  valid_set's l1: 170.282
```

```
[10000] valid_set's l1: 170.28

        -170.2803         = Validation score    (-mean_absolute_error)
        43.61s    = Training    runtime
        0.27s     = Validation runtime
Fitting model: WeightedEnsemble_L2 … Training model for up to 360.0s of the
332.71s of remaining time.
        -163.223          = Validation score    (-mean_absolute_error)
        0.45s     = Training    runtime
        0.0s      = Validation runtime
AutoGluon training complete, total runtime = 267.78s … Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_80_A/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -187.8065158485432
        Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -187.8065158485432,
    "root_mean_squared_error": -414.32909739992886,
    "mean_squared_error": -171668.60095223974,
    "r2": 0.8754434810346822,
    "pearsonr": 0.9358470155799331,
    "median_absolute_error": -12.917183456420899
}

Evaluation on test data:
-187.8065158485432
```

[10]:
```
loc = "B"
predictors[1] = fit_predictor_for_location(loc)
```

```
Values in column 'sample_weight' used as sample weights instead of predictive
features. Evaluation will report weighted metrics, so ensure same column exists
in test data.
Beginning AutoGluon training … Time limit = 600s
AutoGluon will save models to "AutogluonModels/submission_80_B/"
AutoGluon Version:  0.8.2
Python Version:     3.10.12
Operating System:   Linux
Platform Machine:   x86_64
Platform Version:   #1 SMP Debian 5.10.191-1 (2023-08-16)
Disk Space Avail:   310.58 GB / 315.93 GB (98.3%)
Train Data Rows:    30768
Train Data Columns: 46
Label Column: y
```

Preprocessing data …
AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).
        Label info (max, min, mean, stddev): (1152.3, -0.0, 97.74541, 195.0957)
        If 'regression' is not the correct problem_type, please manually specify the problem_type parameter during predictor init (You may specify problem_type as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data …
Fitting AutoMLPipelineFeatureGenerator…
        Available Memory: 130675.17 MB
        Train Data (Original) Memory Usage: 12.62 MB (0.0% of available memory)
        Inferring data type of each feature based on column values. Set feature_metadata_in to manually specify special dtypes of the features.
        Stage 1 Generators:
                Fitting AsTypeFeatureGenerator…
                        Note: Converting 3 features to boolean dtype as they only contain 2 unique values.
        Stage 2 Generators:
                Fitting FillNaFeatureGenerator…
        Stage 3 Generators:
                Fitting IdentityFeatureGenerator…
        Stage 4 Generators:
                Fitting DropUniqueFeatureGenerator…
        Stage 5 Generators:
                Fitting DropDuplicatesFeatureGenerator…

Training model for location B…
Train data sample weight sum: 30768
Train data number of rows: 30768
Test data sample weight sum: 2051
Test data number of rows: 2051

        Useless Original Features (Count: 2): ['elevation:m', 'location']
                These features carry no predictive signal and should be manually investigated.
                This is typically a feature which has the same value for all rows.
                These features do not need to be present at inference time.
        Types of features in original data (raw dtype, special dtypes):
                ('float', []) : 42 | ['absolute_humidity_2m:gm3', 'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J', 'clear_sky_rad:W', …]
                ('int', [])   : 1 | ['estimated_diff_hours']
        Types of features in processed data (raw dtype, special dtypes):
                ('float', [])   : 39 | ['absolute_humidity_2m:gm3', 'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J', 'clear_sky_rad:W', …]
                ('int', [])   : 1 | ['estimated_diff_hours']
                ('int', ['bool']) : 3 | ['is_day:idx', 'is_in_shadow:idx',

14

```
'wind_speed_w_1000hPa:ms']
        0.2s = Fit runtime
        43 features in original data used to generate 43 features in processed
data.
        Train Data (Processed) Memory Usage: 9.94 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.18s …
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
        This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
        To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with
holdout_frac=0.0812532501300052, Train Rows: 28268, Val Rows: 2500
User-specified model hyperparameters to be fit:
{
        'NN_TORCH': {},
        'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {},
'GBMLarge'],
        'CAT': {},
        'XGB': {},
        'FASTAI': {},
        'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
        'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
        'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Fitting 11 L1 models …
Fitting model: KNeighborsUnif … Training model for up to 599.82s of the
599.81s of remaining time.
        -57.0973          = Validation score   (-mean_absolute_error)
        0.03s    = Training    runtime
        0.04s    = Validation runtime
Fitting model: KNeighborsDist … Training model for up to 599.73s of the
599.73s of remaining time.
        -56.8969          = Validation score   (-mean_absolute_error)
        0.03s    = Training    runtime
        0.04s    = Validation runtime
Fitting model: LightGBMXT … Training model for up to 599.65s of the 599.65s of
remaining time.
```

```
[1000]   valid_set's l1: 35.5751
[2000]   valid_set's l1: 33.3902
[3000]   valid_set's l1: 32.2742
[4000]   valid_set's l1: 31.5407
[5000]   valid_set's l1: 31.0096
[6000]   valid_set's l1: 30.6243
[7000]   valid_set's l1: 30.3162
[8000]   valid_set's l1: 30.0585
[9000]   valid_set's l1: 29.8764
[10000]  valid_set's l1: 29.726

        -29.7259       = Validation score   (-mean_absolute_error)
        13.01s    = Training    runtime
        0.17s     = Validation runtime
Fitting model: LightGBM … Training model for up to 586.23s of the 586.22s of
remaining time.

[1000]   valid_set's l1: 33.0342
[2000]   valid_set's l1: 31.7436
[3000]   valid_set's l1: 31.1409
[4000]   valid_set's l1: 30.8249
[5000]   valid_set's l1: 30.6002
[6000]   valid_set's l1: 30.4238
[7000]   valid_set's l1: 30.3416
[8000]   valid_set's l1: 30.2763
[9000]   valid_set's l1: 30.2196
[10000]  valid_set's l1: 30.185

        -30.185   = Validation score   (-mean_absolute_error)
        15.69s    = Training    runtime
        0.16s     = Validation runtime
Fitting model: RandomForestMSE … Training model for up to 570.13s of the
570.13s of remaining time.
        -35.3114       = Validation score   (-mean_absolute_error)
        8.75s     = Training    runtime
        0.09s     = Validation runtime
Fitting model: CatBoost … Training model for up to 560.95s of the 560.95s of
remaining time.
        -32.7391       = Validation score   (-mean_absolute_error)
        113.3s    = Training    runtime
        0.01s     = Validation runtime
Fitting model: ExtraTreesMSE … Training model for up to 447.61s of the 447.6s
of remaining time.
        -36.4349       = Validation score   (-mean_absolute_error)
        1.81s     = Training    runtime
        0.09s     = Validation runtime
Fitting model: NeuralNetFastAI … Training model for up to 445.3s of the
445.29s of remaining time.
        -40.4352       = Validation score   (-mean_absolute_error)
```

```
        25.68s   = Training   runtime
        0.03s    = Validation runtime
Fitting model: XGBoost … Training model for up to 419.55s of the 419.55s of
remaining time.
        -33.3765          = Validation score   (-mean_absolute_error)
        24.41s   = Training   runtime
        0.21s    = Validation runtime
Fitting model: NeuralNetTorch … Training model for up to 394.78s of the
394.77s of remaining time.
        -34.0567          = Validation score   (-mean_absolute_error)
        98.84s   = Training   runtime
        0.04s    = Validation runtime
Fitting model: LightGBMLarge … Training model for up to 295.89s of the 295.89s
of remaining time.

[1000]  valid_set's l1: 30.2342
[2000]  valid_set's l1: 29.3138
[3000]  valid_set's l1: 29.0572
[4000]  valid_set's l1: 28.983
[5000]  valid_set's l1: 28.955
[6000]  valid_set's l1: 28.9422
[7000]  valid_set's l1: 28.9365
[8000]  valid_set's l1: 28.9338
[9000]  valid_set's l1: 28.9325
[10000] valid_set's l1: 28.9318

        -28.9318          = Validation score   (-mean_absolute_error)
        44.72s   = Training   runtime
        0.29s    = Validation runtime
Fitting model: WeightedEnsemble_L2 … Training model for up to 360.0s of the
249.46s of remaining time.
        -28.419  = Validation score   (-mean_absolute_error)
        0.44s    = Training   runtime
        0.0s     = Validation runtime
AutoGluon training complete, total runtime = 351.02s … Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_80_B/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -37.12180427631004
        Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -37.12180427631004,
    "root_mean_squared_error": -81.58609851989235,
    "mean_squared_error": -6656.291471697579,
    "r2": 0.7859139269118017,
```

```
    "pearsonr": 0.9104600209448434,
    "median_absolute_error": -8.015130996704102
}

Evaluation on test data:
-37.12180427631004
```

```
[ ]:  loc = "C"
      predictors[2] = fit_predictor_for_location(loc)
```

Values in column 'sample_weight' used as sample weights instead of predictive
features. Evaluation will report weighted metrics, so ensure same column exists
in test data.
Beginning AutoGluon training … Time limit = 600s
AutoGluon will save models to "AutogluonModels/submission_80_C/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.191-1 (2023-08-16)
Disk Space Avail: 309.67 GB / 315.93 GB (98.0%)
Train Data Rows: 24492
Train Data Columns: 46
Label Column: y
Preprocessing data …
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and label-values can't be converted to int).
        Label info (max, min, mean, stddev): (999.6, 0.0, 78.11911, 167.50151)
        If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data …
Fitting AutoMLPipelineFeatureGenerator…
        Available Memory: 130353.95 MB
        Train Data (Original) Memory Usage: 10.04 MB (0.0% of available memory)
        Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
        Stage 1 Generators:
                Fitting AsTypeFeatureGenerator…
                        Note: Converting 2 features to boolean dtype as they
only contain 2 unique values.
        Stage 2 Generators:
                Fitting FillNaFeatureGenerator…
        Stage 3 Generators:
                Fitting IdentityFeatureGenerator…
        Stage 4 Generators:
                Fitting DropUniqueFeatureGenerator…
        Stage 5 Generators:
                Fitting DropDuplicatesFeatureGenerator…

```

Useless Original Features (Count: 2): ['elevation:m', 'location']
            These features carry no predictive signal and should be manually
investigated.
            This is typically a feature which has the same value for all
rows.
            These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
            ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', …]
            ('int', [])   :  1 | ['estimated_diff_hours']
    Types of features in processed data (raw dtype, special dtypes):
            ('float', [])     : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', …]
            ('int', [])       :  1 | ['estimated_diff_hours']
Training model for location C…
Train data sample weight sum: 24492
Train data number of rows: 24492
Test data sample weight sum: 1579
Test data number of rows: 1579

            ('int', ['bool']) :  2 | ['is_day:idx', 'is_in_shadow:idx']
    0.1s = Fit runtime
    43 features in original data used to generate 43 features in processed
data.
    Train Data (Processed) Memory Usage: 8.08 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.16s …
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with holdout_frac=0.1, Train
Rows: 22042, Val Rows: 2450
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {},
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],

```
        'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
        'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Fitting 11 L1 models …
Fitting model: KNeighborsUnif … Training model for up to 599.84s of the
599.83s of remaining time.
        -33.2822   = Validation score   (-mean_absolute_error)
        0.03s   = Training   runtime
        0.05s   = Validation runtime
Fitting model: KNeighborsDist … Training model for up to 599.74s of the
599.74s of remaining time.
        -33.3446   = Validation score   (-mean_absolute_error)
        0.03s   = Training   runtime
        0.1s   = Validation runtime
Fitting model: LightGBMXT … Training model for up to 599.6s of the 599.59s of
remaining time.

[1000] valid_set's l1: 18.9213
[2000] valid_set's l1: 18.3545
[3000] valid_set's l1: 18.1711
[4000] valid_set's l1: 18.08
[5000] valid_set's l1: 18.0196
[6000] valid_set's l1: 17.9721
[7000] valid_set's l1: 17.9335
[8000] valid_set's l1: 17.9153
[9000] valid_set's l1: 17.9061
[10000] valid_set's l1: 17.8961

        -17.8909   = Validation score   (-mean_absolute_error)
        12.9s   = Training   runtime
        0.18s   = Validation runtime
Fitting model: LightGBM … Training model for up to 586.2s of the 586.2s of
remaining time.

[1000] valid_set's l1: 19.115
[2000] valid_set's l1: 18.8635
[3000] valid_set's l1: 18.8186
[4000] valid_set's l1: 18.7676
[5000] valid_set's l1: 18.7493
[6000] valid_set's l1: 18.7353
[7000] valid_set's l1: 18.7294
[8000] valid_set's l1: 18.7245
[9000] valid_set's l1: 18.7201
[10000] valid_set's l1: 18.7209
```

```
        -18.7199         = Validation score    (-mean_absolute_error)
        13.65s    = Training    runtime
        0.16s     = Validation runtime
Fitting model: RandomForestMSE … Training model for up to 572.1s of the
572.09s of remaining time.
        -20.2022         = Validation score    (-mean_absolute_error)
        4.85s     = Training    runtime
        0.1s      = Validation runtime
Fitting model: CatBoost … Training model for up to 566.99s of the 566.98s of
remaining time.
        -18.5962         = Validation score    (-mean_absolute_error)
        117.36s   = Training    runtime
        0.01s     = Validation runtime
Fitting model: ExtraTreesMSE … Training model for up to 449.59s of the 449.58s
of remaining time.
        -20.1925         = Validation score    (-mean_absolute_error)
        1.09s     = Training    runtime
        0.09s     = Validation runtime
Fitting model: NeuralNetFastAI … Training model for up to 448.24s of the
448.23s of remaining time.
        -20.3136         = Validation score    (-mean_absolute_error)
        20.56s    = Training    runtime
        0.04s     = Validation runtime
Fitting model: XGBoost … Training model for up to 427.6s of the 427.6s of
remaining time.
        -18.7778         = Validation score    (-mean_absolute_error)
        23.62s    = Training    runtime
        0.21s     = Validation runtime
Fitting model: NeuralNetTorch … Training model for up to 403.62s of the
403.61s of remaining time.
```

3 Submit

```python
import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data
```

```python
test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
```

```
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
 ↪"location"], left_on=["ds", "location"])

#test_data_merged
```

```
[ ]: # predict, grouped by location
     predictions = []
     location_map = {
         "A": 0,
         "B": 1,
         "C": 2
     }
     for loc, group in test_data.groupby('location'):
         i = location_map[loc]
         subset = test_data_merged[test_data_merged["location"] == loc].
      ↪reset_index(drop=True)
         #print(subset)
         pred = predictors[i].predict(subset)
         subset["prediction"] = pred
         predictions.append(subset)

         # get past predictions
         past_pred = predictors[i].
      ↪predict(train_data_with_dates[train_data_with_dates["location"] == loc])
         train_data_with_dates.loc[train_data_with_dates["location"] == loc,
      ↪"prediction"] = past_pred
```

```
[ ]: # plot predictions for location A, in addition to train data for A
     for loc, idx in location_map.items():
         fig, ax = plt.subplots(figsize=(20, 10))
         # plot train data
         train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
      ↪y='y', ax=ax, label="train data")

         # plot predictions
         predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

         # plot past predictions
         train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
      ↪y='prediction', ax=ax, label="past predictions")

         # title
         ax.set_title(f"Predictions for location {loc}")
```

```python
# concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df
```

```python
# Save the submission DataFrame to submissions folder, create new name based on
↪last submission, format is submission_<last_submission_number + 1>.csv

# Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
↪index=False)
print("jall1a")
```

```python
# save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)
```

```python
# save this notebook to submissions folder
import subprocess
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
↪join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
↪ipynb"])
```

```python
# feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
↪time_limit=60*10)
```

```python
# feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
↪time_limit=60*10)
```

```python
display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)
```

```
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
↪join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),␣
↪"autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

     # def execute_git_command(directory, command):
     #     """Execute a Git command in the specified directory."""
     #     try:
     #         result = subprocess.check_output(['git', '-C', directory] + command,␣
     ↪stderr=subprocess.STDOUT)
     #         return result.decode('utf-8').strip(), True
     #     except subprocess.CalledProcessError as e:
     #         print(f"Git command failed with message: {e.output.decode('utf-8').
     ↪strip()}")
     #         return e.output.decode('utf-8').strip(), False

     # git_repo_path = "."

     # execute_git_command(git_repo_path, ['config', 'user.email',␣
     ↪'henrikskog01@gmail.com'])
     # execute_git_command(git_repo_path, ['config', 'user.name', hello if hello is␣
     ↪not None else 'Henrik eller Jørgen'])

     # branch_name = new_filename

     # # add datetime to branch name
     # branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

     # commit_msg = "run result"

     # execute_git_command(git_repo_path, ['checkout', '-b',branch_name])

     # # Navigate to your repo and commit changes
     # execute_git_command(git_repo_path, ['add', '.'])
     # execute_git_command(git_repo_path, ['commit', '-m',commit_msg])

     # # Push to remote
     # output, success = execute_git_command(git_repo_path, ['push',␣
     ↪'origin',branch_name])

     # # If the push fails, try setting an upstream branch and push again
     # if not success and 'upstream' in output:
     #     print("Attempting to set upstream and push again...")
     #     execute_git_command(git_repo_path, ['push', '--set-upstream',␣
     ↪'origin',branch_name])
```

```
#        execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])
```