

# autogluon\_each\_location

October 9, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*10
presets = 'best_quality'

do_drop_ds = True

use_groups = False
n_groups = 8

auto_stack = True
num_stack_levels = 1
num_bag_folds = 0
if auto_stack:
    num_stack_levels = None
    num_bag_folds = None

use_tune_data = False
use_test_data = True
tune_and_test_length = 24*30*3 # 3 months from end, this changes the
    ↪evaluations for only test
holdout_frac = None
use_bag_holdout = False # Enable this if there is a large gap between score_val
    ↪and score_test in stack models.

sample_weight = 'sample_weight' #None
weight_evaluation = True #False
sample_weight_estimated = 1 # this changes evaluations for test and tune WTF,
    ↪cant find a fix

run_analysis = False

[2]: import pandas as pd
import numpy as np
```

```

import warnings
warnings.filterwarnings("ignore")

def fix_datetime(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    # with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    # Drop rows where the minute part of the time is not 0
    X = X[X.index.minute == 0].copy()
    return X

def convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_datetime(X_train_observed, "X_train_observed")
    X_train_estimated = fix_datetime(X_train_estimated, "X_train_estimated")
    X_test = fix_datetime(X_test, "X_test")

    # add sample weights, which are 1 for observed and 3 for estimated
    X_train_observed["sample_weight"] = 1
    X_train_estimated["sample_weight"] = sample_weight_estimated
    X_test["sample_weight"] = sample_weight_estimated

    X_train_observed["estimated_diff_hours"] = 0
    X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index - pd.
    to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
    X_test["estimated_diff_hours"] = (X_test.index - pd.
    to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

    X_train_estimated["estimated_diff_hours"] =
    X_train_estimated["estimated_diff_hours"].astype('int64')
    # the filled once will get dropped later anyways, when we drop y nans
    X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].fillna(-50).
    astype('int64')

    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train['ds'] = pd.to_datetime(y_train['time'])
    y_train.drop(columns=['time'], inplace=True)

```

```

y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # fill missng sample_weight with 3
    #X_train["sample_weight"] = X_train["sample_weight"].fillna(0)

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
right_index=True)

    # print number of nans in sample_weight
    print(f"Number of nans in sample_weight: {X_train['sample_weight'].isna().
sum()}")
    # print number of nans in y
    print(f"Number of nans in y: {X_train['y'].isna().sum()}")

    X_train["location"] = location
    X_test["location"] = location

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")

```

```

# Read target training data
y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

# Read estimated training data and add location feature
X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

# Read observed training data and add location feature
X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

# Read estimated test data and add location feature
X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

# Preprocess data
X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
X_test_estimated, y_train, loc)

X_trains.append(X_train)
X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

```

Processing location A...
Number of nans in sample_weight: 0
Number of nans in y: 0
Processing location B...
Number of nans in sample_weight: 0
Number of nans in y: 4
Processing location C...
Number of nans in sample_weight: 0
Number of nans in y: 6059

```

## 1 Feature engineering

```

[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y'], inplace=True)

if not do_drop_ds:
    # add hour datetime feature
    X_train["hour"] = X_train.index.hour
    X_test["hour"] = X_test.index.hour

#print(X_train.head())

```

```

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()

        # Calculate the size of each group for this location
        group_size = len(loc_df) // n_groups

        # Create a new 'group' column for this location
        loc_df['group'] = np.repeat(range(n_groups),
    ↪ repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

        # Append to list of grouped DataFrames
        grouped_dfs.append(loc_df)

    # Concatenate all the grouped DataFrames back together
    X_train = pd.concat(grouped_dfs)
    X_train.sort_index(inplace=True)
    print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

```

[4]: from autogluon.tabular import TabularDataset, TabularPredictor
      from autogluon.timeseries import TimeSeriesDataFrame
      import numpy as np
      train_data = TabularDataset('X_train_raw.csv')

```

```

# set group column of train_data be increasing from 0 to 7 based on time, the
    ↪ first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
train_data['ds'] = pd.to_datetime(train_data['ds'])
train_data = train_data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
split_time = pd.to_datetime(train_data["ds"]).max() - pd.
    ↪ Timedelta(hours=tune_and_test_length)
train_set = TabularDataset(train_data[train_data["ds"] < split_time])
test_set = TabularDataset(train_data[train_data["ds"] >= split_time])
if use_groups:
    test_set = test_set.drop(columns=['group'])

if do_drop_ds:
    train_set = train_set.drop(columns=['ds'])
    test_set = test_set.drop(columns=['ds'])
    train_data = train_data.drop(columns=['ds'])

def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
    ↪ loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

tuning_data = None
if use_tune_data:
    train_data = train_set
    if use_test_data:
        # split test_set in half, use first half for tuning
        tuning_data, test_data = [], []
        for loc in locations:
            loc_test_set = test_set[test_set["location"] == loc]
            loc_tuning_data = loc_test_set.iloc[:len(loc_test_set)//2]
            loc_test_data = loc_test_set.iloc[len(loc_test_set)//2:]
            tuning_data.append(loc_tuning_data)
            test_data.append(loc_test_data)
        tuning_data = pd.concat(tuning_data)
        test_data = pd.concat(test_data)

```

```

        print("Shapes of tuning and test", tuning_data.shape[0], test_data.
↪shape[0], tuning_data.shape[0] + test_data.shape[0])

    else:
        tuning_data = test_set
        print("Shape of tuning", tuning_data.shape[0])

        # ensure sample weights for your tuning data sum to the number of rows in
↪the tuning data.
        tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        train_data = train_set
        test_data = test_set
        print("Shape of test", test_data.shape[0])

    # ensure sample weights for your training (or tuning) data sum to the number of
↪rows in the training (or tuning) data.
    train_data = normalize_sample_weights_per_location(train_data)
    if use_test_data:
        test_data = normalize_sample_weights_per_location(test_data)

```

Shape of test 5791

```

[5]: if run_analysis:
        import autogluon.eda.auto as auto
        auto.dataset_overview(train_data=train_data, test_data=test_data,
↪label="y", sample=None)

```

```

[6]: if run_analysis:
        auto.target_analysis(train_data=train_data, label="y")

```

## 2 Starting

```

[7]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
↪filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

```

```

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)

```

Last submission number: 80  
 Now creating submission number: 81  
 New filename: submission\_81

```
[8]: predictors = [None, None, None]
```

```

[9]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪ train and tune data and test data
    print("Train data sample weight sum:", train_data[train_data["location"] ==
    ↪ loc]["sample_weight"].sum())
    print("Train data number of rows:", train_data[train_data["location"] ==
    ↪ loc].shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
    ↪ tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:", tuning_data[tuning_data["location"]
    ↪ == loc].shape[0])
    if use_test_data:
        print("Test data sample weight sum:", test_data[test_data["location"]
    ↪ == loc]["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"] ==
    ↪ loc].shape[0])
    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        sample_weight=sample_weight,
        weight_evaluation=weight_evaluation,
        groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc],
        time_limit=time_limit,
        #presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
    ↪ somethin, will be overwritten anyways
        tuning_data=tuning_data[tuning_data["location"] == loc] if
    ↪ use_tune_data else None,

```



```

        use_bag_holdout=use_bag_holdout,
        holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
        ↪drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 300s

AutoGluon will save models to "AutogluonModels/submission\_81\_A/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 309.06 GB / 315.93 GB (97.8%)

Train Data Rows: 31900

Train Data Columns: 46

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (5733.42, 0.0, 633.132, 1165.64686)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Training model for location A...

Train data sample weight sum: 31900

Train data number of rows: 31900

Test data sample weight sum: 2161

Test data number of rows: 2161

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 132350.31 MB

Train Data (Original) Memory Usage: 13.08 MB (0.0% of available memory)  
Inferring data type of each feature based on column values. Set  
feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 3 features to boolean dtype as they  
only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

Stage 4 Generators:

Fitting DropUniqueFeatureGenerator...

Stage 5 Generators:

Fitting DropDuplicatesFeatureGenerator...

Useless Original Features (Count: 2): ['elevation:m', 'location']

These features carry no predictive signal and should be manually  
investigated.

This is typically a feature which has the same value for all  
rows.

These features do not need to be present at inference time.

Types of features in original data (raw dtype, special dtypes):

('float', []) : 42 | ['absolute\_humidity\_2m:gm3',  
'air\_density\_2m:kgm3', 'ceiling\_height\_agl:m', 'clear\_sky\_energy\_1h:J',  
'clear\_sky\_rad:W', ...]

('int', []) : 1 | ['estimated\_diff\_hours']

Types of features in processed data (raw dtype, special dtypes):

('float', []) : 39 | ['absolute\_humidity\_2m:gm3',  
'air\_density\_2m:kgm3', 'ceiling\_height\_agl:m', 'clear\_sky\_energy\_1h:J',  
'clear\_sky\_rad:W', ...]

('int', []) : 1 | ['estimated\_diff\_hours']

('int', ['bool']) : 3 | ['is\_day:idx', 'is\_in\_shadow:idx',  
'wind\_speed\_w\_1000hPa:ms']

0.2s = Fit runtime

43 features in original data used to generate 43 features in processed  
data.

Train Data (Processed) Memory Usage: 10.3 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.19s ...

AutoGluon will gauge predictive performance using evaluation metric:

'mean\_absolute\_error'

This metric's sign has been flipped to adhere to being higher\_is\_better.  
The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval\_metric parameter of Predictor()

Automatically generating train/validation split with

holdout\_frac=0.07836990595611286, Train Rows: 29400, Val Rows: 2500

User-specified model hyperparameters to be fit:

```
{  
    'NN_TORCH': {},
```

```

    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif ... Training model for up to 299.81s of the 299.8s of remaining time.

```

-285.3795      = Validation score    (-mean_absolute_error)
0.04s         = Training    runtime
0.07s         = Validation runtime

```

Fitting model: KNeighborsDist ... Training model for up to 299.69s of the 299.69s of remaining time.

```

-288.0059      = Validation score    (-mean_absolute_error)
0.04s         = Training    runtime
0.04s         = Validation runtime

```

Fitting model: LightGBMXT ... Training model for up to 299.6s of the 299.6s of remaining time.

```

[1000] valid_set's l1: 178.37
[2000] valid_set's l1: 174.975
[3000] valid_set's l1: 173.514
[4000] valid_set's l1: 172.456
[5000] valid_set's l1: 172.017
[6000] valid_set's l1: 171.584
[7000] valid_set's l1: 171.168
[8000] valid_set's l1: 170.818
[9000] valid_set's l1: 170.597
[10000] valid_set's l1: 170.345

```

```

-170.3454      = Validation score    (-mean_absolute_error)
13.7s         = Training    runtime
0.16s         = Validation runtime

```

Fitting model: LightGBM ... Training model for up to 285.46s of the 285.46s of remaining time.

```

[1000] valid_set's l1: 182.44
[2000] valid_set's l1: 180.615
[3000] valid_set's l1: 180.212

-180.107          = Validation score  (-mean_absolute_error)
4.84s            = Training  runtime
0.04s           = Validation runtime
Fitting model: RandomForestMSE ... Training model for up to 280.51s of the
280.51s of remaining time.
-187.2246         = Validation score  (-mean_absolute_error)
7.52s            = Training  runtime
0.09s           = Validation runtime
Fitting model: CatBoost ... Training model for up to 272.41s of the 272.41s of
remaining time.
-181.3073         = Validation score  (-mean_absolute_error)
110.67s          = Training  runtime
0.01s           = Validation runtime
Fitting model: ExtraTreesMSE ... Training model for up to 161.69s of the 161.69s
of remaining time.
-186.6021         = Validation score  (-mean_absolute_error)
1.75s           = Training  runtime
0.09s           = Validation runtime
Fitting model: NeuralNetFastAI ... Training model for up to 159.37s of the
159.36s of remaining time.
-192.4111         = Validation score  (-mean_absolute_error)
27.52s          = Training  runtime
0.04s           = Validation runtime
Fitting model: XGBoost ... Training model for up to 131.76s of the 131.76s of
remaining time.
-186.0713         = Validation score  (-mean_absolute_error)
2.36s           = Training  runtime
0.01s           = Validation runtime
Fitting model: NeuralNetTorch ... Training model for up to 129.36s of the
129.36s of remaining time.
-176.1157         = Validation score  (-mean_absolute_error)
52.9s           = Training  runtime
0.04s           = Validation runtime
Fitting model: LightGBMLarge ... Training model for up to 76.42s of the 76.42s
of remaining time.

[1000] valid_set's l1: 172.273
[2000] valid_set's l1: 170.86
[3000] valid_set's l1: 170.486
[4000] valid_set's l1: 170.39
[5000] valid_set's l1: 170.319
[6000] valid_set's l1: 170.298
[7000] valid_set's l1: 170.29
[8000] valid_set's l1: 170.284
[9000] valid_set's l1: 170.282

```

```

[10000] valid_set's l1: 170.28
      -170.2803      = Validation score      (-mean_absolute_error)
      46.18s      = Training      runtime
      0.26s      = Validation runtime
Fitting model: WeightedEnsemble_L2 ... Training model for up to 299.81s of the
28.81s of remaining time.
      -163.223      = Validation score      (-mean_absolute_error)
      0.46s      = Training      runtime
      0.0s      = Validation runtime
AutoGluon training complete, total runtime = 271.69s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_81_A/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -187.8065158485432
      Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
  "mean_absolute_error": -187.8065158485432,
  "root_mean_squared_error": -414.32909739992886,
  "mean_squared_error": -171668.60095223974,
  "r2": 0.8754434810346822,
  "pearsonr": 0.9358470155799331,
  "median_absolute_error": -12.917183456420899
}
Evaluation on test data:
-187.8065158485432

```

```

[10]: loc = "B"
      predictors[1] = fit_predictor_for_location(loc)

```

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

```

Beginning AutoGluon training ... Time limit = 300s
AutoGluon will save models to "AutogluonModels/submission_81_B/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 308.09 GB / 315.93 GB (97.5%)
Train Data Rows: 30768
Train Data Columns: 46
Label Column: y

```

```

Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (1152.3, -0.0, 97.74541, 195.0957)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                130566.93 MB
    Train Data (Original) Memory Usage: 12.62 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 3 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...

Training model for location B..
Train data sample weight sum: 30768
Train data number of rows: 30768
Test data sample weight sum: 2051
Test data number of rows: 2051

    Useless Original Features (Count: 2): ['elevation:m', 'location']
        These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', [])   : 1 | ['estimated_diff_hours']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', [])   : 1 | ['estimated_diff_hours']
        ('int', ['bool']) : 3 | ['is_day:idx', 'is_in_shadow:idx',

```

```

'wind_speed_w_1000hPa:ms']
    0.1s = Fit runtime
    43 features in original data used to generate 43 features in processed
data.
    Train Data (Processed) Memory Usage: 9.94 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.18s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with
holdout_frac=0.0812532501300052, Train Rows: 28268, Val Rows: 2500
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Fitting 11 L1 models ...
Fitting model: KNeighborsUnif ... Training model for up to 299.82s of the
299.82s of remaining time.
    -57.0973          = Validation score    (-mean_absolute_error)
    0.03s           = Training    runtime
    0.04s           = Validation runtime
Fitting model: KNeighborsDist ... Training model for up to 299.74s of the
299.73s of remaining time.
    -56.8969          = Validation score    (-mean_absolute_error)
    0.04s           = Training    runtime
    0.04s           = Validation runtime
Fitting model: LightGBMXT ... Training model for up to 299.65s of the 299.65s of
remaining time.

```

```

[1000] valid_set's l1: 35.5751
[2000] valid_set's l1: 33.3902
[3000] valid_set's l1: 32.2742
[4000] valid_set's l1: 31.5407
[5000] valid_set's l1: 31.0096
[6000] valid_set's l1: 30.6243
[7000] valid_set's l1: 30.3162
[8000] valid_set's l1: 30.0585
[9000] valid_set's l1: 29.8764
[10000] valid_set's l1: 29.726

```

```

-29.7259      = Validation score  (-mean_absolute_error)
13.51s       = Training  runtime
0.17s        = Validation runtime

```

Fitting model: LightGBM ... Training model for up to 285.72s of the 285.72s of remaining time.

```

[1000] valid_set's l1: 33.0342
[2000] valid_set's l1: 31.7436
[3000] valid_set's l1: 31.1409
[4000] valid_set's l1: 30.8249
[5000] valid_set's l1: 30.6002
[6000] valid_set's l1: 30.4238
[7000] valid_set's l1: 30.3416
[8000] valid_set's l1: 30.2763
[9000] valid_set's l1: 30.2196
[10000] valid_set's l1: 30.185

```

```

-30.185      = Validation score  (-mean_absolute_error)
14.18s       = Training  runtime
0.16s        = Validation runtime

```

Fitting model: RandomForestMSE ... Training model for up to 271.13s of the 271.13s of remaining time.

```

-35.3114      = Validation score  (-mean_absolute_error)
8.8s          = Training  runtime
0.1s          = Validation runtime

```

Fitting model: CatBoost ... Training model for up to 261.89s of the 261.88s of remaining time.

```

-32.7391      = Validation score  (-mean_absolute_error)
113.36s       = Training  runtime
0.01s         = Validation runtime

```

Fitting model: ExtraTreesMSE ... Training model for up to 148.49s of the 148.48s of remaining time.

```

-36.4349      = Validation score  (-mean_absolute_error)
1.81s         = Training  runtime
0.09s         = Validation runtime

```

Fitting model: NeuralNetFastAI ... Training model for up to 146.17s of the 146.17s of remaining time.

```

-40.4352      = Validation score  (-mean_absolute_error)

```



```

    26.19s    = Training    runtime
    0.04s     = Validation runtime
Fitting model: XGBoost ... Training model for up to 119.92s of the 119.91s of
remaining time.
    -33.3765          = Validation score    (-mean_absolute_error)
    24.77s    = Training    runtime
    0.21s     = Validation runtime
Fitting model: NeuralNetTorch ... Training model for up to 94.79s of the 94.78s
of remaining time.
    Ran out of time, stopping training early. (Stopping on epoch 142)
    -34.0567          = Validation score    (-mean_absolute_error)
    94.46s    = Training    runtime
    0.04s     = Validation runtime
Fitting model: LightGBMLarge ... Training model for up to 0.28s of the 0.27s of
remaining time.
    Ran out of time, early stopping on iteration 13. Best iteration is:
    [13]    valid_set's l1: 98.5649
    -98.5649          = Validation score    (-mean_absolute_error)
    0.31s    = Training    runtime
    0.0s     = Validation runtime
Fitting model: WeightedEnsemble_L2 ... Training model for up to 299.82s of the
-0.05s of remaining time.
    -29.0846          = Validation score    (-mean_absolute_error)
    0.46s    = Training    runtime
    0.0s     = Validation runtime
AutoGluon training complete, total runtime = 300.54s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_81_B/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -38.66709888168536
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -38.66709888168536,
    "root_mean_squared_error": -84.1069847977034,
    "mean_squared_error": -7073.984891761112,
    "r2": 0.7724796678448118,
    "pearsonr": 0.9044476820459424,
    "median_absolute_error": -8.309171676635742
}
Evaluation on test data:
-38.66709888168536

```

```
[ ]: loc = "C"
predictors[2] = fit_predictor_for_location(loc)
```

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 300s

AutoGluon will save models to "AutogluonModels/submission\_81\_C/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 307.31 GB / 315.93 GB (97.3%)

Train Data Rows: 24492

Train Data Columns: 46

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and label-values can't be converted to int).

Label info (max, min, mean, stddev): (999.6, 0.0, 78.11911, 167.50151)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 130339.77 MB

Train Data (Original) Memory Usage: 10.04 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set

feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 2 features to boolean dtype as they only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

Stage 4 Generators:

Fitting DropUniqueFeatureGenerator...

Stage 5 Generators:

Fitting DropDuplicatesFeatureGenerator...

Useless Original Features (Count: 2): ['elevation:m', 'location']

These features carry no predictive signal and should be manually investigated.

This is typically a feature which has the same value for all rows.

These features do not need to be present at inference time.

```

Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', []) : 1 | ['estimated_diff_hours']
Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', []) : 1 | ['estimated_diff_hours']
    ('int', ['bool']) : 2 | ['is_day:idx', 'is_in_shadow:idx']
0.1s = Fit runtime
43 features in original data used to generate 43 features in processed
data.

```

Train Data (Processed) Memory Usage: 8.08 MB (0.0% of available memory)

```

Training model for location C...
Train data sample weight sum: 24492
Train data number of rows: 24492
Test data sample weight sum: 1579
Test data number of rows: 1579

```

```

Data preprocessing and feature engineering runtime = 0.16s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'

```

This metric's sign has been flipped to adhere to being higher\_is\_better.  
The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval\_metric parameter of Predictor()  
Automatically generating train/validation split with holdout\_frac=0.1, Train  
Rows: 22042, Val Rows: 2450

User-specified model hyperparameters to be fit:

```

{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},

```

```

{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Fitting 11 L1 models ...
Fitting model: KNeighborsUnif ... Training model for up to 299.84s of the
299.84s of remaining time.
    -33.2822          = Validation score    (-mean_absolute_error)
    0.04s           = Training    runtime
    0.42s           = Validation runtime
Fitting model: KNeighborsDist ... Training model for up to 299.38s of the
299.37s of remaining time.
    -33.3446          = Validation score    (-mean_absolute_error)
    0.03s           = Training    runtime
    0.43s           = Validation runtime
Fitting model: LightGBMXT ... Training model for up to 298.9s of the 298.9s of
remaining time.

[1000] valid_set's l1: 18.9213
[2000] valid_set's l1: 18.3545
[3000] valid_set's l1: 18.1711
[4000] valid_set's l1: 18.08
[5000] valid_set's l1: 18.0196
[6000] valid_set's l1: 17.9721
[7000] valid_set's l1: 17.9335
[8000] valid_set's l1: 17.9153
[9000] valid_set's l1: 17.9061
[10000] valid_set's l1: 17.8961

    -17.8909          = Validation score    (-mean_absolute_error)
    12.67s           = Training    runtime
    0.17s           = Validation runtime
Fitting model: LightGBM ... Training model for up to 285.74s of the 285.73s of
remaining time.

[1000] valid_set's l1: 19.115
[2000] valid_set's l1: 18.8635
[3000] valid_set's l1: 18.8186
[4000] valid_set's l1: 18.7676
[5000] valid_set's l1: 18.7493
[6000] valid_set's l1: 18.7353
[7000] valid_set's l1: 18.7294
[8000] valid_set's l1: 18.7245
[9000] valid_set's l1: 18.7201
[10000] valid_set's l1: 18.7209

    -18.7199          = Validation score    (-mean_absolute_error)
    13.13s           = Training    runtime
    0.15s           = Validation runtime
Fitting model: RandomForestMSE ... Training model for up to 272.23s of the
272.22s of remaining time.
    -20.2022          = Validation score    (-mean_absolute_error)

```

```

4.64s    = Training    runtime
0.1s     = Validation runtime
Fitting model: CatBoost ... Training model for up to 267.33s of the 267.33s of
remaining time.
-18.5962          = Validation score    (-mean_absolute_error)
112.02s = Training    runtime
0.01s    = Validation runtime
Fitting model: ExtraTreesMSE ... Training model for up to 155.27s of the 155.27s
of remaining time.
-20.1925          = Validation score    (-mean_absolute_error)
1.0s    = Training    runtime
0.08s    = Validation runtime
Fitting model: NeuralNetFastAI ... Training model for up to 154.01s of the
154.0s of remaining time.
-20.3136          = Validation score    (-mean_absolute_error)
20.87s  = Training    runtime
0.04s    = Validation runtime
Fitting model: XGBoost ... Training model for up to 133.07s of the 133.06s of
remaining time.
-18.7778          = Validation score    (-mean_absolute_error)
24.52s  = Training    runtime
0.21s    = Validation runtime
Fitting model: NeuralNetTorch ... Training model for up to 108.19s of the
108.18s of remaining time.

```

### 3 Submit

```

[ ]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data

[ ]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
↪ "location"], left_on=["ds", "location"])

#test_data_merged

[ ]: # predict, grouped by location
predictions = []

```

```

location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in test_data.groupby('location'):
    i = location_map[loc]
    subset = test_data_merged[test_data_merged["location"] == loc].
    ↪reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)

    # get past predictions
    past_pred = predictors[i].
    ↪predict(train_data_with_dates[train_data_with_dates["location"] == loc])
    train_data_with_dates.loc[train_data_with_dates["location"] == loc,
    ↪"prediction"] = past_pred

```

```

[ ]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='y', ax=ax, label="train data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='prediction', ax=ax, label="past predictions")

    # title
    ax.set_title(f"Predictions for location {loc}")

```

```

[ ]: # concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df

```

```

[ ]: # Save the submission DataFrame to submissions folder, create new name based on
    ↪last submission, format is submission_<last_submission_number + 1>.csv

# Save the submission

```

```
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
    ↪index=False)
print("jall1a")
```

```
[ ]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)
```

```
[ ]: # save this notebook to submissions folder
import subprocess
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
    ↪ipynb"])
```

```
[ ]: # feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
    ↪time_limit=60*10)
```

```
[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
    ↪time_limit=60*10)
```

```
[ ]: display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)

subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↪"autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
```

```

#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
# ↪ stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8')}.
# ↪ strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
# ↪ 'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', hello if hello is
# ↪ not None else 'Henrik eller Jørgen'])

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push',
# ↪ 'origin', branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream',
# ↪ 'origin', branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])

```