

autogluon_each_location

October 9, 2023

```
[6]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*30
presets = 'best_quality'

do_drop_ds = True

use_groups = False
n_groups = 8

auto_stack = True
num_stack_levels = 1
num_bag_folds = 0
if auto_stack:
    num_stack_levels = None
    num_bag_folds = None

use_tune_data = False
use_test_data = True
tune_and_test_length = 24*30*3 # 3 months from end, this changes the
    ↪evaluations for only test
holdout_frac = None
use_bag_holdout = False # Enable this if there is a large gap between score_val
    ↪and score_test in stack models.

sample_weight = 'sample_weight' #None
weight_evaluation = True #False
sample_weight_estimated = 3 # this changes evaluations for test and tune WTF,
    ↪cant find a fix

run_analysis = False

[7]: import pandas as pd
import numpy as np
```

```

import warnings
warnings.filterwarnings("ignore")

def fix_datetime(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    # with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    # Drop rows where the minute part of the time is not 0
    X = X[X.index.minute == 0].copy()
    return X

def convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_datetime(X_train_observed, "X_train_observed")
    X_train_estimated = fix_datetime(X_train_estimated, "X_train_estimated")
    X_test = fix_datetime(X_test, "X_test")

    # add sample weights, which are 1 for observed and 3 for estimated
    X_train_observed["sample_weight"] = 1
    X_train_estimated["sample_weight"] = sample_weight_estimated
    X_test["sample_weight"] = sample_weight_estimated

    X_train_observed["estimated_diff_hours"] = 0
    X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index - pd.
    to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
    X_test["estimated_diff_hours"] = (X_test.index - pd.
    to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

    X_train_estimated["estimated_diff_hours"] =
    X_train_estimated["estimated_diff_hours"].astype('int64')
    # the filled once will get dropped later anyways, when we drop y nans
    X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].fillna(-50).
    astype('int64')

    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train['ds'] = pd.to_datetime(y_train['time'])
    y_train.drop(columns=['time'], inplace=True)

```

```

y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # fill missng sample_weight with 3
    #X_train["sample_weight"] = X_train["sample_weight"].fillna(0)

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
right_index=True)

    # print number of nans in sample_weight
    print(f"Number of nans in sample_weight: {X_train['sample_weight'].isna().
sum()}")
    # print number of nans in y
    print(f"Number of nans in y: {X_train['y'].isna().sum()}")

    X_train["location"] = location
    X_test["location"] = location

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")

```

```

# Read target training data
y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

# Read estimated training data and add location feature
X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

# Read observed training data and add location feature
X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

# Read estimated test data and add location feature
X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

# Preprocess data
X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
X_test_estimated, y_train, loc)

X_trains.append(X_train)
X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

```

Processing location A...
Number of nans in sample_weight: 0
Number of nans in y: 0
Processing location B...
Number of nans in sample_weight: 0
Number of nans in y: 4
Processing location C...
Number of nans in sample_weight: 0
Number of nans in y: 6059

```

1 Feature engineering

```

[8]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y'], inplace=True)

if not do_drop_ds:
    # add hour datetime feature
    X_train["hour"] = X_train.index.hour
    X_test["hour"] = X_test.index.hour

#print(X_train.head())

```

```

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()

        # Calculate the size of each group for this location
        group_size = len(loc_df) // n_groups

        # Create a new 'group' column for this location
        loc_df['group'] = np.repeat(range(n_groups),
    ↪ repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

        # Append to list of grouped DataFrames
        grouped_dfs.append(loc_df)

    # Concatenate all the grouped DataFrames back together
    X_train = pd.concat(grouped_dfs)
    X_train.sort_index(inplace=True)
    print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

```

[9]: from autogluon.tabular import TabularDataset, TabularPredictor
from autogluon.timeseries import TimeSeriesDataFrame
import numpy as np
train_data = TabularDataset('X_train_raw.csv')

```

```

# set group column of train_data be increasing from 0 to 7 based on time, the
    ↪ first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
train_data['ds'] = pd.to_datetime(train_data['ds'])
train_data = train_data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
split_time = pd.to_datetime(train_data["ds"]).max() - pd.
    ↪ Timedelta(hours=tune_and_test_length)
train_set = TabularDataset(train_data[train_data["ds"] < split_time])
test_set = TabularDataset(train_data[train_data["ds"] >= split_time])
if use_groups:
    test_set = test_set.drop(columns=['group'])

if do_drop_ds:
    train_set = train_set.drop(columns=['ds'])
    test_set = test_set.drop(columns=['ds'])
    train_data = train_data.drop(columns=['ds'])

def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
    ↪ loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

tuning_data = None
if use_tune_data:
    train_data = train_set
    if use_test_data:
        # split test_set in half, use first half for tuning
        tuning_data, test_data = [], []
        for loc in locations:
            loc_test_set = test_set[test_set["location"] == loc]
            loc_tuning_data = loc_test_set.iloc[:len(loc_test_set)//2]
            loc_test_data = loc_test_set.iloc[len(loc_test_set)//2:]
            tuning_data.append(loc_tuning_data)
            test_data.append(loc_test_data)
        tuning_data = pd.concat(tuning_data)
        test_data = pd.concat(test_data)

```

```

        print("Shapes of tuning and test", tuning_data.shape[0], test_data.
↪shape[0], tuning_data.shape[0] + test_data.shape[0])

    else:
        tuning_data = test_set
        print("Shape of tuning", tuning_data.shape[0])

        # ensure sample weights for your tuning data sum to the number of rows in
↪the tuning data.
        tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        train_data = train_set
        test_data = test_set
        print("Shape of test", test_data.shape[0])

    # ensure sample weights for your training (or tuning) data sum to the number of
↪rows in the training (or tuning) data.
    train_data = normalize_sample_weights_per_location(train_data)
    if use_test_data:
        test_data = normalize_sample_weights_per_location(test_data)

```

Shape of test 5791

```

[10]: if run_analysis:
        import autogluon.eda.auto as auto
        auto.dataset_overview(train_data=train_data, test_data=test_data,
↪label="y", sample=None)

```

```

[11]: if run_analysis:
        auto.target_analysis(train_data=train_data, label="y")

```

2 Starting

```

[12]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
↪filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

```

```

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)

```

Last submission number: 82
 Now creating submission number: 83
 New filename: submission_83

```
[13]: predictors = [None, None, None]
```

```

[14]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪ train and tune data and test data
    print("Train data sample weight sum:", train_data[train_data["location"] ==
    ↪ loc]["sample_weight"].sum())
    print("Train data number of rows:", train_data[train_data["location"] ==
    ↪ loc].shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
    ↪ tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:", tuning_data[tuning_data["location"]
    ↪ == loc].shape[0])
    if use_test_data:
        print("Test data sample weight sum:", test_data[test_data["location"]
    ↪ == loc]["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"] ==
    ↪ loc].shape[0])
    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        sample_weight=sample_weight,
        weight_evaluation=weight_evaluation,
        groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc],
        time_limit=time_limit,
        #presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
    ↪ somethin, will be overwritten anyways
        tuning_data=tuning_data[tuning_data["location"] == loc] if
    ↪ use_tune_data else None,

```



```

        use_bag_holdout=use_bag_holdout,
        holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
        drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Training model for location A...

Train data sample weight sum: 31899.999999999996

Values in column 'sample_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission_83_A/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 304.26 GB / 315.93 GB (96.3%)

Train Data Rows: 31900

Train Data Columns: 46

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (5733.42, 0.0, 633.132, 1165.64686)

If 'regression' is not the correct problem_type, please manually specify the problem_type parameter during predictor init (You may specify problem_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 132461.02 MB

Train Data (Original) Memory Usage: 13.08 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set feature_metadata_in to manually specify special dtypes of the features.

```

Stage 1 Generators:
    Fitting AsTypeFeatureGenerator...
    Note: Converting 3 features to boolean dtype as they
only contain 2 unique values.
Stage 2 Generators:
    Fitting FillNaFeatureGenerator...
Stage 3 Generators:
    Fitting IdentityFeatureGenerator...
Stage 4 Generators:
    Fitting DropUniqueFeatureGenerator...

Train data number of rows: 31900
Test data sample weight sum: 2161
Test data number of rows: 2161

Stage 5 Generators:
    Fitting DropDuplicatesFeatureGenerator...
Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
    This is typically a feature which has the same value for all
rows.
    These features do not need to be present at inference time.
Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', []) : 1 | ['estimated_diff_hours']
Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', []) : 1 | ['estimated_diff_hours']
    ('int', ['bool']) : 3 | ['is_day:idx', 'is_in_shadow:idx',
'wind_speed_w_1000hPa:ms']
0.2s = Fit runtime
43 features in original data used to generate 43 features in processed
data.
Train Data (Processed) Memory Usage: 10.3 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.2s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with
holdout_frac=0.07836990595611286, Train Rows: 29400, Val Rows: 2500
User-specified model hyperparameters to be fit:
{

```

```

'NN_TORCH': {},
'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
'CAT': {},
'XGB': {},
'FASTAI': {},
'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif ... Training model for up to 1799.8s of the
1799.79s of remaining time.

```

-256.4362      = Validation score    (-mean_absolute_error)
0.04s         = Training    runtime
0.08s         = Validation runtime

```

Fitting model: KNeighborsDist ... Training model for up to 1799.67s of the
1799.66s of remaining time.

```

-258.904      = Validation score    (-mean_absolute_error)
0.04s         = Training    runtime
0.04s         = Validation runtime

```

Fitting model: LightGBMXT ... Training model for up to 1799.58s of the 1799.58s
of remaining time.

```

[1000] valid_set's l1: 160.225
[2000] valid_set's l1: 156.638
[3000] valid_set's l1: 154.856
[4000] valid_set's l1: 153.333
[5000] valid_set's l1: 152.53
[6000] valid_set's l1: 152.107
[7000] valid_set's l1: 151.576
[8000] valid_set's l1: 150.955
[9000] valid_set's l1: 150.397
[10000] valid_set's l1: 150.014

```

```

-149.9836     = Validation score    (-mean_absolute_error)
14.7s         = Training    runtime
0.19s         = Validation runtime

```

Fitting model: LightGBM ... Training model for up to 1784.31s of the 1784.31s of
remaining time.

```
[1000] valid_set's l1: 164.748
[2000] valid_set's l1: 163.139
[3000] valid_set's l1: 161.84
[4000] valid_set's l1: 161.376
[5000] valid_set's l1: 161.17
[6000] valid_set's l1: 161.09
[7000] valid_set's l1: 161.039
[8000] valid_set's l1: 161.06
```

```
-160.9804      = Validation score  (-mean_absolute_error)
13.12s      = Training  runtime
0.12s      = Validation runtime
```

Fitting model: RandomForestMSE ... Training model for up to 1770.89s of the 1770.88s of remaining time.

```
-168.3222      = Validation score  (-mean_absolute_error)
8.31s      = Training  runtime
0.09s      = Validation runtime
```

Fitting model: CatBoost ... Training model for up to 1762.02s of the 1762.01s of remaining time.

```
-164.7748      = Validation score  (-mean_absolute_error)
114.88s     = Training  runtime
0.01s      = Validation runtime
```

Fitting model: ExtraTreesMSE ... Training model for up to 1647.1s of the 1647.09s of remaining time.

```
-168.0615      = Validation score  (-mean_absolute_error)
1.76s      = Training  runtime
0.09s      = Validation runtime
```

Fitting model: NeuralNetFastAI ... Training model for up to 1644.77s of the 1644.76s of remaining time.

```
-173.685       = Validation score  (-mean_absolute_error)
27.5s      = Training  runtime
0.04s      = Validation runtime
```

Fitting model: XGBoost ... Training model for up to 1617.19s of the 1617.18s of remaining time.

```
-166.9422      = Validation score  (-mean_absolute_error)
2.95s      = Training  runtime
0.02s      = Validation runtime
```

Fitting model: NeuralNetTorch ... Training model for up to 1614.2s of the 1614.19s of remaining time.

```
-157.4766      = Validation score  (-mean_absolute_error)
53.37s     = Training  runtime
0.04s      = Validation runtime
```

Fitting model: LightGBMLarge ... Training model for up to 1560.78s of the 1560.77s of remaining time.

```
[1000] valid_set's l1: 152.501
[2000] valid_set's l1: 151.16
[3000] valid_set's l1: 150.868
[4000] valid_set's l1: 150.796
```

```

[5000] valid_set's l1: 150.752
[6000] valid_set's l1: 150.73
[7000] valid_set's l1: 150.721
[8000] valid_set's l1: 150.718
[9000] valid_set's l1: 150.715
[10000] valid_set's l1: 150.715

-150.7145      = Validation score    (-mean_absolute_error)
44.88s        = Training    runtime
0.27s         = Validation runtime

Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
1514.44s of remaining time.
-144.7512      = Validation score    (-mean_absolute_error)
0.47s          = Training    runtime
0.0s           = Validation runtime

AutoGluon training complete, total runtime = 286.08s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_83_A/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -188.93721938998613
      Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
  "mean_absolute_error": -188.93721938998613,
  "root_mean_squared_error": -412.50206821119434,
  "mean_squared_error": -170157.9562785128,
  "r2": 0.876539550093958,
  "pearsonr": 0.9364619268059786,
  "median_absolute_error": -11.871703147888184
}

Evaluation on test data:
-188.93721938998613

```

```

[15]: loc = "B"
      predictors[1] = fit_predictor_for_location(loc)

```

Values in column 'sample_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

```

Beginning AutoGluon training ... Time limit = 1800s
AutoGluon will save models to "AutogluonModels/submission_83_B/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64

```

```

Platform Version:   #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail:  303.27 GB / 315.93 GB (96.0%)
Train Data Rows:   30768
Train Data Columns: 46
Label Column:      y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (1152.3, -0.0, 97.74541, 195.0957)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                  130687.49 MB
    Train Data (Original) Memory Usage: 12.62 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 3 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...

Training model for location B...
Train data sample weight sum: 30768.0
Train data number of rows: 30768
Test data sample weight sum: 2051
Test data number of rows: 2051

    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', [])   : 1 | ['estimated_diff_hours']
    Types of features in processed data (raw dtype, special dtypes):

```

```

('float', [])      : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
('int', [])        : 1 | ['estimated_diff_hours']
('int', ['bool'])  : 3 | ['is_day:idx', 'is_in_shadow:idx',
'wind_speed_w_1000hPa:ms']
0.2s = Fit runtime
43 features in original data used to generate 43 features in processed
data.

Train Data (Processed) Memory Usage: 9.94 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.2s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'

This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with
holdout_frac=0.0812532501300052, Train Rows: 28268, Val Rows: 2500
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Fitting 11 L1 models ...
Fitting model: KNeighborsUnif ... Training model for up to 1799.8s of the
1799.8s of remaining time.
-53.446 = Validation score    (-mean_absolute_error)
0.04s   = Training runtime
0.04s   = Validation runtime
Fitting model: KNeighborsDist ... Training model for up to 1799.72s of the
1799.71s of remaining time.
-53.26  = Validation score    (-mean_absolute_error)

```

```

0.04s      = Training    runtime
0.04s      = Validation runtime
Fitting model: LightGBMXT ... Training model for up to 1799.63s of the 1799.63s
of remaining time.

[1000]  valid_set's l1: 33.1184
[2000]  valid_set's l1: 31.2902
[3000]  valid_set's l1: 30.3861
[4000]  valid_set's l1: 29.7745
[5000]  valid_set's l1: 29.2881
[6000]  valid_set's l1: 28.9165
[7000]  valid_set's l1: 28.69
[8000]  valid_set's l1: 28.5065
[9000]  valid_set's l1: 28.3414
[10000] valid_set's l1: 28.1891

-28.1891      = Validation score  (-mean_absolute_error)
13.87s       = Training    runtime
0.18s        = Validation runtime
Fitting model: LightGBM ... Training model for up to 1785.29s of the 1785.29s of
remaining time.

[1000]  valid_set's l1: 30.9525
[2000]  valid_set's l1: 29.5944
[3000]  valid_set's l1: 28.9858
[4000]  valid_set's l1: 28.7098
[5000]  valid_set's l1: 28.5117
[6000]  valid_set's l1: 28.3986
[7000]  valid_set's l1: 28.3279
[8000]  valid_set's l1: 28.2614
[9000]  valid_set's l1: 28.2242
[10000] valid_set's l1: 28.2037

-28.2036      = Validation score  (-mean_absolute_error)
14.2s         = Training    runtime
0.17s         = Validation runtime
Fitting model: RandomForestMSE ... Training model for up to 1770.63s of the
1770.62s of remaining time.

-32.8223      = Validation score  (-mean_absolute_error)
9.33s         = Training    runtime
0.1s          = Validation runtime
Fitting model: CatBoost ... Training model for up to 1760.83s of the 1760.82s of
remaining time.

-30.4129      = Validation score  (-mean_absolute_error)
115.23s       = Training    runtime
0.01s         = Validation runtime
Fitting model: ExtraTreesMSE ... Training model for up to 1645.56s of the
1645.55s of remaining time.

-33.619       = Validation score  (-mean_absolute_error)
1.89s         = Training    runtime

```



```

    0.09s      = Validation runtime
Fitting model: NeuralNetFastAI ... Training model for up to 1643.17s of the
1643.16s of remaining time.
    -37.5693      = Validation score      (-mean_absolute_error)
    26.48s      = Training      runtime
    0.04s      = Validation runtime
Fitting model: XGBoost ... Training model for up to 1616.62s of the 1616.61s of
remaining time.
    -31.0815      = Validation score      (-mean_absolute_error)
    24.5s      = Training      runtime
    0.21s      = Validation runtime
Fitting model: NeuralNetTorch ... Training model for up to 1591.76s of the
1591.75s of remaining time.
    -31.5314      = Validation score      (-mean_absolute_error)
    103.34s     = Training      runtime
    0.04s      = Validation runtime
Fitting model: LightGBMLarge ... Training model for up to 1488.37s of the
1488.36s of remaining time.

[1000] valid_set's l1: 28.0771
[2000] valid_set's l1: 27.2322
[3000] valid_set's l1: 27.0078
[4000] valid_set's l1: 26.931
[5000] valid_set's l1: 26.9056
[6000] valid_set's l1: 26.8914
[7000] valid_set's l1: 26.8859
[8000] valid_set's l1: 26.8825
[9000] valid_set's l1: 26.8809
[10000] valid_set's l1: 26.88

    -26.88      = Validation score      (-mean_absolute_error)
    51.57s     = Training      runtime
    0.39s      = Validation runtime
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
1435.05s of remaining time.
    -26.5096      = Validation score      (-mean_absolute_error)
    0.45s      = Training      runtime
    0.0s      = Validation runtime
AutoGluon training complete, total runtime = 365.44s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_83_B/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -37.193940933616126
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{

```

```

    "mean_absolute_error": -37.193940933616126,
    "root_mean_squared_error": -81.67094023712487,
    "mean_squared_error": -6670.1424792160215,
    "r2": 0.7854684374345835,
    "pearsonr": 0.9089080478259688,
    "median_absolute_error": -8.027046203613281
}

```

Evaluation on test data:
-37.193940933616126

```

[16]: loc = "C"
      predictors[2] = fit_predictor_for_location(loc)

```

Values in column 'sample_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission_83_C/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 302.33 GB / 315.93 GB (95.7%)

Train Data Rows: 24492

Train Data Columns: 46

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and label-values can't be converted to int).

Label info (max, min, mean, stddev): (999.6, 0.0, 78.11911, 167.50151)

If 'regression' is not the correct problem_type, please manually specify the problem_type parameter during predictor init (You may specify problem_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 130450.71 MB

Train Data (Original) Memory Usage: 10.04 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set feature_metadata_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 2 features to boolean dtype as they only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

```

Stage 4 Generators:
    Fitting DropUniqueFeatureGenerator...
Stage 5 Generators:
    Fitting DropDuplicatesFeatureGenerator...

Training model for location C...
Train data sample weight sum: 24492.000000000004
Train data number of rows: 24492
Test data sample weight sum: 1579
Test data number of rows: 1579

    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
    This is typically a feature which has the same value for all
rows.
    These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['estimated_diff_hours']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['estimated_diff_hours']
        ('int', ['bool']) : 2 | ['is_day:idx', 'is_in_shadow:idx']
0.1s = Fit runtime
43 features in original data used to generate 43 features in processed
data.
    Train Data (Processed) Memory Usage: 8.08 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.18s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with holdout_frac=0.1, Train
Rows: 22042, Val Rows: 2450
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',

```

```

'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif ... Training model for up to 1799.82s of the
1799.82s of remaining time.

```

-30.0314          = Validation score    (-mean_absolute_error)
0.04s            = Training    runtime
0.03s            = Validation runtime

```

Fitting model: KNeighborsDist ... Training model for up to 1799.74s of the
1799.74s of remaining time.

```

-30.0815          = Validation score    (-mean_absolute_error)
0.03s            = Training    runtime
0.03s            = Validation runtime

```

Fitting model: LightGBMXT ... Training model for up to 1799.66s of the 1799.66s
of remaining time.

```

[1000] valid_set's l1: 17.286
[2000] valid_set's l1: 16.8066
[3000] valid_set's l1: 16.5904
[4000] valid_set's l1: 16.4733
[5000] valid_set's l1: 16.4109
[6000] valid_set's l1: 16.3809
[7000] valid_set's l1: 16.3455
[8000] valid_set's l1: 16.3286
[9000] valid_set's l1: 16.3022
[10000] valid_set's l1: 16.2884

```

```

-16.2876          = Validation score    (-mean_absolute_error)
14.45s           = Training    runtime
0.19s            = Validation runtime

```

Fitting model: LightGBM ... Training model for up to 1784.76s of the 1784.75s of
remaining time.

```

[1000] valid_set's l1: 17.3864

```

```

-17.2471          = Validation score    (-mean_absolute_error)
3.11s             = Training    runtime
0.03s             = Validation runtime

```

Fitting model: RandomForestMSE ... Training model for up to 1781.57s of the
1781.56s of remaining time.

```

-18.3698          = Validation score    (-mean_absolute_error)
5.22s            = Training   runtime
0.1s             = Validation runtime
Fitting model: CatBoost ... Training model for up to 1776.08s of the 1776.08s of
remaining time.
-17.1912          = Validation score    (-mean_absolute_error)
125.01s          = Training   runtime
0.01s            = Validation runtime
Fitting model: ExtraTreesMSE ... Training model for up to 1651.02s of the
1651.02s of remaining time.
-18.278           = Validation score    (-mean_absolute_error)
1.11s            = Training   runtime
0.08s            = Validation runtime
Fitting model: NeuralNetFastAI ... Training model for up to 1649.65s of the
1649.65s of remaining time.
-18.5803          = Validation score    (-mean_absolute_error)
20.55s           = Training   runtime
0.03s            = Validation runtime
Fitting model: XGBoost ... Training model for up to 1629.04s of the 1629.03s of
remaining time.
-17.1699          = Validation score    (-mean_absolute_error)
24.02s           = Training   runtime
0.22s            = Validation runtime
Fitting model: NeuralNetTorch ... Training model for up to 1604.65s of the
1604.64s of remaining time.
-17.2237          = Validation score    (-mean_absolute_error)
79.07s           = Training   runtime
0.04s            = Validation runtime
Fitting model: LightGBMLarge ... Training model for up to 1525.53s of the
1525.53s of remaining time.

[1000]  valid_set's l1: 16.652
[2000]  valid_set's l1: 16.5555
[3000]  valid_set's l1: 16.5324
[4000]  valid_set's l1: 16.5278
[5000]  valid_set's l1: 16.5262
[6000]  valid_set's l1: 16.5256
[7000]  valid_set's l1: 16.5254
[8000]  valid_set's l1: 16.5254
[9000]  valid_set's l1: 16.5254
[10000] valid_set's l1: 16.5254

-16.5254          = Validation score    (-mean_absolute_error)
46.16s           = Training   runtime
0.35s            = Validation runtime
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
1477.7s of remaining time.
-15.5368          = Validation score    (-mean_absolute_error)
0.45s            = Training   runtime

```

```

0.0s      = Validation runtime
AutoGluon training complete, total runtime = 322.79s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_83_C/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -30.79885044646661
      Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -30.79885044646661,
    "root_mean_squared_error": -63.830887679445,
    "mean_squared_error": -4074.382221945923,
    "r2": 0.7863794094990258,
    "pearsonr": 0.8938476674072768,
    "median_absolute_error": -2.830535888671875
}

Evaluation on test data:
-30.79885044646661

```

3 Submit

```

[17]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data

```

```

Loaded data from: X_train_raw.csv | Columns = 48 / 48 | Rows = 92951 -> 92951
Loaded data from: X_test_raw.csv | Columns = 47 / 47 | Rows = 2160 -> 2160

```

```

[18]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
↪ "location"], left_on=["ds", "location"])

#test_data_merged

```

```

Loaded data from: test.csv | Columns = 4 / 4 | Rows = 2160 -> 2160

```

```
[19]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in test_data.groupby('location'):
    i = location_map[loc]
    subset = test_data_merged[test_data_merged["location"] == loc].
    ↪reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)



    # get past predictions
    past_pred = predictors[i].
    ↪predict(train_data_with_dates[train_data_with_dates["location"] == loc])
    train_data_with_dates.loc[train_data_with_dates["location"] == loc,
    ↪"prediction"] = past_pred
```


```
[20]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='y', ax=ax, label="train data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='prediction', ax=ax, label="past predictions")

    # title
    ax.set_title(f"Predictions for location {loc}")
```





```
[21]: # concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df
```

```
[21]:
```

	id	prediction
0	0	-2.093352
1	1	-1.161028
2	2	1.035955
3	3	39.386139
4	4	346.093262
..
715	2155	50.987034
716	2156	33.331429
717	2157	10.217902
718	2158	2.306329
719	2159	1.580828

[2160 rows x 2 columns]

```
[22]: # Save the submission DataFrame to submissions folder, create new name based on
↳ last submission, format is submission_<last_submission_number + 1>.csv

# Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
↳ index=False)
print("jallia")
```

Saving submission to submissions/submission_83.csv
jall1a

```
[23]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)
```

<IPython.core.display.Javascript object>

```
[24]: # save this notebook to submissions folder
import subprocess
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f'hei.pdf'), "autogluon_each_location.ipynb"])
```

[NbConvertApp] Converting notebook autogluon_each_location.ipynb to pdf
/opt/conda/lib/python3.10/site-packages/nbconvert/utils/pandoc.py:51:
RuntimeWarning: You are using an unsupported version of pandoc (2.9.2.1).
Your version must be at least (2.14.2) but less than (4.0.0).
Refer to <https://pandoc.org/installing.html>.
Continuing with doubts...
check_pandoc_version()
[NbConvertApp] Writing 110252 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 89019 bytes to notebook_pdfs/hei.pdf

```
[24]: CompletedProcess(args=['jupyter', 'nbconvert', '--to', 'pdf', '--output',
    'notebook_pdfs/hei.pdf', 'autogluon_each_location.ipynb'], returncode=0)
```

```
[25]: # feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
    ↪time_limit=60*10)
```

These features in provided data are not utilized by the predictor and will be

ignored: ['ds', 'elevation:m', 'sample_weight', 'location', 'prediction']
Computing feature importance via permutation shuffling for 43 features using
4394 rows with 10 shuffle sets... Time limit: 600s...

602.15s = Expected runtime (60.21s per shuffle set)

311.81s = Actual runtime (Completed 10 of 10 shuffle sets)

[25]:	importance	stddev	p_value	n \
direct_rad:W	1.517037e+02	2.289968	3.273217e-18	10
clear_sky_rad:W	8.582451e+01	1.703378	3.841559e-17	10
diffuse_rad:W	7.504888e+01	2.161663	1.095237e-15	10
sun_azimuth:d	5.999013e+01	3.158821	2.479818e-13	10
sun_elevation:d	4.082302e+01	1.096936	5.863730e-16	10
direct_rad_1h:J	2.890493e+01	0.671130	1.575765e-16	10
clear_sky_energy_1h:J	2.494764e+01	1.374852	3.732751e-13	10
diffuse_rad_1h:J	1.564581e+01	0.809628	2.121050e-13	10
effective_cloud_cover:p	1.410651e+01	0.829709	6.696896e-13	10
total_cloud_cover:p	1.408536e+01	0.635340	6.176382e-14	10
wind_speed_u_10m:ms	1.007212e+01	1.198845	3.665292e-10	10
cloud_base_agl:m	7.394463e+00	0.472747	1.415201e-12	10
is_day:idx	6.217252e+00	0.317781	1.898621e-13	10
snow_water:kgm2	6.144440e+00	0.759398	5.122269e-10	10
fresh_snow_24h:cm	5.312308e+00	0.572670	1.516854e-10	10
relative_humidity_1000hPa:p	5.234821e+00	0.633852	4.268981e-10	10
visibility:m	5.207509e+00	0.492373	4.705255e-11	10
pressure_50m:hPa	4.911803e+00	0.681923	1.437898e-09	10
is_in_shadow:idx	4.868039e+00	0.260064	2.823661e-13	10
ceiling_height_agl:m	4.435441e+00	0.531354	3.882061e-10	10
pressure_100m:hPa	4.334008e+00	0.611214	1.652148e-09	10
wind_speed_10m:ms	4.199136e+00	0.568926	1.158211e-09	10
wind_speed_v_10m:ms	4.010390e+00	0.698641	1.065526e-08	10
sfc_pressure:hPa	3.930014e+00	0.564905	1.955842e-09	10
msl_pressure:hPa	3.814023e+00	0.395453	1.071865e-10	10
air_density_2m:kgm3	1.902994e+00	0.710650	7.007531e-06	10
t_1000hPa:K	1.902247e+00	1.067549	1.598663e-04	10
estimated_diff_hours	1.826127e+00	0.185571	8.958305e-11	10
fresh_snow_6h:cm	1.756706e+00	0.210241	3.848165e-10	10
fresh_snow_12h:cm	1.626520e+00	0.309084	2.282256e-08	10
super_cooled_liquid_water:kgm2	1.439939e+00	0.356188	2.241701e-07	10
snow_depth:cm	1.377273e+00	0.412376	1.133366e-06	10
precip_5min:mm	1.014655e+00	0.413575	1.413010e-05	10
dew_point_2m:K	9.633977e-01	0.399315	1.613398e-05	10
fresh_snow_3h:cm	9.372222e-01	0.194065	4.825074e-08	10
dew_or_rime:idx	6.006628e-01	0.177677	1.023541e-06	10
precip_type_5min:idx	5.757253e-01	0.247069	2.120778e-05	10
fresh_snow_1h:cm	4.291272e-01	0.229121	1.113596e-04	10
rain_water:kgm2	1.917926e-01	0.115546	2.641963e-04	10
prob_rime:p	1.070263e-01	0.155762	2.892224e-02	10

absolute_humidity_2m:gm3	2.790905e-02	0.160047	2.973783e-01	10
wind_speed_w_1000hPa:ms	-1.085141e-10	0.000000	5.000000e-01	10
snow_melt_10min:mm	-1.138008e-01	0.137336	9.861026e-01	10

	p99_high	p99_low
direct_rad:W	1.540571e+02	1.493504e+02
clear_sky_rad:W	8.757506e+01	8.407397e+01
diffuse_rad:W	7.727039e+01	7.282736e+01
sun_azimuth:d	6.323642e+01	5.674385e+01
sun_elevation:d	4.195033e+01	3.969571e+01
direct_rad_1h:J	2.959464e+01	2.821521e+01
clear_sky_energy_1h:J	2.636056e+01	2.353472e+01
diffuse_rad_1h:J	1.647785e+01	1.481376e+01
effective_cloud_cover:p	1.495919e+01	1.325383e+01
total_cloud_cover:p	1.473829e+01	1.343243e+01
wind_speed_u_10m:ms	1.130416e+01	8.840079e+00
cloud_base_agl:m	7.880300e+00	6.908627e+00
is_day:idx	6.543833e+00	5.890672e+00
snow_water:kgm2	6.924864e+00	5.364017e+00
fresh_snow_24h:cm	5.900834e+00	4.723782e+00
relative_humidity_1000hPa:p	5.886223e+00	4.583418e+00
visibility:m	5.713515e+00	4.701503e+00
pressure_50m:hPa	5.612607e+00	4.210999e+00
is_in_shadow:idx	5.135303e+00	4.600774e+00
ceiling_height_agl:m	4.981508e+00	3.889375e+00
pressure_100m:hPa	4.962146e+00	3.705871e+00
wind_speed_10m:ms	4.783815e+00	3.614458e+00
wind_speed_v_10m:ms	4.728375e+00	3.292405e+00
sfc_pressure:hPa	4.510561e+00	3.349468e+00
msl_pressure:hPa	4.220426e+00	3.407621e+00
air_density_2m:kgm3	2.633321e+00	1.172667e+00
t_1000hPa:K	2.999354e+00	8.051388e-01
estimated_diff_hours	2.016836e+00	1.635417e+00
fresh_snow_6h:cm	1.972768e+00	1.540643e+00
fresh_snow_12h:cm	1.944162e+00	1.308877e+00
super_cooled_liquid_water:kgm2	1.805989e+00	1.073889e+00
snow_depth:cm	1.801067e+00	9.534785e-01
precip_5min:mm	1.439682e+00	5.896286e-01
dew_point_2m:K	1.373769e+00	5.530260e-01
fresh_snow_3h:cm	1.136661e+00	7.377835e-01
dew_or_rime:idx	7.832597e-01	4.180659e-01
precip_type_5min:idx	8.296352e-01	3.218155e-01
fresh_snow_1h:cm	6.645924e-01	1.936620e-01
rain_water:kgm2	3.105377e-01	7.304745e-02
prob_rime:p	2.671008e-01	-5.304827e-02
absolute_humidity_2m:gm3	1.923874e-01	-1.365693e-01
wind_speed_w_1000hPa:ms	-1.085141e-10	-1.085141e-10

snow_melt_10min:mm 2.733802e-02 -2.549395e-01

```
[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
    ↪time_limit=60*10)
```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'sample_weight', 'location', 'prediction']
Computing feature importance via permutation shuffling for 43 features using 5000 rows with 10 shuffle sets... Time limit: 600s...
628.12s = Expected runtime (62.81s per shuffle set)

```
[ ]: display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)

subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↪"autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
            ↪stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
            ↪strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
    ↪'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', 'hello if hello is
    ↪not None else 'Henrik eller Jørgen'])

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"
```

```

# execute_git_command git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command git_repo_path, ['add', '.'])
# execute_git_command git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote
# output, success = execute_git_command git_repo_path, ['push',
↳ 'origin', branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command git_repo_path, ['push', '--set-upstream',
↳ 'origin', branch_name])
#     execute_git_command git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command git_repo_path, ['checkout', 'main'])

```