

## autogluon\_each\_location

October 9, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*30
presets = 'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = [] # "hour", "dayofweek", "day", "month", "year"
use_estimated_diff_attr = False
use_is_estimated_attr = True

use_groups = True
n_groups = 8

auto_stack = True
num_stack_levels = 1
num_bag_folds = 0
if auto_stack:
    num_stack_levels = None
    num_bag_folds = None

use_tune_data = False
use_test_data = True
tune_and_test_length = 24*30*3 # 3 months from end, this changes the
    ↪ evaluations for only test
holdout_frac = None
use_bag_holdout = False # Enable this if there is a large gap between score_val
    ↪ and score_test in stack models.

sample_weight = 'sample_weight' #None
weight_evaluation = True #False
sample_weight_estimated = 1 # this changes evaluations for test and tune WTF,
    ↪ cant find a fix

run_analysis = False
```

```

[2]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings("ignore")

def fix_datetime(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    ↪with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    # Drop rows where the minute part of the time is not 0
    X = X[X.index.minute == 0].copy()
    return X

def convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_datetime(X_train_observed, "X_train_observed")
    X_train_estimated = fix_datetime(X_train_estimated, "X_train_estimated")
    X_test = fix_datetime(X_test, "X_test")

    # add sample weights, which are 1 for observed and 3 for estimated
    X_train_observed["sample_weight"] = 1
    X_train_estimated["sample_weight"] = sample_weight_estimated
    X_test["sample_weight"] = sample_weight_estimated

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
        ↪pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
        X_test["estimated_diff_hours"] = (X_test.index - pd.
        ↪to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

        X_train_estimated["estimated_diff_hours"] =
        ↪X_train_estimated["estimated_diff_hours"].astype('int64')
        # the filled once will get dropped later anyways, when we drop y nans
        X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
        ↪fillna(-50).astype('int64')

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0

```

```

X_train_estimated["is_estimated"] = 1
X_test["is_estimated"] = 1

X_train_estimated.drop(columns=['date_calc'], inplace=True)
X_test.drop(columns=['date_calc'], inplace=True)

y_train['ds'] = pd.to_datetime(y_train['time'])
y_train.drop(columns=['time'], inplace=True)
y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
↳location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
↳convert_to_datetime(X_train_observed, X_train_estimated, X_test, y_train)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # fill missng sample_weight with 3
    #X_train["sample_weight"] = X_train["sample_weight"].fillna(0)

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

    # print number of nans in sample_weight
    print(f"Number of nans in sample_weight: {X_train['sample_weight'].isna().
↳sum()}")

    # print number of nans in y
    print(f"Number of nans in y: {X_train['y'].isna().sum()}")

    X_train["location"] = location
    X_test["location"] = location

```

```

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
    ↪X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

```

Processing location A...
Number of nans in sample_weight: 0
Number of nans in y: 0
Processing location B...
Number of nans in sample_weight: 0
Number of nans in y: 4
Processing location C...
Number of nans in sample_weight: 0
Number of nans in y: 6059

```

# 1 Feature engineering

```
[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y'], inplace=True)

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

print(X_train.head())

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()

        # Calculate the size of each group for this location
        group_size = len(loc_df) // n_groups

        # Create a new 'group' column for this location
        loc_df['group'] = np.repeat(range(n_groups),
    ↪ repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

        # Append to list of grouped DataFrames
        grouped_dfs.append(loc_df)

    # Concatenate all the grouped DataFrames back together
    X_train = pd.concat(grouped_dfs)
    X_train.sort_index(inplace=True)
    print(X_train["group"].head())
```

```
to_drop = ["snow_drift:idx", "snow_density:kgm3"]
```

```
X_train.drop(columns=to_drop, inplace=True)
```

```
X_test.drop(columns=to_drop, inplace=True)
```

```
X_train.to_csv('X_train_raw.csv', index=True)
```

```
X_test.to_csv('X_test_raw.csv', index=True)
```

```

absolute_humidity_2m:gm3  air_density_2m:kgm3  \
ds
2019-06-02 22:00:00          7.7          1.230
2019-06-02 23:00:00          7.7          1.225
2019-06-03 00:00:00          7.7          1.221
2019-06-03 01:00:00          8.2          1.218
2019-06-03 02:00:00          8.8          1.219

```

```

ceiling_height_agl:m  clear_sky_energy_1h:J  \
ds
2019-06-02 22:00:00      1744.900024          0.000000
2019-06-02 23:00:00      1703.599976          0.000000
2019-06-03 00:00:00      1668.099976          0.000000
2019-06-03 01:00:00      1388.400024          0.000000
2019-06-03 02:00:00      1108.500000      6546.899902

```

```

clear_sky_rad:W  cloud_base_agl:m  dew_or_rime:idx  \
ds
2019-06-02 22:00:00          0.0      1744.900024          0.0
2019-06-02 23:00:00          0.0      1703.599976          0.0
2019-06-03 00:00:00          0.0      1668.099976          0.0
2019-06-03 01:00:00          0.0      1388.400024          0.0
2019-06-03 02:00:00          9.8      1108.500000          0.0

```

```

dew_point_2m:K  diffuse_rad:W  diffuse_rad_1h:J  ...  \
ds
2019-06-02 22:00:00      280.299988          0.0          0.000000  ...
2019-06-02 23:00:00      280.299988          0.0          0.000000  ...
2019-06-03 00:00:00      280.200012          0.0          0.000000  ...
2019-06-03 01:00:00      281.299988          0.0          0.000000  ...
2019-06-03 02:00:00      282.299988          4.3      7743.299805  ...

```

```

total_cloud_cover:p  visibility:m  wind_speed_10m:ms  \
ds
2019-06-02 22:00:00          100.0  39640.101562          3.7
2019-06-02 23:00:00          100.0  41699.898438          3.5
2019-06-03 00:00:00          100.0  20473.000000          3.2

```

2019-06-03 01:00:00	100.0	2104.600098	2.8
2019-06-03 02:00:00	100.0	2681.600098	2.7

	wind_speed_u_10m:ms	wind_speed_v_10m:ms	\
ds			
2019-06-02 22:00:00	-3.6	-0.8	
2019-06-02 23:00:00	-3.5	0.0	
2019-06-03 00:00:00	-3.1	0.7	
2019-06-03 01:00:00	-2.7	0.8	
2019-06-03 02:00:00	-2.5	1.0	

	wind_speed_w_1000hPa:ms	sample_weight	is_estimated	\
ds				
2019-06-02 22:00:00	-0.0	1	0	
2019-06-02 23:00:00	-0.0	1	0	
2019-06-03 00:00:00	-0.0	1	0	
2019-06-03 01:00:00	-0.0	1	0	
2019-06-03 02:00:00	-0.0	1	0	

	y	location
ds		
2019-06-02 22:00:00	0.00	A
2019-06-02 23:00:00	0.00	A
2019-06-03 00:00:00	0.00	A
2019-06-03 01:00:00	0.00	A
2019-06-03 02:00:00	19.36	A

[5 rows x 49 columns]

ds	
2019-01-01 00:00:00	0
2019-01-01 01:00:00	0
2019-01-01 02:00:00	0
2019-01-01 03:00:00	0
2019-01-01 04:00:00	0

Name: group, dtype: int64

```
[4]: from autogluon.tabular import TabularDataset, TabularPredictor
from autogluon.timeseries import TimeSeriesDataFrame
import numpy as np
train_data = TabularDataset('X_train_raw.csv')
# set group column of train_data be increasing from 0 to 7 based on time, the
# first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
train_data['ds'] = pd.to_datetime(train_data['ds'])
train_data = train_data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
```

```

#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
split_time = pd.to_datetime(train_data["ds"]).max() - pd.
    ↳ Timedelta(hours=tune_and_test_length)
train_set = TabularDataset(train_data[train_data["ds"] < split_time])
test_set = TabularDataset(train_data[train_data["ds"] >= split_time])
if use_groups:
    test_set = test_set.drop(columns=['group'])

if do_drop_ds:
    train_set = train_set.drop(columns=['ds'])
    test_set = test_set.drop(columns=['ds'])
    train_data = train_data.drop(columns=['ds'])

def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
    ↳ loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

tuning_data = None
if use_tune_data:
    train_data = train_set
    if use_test_data:
        # split test_set in half, use first half for tuning
        tuning_data, test_data = [], []
        for loc in locations:
            loc_test_set = test_set[test_set["location"] == loc]
            loc_tuning_data = loc_test_set.iloc[:len(loc_test_set)//2]
            loc_test_data = loc_test_set.iloc[len(loc_test_set)//2:]
            tuning_data.append(loc_tuning_data)
            test_data.append(loc_test_data)
        tuning_data = pd.concat(tuning_data)
        test_data = pd.concat(test_data)
        print("Shapes of tuning and test", tuning_data.shape[0], test_data.
    ↳ shape[0], tuning_data.shape[0] + test_data.shape[0])

    else:
        tuning_data = test_set
        print("Shape of tuning", tuning_data.shape[0])

```



```

    # ensure sample weights for your tuning data sum to the number of rows in
    ↪ the tuning data.
    tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        train_data = train_set
        test_data = test_set
        print("Shape of test", test_data.shape[0])

# ensure sample weights for your training (or tuning) data sum to the number of
↪ rows in the training (or tuning) data.
train_data = normalize_sample_weights_per_location(train_data)
if use_test_data:
    test_data = normalize_sample_weights_per_location(test_data)

```

Shape of test 5791

```

[5]: if run_analysis:
        import autogluon.eda.auto as auto
        auto.dataset_overview(train_data=train_data, test_data=test_data,
        ↪ label="y", sample=None)

```

```

[6]: if run_analysis:
        auto.target_analysis(train_data=train_data, label="y")

```

## 2 Starting

```

[7]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    ↪ filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)

```

Last submission number: 84  
Now creating submission number: 85  
New filename: submission\_85

```
[8]: predictors = [None, None, None]
```

```
[9]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪ train and tune data and test data
    print("Train data sample weight sum:", train_data[train_data["location"] == loc]
    ↪ ["sample_weight"].sum())
    print("Train data number of rows:", train_data[train_data["location"] == loc]
    ↪ .shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
        ↪ tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:", tuning_data[tuning_data["location"] == loc]
        ↪ .shape[0])
    if use_test_data:
        print("Test data sample weight sum:", test_data[test_data["location"] == loc]
        ↪ ["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"] == loc]
        ↪ .shape[0])
    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        sample_weight=sample_weight,
        weight_evaluation=weight_evaluation,
        groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc],
        time_limit=time_limit,
        #presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
        ↪ somethin, will be overwritten anyways
        tuning_data=tuning_data[tuning_data["location"] == loc] if
        ↪ use_tune_data else None,
        use_bag_holdout=use_bag_holdout,
        holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
```

```

        t = test_data[test_data["location"] == loc]#.
        ↪drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Warning: path already exists! This predictor may overwrite an existing predictor! path="AutogluonModels/submission\_85\_A"

Training model for location A...

Train data sample weight sum: 31900

Train data number of rows: 31900

Test data sample weight sum: 2161

Test data number of rows: 2161

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Values in column 'group' used as split folds instead of being automatically set.

Bagged models will have 8 splits.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission\_85\_A/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 307.23 GB / 315.93 GB (97.2%)

Train Data Rows: 31900

Train Data Columns: 47

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (5733.42, 0.0, 633.132, 1165.64686)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 132348.32 MB

Train Data (Original) Memory Usage: 13.08 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set feature\_metadata\_in to manually specify special dtypes of the features.

```

Stage 1 Generators:
    Fitting AsTypeFeatureGenerator...
    Note: Converting 4 features to boolean dtype as they
only contain 2 unique values.
Stage 2 Generators:
    Fitting FillNaFeatureGenerator...
Stage 3 Generators:
    Fitting IdentityFeatureGenerator...
Stage 4 Generators:
    Fitting DropUniqueFeatureGenerator...
Stage 5 Generators:
    Fitting DropDuplicatesFeatureGenerator...
Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
    This is typically a feature which has the same value for all
rows.
    These features do not need to be present at inference time.
Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', [])   : 1 | ['is_estimated']
Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', ['bool']) : 4 | ['is_day:idx', 'is_in_shadow:idx',
'wind_speed_w_1000hPa:ms', 'is_estimated']
0.2s = Fit runtime
43 features in original data used to generate 43 features in processed
data.
Train Data (Processed) Memory Usage: 10.08 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.25s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',

```

```

'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1799.75s of the 1799.75s of remaining time.

```

-299.6339      = Validation score    (-mean_absolute_error)
0.04s         = Training    runtime
0.4s          = Validation runtime

```

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 1799.23s of the 1799.23s of remaining time.

```

-300.6895      = Validation score    (-mean_absolute_error)
0.04s          = Training    runtime
0.4s           = Validation runtime

```

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 1798.74s of the 1798.74s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-208.9611      = Validation score    (-mean_absolute_error)
1.82s          = Training    runtime
0.18s          = Validation runtime

```

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 1791.54s of the 1791.53s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-208.8242      = Validation score    (-mean_absolute_error)
2.68s          = Training    runtime
0.14s          = Validation runtime

```

Fitting model: RandomForestMSE\_BAG\_L1 ... Training model for up to 1787.7s of the 1787.7s of remaining time.

```

-192.1045      = Validation score    (-mean_absolute_error)
8.17s          = Training    runtime
1.24s          = Validation runtime

```

Fitting model: CatBoost\_BAG\_L1 ... Training model for up to 1776.24s of the 1776.24s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-212.9408      = Validation score    (-mean_absolute_error)
46.78s         = Training    runtime

```

```

    0.06s      = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1728.4s of the
1728.39s of remaining time.
    -191.928      = Validation score    (-mean_absolute_error)
    1.83s      = Training    runtime
    1.17s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1723.32s of
the 1723.32s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -212.2221      = Validation score    (-mean_absolute_error)
    38.7s      = Training    runtime
    0.51s      = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 1683.54s of the
1683.53s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -207.9489      = Validation score    (-mean_absolute_error)
    2.25s      = Training    runtime
    0.11s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1679.48s of
the 1679.48s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -199.6128      = Validation score    (-mean_absolute_error)
    59.51s      = Training    runtime
    0.41s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1618.63s of the
1618.62s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -208.4925      = Validation score    (-mean_absolute_error)
    5.47s      = Training    runtime
    0.2s      = Validation runtime
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
1611.85s of remaining time.
    -187.6368      = Validation score    (-mean_absolute_error)
    0.84s      = Training    runtime
    0.0s      = Validation runtime
AutoGluon training complete, total runtime = 189.05s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_85_A/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -189.8488279945972
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.

```

Evaluations on test data:

```
{
  "mean_absolute_error": -189.8488279945972,
  "root_mean_squared_error": -417.9352031434312,
  "mean_squared_error": -174669.83402654107,
  "r2": 0.873265895021486,
  "pearsonr": 0.9349789381714066,
  "median_absolute_error": -2.470029582977295
}
```

Evaluation on test data:

-189.8488279945972

```
[10]: loc = "B"
      predictors[1] = fit_predictor_for_location(loc)
```

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Values in column 'group' used as split folds instead of being automatically set. Bagged models will have 8 splits.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission\_85\_B/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 307.18 GB / 315.93 GB (97.2%)

Train Data Rows: 30768

Train Data Columns: 47

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (1152.3, -0.0, 97.74541, 195.0957)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 131098.62 MB

Train Data (Original) Memory Usage: 12.62 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set

feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 4 features to boolean dtype as they only contain 2 unique values.

```

Stage 2 Generators:
    Fitting FillNaFeatureGenerator...
Stage 3 Generators:
    Fitting IdentityFeatureGenerator...
Stage 4 Generators:
    Fitting DropUniqueFeatureGenerator...

Training model for location B...
Train data sample weight sum: 30768
Train data number of rows: 30768
Test data sample weight sum: 2051
Test data number of rows: 2051

Stage 5 Generators:
    Fitting DropDuplicatesFeatureGenerator...
Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
    This is typically a feature which has the same value for all
rows.
    These features do not need to be present at inference time.
Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', [])   : 1 | ['is_estimated']
Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
    ('int', ['bool']) : 4 | ['is_day:idx', 'is_in_shadow:idx',
'wind_speed_w_1000hPa:ms', 'is_estimated']
    0.2s = Fit runtime
    43 features in original data used to generate 43 features in processed
data.
    Train Data (Processed) Memory Usage: 9.72 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.19s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},

```



```

'FASTAI': {},
'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1799.81s of the 1799.81s of remaining time.

```

-57.5698          = Validation score    (-mean_absolute_error)
0.03s            = Training    runtime
0.42s            = Validation runtime

```

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 1799.13s of the 1799.12s of remaining time.

```

-57.4932          = Validation score    (-mean_absolute_error)
0.03s            = Training    runtime
0.42s            = Validation runtime

```

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 1798.62s of the 1798.61s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-47.7643          = Validation score    (-mean_absolute_error)
2.56s            = Training    runtime
0.16s            = Validation runtime

```

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 1794.74s of the 1794.74s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-48.2753          = Validation score    (-mean_absolute_error)
3.98s            = Training    runtime
0.22s            = Validation runtime

```

Fitting model: RandomForestMSE\_BAG\_L1 ... Training model for up to 1789.38s of the 1789.37s of remaining time.

```

-36.2588          = Validation score    (-mean_absolute_error)
9.71s            = Training    runtime
1.18s            = Validation runtime

```

Fitting model: CatBoost\_BAG\_L1 ... Training model for up to 1777.96s of the 1777.95s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-48.773 = Validation score (-mean_absolute_error)
5.16s   = Training runtime
0.06s   = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1771.63s of the
1771.63s of remaining time.
-37.0616 = Validation score (-mean_absolute_error)
1.94s    = Training runtime
1.19s    = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1767.91s of
the 1767.91s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-47.4511 = Validation score (-mean_absolute_error)
36.52s   = Training runtime
0.51s    = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 1730.13s of the
1730.12s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-46.8022 = Validation score (-mean_absolute_error)
4.31s    = Training runtime
0.12s    = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1724.48s of
the 1724.48s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-41.8723 = Validation score (-mean_absolute_error)
34.2s    = Training runtime
0.37s    = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1689.06s of the
1689.05s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-48.8925 = Validation score (-mean_absolute_error)
6.87s    = Training runtime
0.28s    = Validation runtime
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
1680.72s of remaining time.
-36.2588 = Validation score (-mean_absolute_error)
0.81s    = Training runtime
0.0s     = Validation runtime
AutoGluon training complete, total runtime = 120.13s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_85_B/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -35.73202831478846

```

Note: Scores are always higher\_is\_better. This metric score can be multiplied by -1 to get the metric value.

Evaluations on test data:

```
{
  "mean_absolute_error": -35.73202831478846,
  "root_mean_squared_error": -78.04295480779328,
  "mean_squared_error": -6090.702795131265,
  "r2": 0.8041049360141048,
  "pearsonr": 0.9152552872295632,
  "median_absolute_error": -6.592377662658691
}
```

Evaluation on test data:

-35.73202831478846

```
[ ]: loc = "C"
      predictors[2] = fit_predictor_for_location(loc)
```

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Values in column 'group' used as split folds instead of being automatically set. Bagged models will have 8 splits.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission\_85\_C/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 306.39 GB / 315.93 GB (97.0%)

Train Data Rows: 24492

Train Data Columns: 47

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and label-values can't be converted to int).

Label info (max, min, mean, stddev): (999.6, 0.0, 78.11911, 167.50151)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 130473.17 MB

Train Data (Original) Memory Usage: 10.04 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set

feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

```

        Note: Converting 3 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
        These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]

Training model for location C...
Train data sample weight sum: 24492
Train data number of rows: 24492
Test data sample weight sum: 1579
Test data number of rows: 1579

        ('int', ['bool']) : 3 | ['is_day:idx', 'is_in_shadow:idx',
'is_estimated']
    0.1s = Fit runtime
    43 features in original data used to generate 43 features in processed
data.
    Train Data (Processed) Memory Usage: 7.91 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.16s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge'],

```

```

    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1799.84s of the 1799.84s of remaining time.

```

-32.6988      = Validation score    (-mean_absolute_error)
0.03s        = Training    runtime
0.29s        = Validation runtime

```

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 1799.48s of the 1799.47s of remaining time.

```

-32.7258      = Validation score    (-mean_absolute_error)
0.03s        = Training    runtime
0.28s        = Validation runtime

```

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 1799.11s of the 1799.1s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

```

-23.9492      = Validation score    (-mean_absolute_error)
2.43s        = Training    runtime
0.13s        = Validation runtime

```

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 1795.51s of the 1795.5s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

```

-24.2438      = Validation score    (-mean_absolute_error)
2.01s        = Training    runtime
0.11s        = Validation runtime

```

Fitting model: RandomForestMSE\_BAG\_L1 ... Training model for up to 1792.21s of the 1792.21s of remaining time.

```

-20.7654      = Validation score    (-mean_absolute_error)
5.44s        = Training    runtime
0.76s        = Validation runtime

```

Fitting model: CatBoost\_BAG\_L1 ... Training model for up to 1785.68s of the 1785.67s of remaining time.

```

Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-24.4461      = Validation score    (-mean_absolute_error)
96.28s       = Training    runtime
0.06s       = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1688.18s of the
1688.18s of remaining time.
-20.7137      = Validation score    (-mean_absolute_error)
1.14s       = Training    runtime
0.78s       = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1685.9s of
the 1685.9s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

### 3 Submit

```

[ ]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data

[ ]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time"],
↪    left_on=["ds", "location"])

#test_data_merged

[ ]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in test_data.groupby('location'):
    i = location_map[loc]
    subset = test_data_merged[test_data_merged["location"] == loc].
↪    reset_index(drop=True)
    #print(subset)

```

```

    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)

    # get past predictions
    past_pred = predictors[i].
    ↪predict(train_data_with_dates[train_data_with_dates["location"] == loc])
        train_data_with_dates.loc[train_data_with_dates["location"] == loc,
    ↪"prediction"] = past_pred

```

```

[ ]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='y', ax=ax, label="train data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='prediction', ax=ax, label="past predictions")

    # title
    ax.set_title(f"Predictions for location {loc}")

```

```

[ ]: # concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df

```

```

[ ]: # Save the submission DataFrame to submissions folder, create new name based on
    ↪last submission, format is submission_<last_submission_number + 1>.csv

# Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
    ↪index=False)
print("jallia")

```

```

[ ]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

```

```
display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)
```

```
[ ]: # save this notebook to submissions folder
import subprocess
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
    ↪ipynb"])
```

```
[ ]: # feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
    ↪time_limit=60*10)
```

```
[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
    ↪time_limit=60*10)
```

```
[ ]: display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)

subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↪"autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
#             ↪stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
#             ↪strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."
```



```

# execute_git_command(git_repo_path, ['config', 'user.email',
↳ 'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', hello if hello is
↳ not None else 'Henrik eller Jørgen'])

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push',
↳ 'origin', branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream',
↳ 'origin', branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])

```