

## autogluon\_each\_location

October 16, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*30
presets = 'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = [] # ["hour", "year"]
use_estimated_diff_attr = False
use_is_estimated_attr = True

use_groups = False
n_groups = 8

auto_stack = True
num_stack_levels = 1
num_bag_folds = 8

use_tune_data = False
use_test_data = True
tune_and_test_length = 24*30*3 # 3 months from end
holdout_frac = None
use_bag_holdout = False # Enable this if there is a large gap between score_val
    ↪ and score_test in stack models.

sample_weight = 'sample_weight' #None
weight_evaluation = True
sample_weight_estimated = 3

run_analysis = True

[2]: import pandas as pd
import numpy as np
```

```

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift all columns with "1h" in them by 1 hour, so that for index 16:00,
    # we have the values from 17:00
    # but only for the columns with "1h" in the name
    #X_shifted = X.filter(regex="\dh").shift(-1, axis=1)
    #print(f"Number of columns with 1h in name: {X_shifted.columns}")

    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

    X_shifted = X[X.index.minute==0][columns].copy()
    # loop through all rows and check if index + 1 hour is in the index, if so
    # get that value, else nan
    count1 = 0
    count2 = 0
    for i in range(len(X_shifted)):
        if X_shifted.index[i] + pd.Timedelta('1 hour') in X.index:
            count1 += 1
            X_shifted.iloc[i] = X.loc[X_shifted.index[i] + pd.Timedelta('1
            hour')][columns]
        else:
            count2 += 1
            X_shifted.iloc[i] = np.nan

    print("COUNT1", count1)
    print("COUNT2", count2)

    X_old_unshifted = X[X.index.minute==0][columns]
    # rename X_old_unshifted columns to have _not_shifted at the end
    X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
    columns]

    # put the shifted columns back into the original dataframe
    #X[columns] = X_shifted[columns]

    date_calc = None
    if "date_calc" in X.columns:

```

```

        date_calc = X[X.index.minute == 0]['date_calc']

    # resample to hourly
    X = X.resample('H').mean()

    X[columns] = X_shifted[columns]
    #X[X_old_unshifted.columns] = X_old_unshifted

    if date_calc is not None:
        X['date_calc'] = date_calc

    return X

def fix_X(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    # with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    X = feature_engineering(X)

    return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    # add sample weights, which are 1 for observed and 3 for estimated
    X_train_observed["sample_weight"] = 1
    X_train_estimated["sample_weight"] = sample_weight_estimated
    X_test["sample_weight"] = sample_weight_estimated

    y_train['ds'] = pd.to_datetime(y_train['time'])
    y_train.drop(columns=['time'], inplace=True)
    y_train.sort_values(by='ds', inplace=True)
    y_train.set_index('ds', inplace=True)

```

```

    return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
    ↪location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
    ↪handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
    ↪pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
        X_test["estimated_diff_hours"] = (X_test.index - pd.
    ↪to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

        X_train_estimated["estimated_diff_hours"] =
    ↪X_train_estimated["estimated_diff_hours"].astype('int64')
        # the filled once will get dropped later anyways, when we drop y nans
        X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
    ↪fillna(-50).astype('int64')

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0
        X_train_estimated["is_estimated"] = 1
        X_test["is_estimated"] = 1

    # drop date_calc
    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
    ↪right_index=True)

```

```

    # print number of nans in sample_weight
    print(f"Number of nans in sample_weight: {X_train['sample_weight'].isna().
↪sum()}")
    # print number of nans in y
    print(f"Number of nans in y: {X_train['y'].isna().sum()}")

    X_train["location"] = location
    X_test["location"] = location

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↪X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

Processing location A...

COUNT1 29667

COUNT2 1

COUNT1 4392

COUNT2 2

```

COUNT1 702
COUNT2 18
Number of nans in sample_weight: 0
Number of nans in y: 0
Processing location B...
COUNT1 29232
COUNT2 1
COUNT1 4392
COUNT2 2
COUNT1 702
COUNT2 18
Number of nans in sample_weight: 0
Number of nans in y: 4
Processing location C...
COUNT1 29206
COUNT2 1
COUNT1 4392
COUNT2 2
COUNT1 702
COUNT2 18
Number of nans in sample_weight: 0
Number of nans in y: 6059

```

## 1 Feature engineering

```

[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y'], inplace=True)

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

print(X_train.head())

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

```

```

# Loop through each unique location
for loc in locations:
    loc_df = X_train[X_train['location'] == loc]

    # Sort the DataFrame for this location by the time column
    loc_df = loc_df.sort_index()

    # Calculate the size of each group for this location
    group_size = len(loc_df) // n_groups

    # Create a new 'group' column for this location
    loc_df['group'] = np.repeat(range(n_groups),
    ↪repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

    # Append to list of grouped DataFrames
    grouped_dfs.append(loc_df)

# Concatenate all the grouped DataFrames back together
X_train = pd.concat(grouped_dfs)
X_train.sort_index(inplace=True)
print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
    ↪"dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

	absolute_humidity_2m:gm3	air_density_2m:kgm3	\
ds			
2019-06-02 22:00:00	7.700	1.22825	
2019-06-02 23:00:00	7.700	1.22350	
2019-06-03 00:00:00	7.875	1.21975	
2019-06-03 01:00:00	8.425	1.21800	
2019-06-03 02:00:00	8.950	1.21800	

	ceiling_height_agl:m	clear_sky_energy_1h:J	\
ds			
2019-06-02 22:00:00	1728.949951	0.000000	
2019-06-02 23:00:00	1689.824951	0.000000	

2019-06-03 00:00:00	1563.224976	0.000000
2019-06-03 01:00:00	1283.425049	6546.899902
2019-06-03 02:00:00	1003.500000	102225.898438

	clear_sky_rad:W	cloud_base_agl:m	dew_or_rime:idx	\
ds				
2019-06-02 22:00:00	0.00	1728.949951	0.0	
2019-06-02 23:00:00	0.00	1689.824951	0.0	
2019-06-03 00:00:00	0.00	1563.224976	0.0	
2019-06-03 01:00:00	0.75	1283.425049	0.0	
2019-06-03 02:00:00	23.10	1003.500000	0.0	

	dew_point_2m:K	diffuse_rad:W	diffuse_rad_1h:J	...	\
ds				...	
2019-06-02 22:00:00	280.299988	0.000	0.000000	...	
2019-06-02 23:00:00	280.299988	0.000	0.000000	...	
2019-06-03 00:00:00	280.649994	0.000	0.000000	...	
2019-06-03 01:00:00	281.674988	0.300	7743.299805	...	
2019-06-03 02:00:00	282.500000	11.975	60137.601562	...	

	total_cloud_cover:p	visibility:m	wind_speed_10m:ms	\
ds				
2019-06-02 22:00:00	100.000000	40386.476562	3.600	
2019-06-02 23:00:00	100.000000	33770.648438	3.350	
2019-06-03 00:00:00	100.000000	13595.500000	3.050	
2019-06-03 01:00:00	100.000000	2321.850098	2.725	
2019-06-03 02:00:00	99.224998	11634.799805	2.550	

	wind_speed_u_10m:ms	wind_speed_v_10m:ms	\
ds			
2019-06-02 22:00:00	-3.575	-0.500	
2019-06-02 23:00:00	-3.350	0.275	
2019-06-03 00:00:00	-2.950	0.750	
2019-06-03 01:00:00	-2.600	0.875	
2019-06-03 02:00:00	-2.350	0.925	

	wind_speed_w_1000hPa:ms	sample_weight	is_estimated	\
ds				
2019-06-02 22:00:00	0.0	1	0	
2019-06-02 23:00:00	0.0	1	0	
2019-06-03 00:00:00	0.0	1	0	
2019-06-03 01:00:00	0.0	1	0	
2019-06-03 02:00:00	0.0	1	0	

	y	location
ds		
2019-06-02 22:00:00	0.00	A
2019-06-02 23:00:00	0.00	A



2019-06-03 00:00:00	0.00	A
2019-06-03 01:00:00	0.00	A
2019-06-03 02:00:00	19.36	A

[5 rows x 49 columns]

```
[4]: from autogluon.tabular import TabularDataset, TabularPredictor
from autogluon.timeseries import TimeSeriesDataFrame
import numpy as np
train_data = TabularDataset('X_train_raw.csv')
# set group column of train_data be increasing from 0 to 7 based on time, the
# first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
train_data['ds'] = pd.to_datetime(train_data['ds'])
train_data = train_data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
split_time = pd.to_datetime(train_data["ds"]).max() - pd.
    Timedelta(hours=tune_and_test_length)
train_set = TabularDataset(train_data[train_data["ds"] < split_time])
test_set = TabularDataset(train_data[train_data["ds"] >= split_time])
if use_groups:
    test_set = test_set.drop(columns=['group'])

if do_drop_ds:
    train_set = train_set.drop(columns=['ds'])
    test_set = test_set.drop(columns=['ds'])
    train_data = train_data.drop(columns=['ds'])

def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
        loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

tuning_data = None
if use_tune_data:
    train_data = train_set
    if use_test_data:
```

```

# split test_set in half, use first half for tuning
tuning_data, test_data = [], []
for loc in locations:
    loc_test_set = test_set[test_set["location"] == loc]
    loc_tuning_data = loc_test_set.iloc[:len(loc_test_set)//2]
    loc_test_data = loc_test_set.iloc[len(loc_test_set)//2:]
    tuning_data.append(loc_tuning_data)
    test_data.append(loc_test_data)
tuning_data = pd.concat(tuning_data)
test_data = pd.concat(test_data)
print("Shapes of tuning and test", tuning_data.shape[0], test_data.
↪shape[0], tuning_data.shape[0] + test_data.shape[0])

else:
    tuning_data = test_set
    print("Shape of tuning", tuning_data.shape[0])

# ensure sample weights for your tuning data sum to the number of rows in
↪the tuning data.
tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        train_data = train_set
        test_data = test_set
        print("Shape of test", test_data.shape[0])

# ensure sample weights for your training (or tuning) data sum to the number of
↪rows in the training (or tuning) data.
train_data = normalize_sample_weights_per_location(train_data)
if use_test_data:
    test_data = normalize_sample_weights_per_location(test_data)

```

Shape of test 5791

```

[5]: if run_analysis:
    import autogluon.eda.auto as auto
    auto.dataset_overview(train_data=train_data, test_data=test_data,
↪label="y", sample=None)

```

train\_data dataset summary

	count	unique	top	freq	mean	\
absolute_humidity_2m:gm3	87160	757			6.138632	
air_density_2m:kgm3	87160	1370			1.253802	
ceiling_height_agl:m	72139	59833			2864.542561	
clear_sky_energy_1h:J	87157	45557			518241.825283	

clear_sky_rad:W	87160	19511	143.951884
cloud_base_agl:m	81279	61233	1736.160546
dew_point_2m:K	87160	2001	275.537267
diffuse_rad:W	87160	10980	39.210491
diffuse_rad_1h:J	87157	45515	141514.615214
direct_rad:W	87160	13914	50.139922
direct_rad_1h:J	87157	39280	180360.846534
effective_cloud_cover:p	87160	5652	67.118857
elevation:m	87160	3	11.411014
fresh_snow_1h:cm	87157	39	0.008136
fresh_snow_3h:cm	87157	68	0.024312
fresh_snow_6h:cm	87157	94	0.048424
is_day:idx	87160	5	0.482965
is_estimated	87232	2	0.05968
is_in_shadow:idx	87160	5	0.564895
location	87232	3	A 31924
msl_pressure:hPa	87160	3693	1009.291473
precip_5min:mm	87160	270	0.005788
precip_type_5min:idx	87160	15	0.084236
pressure_100m:hPa	87160	3705	995.621975
pressure_50m:hPa	87160	3758	1001.745166
rain_water:kgm2	87160	39	0.010136
relative_humidity_1000hPa:p	87160	3787	73.860635
sample_weight	87232	6	1.0
sfc_pressure:hPa	87160	3780	1007.89561
snow_depth:cm	87160	483	0.197251
snow_melt_10min:mm	87160	63	0.000245
snow_water:kgm2	87160	161	0.09109
sun_azimuth:d	87160	82801	179.660078
sun_elevation:d	87160	71854	-1.225457
super_cooled_liquid_water:kgm2	87160	53	0.058341
t_1000hPa:K	87160	1986	279.712685
total_cloud_cover:p	87160	5546	73.819247
visibility:m	87160	85645	33233.674454
wind_speed_10m:ms	87160	594	3.025581
wind_speed_u_10m:ms	87160	988	0.664335
wind_speed_v_10m:ms	87160	848	0.694845
y	87232	10750	287.954185

	std	min	25% \
absolute_humidity_2m:gm3	2.73761	0.5	4.1
air_density_2m:kgm3	0.036657	1.13925	1.22875
ceiling_height_agl:m	2531.428872	27.8	1085.2625
clear_sky_energy_1h:J	828416.074463	0.0	0.0
clear_sky_rad:W	230.149085	0.0	0.0
cloud_base_agl:m	1797.954658	27.8	598.2875
dew_point_2m:K	6.846723	247.425	271.0
diffuse_rad:W	60.603659	0.0	0.0

diffuse_rad_1h:J	216225.961408	0.0	0.0
direct_rad:W	113.07899	0.0	0.0
direct_rad_1h:J	402277.99935	0.0	0.0
effective_cloud_cover:p	34.037938	0.0	42.41875
elevation:m	7.881548	6.0	6.0
fresh_snow_1h:cm	0.107503	0.0	0.0
fresh_snow_3h:cm	0.267725	0.0	0.0
fresh_snow_6h:cm	0.457725	0.0	0.0
is_day:idx	0.485944	0.0	0.0
is_estimated	0.236894	0.0	0.0
is_in_shadow:idx	0.483131	0.0	0.0
location			
msl_pressure:hPa	12.998509	944.375	1001.275
precip_5min:mm	0.029771	0.0	0.0
precip_type_5min:idx	0.325388	0.0	0.0
pressure_100m:hPa	12.924683	929.975	987.69995
pressure_50m:hPa	12.982562	935.75	993.75
rain_water:kgm2	0.042332	0.0	0.0
relative_humidity_1000hPa:p	14.160229	19.575	64.425
sample_weight	0.422269	0.876118	0.876118
sfc_pressure:hPa	13.042592	941.55	999.85
snow_depth:cm	1.284395	0.0	0.0
snow_melt_10min:mm	0.003958	0.0	0.0
snow_water:kgm2	0.240712	0.0	0.0
sun_azimuth:d	97.308971	6.983	94.72475
sun_elevation:d	24.168008	-49.932	-18.737563
super_cooled_liquid_water:kgm2	0.106882	0.0	0.0
t_1000hPa:K	6.559438	258.025	275.15
total_cloud_cover:p	33.768818	0.0	53.725
visibility:m	18089.724083	132.375	16688.61875
wind_speed_10m:ms	1.752114	0.025	1.65
wind_speed_u_10m:ms	2.779236	-7.225	-1.35
wind_speed_v_10m:ms	1.881059	-8.4	-0.55
y	766.111697	0.0	0.0

	50%	75%	max	dtypes \
absolute_humidity_2m:gm3	5.6	8.0	17.35	float64
air_density_2m:kgm3	1.2525	1.27675	1.441	float64
ceiling_height_agl:m	1859.4751	3925.32485	12285.775	float64
clear_sky_energy_1h:J	4312.0	777195.2	3006697.2	float64
clear_sky_rad:W	1.6	216.2	835.65	float64
cloud_base_agl:m	1178.425	2081.0375	11673.725	float64
dew_point_2m:K	275.4	280.8	293.625	float64
diffuse_rad:W	0.875	64.325	334.75	float64
diffuse_rad_1h:J	9534.0	233004.8	1182265.4	float64
direct_rad:W	0.0	28.925	683.4	float64
direct_rad_1h:J	0.0	111408.5	2445897.0	float64
effective_cloud_cover:p	79.675	98.475	100.0	float64

elevation:m	7.0	24.0	24.0	float64
fresh_snow_1h:cm	0.0	0.0	7.1	float64
fresh_snow_3h:cm	0.0	0.0	20.6	float64
fresh_snow_6h:cm	0.0	0.0	34.0	float64
is_day:idx	0.25	1.0	1.0	float64
is_estimated	0.0	0.0	1.0	int64
is_in_shadow:idx	1.0	1.0	1.0	float64
location				object
msl_pressure:hPa	1010.275	1018.35	1044.1	float64
precip_5min:mm	0.0	0.0	0.6225	float64
precip_type_5min:idx	0.0	0.0	5.0	float64
pressure_100m:hPa	996.7	1004.7	1030.875	float64
pressure_50m:hPa	1002.8	1010.825	1037.25	float64
rain_water:kgm2	0.0	0.0	1.1	float64
relative_humidity_1000hPa:p	76.2	85.25	100.0	float64
sample_weight	0.899142	0.907248	2.721744	float64
sfc_pressure:hPa	1008.925	1017.0	1043.725	float64
snow_depth:cm	0.0	0.0	18.2	float64
snow_melt_10min:mm	0.0	0.0	0.18	float64
snow_water:kgm2	0.0	0.1	5.65	float64
sun_azimuth:d	180.007	264.513	348.48752	float64
sun_elevation:d	-0.8645	15.234063	49.94375	float64
super_cooled_liquid_water:kgm2	0.0	0.1	1.375	float64
t_1000hPa:K	279.075	284.25	303.25	float64
total_cloud_cover:p	92.85	99.9	100.0	float64
visibility:m	37320.0515	48663.7615	75489.33	float64
wind_speed_10m:ms	2.7	4.05	13.275	float64
wind_speed_u_10m:ms	0.3	2.475	11.2	float64
wind_speed_v_10m:ms	0.725	1.875	8.825	float64
y	0.0	176.4	5733.42	float64

	missing_count	missing_ratio	raw_type	\
absolute_humidity_2m:gm3	72	0.000825	float	
air_density_2m:kgm3	72	0.000825	float	
ceiling_height_agl:m	15093	0.173021	float	
clear_sky_energy_1h:J	75	0.00086	float	
clear_sky_rad:W	72	0.000825	float	
cloud_base_agl:m	5953	0.068243	float	
dew_point_2m:K	72	0.000825	float	
diffuse_rad:W	72	0.000825	float	
diffuse_rad_1h:J	75	0.00086	float	
direct_rad:W	72	0.000825	float	
direct_rad_1h:J	75	0.00086	float	
effective_cloud_cover:p	72	0.000825	float	
elevation:m	72	0.000825	float	
fresh_snow_1h:cm	75	0.00086	float	
fresh_snow_3h:cm	75	0.00086	float	
fresh_snow_6h:cm	75	0.00086	float	

is_day:idx	72	0.000825	float
is_estimated			int
is_in_shadow:idx	72	0.000825	float
location			object
msl_pressure:hPa	72	0.000825	float
precip_5min:mm	72	0.000825	float
precip_type_5min:idx	72	0.000825	float
pressure_100m:hPa	72	0.000825	float
pressure_50m:hPa	72	0.000825	float
rain_water:kgm2	72	0.000825	float
relative_humidity_1000hPa:p	72	0.000825	float
sample_weight			float
sfc_pressure:hPa	72	0.000825	float
snow_depth:cm	72	0.000825	float
snow_melt_10min:mm	72	0.000825	float
snow_water:kgm2	72	0.000825	float
sun_azimuth:d	72	0.000825	float
sun_elevation:d	72	0.000825	float
super_cooled_liquid_water:kgm2	72	0.000825	float
t_1000hPa:K	72	0.000825	float
total_cloud_cover:p	72	0.000825	float
visibility:m	72	0.000825	float
wind_speed_10m:ms	72	0.000825	float
wind_speed_u_10m:ms	72	0.000825	float
wind_speed_v_10m:ms	72	0.000825	float
y			float

	variable_type	special_types
absolute_humidity_2m:gm3	numeric	
air_density_2m:kgm3	numeric	
ceiling_height_agl:m	numeric	
clear_sky_energy_1h:J	numeric	
clear_sky_rad:W	numeric	
cloud_base_agl:m	numeric	
dew_point_2m:K	numeric	
diffuse_rad:W	numeric	
diffuse_rad_1h:J	numeric	
direct_rad:W	numeric	
direct_rad_1h:J	numeric	
effective_cloud_cover:p	numeric	
elevation:m	category	
fresh_snow_1h:cm	numeric	
fresh_snow_3h:cm	numeric	
fresh_snow_6h:cm	numeric	
is_day:idx	category	
is_estimated	category	
is_in_shadow:idx	category	
location	category	

msl_pressure:hPa	numeric
precip_5min:mm	numeric
precip_type_5min:idx	category
pressure_100m:hPa	numeric
pressure_50m:hPa	numeric
rain_water:kgm2	numeric
relative_humidity_1000hPa:p	numeric
sample_weight	category
sfc_pressure:hPa	numeric
snow_depth:cm	numeric
snow_melt_10min:mm	numeric
snow_water:kgm2	numeric
sun_azimuth:d	numeric
sun_elevation:d	numeric
super_cooled_liquid_water:kgm2	numeric
t_1000hPa:K	numeric
total_cloud_cover:p	numeric
visibility:m	numeric
wind_speed_10m:ms	numeric
wind_speed_u_10m:ms	numeric
wind_speed_v_10m:ms	numeric
y	numeric

#### test\_data dataset summary

	count	unique	top	freq	mean	\
absolute_humidity_2m:gm3	5791	289			4.192639	
air_density_2m:kgm3	5791	640			1.280018	
ceiling_height_agl:m	4395	4247			3278.267059	
clear_sky_energy_1h:J	5788	3059		469132.824948		
clear_sky_rad:W	5791	2046			130.246477	
cloud_base_agl:m	4934	4719			1733.271034	
dew_point_2m:K	5791	948			270.733081	
diffuse_rad:W	5791	2237			42.175259	
diffuse_rad_1h:J	5788	3065		152461.828645		
direct_rad:W	5791	1829			51.829421	
direct_rad_1h:J	5788	2676		186526.762509		
effective_cloud_cover:p	5791	2100			66.598541	
elevation:m	5791	3			11.262131	
fresh_snow_1h:cm	5788	23			0.032308	
fresh_snow_3h:cm	5788	42			0.100259	
fresh_snow_6h:cm	5788	60			0.204492	
is_day:idx	5791	5			0.488387	
is_estimated	5791	1			1.0	
is_in_shadow:idx	5791	5			0.555085	
location	5791	3	A	2161		
msl_pressure:hPa	5791	2040			1012.678587	
precip_5min:mm	5791	63			0.003687	
precip_type_5min:idx	5791	12			0.086039	

pressure_100m:hPa	5791	2124	998.781639
pressure_50m:hPa	5791	2134	1005.02648
rain_water:kgm2	5791	7	0.000984
relative_humidity_1000hPa:p	5791	2051	70.810205
sample_weight	5791	1	1.0
sfc_pressure:hPa	5791	2148	1011.29959
snow_depth:cm	5791	78	0.131661
snow_melt_10min:mm	5791	38	0.000695
snow_water:kgm2	5791	68	0.078393
sun_azimuth:d	5791	5681	179.475343
sun_elevation:d	5791	5093	-0.927197
super_cooled_liquid_water:kgm2	5791	31	0.035175
t_1000hPa:K	5791	825	275.185991
total_cloud_cover:p	5791	1838	71.785616
visibility:m	5791	5784	29884.461577
wind_speed_10m:ms	5791	424	3.227599
wind_speed_u_10m:ms	5791	672	0.668019
wind_speed_v_10m:ms	5791	483	0.538344
y	5791	2304	272.991992

	std	min	25% \
absolute_humidity_2m:gm3	1.300644	1.1	3.35
air_density_2m:kgm3	0.024372	1.219	1.26375
ceiling_height_agl:m	2590.751931	27.925	1149.0625
clear_sky_energy_1h:J	689638.596662	0.0	0.0
clear_sky_rad:W	191.578221	0.0	0.0
cloud_base_agl:m	1987.046511	27.5	525.4375
dew_point_2m:K	4.634046	255.05	268.33749
diffuse_rad:W	59.158733	0.0	0.0
diffuse_rad_1h:J	211011.771342	0.0	0.0
direct_rad:W	110.450287	0.0	0.0
direct_rad_1h:J	393513.65175	0.0	0.0
effective_cloud_cover:p	37.583548	0.0	33.6375
elevation:m	7.8114	6.0	6.0
fresh_snow_1h:cm	0.170919	0.0	0.0
fresh_snow_3h:cm	0.425766	0.0	0.0
fresh_snow_6h:cm	0.738932	0.0	0.0
is_day:idx	0.486436	0.0	0.0
is_estimated	0.0	1.0	1.0
is_in_shadow:idx	0.483636	0.0	0.0
location			
msl_pressure:hPa	13.953847	975.3	1003.875
precip_5min:mm	0.017701	0.0	0.0
precip_type_5min:idx	0.393918	0.0	0.0
pressure_100m:hPa	13.825369	962.4	989.9
pressure_50m:hPa	13.873049	968.45	996.087475
rain_water:kgm2	0.009596	0.0	0.0
relative_humidity_1000hPa:p	14.940249	21.325	60.75



sample_weight	0.0	1.0	1.0
sfc_pressure:hPa	13.921629	974.55	1002.25
snow_depth:cm	0.635847	0.0	0.0
snow_melt_10min:mm	0.007333	0.0	0.0
snow_water:kgm2	0.189057	0.0	0.0
sun_azimuth:d	96.891969	14.913	94.264625
sun_elevation:d	20.775858	-44.28175	-17.109625
super_cooled_liquid_water:kgm2	0.084895	0.0	0.0
t_1000hPa:K	3.823552	261.975	272.8
total_cloud_cover:p	37.578218	0.0	41.8
visibility:m	14669.627165	1215.4	18727.05
wind_speed_10m:ms	1.869023	0.05	1.725
wind_speed_u_10m:ms	3.12501	-7.15	-1.75
wind_speed_v_10m:ms	1.838513	-5.3	-0.8
y	770.841016	-0.0	0.0

	50%	75%	max	dtypes \
absolute_humidity_2m:gm3	4.3	5.05	7.7	float64
air_density_2m:kgm3	1.279	1.29375	1.37175	float64
ceiling_height_agl:m	2618.95	4661.025	12294.901	float64
clear_sky_energy_1h:J	11008.5	791394.0	2554290.5	float64
clear_sky_rad:W	2.675	221.925	710.5	float64
cloud_base_agl:m	904.825	2014.962525	10674.3	float64
dew_point_2m:K	271.6	273.9	280.4	float64
diffuse_rad:W	1.775	78.4875	311.95	float64
diffuse_rad_1h:J	18860.9	279202.425	1071799.5	float64
direct_rad:W	0.0	34.0875	530.15	float64
direct_rad_1h:J	0.0	129529.5	1895533.0	float64
effective_cloud_cover:p	85.375	99.975	100.0	float64
elevation:m	7.0	24.0	24.0	float64
fresh_snow_1h:cm	0.0	0.0	2.6	float64
fresh_snow_3h:cm	0.0	0.0	5.2	float64
fresh_snow_6h:cm	0.0	0.0	7.5	float64
is_day:idx	0.25	1.0	1.0	float64
is_estimated	1.0	1.0	1.0	int64
is_in_shadow:idx	1.0	1.0	1.0	float64
location				object
msl_pressure:hPa	1011.625	1023.8125	1041.3501	float64
precip_5min:mm	0.0	0.0	0.2475	float64
precip_type_5min:idx	0.0	0.0	3.0	float64
pressure_100m:hPa	997.9	1009.875	1028.05	float64
pressure_50m:hPa	1004.1	1016.1625	1034.45	float64
rain_water:kgm2	0.0	0.0	0.175	float64
relative_humidity_1000hPa:p	73.1	82.075	98.0	float64
sample_weight	1.0	1.0	1.0	int64
sfc_pressure:hPa	1010.35	1022.5125	1040.8501	float64
snow_depth:cm	0.0	0.0	4.9	float64
snow_melt_10min:mm	0.0	0.0	0.14	float64

snow_water:kgm2	0.0	0.1	2.15	float64
sun_azimuth:d	179.52899	263.49875	347.37848	float64
sun_elevation:d	-0.79825	15.30325	41.13025	float64
super_cooled_liquid_water:kgm2	0.0	0.0	0.75	float64
t_1000hPa:K	275.175	277.525	285.1	float64
total_cloud_cover:p	96.65	100.0	100.0	float64
visibility:m	31311.025	40438.6635	66178.45	float64
wind_speed_10m:ms	2.9	4.45	10.2	float64
wind_speed_u_10m:ms	0.3	2.9	9.95	float64
wind_speed_v_10m:ms	0.625	1.825	7.15	float64
y	0.0	142.906699	5172.64	float64

	missing_count	missing_ratio	raw_type	\
absolute_humidity_2m:gm3			float	
air_density_2m:kgm3			float	
ceiling_height_agl:m	1396	0.241064	float	
clear_sky_energy_1h:J	3	0.000518	float	
clear_sky_rad:W			float	
cloud_base_agl:m	857	0.147988	float	
dew_point_2m:K			float	
diffuse_rad:W			float	
diffuse_rad_1h:J	3	0.000518	float	
direct_rad:W			float	
direct_rad_1h:J	3	0.000518	float	
effective_cloud_cover:p			float	
elevation:m			float	
fresh_snow_1h:cm	3	0.000518	float	
fresh_snow_3h:cm	3	0.000518	float	
fresh_snow_6h:cm	3	0.000518	float	
is_day:idx			float	
is_estimated			int	
is_in_shadow:idx			float	
location			object	
msl_pressure:hPa			float	
precip_5min:mm			float	
precip_type_5min:idx			float	
pressure_100m:hPa			float	
pressure_50m:hPa			float	
rain_water:kgm2			float	
relative_humidity_1000hPa:p			float	
sample_weight			int	
sfc_pressure:hPa			float	
snow_depth:cm			float	
snow_melt_10min:mm			float	
snow_water:kgm2			float	
sun_azimuth:d			float	
sun_elevation:d			float	
super_cooled_liquid_water:kgm2			float	

t_1000hPa:K	float
total_cloud_cover:p	float
visibility:m	float
wind_speed_10m:ms	float
wind_speed_u_10m:ms	float
wind_speed_v_10m:ms	float
y	float


	variable_type	special_types
absolute_humidity_2m:gm3	numeric	
air_density_2m:kgm3	numeric	
ceiling_height_agl:m	numeric	
clear_sky_energy_1h:J	numeric	
clear_sky_rad:W	numeric	
cloud_base_agl:m	numeric	
dew_point_2m:K	numeric	
diffuse_rad:W	numeric	
diffuse_rad_1h:J	numeric	
direct_rad:W	numeric	
direct_rad_1h:J	numeric	
effective_cloud_cover:p	numeric	
elevation:m	category	
fresh_snow_1h:cm	numeric	
fresh_snow_3h:cm	numeric	
fresh_snow_6h:cm	numeric	
is_day:idx	category	
is_estimated	category	
is_in_shadow:idx	category	
location	category	
msl_pressure:hPa	numeric	
precip_5min:mm	numeric	
precip_type_5min:idx	category	
pressure_100m:hPa	numeric	
pressure_50m:hPa	numeric	
rain_water:kgm2	category	
relative_humidity_1000hPa:p	numeric	
sample_weight	category	
sfc_pressure:hPa	numeric	
snow_depth:cm	numeric	
snow_melt_10min:mm	numeric	
snow_water:kgm2	numeric	
sun_azimuth:d	numeric	
sun_elevation:d	numeric	
super_cooled_liquid_water:kgm2	numeric	
t_1000hPa:K	numeric	
total_cloud_cover:p	numeric	
visibility:m	numeric	
wind_speed_10m:ms	numeric	

```
wind_speed_u_10m:ms          numeric
wind_speed_v_10m:ms          numeric
y                             numeric
```

## Types warnings summary

```
train_data test_data warnings
sample_weight    float      int  warning
```

### 1.0.1 Feature Distance




```
[6]: if run_analysis:
      auto.target_analysis(train_data=train_data, label="y")
```

### 1.1 Target variable analysis

```
count      mean      std  min  25%  50%      75%      max  dtypes  \
y  10000  299.128516  787.495283  0.0  0.0  0.0  183.7125  5596.36  float64

unique missing_count missing_ratio raw_type special_types
y      2419                                float
```




### 1.1.1 Distribution fits for target variable


- none of the [attempted](#) distribution fits satisfy specified minimum p-value threshold: 0.01

### 1.1.2 Target variable correlations


`train_data` - spearman correlation matrix; focus: absolute correlation for  $y \geq 0.5$   
(sample size: 10000)




Feature interaction between `clear_sky_rad:W/y` in `train_data` (sample size: 10000)




Feature interaction between clear\_sky\_energy\_1h:J/y in train\_data (sample size: 10000)



Feature interaction between diffuse\_rad:W/y in train\_data (sample size: 10000)




Feature interaction between diffuse\_rad\_1h:J/y in train\_data (sample size: 10000)




Feature interaction between direct\_rad\_1h:J/y in train\_data (sample size: 10000)







Feature interaction between direct\_rad:W/y in train\_data (sample size: 10000)




Feature interaction between sun\_elevation:d/y in train\_data (sample size: 10000)



Feature interaction between `is_day:idx/y` in `train_data` (sample size: 10000)



Feature interaction between `is_in_shadow:idx/y` in `train_data` (sample size: 10000)



## 2 Starting

```
[7]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

```
Last submission number: 89
Now creating submission number: 90
New filename: submission_90
```

```
[8]: predictors = [None, None, None]
```

```

[9]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪train and tune data and test data
    print("Train data sample weight sum:", train_data[train_data["location"] ==
    ↪loc]["sample_weight"].sum())
    print("Train data number of rows:", train_data[train_data["location"] ==
    ↪loc].shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
    ↪tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:", tuning_data[tuning_data["location"]
    ↪== loc].shape[0])
    if use_test_data:
        print("Test data sample weight sum:", test_data[test_data["location"]
    ↪== loc]["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"] ==
    ↪loc].shape[0])
    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        sample_weight=sample_weight,
        weight_evaluation=weight_evaluation,
        groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc],
        time_limit=time_limit,
        presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
    ↪somethin, will be overwritten anyways
        tuning_data=tuning_data[tuning_data["location"] == loc] if
    ↪use_tune_data else None,
        use_bag_holdout=use_bag_holdout,
        holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
    ↪drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

```

```

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Warning: path already exists! This predictor may overwrite an existing predictor! path="AutogluonModels/submission\_90\_A"

Presets specified: ['best\_quality']

Stack configuration (auto\_stack=True): num\_stack\_levels=1, num\_bag\_folds=8, num\_bag\_sets=20

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission\_90\_A/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 298.11 GB / 315.93 GB (94.4%)

Train Data Rows: 31924

Train Data Columns: 41

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (5733.42, 0.0, 632.68576, 1165.32372)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 131761.1 MB

Train Data (Original) Memory Usage: 11.81 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 2 features to boolean dtype as they only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

Stage 4 Generators:

Fitting DropUniqueFeatureGenerator...

Training model for location A...

Train data sample weight sum: 31924.000000000007

Train data number of rows: 31924

Test data sample weight sum: 2161

Test data number of rows: 2161

Stage 5 Generators:

Fitting DropDuplicatesFeatureGenerator...

Useless Original Features (Count: 1): ['location']

These features carry no predictive signal and should be manually investigated.

This is typically a feature which has the same value for all rows.

These features do not need to be present at inference time.

Types of features in original data (raw dtype, special dtypes):

('float', []) : 38 | ['absolute\_humidity\_2m:gm3',  
'air\_density\_2m:kgm3', 'ceiling\_height\_agl:m', 'clear\_sky\_energy\_1h:J',  
'clear\_sky\_rad:W', ...]

('int', []) : 1 | ['is\_estimated']

Types of features in processed data (raw dtype, special dtypes):

('float', []) : 37 | ['absolute\_humidity\_2m:gm3',  
'air\_density\_2m:kgm3', 'ceiling\_height\_agl:m', 'clear\_sky\_energy\_1h:J',  
'clear\_sky\_rad:W', ...]

('int', ['bool']) : 2 | ['elevation:m', 'is\_estimated']

0.2s = Fit runtime

39 features in original data used to generate 39 features in processed data.

Train Data (Processed) Memory Usage: 9.51 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.18s ...

AutoGluon will gauge predictive performance using evaluation metric:

'mean\_absolute\_error'

This metric's sign has been flipped to adhere to being higher\_is\_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval\_metric parameter of Predictor()

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge',
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
```

```

'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1199.58s of the 1799.81s of remaining time.

```

-207.0745      = Validation score    (-mean_absolute_error)
0.04s         = Training    runtime
0.42s         = Validation runtime

```

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 1198.94s of the 1799.18s of remaining time.

```

-208.1423      = Validation score    (-mean_absolute_error)
0.04s         = Training    runtime
0.4s          = Validation runtime

```

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 1198.44s of the 1798.68s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

```

-139.7504      = Validation score    (-mean_absolute_error)
38.16s        = Training    runtime
21.85s        = Validation runtime

```

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 1151.98s of the 1752.21s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

```

-148.2262      = Validation score    (-mean_absolute_error)
41.99s        = Training    runtime
14.42s        = Validation runtime

```

Fitting model: RandomForestMSE\_BAG\_L1 ... Training model for up to 1106.75s of the 1706.98s of remaining time.

```

-162.4541      = Validation score    (-mean_absolute_error)
10.04s        = Training    runtime
1.2s          = Validation runtime

```

Fitting model: CatBoost\_BAG\_L1 ... Training model for up to 1093.46s of the 1693.7s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

```

-155.3753      = Validation score    (-mean_absolute_error)
204.38s       = Training    runtime
0.13s         = Validation runtime

```

Fitting model: ExtraTreesMSE\_BAG\_L1 ... Training model for up to 888.02s of the 1488.25s of remaining time.

```

-163.5846      = Validation score    (-mean_absolute_error)
2.01s         = Training   runtime
1.2s          = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 882.7s of the
1482.93s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -169.7354      = Validation score    (-mean_absolute_error)
    39.32s         = Training   runtime
    0.64s          = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 842.26s of the
1442.5s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -158.1199      = Validation score    (-mean_absolute_error)
    63.81s         = Training   runtime
    3.27s          = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 775.1s of the
1375.34s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -153.7348      = Validation score    (-mean_absolute_error)
    132.53s        = Training   runtime
    0.36s          = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 641.4s of the
1241.64s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -144.4866      = Validation score    (-mean_absolute_error)
    133.11s        = Training   runtime
    19.94s         = Validation runtime
Completed 1/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
1101.5s of remaining time.
    -136.91        = Validation score    (-mean_absolute_error)
    0.86s          = Training   runtime
    0.0s           = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 1100.62s of the
1100.6s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -139.2783      = Validation score    (-mean_absolute_error)
    3.0s           = Training   runtime
    0.17s          = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 1096.36s of the
1096.34s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with

```



```

ParallelLocalFoldFittingStrategy
    -136.5524      = Validation score    (-mean_absolute_error)
    2.29s         = Training runtime
    0.09s         = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ... Training model for up to 1092.84s of
the 1092.82s of remaining time.
    -135.9141      = Validation score    (-mean_absolute_error)
    15.92s        = Training runtime
    1.27s         = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 1073.58s of the
1073.56s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.9547      = Validation score    (-mean_absolute_error)
    4.83s         = Training runtime
    0.06s         = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L2 ... Training model for up to 1067.58s of the
1067.56s of remaining time.
    -135.6456      = Validation score    (-mean_absolute_error)
    2.61s         = Training runtime
    1.25s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 1061.6s of
the 1061.59s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0584      = Validation score    (-mean_absolute_error)
    39.72s        = Training runtime
    0.67s         = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 1020.58s of the
1020.57s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.1474      = Validation score    (-mean_absolute_error)
    3.02s         = Training runtime
    0.12s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 1016.3s of the
1016.28s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -137.1077      = Validation score    (-mean_absolute_error)
    53.6s         = Training runtime
    0.54s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 961.31s of the
961.29s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -137.4687      = Validation score    (-mean_absolute_error)
    7.9s          = Training runtime

```

```

    0.25s      = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 952.03s of the
952.01s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -138.7235      = Validation score    (-mean_absolute_error)
    6.69s      = Training    runtime
    0.37s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 947.09s of the
947.07s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.2762      = Validation score    (-mean_absolute_error)
    4.75s      = Training    runtime
    0.19s      = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 943.34s of the
943.33s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.7681      = Validation score    (-mean_absolute_error)
    9.76s      = Training    runtime
    0.12s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 936.98s of
the 936.97s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.066      = Validation score    (-mean_absolute_error)
    79.58s      = Training    runtime
    1.33s      = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 895.67s of the
895.65s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.3834      = Validation score    (-mean_absolute_error)
    5.99s      = Training    runtime
    0.26s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 891.22s of the
891.21s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.3568      = Validation score    (-mean_absolute_error)
    112.01s     = Training    runtime
    1.05s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 831.46s of the
831.44s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

-136.6957      = Validation score    (-mean_absolute_error)
15.37s      = Training    runtime
0.46s      = Validation runtime
Repeating k-fold bagging: 3/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 822.71s of the
822.69s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-138.6952      = Validation score    (-mean_absolute_error)
10.4s      = Training    runtime
0.54s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 817.7s of the
817.68s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.1345      = Validation score    (-mean_absolute_error)
7.01s      = Training    runtime
0.28s      = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 814.11s of the
814.09s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.6666      = Validation score    (-mean_absolute_error)
14.61s      = Training    runtime
0.18s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 807.97s of
the 807.96s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-134.7458      = Validation score    (-mean_absolute_error)
119.65s     = Training    runtime
1.97s      = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 766.57s of the
766.55s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-135.2141      = Validation score    (-mean_absolute_error)
8.89s      = Training    runtime
0.39s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 762.35s of the
762.33s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-135.9574      = Validation score    (-mean_absolute_error)
175.66s     = Training    runtime
1.69s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 697.32s of the
697.31s of remaining time.

```

```

    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.3156      = Validation score    (-mean_absolute_error)
    23.51s      = Training    runtime
    0.66s      = Validation runtime
Repeating k-fold bagging: 4/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 687.88s of the
687.87s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -138.5887      = Validation score    (-mean_absolute_error)
    14.25s      = Training    runtime
    0.73s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 682.66s of the
682.65s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0618      = Validation score    (-mean_absolute_error)
    9.28s      = Training    runtime
    0.37s      = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 679.21s of the
679.2s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.5893      = Validation score    (-mean_absolute_error)
    19.79s      = Training    runtime
    0.24s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 672.76s of
the 672.74s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -134.5501      = Validation score    (-mean_absolute_error)
    160.12s      = Training    runtime
    2.7s      = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 630.86s of the
630.85s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -134.9661      = Validation score    (-mean_absolute_error)
    12.22s      = Training    runtime
    0.54s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 626.23s of the
626.22s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.7212      = Validation score    (-mean_absolute_error)
    239.1s      = Training    runtime
    2.18s      = Validation runtime

```

Fitting model: LightGBMLarge\_BAG\_L2 ... Training model for up to 561.4s of the 561.39s of remaining time.

Fitting 8 child models (S4F1 - S4F8) | Fitting with ParallelLocalFoldFittingStrategy

-136.1661 = Validation score (-mean\_absolute\_error)  
30.73s = Training runtime  
0.85s = Validation runtime

Repeating k-fold bagging: 5/20

Fitting model: LightGBMXT\_BAG\_L2 ... Training model for up to 552.9s of the 552.88s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-138.5468 = Validation score (-mean\_absolute\_error)  
18.09s = Training runtime  
0.92s = Validation runtime

Fitting model: LightGBM\_BAG\_L2 ... Training model for up to 547.72s of the 547.7s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-136.0882 = Validation score (-mean\_absolute\_error)  
11.64s = Training runtime  
0.47s = Validation runtime

Fitting model: CatBoost\_BAG\_L2 ... Training model for up to 544.12s of the 544.1s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-136.5497 = Validation score (-mean\_absolute\_error)  
25.51s = Training runtime  
0.3s = Validation runtime

Fitting model: NeuralNetFastAI\_BAG\_L2 ... Training model for up to 537.1s of the 537.08s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-134.5143 = Validation score (-mean\_absolute\_error)  
200.32s = Training runtime  
3.36s = Validation runtime

Fitting model: XGBoost\_BAG\_L2 ... Training model for up to 495.48s of the 495.46s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-134.9755 = Validation score (-mean\_absolute\_error)  
15.13s = Training runtime  
0.66s = Validation runtime

Fitting model: NeuralNetTorch\_BAG\_L2 ... Training model for up to 491.1s of the 491.09s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-135.6579 = Validation score (-mean\_absolute\_error)

```

300.76s = Training runtime
2.73s   = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 428.2s of the
428.18s of remaining time.
Fitting 8 child models (S5F1 - S5F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.106 = Validation score (-mean_absolute_error)
38.7s    = Training runtime
1.07s    = Validation runtime
Repeating k-fold bagging: 6/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 418.88s of the
418.87s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-138.5034 = Validation score (-mean_absolute_error)
21.42s    = Training runtime
1.08s     = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 414.16s of the
414.15s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.0657 = Validation score (-mean_absolute_error)
14.05s    = Training runtime
0.56s     = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 410.49s of the
410.47s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.5486 = Validation score (-mean_absolute_error)
30.66s    = Training runtime
0.36s     = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 403.95s of
the 403.93s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-134.4519 = Validation score (-mean_absolute_error)
240.63s   = Training runtime
4.11s     = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 362.21s of the
362.19s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-134.9381 = Validation score (-mean_absolute_error)
17.98s    = Training runtime
0.8s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 357.82s of the
357.8s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with

```

```

ParallelLocalFoldFittingStrategy
    -135.6084      = Validation score    (-mean_absolute_error)
    366.09s      = Training    runtime
    3.24s        = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 291.08s of the
291.06s of remaining time.
    Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0691      = Validation score    (-mean_absolute_error)
    47.02s        = Training    runtime
    1.3s          = Validation runtime
Repeating k-fold bagging: 7/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 281.45s of the
281.43s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -138.4553      = Validation score    (-mean_absolute_error)
    25.67s        = Training    runtime
    1.29s        = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 275.97s of the
275.96s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0951      = Validation score    (-mean_absolute_error)
    16.48s        = Training    runtime
    0.65s        = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 272.33s of the
272.32s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.5204      = Validation score    (-mean_absolute_error)
    35.88s        = Training    runtime
    0.42s        = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 265.75s of
the 265.73s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -134.341       = Validation score    (-mean_absolute_error)
    281.2s        = Training    runtime
    4.77s        = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 223.79s of the
223.78s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.0242      = Validation score    (-mean_absolute_error)
    21.05s        = Training    runtime
    0.93s        = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 219.29s of the

```

```

219.27s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.5437      = Validation score    (-mean_absolute_error)
    429.87s      = Training    runtime
    3.8s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 154.06s of the
154.05s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.067      = Validation score    (-mean_absolute_error)
    54.33s       = Training    runtime
    1.48s        = Validation runtime
Completed 7/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L3 ... Training model for up to 360.0s of the
145.45s of remaining time.
    -132.8711     = Validation score    (-mean_absolute_error)
    0.73s        = Training    runtime
    0.0s         = Validation runtime
AutoGluon training complete, total runtime = 1655.32s ... Best model:
"WeightedEnsemble_L3"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_90_A/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -186.24721157826426
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -186.24721157826426,
    "root_mean_squared_error": -420.62416270646276,
    "mean_squared_error": -176924.68625251285,
    "r2": 0.8716298559177108,
    "pearsonr": 0.9358825898767282,
    "median_absolute_error": -3.769514799118042
}

Evaluation on test data:
-186.24721157826426

```

```

[ ]: loc = "B"
predictors[1] = fit_predictor_for_location(loc)

```

```

Warning: path already exists! This predictor may overwrite an existing
predictor! path="AutogluonModels/submission_90_B"
Presets specified: ['best_quality']
Stack configuration (auto_stack=True): num_stack_levels=1, num_bag_folds=8,
num_bag_sets=20

```



Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission\_90\_B/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 297.94 GB / 315.93 GB (94.3%)

Train Data Rows: 30792

Train Data Columns: 41

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (1152.3, -0.0, 97.67477, 195.03642)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 130207.91 MB

Train Data (Original) Memory Usage: 11.39 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 2 features to boolean dtype as they only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

Stage 4 Generators:

Fitting DropUniqueFeatureGenerator...

Stage 5 Generators:

Fitting DropDuplicatesFeatureGenerator...

Training model for location B...

Train data sample weight sum: 30792.0

Train data number of rows: 30792

Test data sample weight sum: 2051

Test data number of rows: 2051

Useless Original Features (Count: 1): ['location']

These features carry no predictive signal and should be manually investigated.

This is typically a feature which has the same value for all rows.

These features do not need to be present at inference time.

Types of features in original data (raw dtype, special dtypes):

```
('float', []) : 38 | ['absolute_humidity_2m:gm3',  
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',  
'clear_sky_rad:W', ...]
```

```
('int', []) : 1 | ['is_estimated']
```

Types of features in processed data (raw dtype, special dtypes):

```
('float', []) : 37 | ['absolute_humidity_2m:gm3',  
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',  
'clear_sky_rad:W', ...]
```

```
('int', ['bool']) : 2 | ['elevation:m', 'is_estimated']
```

0.2s = Fit runtime

39 features in original data used to generate 39 features in processed data.

Train Data (Processed) Memory Usage: 9.18 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.18s ...

AutoGluon will gauge predictive performance using evaluation metric:

'mean\_absolute\_error'

This metric's sign has been flipped to adhere to being higher\_is\_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval\_metric parameter of Predictor()

User-specified model hyperparameters to be fit:

```
{  
    'NN_TORCH': {},  
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],  
'GBMLarge'],  
    'CAT': {},  
    'XGB': {},  
    'FASTAI': {},  
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',  
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':  
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},  
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',  
'problem_types': ['regression', 'quantile']}}],  
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',  
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':  
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},  
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',  
'problem_types': ['regression', 'quantile']}}],  
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},  
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],  
}
```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1199.58s of the 1799.81s of remaining time.

```

-45.0801          = Validation score    (-mean_absolute_error)
0.03s            = Training   runtime
0.43s            = Validation runtime
Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 1199.06s of
the 1799.29s of remaining time.
-44.9883          = Validation score    (-mean_absolute_error)
0.03s            = Training   runtime
0.43s            = Validation runtime
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1198.55s of the
1798.78s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-26.235           = Validation score    (-mean_absolute_error)
38.79s           = Training   runtime
20.81s           = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1155.96s of the
1756.2s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-27.6872          = Validation score    (-mean_absolute_error)
44.86s           = Training   runtime
16.33s           = Validation runtime
Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1107.7s of
the 1707.94s of remaining time.
-31.8894          = Validation score    (-mean_absolute_error)
11.38s           = Training   runtime
1.17s            = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1093.5s of the
1693.74s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-29.4372          = Validation score    (-mean_absolute_error)
206.59s          = Training   runtime
0.11s            = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 885.71s of the
1485.94s of remaining time.
-32.6592          = Validation score    (-mean_absolute_error)
1.93s            = Training   runtime
1.16s            = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 880.85s of
the 1481.08s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-35.7575          = Validation score    (-mean_absolute_error)
38.29s           = Training   runtime
0.6s             = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 841.26s of the
1441.5s of remaining time.

```

```

    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -30.4775      = Validation score    (-mean_absolute_error)
    105.3s       = Training    runtime
    24.76s       = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 731.93s of the
1332.17s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -28.9112      = Validation score    (-mean_absolute_error)
    198.41s      = Training    runtime
    0.46s        = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 532.27s of the
1132.5s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -26.5707      = Validation score    (-mean_absolute_error)
    139.44s      = Training    runtime
    26.82s       = Validation runtime
Completed 1/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
985.81s of remaining time.
    -25.5137      = Validation score    (-mean_absolute_error)
    0.82s         = Training    runtime
    0.0s          = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 984.97s of the
984.95s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -23.1597      = Validation score    (-mean_absolute_error)
    10.07s        = Training    runtime
    0.69s         = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 973.33s of the
973.31s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.8809      = Validation score    (-mean_absolute_error)
    2.81s         = Training    runtime
    0.12s         = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ... Training model for up to 969.32s of
the 969.3s of remaining time.
    -21.959       = Validation score    (-mean_absolute_error)
    15.03s        = Training    runtime
    1.19s         = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 952.57s of the
952.55s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with

```

```

ParallelLocalFoldFittingStrategy
    -22.9341          = Validation score    (-mean_absolute_error)
    37.75s           = Training   runtime
    0.06s            = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L2 ... Training model for up to 913.65s of the
913.64s of remaining time.
    -21.9244          = Validation score    (-mean_absolute_error)
    2.39s            = Training   runtime
    1.22s            = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 909.48s of
the 909.47s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.1783          = Validation score    (-mean_absolute_error)
    38.24s           = Training   runtime
    0.63s            = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 869.92s of the
869.91s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.6477          = Validation score    (-mean_absolute_error)
    3.04s            = Training   runtime
    0.13s            = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 865.47s of the
865.46s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.2668          = Validation score    (-mean_absolute_error)
    89.3s            = Training   runtime
    0.6s             = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 774.85s of the
774.84s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.4653          = Validation score    (-mean_absolute_error)
    98.41s           = Training   runtime
    1.73s            = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 671.26s of the
671.25s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -23.0488          = Validation score    (-mean_absolute_error)
    17.41s           = Training   runtime
    1.12s            = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 662.56s of the
662.55s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with

```

```

ParallelLocalFoldFittingStrategy
    -22.8141          = Validation score    (-mean_absolute_error)
    5.29s            = Training    runtime
    0.23s            = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 658.7s of the
658.69s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.8401          = Validation score    (-mean_absolute_error)
    48.76s            = Training    runtime
    0.12s            = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 646.5s of the
646.48s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.0293          = Validation score    (-mean_absolute_error)
    76.47s            = Training    runtime
    1.31s            = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 606.94s of the
606.92s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.5148          = Validation score    (-mean_absolute_error)
    6.21s             = Training    runtime
    0.26s             = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 602.47s of the
602.45s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.0278          = Validation score    (-mean_absolute_error)
    189.07s           = Training    runtime
    1.11s             = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 501.31s of the
501.3s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.2095          = Validation score    (-mean_absolute_error)
    204.76s           = Training    runtime
    4.79s             = Validation runtime
Repeating k-fold bagging: 3/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 389.89s of the
389.88s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -23.0258          = Validation score    (-mean_absolute_error)
    24.06s            = Training    runtime
    1.58s            = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 381.85s of the

```

```

381.83s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.7665      = Validation score    (-mean_absolute_error)
    8.3s         = Training    runtime
    0.35s        = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 377.65s of the
377.63s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.7782      = Validation score    (-mean_absolute_error)
    62.72s        = Training    runtime
    0.19s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 362.45s of
the 362.43s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.9458      = Validation score    (-mean_absolute_error)
    115.45s       = Training    runtime
    1.97s         = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 322.19s of the
322.17s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.4749      = Validation score    (-mean_absolute_error)
    9.33s         = Training    runtime
    0.39s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 317.71s of the
317.69s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.9414      = Validation score    (-mean_absolute_error)
    270.03s       = Training    runtime
    1.62s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 235.42s of the
235.4s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.1519      = Validation score    (-mean_absolute_error)
    296.67s       = Training    runtime
    6.16s         = Validation runtime
Completed 3/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L3 ... Training model for up to 360.0s of the
139.58s of remaining time.
    -21.4202      = Validation score    (-mean_absolute_error)
    0.7s          = Training    runtime
    0.0s          = Validation runtime
AutoGluon training complete, total runtime = 1661.17s ... Best model:

```

```
"WeightedEnsemble_L3"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_90_B/")
```

```
[ ]: loc = "C"
predictors[2] = fit_predictor_for_location(loc)
```

### 3 Submit

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data
```

```
[ ]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
↪ "location"], left_on=["ds", "location"])

#test_data_merged
```

```
[ ]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in test_data.groupby('location'):
    i = location_map[loc]
    subset = test_data_merged[test_data_merged["location"] == loc].
↪reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)

# get past predictions
past_pred = predictors[i].
↪predict(train_data_with_dates[train_data_with_dates["location"] == loc])
```



```

train_data_with_dates.loc[train_data_with_dates["location"] == loc,
↪ "prediction"] = past_pred

```

```

[ ]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
↪ y='y', ax=ax, label="train data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
↪ y='prediction', ax=ax, label="past predictions")

    # title
    ax.set_title(f"Predictions for location {loc}")

```

```

[ ]: # concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df

```

```

[ ]: # Save the submission DataFrame to submissions folder, create new name based on
↪ last submission, format is submission_<last_submission_number + 1>.csv

# Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
↪ index=False)
print("jall1a")

```

```

[ ]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)

```

```

[ ]: # save this notebook to submissions folder
import subprocess

```

```
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↳join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
    ↳ipynb"])
```

```
[ ]: # feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
    ↳time_limit=60*10)
```

```
[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
    ↳time_limit=60*10)
```

```
[ ]: display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)

subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↳join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↳"autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
#             ↳stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
#             ↳strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
#     ↳'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', hello if hello is
#     ↳not None else 'Henrik eller Jørgen'])

# branch_name = new_filename
```

```

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push', '↪
'origin', branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream', '↪
'origin', branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])

```