

# autogluon\_each\_location

October 16, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*30
presets = 'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = [] # ["hour", "year"]
use_estimated_diff_attr = False
use_is_estimated_attr = True

use_groups = False
n_groups = 8

auto_stack = True
num_stack_levels = 1
num_bag_folds = 8

use_tune_data = False
use_test_data = True
tune_and_test_length = 24*30*3 # 3 months from end
holdout_frac = None
use_bag_holdout = False # Enable this if there is a large gap between score_val
    ↪ and score_test in stack models.

sample_weight = 'sample_weight' #None
weight_evaluation = True
sample_weight_estimated = 3

run_analysis = True

[2]: import pandas as pd
import numpy as np
```

```

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift all columns with "1h" in them by 1 hour, so that for index 16:00,
    # we have the values from 17:00
    # but only for the columns with "1h" in the name
    #X_shifted = X.filter(regex="\dh").shift(-1, axis=1)
    #print(f"Number of columns with 1h in name: {X_shifted.columns}")

    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

    X_shifted = X[X.index.minute==0][columns].copy()
    # loop through all rows and check if index + 1 hour is in the index, if so
    # get that value, else nan
    count1 = 0
    count2 = 0
    for i in range(len(X_shifted)):
        if X_shifted.index[i] + pd.Timedelta('1 hour') in X.index:
            count1 += 1
            X_shifted.iloc[i] = X.loc[X_shifted.index[i] + pd.Timedelta('1
            hour')][columns]
        else:
            count2 += 1
            X_shifted.iloc[i] = np.nan

    print("COUNT1", count1)
    print("COUNT2", count2)

    X_old_unshifted = X[X.index.minute==0][columns]
    # rename X_old_unshifted columns to have _not_shifted at the end
    X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
    columns]

    # put the shifted columns back into the original dataframe
    #X[columns] = X_shifted[columns]

    date_calc = None
    if "date_calc" in X.columns:

```

```

        date_calc = X[X.index.minute == 0]['date_calc']

        # resample to hourly
        X = X.resample('H').mean()

        X[columns] = X_shifted[columns]
        #X[X_old_unshifted.columns] = X_old_unshifted

        if date_calc is not None:
            X['date_calc'] = date_calc

        return X

def fix_X(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    # with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    X = feature_engineering(X)

    return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    # add sample weights, which are 1 for observed and 3 for estimated
    X_train_observed["sample_weight"] = 1
    X_train_estimated["sample_weight"] = sample_weight_estimated
    X_test["sample_weight"] = sample_weight_estimated

    y_train['ds'] = pd.to_datetime(y_train['time'])
    y_train.drop(columns=['time'], inplace=True)
    y_train.sort_values(by='ds', inplace=True)
    y_train.set_index('ds', inplace=True)

```

```

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
        X_test["estimated_diff_hours"] = (X_test.index - pd.
to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

        X_train_estimated["estimated_diff_hours"] =
X_train_estimated["estimated_diff_hours"].astype('int64')
        # the filled once will get dropped later anyways, when we drop y nans
        X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
fillna(-50).astype('int64')

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0
        X_train_estimated["is_estimated"] = 1
        X_test["is_estimated"] = 1

    # drop date_calc
    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
right_index=True)

```

```

    # print number of nans in sample_weight
    print(f"Number of nans in sample_weight: {X_train['sample_weight'].isna().
↪sum()}")
    # print number of nans in y
    print(f"Number of nans in y: {X_train['y'].isna().sum()}")

    X_train["location"] = location
    X_test["location"] = location

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↪X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

Processing location A...

COUNT1 29667

COUNT2 1

COUNT1 4392

COUNT2 2

```

COUNT1 702
COUNT2 18
Number of nans in sample_weight: 0
Number of nans in y: 0
Processing location B...
COUNT1 29232
COUNT2 1
COUNT1 4392
COUNT2 2
COUNT1 702
COUNT2 18
Number of nans in sample_weight: 0
Number of nans in y: 4
Processing location C...
COUNT1 29206
COUNT2 1
COUNT1 4392
COUNT2 2
COUNT1 702
COUNT2 18
Number of nans in sample_weight: 0
Number of nans in y: 6059

```

## 1 Feature engineering

```

[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y'], inplace=True)

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

print(X_train.head())

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

```

```

# Loop through each unique location
for loc in locations:
    loc_df = X_train[X_train['location'] == loc]

    # Sort the DataFrame for this location by the time column
    loc_df = loc_df.sort_index()

    # Calculate the size of each group for this location
    group_size = len(loc_df) // n_groups

    # Create a new 'group' column for this location
    loc_df['group'] = np.repeat(range(n_groups),
    ↪repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

    # Append to list of grouped DataFrames
    grouped_dfs.append(loc_df)

# Concatenate all the grouped DataFrames back together
X_train = pd.concat(grouped_dfs)
X_train.sort_index(inplace=True)
print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
    ↪"dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

|                     | absolute_humidity_2m:gm3 | air_density_2m:kgm3 | \ |
|---------------------|--------------------------|---------------------|---|
| ds                  |                          |                     |   |
| 2019-06-02 22:00:00 | 7.700                    | 1.22825             |   |
| 2019-06-02 23:00:00 | 7.700                    | 1.22350             |   |
| 2019-06-03 00:00:00 | 7.875                    | 1.21975             |   |
| 2019-06-03 01:00:00 | 8.425                    | 1.21800             |   |
| 2019-06-03 02:00:00 | 8.950                    | 1.21800             |   |

|                     | ceiling_height_agl:m | clear_sky_energy_1h:J | \ |
|---------------------|----------------------|-----------------------|---|
| ds                  |                      |                       |   |
| 2019-06-02 22:00:00 | 1728.949951          | 0.000000              |   |
| 2019-06-02 23:00:00 | 1689.824951          | 0.000000              |   |

|                     |             |               |
|---------------------|-------------|---------------|
| 2019-06-03 00:00:00 | 1563.224976 | 0.000000      |
| 2019-06-03 01:00:00 | 1283.425049 | 6546.899902   |
| 2019-06-03 02:00:00 | 1003.500000 | 102225.898438 |

|                     | clear_sky_rad:W | cloud_base_agl:m | dew_or_rime:idx | \ |
|---------------------|-----------------|------------------|-----------------|---|
| ds                  |                 |                  |                 |   |
| 2019-06-02 22:00:00 | 0.00            | 1728.949951      | 0.0             |   |
| 2019-06-02 23:00:00 | 0.00            | 1689.824951      | 0.0             |   |
| 2019-06-03 00:00:00 | 0.00            | 1563.224976      | 0.0             |   |
| 2019-06-03 01:00:00 | 0.75            | 1283.425049      | 0.0             |   |
| 2019-06-03 02:00:00 | 23.10           | 1003.500000      | 0.0             |   |

|                     | dew_point_2m:K | diffuse_rad:W | diffuse_rad_1h:J | ... | \ |
|---------------------|----------------|---------------|------------------|-----|---|
| ds                  |                |               |                  | ... |   |
| 2019-06-02 22:00:00 | 280.299988     | 0.000         | 0.000000         | ... |   |
| 2019-06-02 23:00:00 | 280.299988     | 0.000         | 0.000000         | ... |   |
| 2019-06-03 00:00:00 | 280.649994     | 0.000         | 0.000000         | ... |   |
| 2019-06-03 01:00:00 | 281.674988     | 0.300         | 7743.299805      | ... |   |
| 2019-06-03 02:00:00 | 282.500000     | 11.975        | 60137.601562     | ... |   |

|                     | total_cloud_cover:p | visibility:m | wind_speed_10m:ms | \ |
|---------------------|---------------------|--------------|-------------------|---|
| ds                  |                     |              |                   |   |
| 2019-06-02 22:00:00 | 100.000000          | 40386.476562 | 3.600             |   |
| 2019-06-02 23:00:00 | 100.000000          | 33770.648438 | 3.350             |   |
| 2019-06-03 00:00:00 | 100.000000          | 13595.500000 | 3.050             |   |
| 2019-06-03 01:00:00 | 100.000000          | 2321.850098  | 2.725             |   |
| 2019-06-03 02:00:00 | 99.224998           | 11634.799805 | 2.550             |   |

|                     | wind_speed_u_10m:ms | wind_speed_v_10m:ms | \ |
|---------------------|---------------------|---------------------|---|
| ds                  |                     |                     |   |
| 2019-06-02 22:00:00 | -3.575              | -0.500              |   |
| 2019-06-02 23:00:00 | -3.350              | 0.275               |   |
| 2019-06-03 00:00:00 | -2.950              | 0.750               |   |
| 2019-06-03 01:00:00 | -2.600              | 0.875               |   |
| 2019-06-03 02:00:00 | -2.350              | 0.925               |   |

|                     | wind_speed_w_1000hPa:ms | sample_weight | is_estimated | \ |
|---------------------|-------------------------|---------------|--------------|---|
| ds                  |                         |               |              |   |
| 2019-06-02 22:00:00 | 0.0                     | 1             | 0            |   |
| 2019-06-02 23:00:00 | 0.0                     | 1             | 0            |   |
| 2019-06-03 00:00:00 | 0.0                     | 1             | 0            |   |
| 2019-06-03 01:00:00 | 0.0                     | 1             | 0            |   |
| 2019-06-03 02:00:00 | 0.0                     | 1             | 0            |   |

|                     | y    | location |
|---------------------|------|----------|
| ds                  |      |          |
| 2019-06-02 22:00:00 | 0.00 | A        |
| 2019-06-02 23:00:00 | 0.00 | A        |



|                     |       |   |
|---------------------|-------|---|
| 2019-06-03 00:00:00 | 0.00  | A |
| 2019-06-03 01:00:00 | 0.00  | A |
| 2019-06-03 02:00:00 | 19.36 | A |

[5 rows x 49 columns]

```
[4]: from autogluon.tabular import TabularDataset, TabularPredictor
from autogluon.timeseries import TimeSeriesDataFrame
import numpy as np
train_data = TabularDataset('X_train_raw.csv')
# set group column of train_data be increasing from 0 to 7 based on time, the
# first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
train_data['ds'] = pd.to_datetime(train_data['ds'])
train_data = train_data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
split_time = pd.to_datetime(train_data["ds"]).max() - pd.
    Timedelta(hours=tune_and_test_length)
train_set = TabularDataset(train_data[train_data["ds"] < split_time])
test_set = TabularDataset(train_data[train_data["ds"] >= split_time])
if use_groups:
    test_set = test_set.drop(columns=['group'])

if do_drop_ds:
    train_set = train_set.drop(columns=['ds'])
    test_set = test_set.drop(columns=['ds'])
    train_data = train_data.drop(columns=['ds'])

def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
        loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

tuning_data = None
if use_tune_data:
    train_data = train_set
    if use_test_data:
```

```

# split test_set in half, use first half for tuning
tuning_data, test_data = [], []
for loc in locations:
    loc_test_set = test_set[test_set["location"] == loc]
    loc_tuning_data = loc_test_set.iloc[:len(loc_test_set)//2]
    loc_test_data = loc_test_set.iloc[len(loc_test_set)//2:]
    tuning_data.append(loc_tuning_data)
    test_data.append(loc_test_data)
tuning_data = pd.concat(tuning_data)
test_data = pd.concat(test_data)
print("Shapes of tuning and test", tuning_data.shape[0], test_data.
↪shape[0], tuning_data.shape[0] + test_data.shape[0])

else:
    tuning_data = test_set
    print("Shape of tuning", tuning_data.shape[0])

# ensure sample weights for your tuning data sum to the number of rows in
↪the tuning data.
tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        train_data = train_set
        test_data = test_set
        print("Shape of test", test_data.shape[0])

# ensure sample weights for your training (or tuning) data sum to the number of
↪rows in the training (or tuning) data.
train_data = normalize_sample_weights_per_location(train_data)
if use_test_data:
    test_data = normalize_sample_weights_per_location(test_data)

```

Shape of test 5791

```

[5]: if run_analysis:
    import autogluon.eda.auto as auto
    auto.dataset_overview(train_data=train_data, test_data=test_data,
↪label="y", sample=None)

```

train\_data dataset summary

|                          | count | unique top | freq | mean \        |
|--------------------------|-------|------------|------|---------------|
| absolute_humidity_2m:gm3 | 87160 | 757        |      | 6.138632      |
| air_density_2m:kgm3      | 87160 | 1370       |      | 1.253802      |
| ceiling_height_agl:m     | 72139 | 59833      |      | 2864.542561   |
| clear_sky_energy_1h:J    | 87157 | 45557      |      | 518241.825283 |

|                                |       |       |               |
|--------------------------------|-------|-------|---------------|
| clear_sky_rad:W                | 87160 | 19511 | 143.951884    |
| cloud_base_agl:m               | 81279 | 61233 | 1736.160546   |
| dew_point_2m:K                 | 87160 | 2001  | 275.537267    |
| diffuse_rad:W                  | 87160 | 10980 | 39.210491     |
| diffuse_rad_1h:J               | 87157 | 45515 | 141514.615214 |
| direct_rad:W                   | 87160 | 13914 | 50.139922     |
| direct_rad_1h:J                | 87157 | 39280 | 180360.846534 |
| effective_cloud_cover:p        | 87160 | 5652  | 67.118857     |
| elevation:m                    | 87160 | 3     | 11.411014     |
| fresh_snow_1h:cm               | 87157 | 39    | 0.008136      |
| fresh_snow_3h:cm               | 87157 | 68    | 0.024312      |
| fresh_snow_6h:cm               | 87157 | 94    | 0.048424      |
| is_day:idx                     | 87160 | 5     | 0.482965      |
| is_estimated                   | 87232 | 2     | 0.05968       |
| is_in_shadow:idx               | 87160 | 5     | 0.564895      |
| location                       | 87232 | 3     | A 31924       |
| msl_pressure:hPa               | 87160 | 3693  | 1009.291473   |
| precip_5min:mm                 | 87160 | 270   | 0.005788      |
| precip_type_5min:idx           | 87160 | 15    | 0.084236      |
| pressure_100m:hPa              | 87160 | 3705  | 995.621975    |
| pressure_50m:hPa               | 87160 | 3758  | 1001.745166   |
| rain_water:kgm2                | 87160 | 39    | 0.010136      |
| relative_humidity_1000hPa:p    | 87160 | 3787  | 73.860635     |
| sample_weight                  | 87232 | 6     | 1.0           |
| sfc_pressure:hPa               | 87160 | 3780  | 1007.89561    |
| snow_depth:cm                  | 87160 | 483   | 0.197251      |
| snow_melt_10min:mm             | 87160 | 63    | 0.000245      |
| snow_water:kgm2                | 87160 | 161   | 0.09109       |
| sun_azimuth:d                  | 87160 | 82801 | 179.660078    |
| sun_elevation:d                | 87160 | 71854 | -1.225457     |
| super_cooled_liquid_water:kgm2 | 87160 | 53    | 0.058341      |
| t_1000hPa:K                    | 87160 | 1986  | 279.712685    |
| total_cloud_cover:p            | 87160 | 5546  | 73.819247     |
| visibility:m                   | 87160 | 85645 | 33233.674454  |
| wind_speed_10m:ms              | 87160 | 594   | 3.025581      |
| wind_speed_u_10m:ms            | 87160 | 988   | 0.664335      |
| wind_speed_v_10m:ms            | 87160 | 848   | 0.694845      |
| y                              | 87232 | 10750 | 287.954185    |

|                          | std           | min     | 25% \     |
|--------------------------|---------------|---------|-----------|
| absolute_humidity_2m:gm3 | 2.73761       | 0.5     | 4.1       |
| air_density_2m:kgm3      | 0.036657      | 1.13925 | 1.22875   |
| ceiling_height_agl:m     | 2531.428872   | 27.8    | 1085.2625 |
| clear_sky_energy_1h:J    | 828416.074463 | 0.0     | 0.0       |
| clear_sky_rad:W          | 230.149085    | 0.0     | 0.0       |
| cloud_base_agl:m         | 1797.954658   | 27.8    | 598.2875  |
| dew_point_2m:K           | 6.846723      | 247.425 | 271.0     |
| diffuse_rad:W            | 60.603659     | 0.0     | 0.0       |

|                                |               |          |             |
|--------------------------------|---------------|----------|-------------|
| diffuse_rad_1h:J               | 216225.961408 | 0.0      | 0.0         |
| direct_rad:W                   | 113.07899     | 0.0      | 0.0         |
| direct_rad_1h:J                | 402277.99935  | 0.0      | 0.0         |
| effective_cloud_cover:p        | 34.037938     | 0.0      | 42.41875    |
| elevation:m                    | 7.881548      | 6.0      | 6.0         |
| fresh_snow_1h:cm               | 0.107503      | 0.0      | 0.0         |
| fresh_snow_3h:cm               | 0.267725      | 0.0      | 0.0         |
| fresh_snow_6h:cm               | 0.457725      | 0.0      | 0.0         |
| is_day:idx                     | 0.485944      | 0.0      | 0.0         |
| is_estimated                   | 0.236894      | 0.0      | 0.0         |
| is_in_shadow:idx               | 0.483131      | 0.0      | 0.0         |
| location                       |               |          |             |
| msl_pressure:hPa               | 12.998509     | 944.375  | 1001.275    |
| precip_5min:mm                 | 0.029771      | 0.0      | 0.0         |
| precip_type_5min:idx           | 0.325388      | 0.0      | 0.0         |
| pressure_100m:hPa              | 12.924683     | 929.975  | 987.69995   |
| pressure_50m:hPa               | 12.982562     | 935.75   | 993.75      |
| rain_water:kgm2                | 0.042332      | 0.0      | 0.0         |
| relative_humidity_1000hPa:p    | 14.160229     | 19.575   | 64.425      |
| sample_weight                  | 0.422269      | 0.876118 | 0.876118    |
| sfc_pressure:hPa               | 13.042592     | 941.55   | 999.85      |
| snow_depth:cm                  | 1.284395      | 0.0      | 0.0         |
| snow_melt_10min:mm             | 0.003958      | 0.0      | 0.0         |
| snow_water:kgm2                | 0.240712      | 0.0      | 0.0         |
| sun_azimuth:d                  | 97.308971     | 6.983    | 94.72475    |
| sun_elevation:d                | 24.168008     | -49.932  | -18.737563  |
| super_cooled_liquid_water:kgm2 | 0.106882      | 0.0      | 0.0         |
| t_1000hPa:K                    | 6.559438      | 258.025  | 275.15      |
| total_cloud_cover:p            | 33.768818     | 0.0      | 53.725      |
| visibility:m                   | 18089.724083  | 132.375  | 16688.61875 |
| wind_speed_10m:ms              | 1.752114      | 0.025    | 1.65        |
| wind_speed_u_10m:ms            | 2.779236      | -7.225   | -1.35       |
| wind_speed_v_10m:ms            | 1.881059      | -8.4     | -0.55       |
| y                              | 766.111697    | 0.0      | 0.0         |

|                          | 50%       | 75%        | max       | dtypes \ |
|--------------------------|-----------|------------|-----------|----------|
| absolute_humidity_2m:gm3 | 5.6       | 8.0        | 17.35     | float64  |
| air_density_2m:kgm3      | 1.2525    | 1.27675    | 1.441     | float64  |
| ceiling_height_agl:m     | 1859.4751 | 3925.32485 | 12285.775 | float64  |
| clear_sky_energy_1h:J    | 4312.0    | 777195.2   | 3006697.2 | float64  |
| clear_sky_rad:W          | 1.6       | 216.2      | 835.65    | float64  |
| cloud_base_agl:m         | 1178.425  | 2081.0375  | 11673.725 | float64  |
| dew_point_2m:K           | 275.4     | 280.8      | 293.625   | float64  |
| diffuse_rad:W            | 0.875     | 64.325     | 334.75    | float64  |
| diffuse_rad_1h:J         | 9534.0    | 233004.8   | 1182265.4 | float64  |
| direct_rad:W             | 0.0       | 28.925     | 683.4     | float64  |
| direct_rad_1h:J          | 0.0       | 111408.5   | 2445897.0 | float64  |
| effective_cloud_cover:p  | 79.675    | 98.475     | 100.0     | float64  |

|                                |            |            |           |         |
|--------------------------------|------------|------------|-----------|---------|
| elevation:m                    | 7.0        | 24.0       | 24.0      | float64 |
| fresh_snow_1h:cm               | 0.0        | 0.0        | 7.1       | float64 |
| fresh_snow_3h:cm               | 0.0        | 0.0        | 20.6      | float64 |
| fresh_snow_6h:cm               | 0.0        | 0.0        | 34.0      | float64 |
| is_day:idx                     | 0.25       | 1.0        | 1.0       | float64 |
| is_estimated                   | 0.0        | 0.0        | 1.0       | int64   |
| is_in_shadow:idx               | 1.0        | 1.0        | 1.0       | float64 |
| location                       |            |            |           | object  |
| msl_pressure:hPa               | 1010.275   | 1018.35    | 1044.1    | float64 |
| precip_5min:mm                 | 0.0        | 0.0        | 0.6225    | float64 |
| precip_type_5min:idx           | 0.0        | 0.0        | 5.0       | float64 |
| pressure_100m:hPa              | 996.7      | 1004.7     | 1030.875  | float64 |
| pressure_50m:hPa               | 1002.8     | 1010.825   | 1037.25   | float64 |
| rain_water:kgm2                | 0.0        | 0.0        | 1.1       | float64 |
| relative_humidity_1000hPa:p    | 76.2       | 85.25      | 100.0     | float64 |
| sample_weight                  | 0.899142   | 0.907248   | 2.721744  | float64 |
| sfc_pressure:hPa               | 1008.925   | 1017.0     | 1043.725  | float64 |
| snow_depth:cm                  | 0.0        | 0.0        | 18.2      | float64 |
| snow_melt_10min:mm             | 0.0        | 0.0        | 0.18      | float64 |
| snow_water:kgm2                | 0.0        | 0.1        | 5.65      | float64 |
| sun_azimuth:d                  | 180.007    | 264.513    | 348.48752 | float64 |
| sun_elevation:d                | -0.8645    | 15.234063  | 49.94375  | float64 |
| super_cooled_liquid_water:kgm2 | 0.0        | 0.1        | 1.375     | float64 |
| t_1000hPa:K                    | 279.075    | 284.25     | 303.25    | float64 |
| total_cloud_cover:p            | 92.85      | 99.9       | 100.0     | float64 |
| visibility:m                   | 37320.0515 | 48663.7615 | 75489.33  | float64 |
| wind_speed_10m:ms              | 2.7        | 4.05       | 13.275    | float64 |
| wind_speed_u_10m:ms            | 0.3        | 2.475      | 11.2      | float64 |
| wind_speed_v_10m:ms            | 0.725      | 1.875      | 8.825     | float64 |
| y                              | 0.0        | 176.4      | 5733.42   | float64 |

|                          | missing_count | missing_ratio | raw_type | \ |
|--------------------------|---------------|---------------|----------|---|
| absolute_humidity_2m:gm3 | 72            | 0.000825      | float    |   |
| air_density_2m:kgm3      | 72            | 0.000825      | float    |   |
| ceiling_height_agl:m     | 15093         | 0.173021      | float    |   |
| clear_sky_energy_1h:J    | 75            | 0.00086       | float    |   |
| clear_sky_rad:W          | 72            | 0.000825      | float    |   |
| cloud_base_agl:m         | 5953          | 0.068243      | float    |   |
| dew_point_2m:K           | 72            | 0.000825      | float    |   |
| diffuse_rad:W            | 72            | 0.000825      | float    |   |
| diffuse_rad_1h:J         | 75            | 0.00086       | float    |   |
| direct_rad:W             | 72            | 0.000825      | float    |   |
| direct_rad_1h:J          | 75            | 0.00086       | float    |   |
| effective_cloud_cover:p  | 72            | 0.000825      | float    |   |
| elevation:m              | 72            | 0.000825      | float    |   |
| fresh_snow_1h:cm         | 75            | 0.00086       | float    |   |
| fresh_snow_3h:cm         | 75            | 0.00086       | float    |   |
| fresh_snow_6h:cm         | 75            | 0.00086       | float    |   |

|                                |    |          |        |
|--------------------------------|----|----------|--------|
| is_day:idx                     | 72 | 0.000825 | float  |
| is_estimated                   |    |          | int    |
| is_in_shadow:idx               | 72 | 0.000825 | float  |
| location                       |    |          | object |
| msl_pressure:hPa               | 72 | 0.000825 | float  |
| precip_5min:mm                 | 72 | 0.000825 | float  |
| precip_type_5min:idx           | 72 | 0.000825 | float  |
| pressure_100m:hPa              | 72 | 0.000825 | float  |
| pressure_50m:hPa               | 72 | 0.000825 | float  |
| rain_water:kgm2                | 72 | 0.000825 | float  |
| relative_humidity_1000hPa:p    | 72 | 0.000825 | float  |
| sample_weight                  |    |          | float  |
| sfc_pressure:hPa               | 72 | 0.000825 | float  |
| snow_depth:cm                  | 72 | 0.000825 | float  |
| snow_melt_10min:mm             | 72 | 0.000825 | float  |
| snow_water:kgm2                | 72 | 0.000825 | float  |
| sun_azimuth:d                  | 72 | 0.000825 | float  |
| sun_elevation:d                | 72 | 0.000825 | float  |
| super_cooled_liquid_water:kgm2 | 72 | 0.000825 | float  |
| t_1000hPa:K                    | 72 | 0.000825 | float  |
| total_cloud_cover:p            | 72 | 0.000825 | float  |
| visibility:m                   | 72 | 0.000825 | float  |
| wind_speed_10m:ms              | 72 | 0.000825 | float  |
| wind_speed_u_10m:ms            | 72 | 0.000825 | float  |
| wind_speed_v_10m:ms            | 72 | 0.000825 | float  |
| y                              |    |          | float  |

|                          | variable_type | special_types |
|--------------------------|---------------|---------------|
| absolute_humidity_2m:gm3 | numeric       |               |
| air_density_2m:kgm3      | numeric       |               |
| ceiling_height_agl:m     | numeric       |               |
| clear_sky_energy_1h:J    | numeric       |               |
| clear_sky_rad:W          | numeric       |               |
| cloud_base_agl:m         | numeric       |               |
| dew_point_2m:K           | numeric       |               |
| diffuse_rad:W            | numeric       |               |
| diffuse_rad_1h:J         | numeric       |               |
| direct_rad:W             | numeric       |               |
| direct_rad_1h:J          | numeric       |               |
| effective_cloud_cover:p  | numeric       |               |
| elevation:m              | category      |               |
| fresh_snow_1h:cm         | numeric       |               |
| fresh_snow_3h:cm         | numeric       |               |
| fresh_snow_6h:cm         | numeric       |               |
| is_day:idx               | category      |               |
| is_estimated             | category      |               |
| is_in_shadow:idx         | category      |               |
| location                 | category      |               |

|                                |          |
|--------------------------------|----------|
| msl_pressure:hPa               | numeric  |
| precip_5min:mm                 | numeric  |
| precip_type_5min:idx           | category |
| pressure_100m:hPa              | numeric  |
| pressure_50m:hPa               | numeric  |
| rain_water:kgm2                | numeric  |
| relative_humidity_1000hPa:p    | numeric  |
| sample_weight                  | category |
| sfc_pressure:hPa               | numeric  |
| snow_depth:cm                  | numeric  |
| snow_melt_10min:mm             | numeric  |
| snow_water:kgm2                | numeric  |
| sun_azimuth:d                  | numeric  |
| sun_elevation:d                | numeric  |
| super_cooled_liquid_water:kgm2 | numeric  |
| t_1000hPa:K                    | numeric  |
| total_cloud_cover:p            | numeric  |
| visibility:m                   | numeric  |
| wind_speed_10m:ms              | numeric  |
| wind_speed_u_10m:ms            | numeric  |
| wind_speed_v_10m:ms            | numeric  |
| y                              | numeric  |

#### test\_data dataset summary

|                          | count | unique | top | freq          | mean        | \ |
|--------------------------|-------|--------|-----|---------------|-------------|---|
| absolute_humidity_2m:gm3 | 5791  | 289    |     |               | 4.192639    |   |
| air_density_2m:kgm3      | 5791  | 640    |     |               | 1.280018    |   |
| ceiling_height_agl:m     | 4395  | 4247   |     |               | 3278.267059 |   |
| clear_sky_energy_1h:J    | 5788  | 3059   |     | 469132.824948 |             |   |
| clear_sky_rad:W          | 5791  | 2046   |     |               | 130.246477  |   |
| cloud_base_agl:m         | 4934  | 4719   |     |               | 1733.271034 |   |
| dew_point_2m:K           | 5791  | 948    |     |               | 270.733081  |   |
| diffuse_rad:W            | 5791  | 2237   |     |               | 42.175259   |   |
| diffuse_rad_1h:J         | 5788  | 3065   |     | 152461.828645 |             |   |
| direct_rad:W             | 5791  | 1829   |     |               | 51.829421   |   |
| direct_rad_1h:J          | 5788  | 2676   |     | 186526.762509 |             |   |
| effective_cloud_cover:p  | 5791  | 2100   |     |               | 66.598541   |   |
| elevation:m              | 5791  | 3      |     |               | 11.262131   |   |
| fresh_snow_1h:cm         | 5788  | 23     |     |               | 0.032308    |   |
| fresh_snow_3h:cm         | 5788  | 42     |     |               | 0.100259    |   |
| fresh_snow_6h:cm         | 5788  | 60     |     |               | 0.204492    |   |
| is_day:idx               | 5791  | 5      |     |               | 0.488387    |   |
| is_estimated             | 5791  | 1      |     |               | 1.0         |   |
| is_in_shadow:idx         | 5791  | 5      |     |               | 0.555085    |   |
| location                 | 5791  | 3      | A   | 2161          |             |   |
| msl_pressure:hPa         | 5791  | 2040   |     |               | 1012.678587 |   |
| precip_5min:mm           | 5791  | 63     |     |               | 0.003687    |   |
| precip_type_5min:idx     | 5791  | 12     |     |               | 0.086039    |   |

|                                |      |      |              |
|--------------------------------|------|------|--------------|
| pressure_100m:hPa              | 5791 | 2124 | 998.781639   |
| pressure_50m:hPa               | 5791 | 2134 | 1005.02648   |
| rain_water:kgm2                | 5791 | 7    | 0.000984     |
| relative_humidity_1000hPa:p    | 5791 | 2051 | 70.810205    |
| sample_weight                  | 5791 | 1    | 1.0          |
| sfc_pressure:hPa               | 5791 | 2148 | 1011.29959   |
| snow_depth:cm                  | 5791 | 78   | 0.131661     |
| snow_melt_10min:mm             | 5791 | 38   | 0.000695     |
| snow_water:kgm2                | 5791 | 68   | 0.078393     |
| sun_azimuth:d                  | 5791 | 5681 | 179.475343   |
| sun_elevation:d                | 5791 | 5093 | -0.927197    |
| super_cooled_liquid_water:kgm2 | 5791 | 31   | 0.035175     |
| t_1000hPa:K                    | 5791 | 825  | 275.185991   |
| total_cloud_cover:p            | 5791 | 1838 | 71.785616    |
| visibility:m                   | 5791 | 5784 | 29884.461577 |
| wind_speed_10m:ms              | 5791 | 424  | 3.227599     |
| wind_speed_u_10m:ms            | 5791 | 672  | 0.668019     |
| wind_speed_v_10m:ms            | 5791 | 483  | 0.538344     |
| y                              | 5791 | 2304 | 272.991992   |

|                             | std           | min    | 25% \      |
|-----------------------------|---------------|--------|------------|
| absolute_humidity_2m:gm3    | 1.300644      | 1.1    | 3.35       |
| air_density_2m:kgm3         | 0.024372      | 1.219  | 1.26375    |
| ceiling_height_agl:m        | 2590.751931   | 27.925 | 1149.0625  |
| clear_sky_energy_1h:J       | 689638.596662 | 0.0    | 0.0        |
| clear_sky_rad:W             | 191.578221    | 0.0    | 0.0        |
| cloud_base_agl:m            | 1987.046511   | 27.5   | 525.4375   |
| dew_point_2m:K              | 4.634046      | 255.05 | 268.33749  |
| diffuse_rad:W               | 59.158733     | 0.0    | 0.0        |
| diffuse_rad_1h:J            | 211011.771342 | 0.0    | 0.0        |
| direct_rad:W                | 110.450287    | 0.0    | 0.0        |
| direct_rad_1h:J             | 393513.65175  | 0.0    | 0.0        |
| effective_cloud_cover:p     | 37.583548     | 0.0    | 33.6375    |
| elevation:m                 | 7.8114        | 6.0    | 6.0        |
| fresh_snow_1h:cm            | 0.170919      | 0.0    | 0.0        |
| fresh_snow_3h:cm            | 0.425766      | 0.0    | 0.0        |
| fresh_snow_6h:cm            | 0.738932      | 0.0    | 0.0        |
| is_day:idx                  | 0.486436      | 0.0    | 0.0        |
| is_estimated                | 0.0           | 1.0    | 1.0        |
| is_in_shadow:idx            | 0.483636      | 0.0    | 0.0        |
| location                    |               |        |            |
| msl_pressure:hPa            | 13.953847     | 975.3  | 1003.875   |
| precip_5min:mm              | 0.017701      | 0.0    | 0.0        |
| precip_type_5min:idx        | 0.393918      | 0.0    | 0.0        |
| pressure_100m:hPa           | 13.825369     | 962.4  | 989.9      |
| pressure_50m:hPa            | 13.873049     | 968.45 | 996.087475 |
| rain_water:kgm2             | 0.009596      | 0.0    | 0.0        |
| relative_humidity_1000hPa:p | 14.940249     | 21.325 | 60.75      |



|                                |              |           |            |
|--------------------------------|--------------|-----------|------------|
| sample_weight                  | 0.0          | 1.0       | 1.0        |
| sfc_pressure:hPa               | 13.921629    | 974.55    | 1002.25    |
| snow_depth:cm                  | 0.635847     | 0.0       | 0.0        |
| snow_melt_10min:mm             | 0.007333     | 0.0       | 0.0        |
| snow_water:kgm2                | 0.189057     | 0.0       | 0.0        |
| sun_azimuth:d                  | 96.891969    | 14.913    | 94.264625  |
| sun_elevation:d                | 20.775858    | -44.28175 | -17.109625 |
| super_cooled_liquid_water:kgm2 | 0.084895     | 0.0       | 0.0        |
| t_1000hPa:K                    | 3.823552     | 261.975   | 272.8      |
| total_cloud_cover:p            | 37.578218    | 0.0       | 41.8       |
| visibility:m                   | 14669.627165 | 1215.4    | 18727.05   |
| wind_speed_10m:ms              | 1.869023     | 0.05      | 1.725      |
| wind_speed_u_10m:ms            | 3.12501      | -7.15     | -1.75      |
| wind_speed_v_10m:ms            | 1.838513     | -5.3      | -0.8       |
| y                              | 770.841016   | -0.0      | 0.0        |

|                             | 50%      | 75%         | max       | dtypes \ |
|-----------------------------|----------|-------------|-----------|----------|
| absolute_humidity_2m:gm3    | 4.3      | 5.05        | 7.7       | float64  |
| air_density_2m:kgm3         | 1.279    | 1.29375     | 1.37175   | float64  |
| ceiling_height_agl:m        | 2618.95  | 4661.025    | 12294.901 | float64  |
| clear_sky_energy_1h:J       | 11008.5  | 791394.0    | 2554290.5 | float64  |
| clear_sky_rad:W             | 2.675    | 221.925     | 710.5     | float64  |
| cloud_base_agl:m            | 904.825  | 2014.962525 | 10674.3   | float64  |
| dew_point_2m:K              | 271.6    | 273.9       | 280.4     | float64  |
| diffuse_rad:W               | 1.775    | 78.4875     | 311.95    | float64  |
| diffuse_rad_1h:J            | 18860.9  | 279202.425  | 1071799.5 | float64  |
| direct_rad:W                | 0.0      | 34.0875     | 530.15    | float64  |
| direct_rad_1h:J             | 0.0      | 129529.5    | 1895533.0 | float64  |
| effective_cloud_cover:p     | 85.375   | 99.975      | 100.0     | float64  |
| elevation:m                 | 7.0      | 24.0        | 24.0      | float64  |
| fresh_snow_1h:cm            | 0.0      | 0.0         | 2.6       | float64  |
| fresh_snow_3h:cm            | 0.0      | 0.0         | 5.2       | float64  |
| fresh_snow_6h:cm            | 0.0      | 0.0         | 7.5       | float64  |
| is_day:idx                  | 0.25     | 1.0         | 1.0       | float64  |
| is_estimated                | 1.0      | 1.0         | 1.0       | int64    |
| is_in_shadow:idx            | 1.0      | 1.0         | 1.0       | float64  |
| location                    |          |             |           | object   |
| msl_pressure:hPa            | 1011.625 | 1023.8125   | 1041.3501 | float64  |
| precip_5min:mm              | 0.0      | 0.0         | 0.2475    | float64  |
| precip_type_5min:idx        | 0.0      | 0.0         | 3.0       | float64  |
| pressure_100m:hPa           | 997.9    | 1009.875    | 1028.05   | float64  |
| pressure_50m:hPa            | 1004.1   | 1016.1625   | 1034.45   | float64  |
| rain_water:kgm2             | 0.0      | 0.0         | 0.175     | float64  |
| relative_humidity_1000hPa:p | 73.1     | 82.075      | 98.0      | float64  |
| sample_weight               | 1.0      | 1.0         | 1.0       | int64    |
| sfc_pressure:hPa            | 1010.35  | 1022.5125   | 1040.8501 | float64  |
| snow_depth:cm               | 0.0      | 0.0         | 4.9       | float64  |
| snow_melt_10min:mm          | 0.0      | 0.0         | 0.14      | float64  |

|                                |           |            |           |         |
|--------------------------------|-----------|------------|-----------|---------|
| snow_water:kgm2                | 0.0       | 0.1        | 2.15      | float64 |
| sun_azimuth:d                  | 179.52899 | 263.49875  | 347.37848 | float64 |
| sun_elevation:d                | -0.79825  | 15.30325   | 41.13025  | float64 |
| super_cooled_liquid_water:kgm2 | 0.0       | 0.0        | 0.75      | float64 |
| t_1000hPa:K                    | 275.175   | 277.525    | 285.1     | float64 |
| total_cloud_cover:p            | 96.65     | 100.0      | 100.0     | float64 |
| visibility:m                   | 31311.025 | 40438.6635 | 66178.45  | float64 |
| wind_speed_10m:ms              | 2.9       | 4.45       | 10.2      | float64 |
| wind_speed_u_10m:ms            | 0.3       | 2.9        | 9.95      | float64 |
| wind_speed_v_10m:ms            | 0.625     | 1.825      | 7.15      | float64 |
| y                              | 0.0       | 142.906699 | 5172.64   | float64 |

|                                | missing_count | missing_ratio | raw_type | \ |
|--------------------------------|---------------|---------------|----------|---|
| absolute_humidity_2m:gm3       |               |               | float    |   |
| air_density_2m:kgm3            |               |               | float    |   |
| ceiling_height_agl:m           | 1396          | 0.241064      | float    |   |
| clear_sky_energy_1h:J          | 3             | 0.000518      | float    |   |
| clear_sky_rad:W                |               |               | float    |   |
| cloud_base_agl:m               | 857           | 0.147988      | float    |   |
| dew_point_2m:K                 |               |               | float    |   |
| diffuse_rad:W                  |               |               | float    |   |
| diffuse_rad_1h:J               | 3             | 0.000518      | float    |   |
| direct_rad:W                   |               |               | float    |   |
| direct_rad_1h:J                | 3             | 0.000518      | float    |   |
| effective_cloud_cover:p        |               |               | float    |   |
| elevation:m                    |               |               | float    |   |
| fresh_snow_1h:cm               | 3             | 0.000518      | float    |   |
| fresh_snow_3h:cm               | 3             | 0.000518      | float    |   |
| fresh_snow_6h:cm               | 3             | 0.000518      | float    |   |
| is_day:idx                     |               |               | float    |   |
| is_estimated                   |               |               | int      |   |
| is_in_shadow:idx               |               |               | float    |   |
| location                       |               |               | object   |   |
| msl_pressure:hPa               |               |               | float    |   |
| precip_5min:mm                 |               |               | float    |   |
| precip_type_5min:idx           |               |               | float    |   |
| pressure_100m:hPa              |               |               | float    |   |
| pressure_50m:hPa               |               |               | float    |   |
| rain_water:kgm2                |               |               | float    |   |
| relative_humidity_1000hPa:p    |               |               | float    |   |
| sample_weight                  |               |               | int      |   |
| sfc_pressure:hPa               |               |               | float    |   |
| snow_depth:cm                  |               |               | float    |   |
| snow_melt_10min:mm             |               |               | float    |   |
| snow_water:kgm2                |               |               | float    |   |
| sun_azimuth:d                  |               |               | float    |   |
| sun_elevation:d                |               |               | float    |   |
| super_cooled_liquid_water:kgm2 |               |               | float    |   |

|                     |       |
|---------------------|-------|
| t_1000hPa:K         | float |
| total_cloud_cover:p | float |
| visibility:m        | float |
| wind_speed_10m:ms   | float |
| wind_speed_u_10m:ms | float |
| wind_speed_v_10m:ms | float |
| y                   | float |


|                                | variable_type | special_types |
|--------------------------------|---------------|---------------|
| absolute_humidity_2m:gm3       | numeric       |               |
| air_density_2m:kgm3            | numeric       |               |
| ceiling_height_agl:m           | numeric       |               |
| clear_sky_energy_1h:J          | numeric       |               |
| clear_sky_rad:W                | numeric       |               |
| cloud_base_agl:m               | numeric       |               |
| dew_point_2m:K                 | numeric       |               |
| diffuse_rad:W                  | numeric       |               |
| diffuse_rad_1h:J               | numeric       |               |
| direct_rad:W                   | numeric       |               |
| direct_rad_1h:J                | numeric       |               |
| effective_cloud_cover:p        | numeric       |               |
| elevation:m                    | category      |               |
| fresh_snow_1h:cm               | numeric       |               |
| fresh_snow_3h:cm               | numeric       |               |
| fresh_snow_6h:cm               | numeric       |               |
| is_day:idx                     | category      |               |
| is_estimated                   | category      |               |
| is_in_shadow:idx               | category      |               |
| location                       | category      |               |
| msl_pressure:hPa               | numeric       |               |
| precip_5min:mm                 | numeric       |               |
| precip_type_5min:idx           | category      |               |
| pressure_100m:hPa              | numeric       |               |
| pressure_50m:hPa               | numeric       |               |
| rain_water:kgm2                | category      |               |
| relative_humidity_1000hPa:p    | numeric       |               |
| sample_weight                  | category      |               |
| sfc_pressure:hPa               | numeric       |               |
| snow_depth:cm                  | numeric       |               |
| snow_melt_10min:mm             | numeric       |               |
| snow_water:kgm2                | numeric       |               |
| sun_azimuth:d                  | numeric       |               |
| sun_elevation:d                | numeric       |               |
| super_cooled_liquid_water:kgm2 | numeric       |               |
| t_1000hPa:K                    | numeric       |               |
| total_cloud_cover:p            | numeric       |               |
| visibility:m                   | numeric       |               |
| wind_speed_10m:ms              | numeric       |               |

|                     |         |
|---------------------|---------|
| wind_speed_u_10m:ms | numeric |
| wind_speed_v_10m:ms | numeric |
| y                   | numeric |

## Types warnings summary

|               | train_data | test_data | warnings |
|---------------|------------|-----------|----------|
| sample_weight | float      | int       | warning  |

### 1.0.1 Feature Distance




```
[6]: if run_analysis:
      auto.target_analysis(train_data=train_data, label="y")
```

### 1.1 Target variable analysis

|   | count | mean       | std        | min | 25% | 50% | 75%      | max     | dtypes  | \ |
|---|-------|------------|------------|-----|-----|-----|----------|---------|---------|---|
| y | 10000 | 299.128516 | 787.495283 | 0.0 | 0.0 | 0.0 | 183.7125 | 5596.36 | float64 |   |

|   | unique | missing_count | missing_ratio | raw_type | special_types |
|---|--------|---------------|---------------|----------|---------------|
| y | 2419   |               |               | float    |               |




### 1.1.1 Distribution fits for target variable


- none of the [attempted](#) distribution fits satisfy specified minimum p-value threshold: 0.01

### 1.1.2 Target variable correlations


`train_data` - spearman correlation matrix; focus: absolute correlation for  $y \geq 0.5$   
(sample size: 10000)




Feature interaction between `clear_sky_rad:W/y` in `train_data` (sample size: 10000)




Feature interaction between clear\_sky\_energy\_1h:J/y in train\_data (sample size: 10000)



Feature interaction between diffuse\_rad:W/y in train\_data (sample size: 10000)




Feature interaction between diffuse\_rad\_1h:J/y in train\_data (sample size: 10000)




Feature interaction between direct\_rad\_1h:J/y in train\_data (sample size: 10000)







Feature interaction between direct\_rad:W/y in train\_data (sample size: 10000)




Feature interaction between sun\_elevation:d/y in train\_data (sample size: 10000)



Feature interaction between `is_day:idx/y` in `train_data` (sample size: 10000)



Feature interaction between `is_in_shadow:idx/y` in `train_data` (sample size: 10000)



## 2 Starting

```
[7]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

```
Last submission number: 89
Now creating submission number: 90
New filename: submission_90
```

```
[8]: predictors = [None, None, None]
```

```

[9]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪train and tune data and test data
    print("Train data sample weight sum:", train_data[train_data["location"] ==
    ↪loc]["sample_weight"].sum())
    print("Train data number of rows:", train_data[train_data["location"] ==
    ↪loc].shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
    ↪tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:", tuning_data[tuning_data["location"]
    ↪== loc].shape[0])
    if use_test_data:
        print("Test data sample weight sum:", test_data[test_data["location"]
    ↪== loc]["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"] ==
    ↪loc].shape[0])
    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        sample_weight=sample_weight,
        weight_evaluation=weight_evaluation,
        groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc],
        time_limit=time_limit,
        presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
    ↪somethin, will be overwritten anyways
        tuning_data=tuning_data[tuning_data["location"] == loc] if
    ↪use_tune_data else None,
        use_bag_holdout=use_bag_holdout,
        holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
    ↪drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

```

```

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Warning: path already exists! This predictor may overwrite an existing predictor! path="AutogluonModels/submission\_90\_A"

Presets specified: ['best\_quality']

Stack configuration (auto\_stack=True): num\_stack\_levels=1, num\_bag\_folds=8, num\_bag\_sets=20

Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission\_90\_A/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 298.11 GB / 315.93 GB (94.4%)

Train Data Rows: 31924

Train Data Columns: 41

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (5733.42, 0.0, 632.68576, 1165.32372)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 131761.1 MB

Train Data (Original) Memory Usage: 11.81 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 2 features to boolean dtype as they only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

Stage 4 Generators:

Fitting DropUniqueFeatureGenerator...

Training model for location A...

Train data sample weight sum: 31924.000000000007

Train data number of rows: 31924

Test data sample weight sum: 2161

Test data number of rows: 2161

Stage 5 Generators:

Fitting DropDuplicatesFeatureGenerator...

Useless Original Features (Count: 1): ['location']

These features carry no predictive signal and should be manually investigated.

This is typically a feature which has the same value for all rows.

These features do not need to be present at inference time.

Types of features in original data (raw dtype, special dtypes):

('float', []) : 38 | ['absolute\_humidity\_2m:gm3',  
'air\_density\_2m:kgm3', 'ceiling\_height\_agl:m', 'clear\_sky\_energy\_1h:J',  
'clear\_sky\_rad:W', ...]

('int', []) : 1 | ['is\_estimated']

Types of features in processed data (raw dtype, special dtypes):

('float', []) : 37 | ['absolute\_humidity\_2m:gm3',  
'air\_density\_2m:kgm3', 'ceiling\_height\_agl:m', 'clear\_sky\_energy\_1h:J',  
'clear\_sky\_rad:W', ...]

('int', ['bool']) : 2 | ['elevation:m', 'is\_estimated']

0.2s = Fit runtime

39 features in original data used to generate 39 features in processed data.

Train Data (Processed) Memory Usage: 9.51 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.18s ...

AutoGluon will gauge predictive performance using evaluation metric:

'mean\_absolute\_error'

This metric's sign has been flipped to adhere to being higher\_is\_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval\_metric parameter of Predictor()

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge',
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
```

```

'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1199.58s of the 1799.81s of remaining time.

```

-207.0745      = Validation score    (-mean_absolute_error)
0.04s         = Training   runtime
0.42s         = Validation runtime

```

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 1198.94s of the 1799.18s of remaining time.

```

-208.1423      = Validation score    (-mean_absolute_error)
0.04s         = Training   runtime
0.4s          = Validation runtime

```

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 1198.44s of the 1798.68s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-139.7504      = Validation score    (-mean_absolute_error)
38.16s        = Training   runtime
21.85s        = Validation runtime

```

Fitting model: LightGBM\_BAG\_L1 ... Training model for up to 1151.98s of the 1752.21s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-148.2262      = Validation score    (-mean_absolute_error)
41.99s        = Training   runtime
14.42s        = Validation runtime

```

Fitting model: RandomForestMSE\_BAG\_L1 ... Training model for up to 1106.75s of the 1706.98s of remaining time.

```

-162.4541      = Validation score    (-mean_absolute_error)
10.04s        = Training   runtime
1.2s          = Validation runtime

```

Fitting model: CatBoost\_BAG\_L1 ... Training model for up to 1093.46s of the 1693.7s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with  
ParallelLocalFoldFittingStrategy

```

-155.3753      = Validation score    (-mean_absolute_error)
204.38s       = Training   runtime
0.13s         = Validation runtime

```

Fitting model: ExtraTreesMSE\_BAG\_L1 ... Training model for up to 888.02s of the 1488.25s of remaining time.

```

-163.5846      = Validation score    (-mean_absolute_error)
2.01s      = Training    runtime
1.2s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 882.7s of the
1482.93s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-169.7354      = Validation score    (-mean_absolute_error)
39.32s      = Training    runtime
0.64s      = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 842.26s of the
1442.5s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-158.1199      = Validation score    (-mean_absolute_error)
63.81s      = Training    runtime
3.27s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 775.1s of the
1375.34s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-153.7348      = Validation score    (-mean_absolute_error)
132.53s      = Training    runtime
0.36s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 641.4s of the
1241.64s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-144.4866      = Validation score    (-mean_absolute_error)
133.11s      = Training    runtime
19.94s      = Validation runtime
Completed 1/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
1101.5s of remaining time.
-136.91      = Validation score    (-mean_absolute_error)
0.86s      = Training    runtime
0.0s      = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 1100.62s of the
1100.6s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-139.2783      = Validation score    (-mean_absolute_error)
3.0s      = Training    runtime
0.17s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 1096.36s of the
1096.34s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with

```



```

ParallelLocalFoldFittingStrategy
    -136.5524      = Validation score    (-mean_absolute_error)
    2.29s         = Training   runtime
    0.09s         = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ... Training model for up to 1092.84s of
the 1092.82s of remaining time.
    -135.9141      = Validation score    (-mean_absolute_error)
    15.92s        = Training   runtime
    1.27s         = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 1073.58s of the
1073.56s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.9547      = Validation score    (-mean_absolute_error)
    4.83s         = Training   runtime
    0.06s         = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L2 ... Training model for up to 1067.58s of the
1067.56s of remaining time.
    -135.6456      = Validation score    (-mean_absolute_error)
    2.61s         = Training   runtime
    1.25s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 1061.6s of
the 1061.59s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0584      = Validation score    (-mean_absolute_error)
    39.72s        = Training   runtime
    0.67s         = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 1020.58s of the
1020.57s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.1474      = Validation score    (-mean_absolute_error)
    3.02s         = Training   runtime
    0.12s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 1016.3s of the
1016.28s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -137.1077      = Validation score    (-mean_absolute_error)
    53.6s         = Training   runtime
    0.54s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 961.31s of the
961.29s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -137.4687      = Validation score    (-mean_absolute_error)
    7.9s          = Training   runtime

```

```

    0.25s      = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 952.03s of the
952.01s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -138.7235      = Validation score      (-mean_absolute_error)
    6.69s      = Training      runtime
    0.37s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 947.09s of the
947.07s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.2762      = Validation score      (-mean_absolute_error)
    4.75s      = Training      runtime
    0.19s      = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 943.34s of the
943.33s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.7681      = Validation score      (-mean_absolute_error)
    9.76s      = Training      runtime
    0.12s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 936.98s of
the 936.97s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.066      = Validation score      (-mean_absolute_error)
    79.58s      = Training      runtime
    1.33s      = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 895.67s of the
895.65s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.3834      = Validation score      (-mean_absolute_error)
    5.99s      = Training      runtime
    0.26s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 891.22s of the
891.21s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.3568      = Validation score      (-mean_absolute_error)
    112.01s      = Training      runtime
    1.05s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 831.46s of the
831.44s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

-136.6957      = Validation score    (-mean_absolute_error)
15.37s      = Training    runtime
0.46s      = Validation runtime
Repeating k-fold bagging: 3/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 822.71s of the
822.69s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-138.6952      = Validation score    (-mean_absolute_error)
10.4s      = Training    runtime
0.54s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 817.7s of the
817.68s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.1345      = Validation score    (-mean_absolute_error)
7.01s      = Training    runtime
0.28s      = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 814.11s of the
814.09s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.6666      = Validation score    (-mean_absolute_error)
14.61s      = Training    runtime
0.18s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 807.97s of
the 807.96s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-134.7458      = Validation score    (-mean_absolute_error)
119.65s     = Training    runtime
1.97s      = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 766.57s of the
766.55s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-135.2141      = Validation score    (-mean_absolute_error)
8.89s      = Training    runtime
0.39s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 762.35s of the
762.33s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-135.9574      = Validation score    (-mean_absolute_error)
175.66s     = Training    runtime
1.69s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 697.32s of the
697.31s of remaining time.

```

```

    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.3156      = Validation score    (-mean_absolute_error)
    23.51s      = Training    runtime
    0.66s      = Validation runtime
Repeating k-fold bagging: 4/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 687.88s of the
687.87s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -138.5887      = Validation score    (-mean_absolute_error)
    14.25s      = Training    runtime
    0.73s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 682.66s of the
682.65s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0618      = Validation score    (-mean_absolute_error)
    9.28s      = Training    runtime
    0.37s      = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 679.21s of the
679.2s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.5893      = Validation score    (-mean_absolute_error)
    19.79s      = Training    runtime
    0.24s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 672.76s of
the 672.74s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -134.5501      = Validation score    (-mean_absolute_error)
    160.12s      = Training    runtime
    2.7s      = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 630.86s of the
630.85s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -134.9661      = Validation score    (-mean_absolute_error)
    12.22s      = Training    runtime
    0.54s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 626.23s of the
626.22s of remaining time.
    Fitting 8 child models (S4F1 - S4F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.7212      = Validation score    (-mean_absolute_error)
    239.1s      = Training    runtime
    2.18s      = Validation runtime

```

Fitting model: LightGBMLarge\_BAG\_L2 ... Training model for up to 561.4s of the 561.39s of remaining time.

Fitting 8 child models (S4F1 - S4F8) | Fitting with ParallelLocalFoldFittingStrategy

-136.1661 = Validation score (-mean\_absolute\_error)

30.73s = Training runtime

0.85s = Validation runtime

Repeating k-fold bagging: 5/20

Fitting model: LightGBMXT\_BAG\_L2 ... Training model for up to 552.9s of the 552.88s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-138.5468 = Validation score (-mean\_absolute\_error)

18.09s = Training runtime

0.92s = Validation runtime

Fitting model: LightGBM\_BAG\_L2 ... Training model for up to 547.72s of the 547.7s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-136.0882 = Validation score (-mean\_absolute\_error)

11.64s = Training runtime

0.47s = Validation runtime

Fitting model: CatBoost\_BAG\_L2 ... Training model for up to 544.12s of the 544.1s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-136.5497 = Validation score (-mean\_absolute\_error)

25.51s = Training runtime

0.3s = Validation runtime

Fitting model: NeuralNetFastAI\_BAG\_L2 ... Training model for up to 537.1s of the 537.08s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-134.5143 = Validation score (-mean\_absolute\_error)

200.32s = Training runtime

3.36s = Validation runtime

Fitting model: XGBoost\_BAG\_L2 ... Training model for up to 495.48s of the 495.46s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-134.9755 = Validation score (-mean\_absolute\_error)

15.13s = Training runtime

0.66s = Validation runtime

Fitting model: NeuralNetTorch\_BAG\_L2 ... Training model for up to 491.1s of the 491.09s of remaining time.

Fitting 8 child models (S5F1 - S5F8) | Fitting with ParallelLocalFoldFittingStrategy

-135.6579 = Validation score (-mean\_absolute\_error)

```

300.76s = Training runtime
2.73s   = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 428.2s of the
428.18s of remaining time.
Fitting 8 child models (S5F1 - S5F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.106 = Validation score (-mean_absolute_error)
38.7s    = Training runtime
1.07s    = Validation runtime
Repeating k-fold bagging: 6/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 418.88s of the
418.87s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-138.5034 = Validation score (-mean_absolute_error)
21.42s    = Training runtime
1.08s     = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 414.16s of the
414.15s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.0657 = Validation score (-mean_absolute_error)
14.05s    = Training runtime
0.56s     = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 410.49s of the
410.47s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-136.5486 = Validation score (-mean_absolute_error)
30.66s    = Training runtime
0.36s     = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 403.95s of
the 403.93s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-134.4519 = Validation score (-mean_absolute_error)
240.63s   = Training runtime
4.11s     = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 362.21s of the
362.19s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
-134.9381 = Validation score (-mean_absolute_error)
17.98s    = Training runtime
0.8s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 357.82s of the
357.8s of remaining time.
Fitting 8 child models (S6F1 - S6F8) | Fitting with

```

```

ParallelLocalFoldFittingStrategy
    -135.6084      = Validation score    (-mean_absolute_error)
    366.09s      = Training    runtime
    3.24s        = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 291.08s of the
291.06s of remaining time.
    Fitting 8 child models (S6F1 - S6F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0691      = Validation score    (-mean_absolute_error)
    47.02s        = Training    runtime
    1.3s          = Validation runtime
Repeating k-fold bagging: 7/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 281.45s of the
281.43s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -138.4553      = Validation score    (-mean_absolute_error)
    25.67s        = Training    runtime
    1.29s        = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 275.97s of the
275.96s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.0951      = Validation score    (-mean_absolute_error)
    16.48s        = Training    runtime
    0.65s         = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 272.33s of the
272.32s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.5204      = Validation score    (-mean_absolute_error)
    35.88s        = Training    runtime
    0.42s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 265.75s of
the 265.73s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -134.341       = Validation score    (-mean_absolute_error)
    281.2s        = Training    runtime
    4.77s        = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 223.79s of the
223.78s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.0242      = Validation score    (-mean_absolute_error)
    21.05s        = Training    runtime
    0.93s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 219.29s of the

```

```

219.27s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -135.5437      = Validation score    (-mean_absolute_error)
    429.87s      = Training    runtime
    3.8s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 154.06s of the
154.05s of remaining time.
    Fitting 8 child models (S7F1 - S7F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -136.067      = Validation score    (-mean_absolute_error)
    54.33s       = Training    runtime
    1.48s        = Validation runtime
Completed 7/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L3 ... Training model for up to 360.0s of the
145.45s of remaining time.
    -132.8711     = Validation score    (-mean_absolute_error)
    0.73s        = Training    runtime
    0.0s         = Validation runtime
AutoGluon training complete, total runtime = 1655.32s ... Best model:
"WeightedEnsemble_L3"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_90_A/")
WARNING: eval_metric='pearsonr' does not support sample weights so they will be
ignored in reported metric.
Evaluation: mean_absolute_error on test data: -186.24721157826426
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -186.24721157826426,
    "root_mean_squared_error": -420.62416270646276,
    "mean_squared_error": -176924.68625251285,
    "r2": 0.8716298559177108,
    "pearsonr": 0.9358825898767282,
    "median_absolute_error": -3.769514799118042
}

Evaluation on test data:
-186.24721157826426

```

```

[ ]: loc = "B"
predictors[1] = fit_predictor_for_location(loc)

```

```

Warning: path already exists! This predictor may overwrite an existing
predictor! path="AutogluonModels/submission_90_B"
Presets specified: ['best_quality']
Stack configuration (auto_stack=True): num_stack_levels=1, num_bag_folds=8,
num_bag_sets=20

```



Values in column 'sample\_weight' used as sample weights instead of predictive features. Evaluation will report weighted metrics, so ensure same column exists in test data.

Beginning AutoGluon training ... Time limit = 1800s

AutoGluon will save models to "AutogluonModels/submission\_90\_B/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 297.94 GB / 315.93 GB (94.3%)

Train Data Rows: 30792

Train Data Columns: 41

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of label-column == float and many unique label-values observed).

Label info (max, min, mean, stddev): (1152.3, -0.0, 97.67477, 195.03642)

If 'regression' is not the correct problem\_type, please manually specify the problem\_type parameter during predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression'])

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 130207.91 MB

Train Data (Original) Memory Usage: 11.39 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 2 features to boolean dtype as they only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

Stage 4 Generators:

Fitting DropUniqueFeatureGenerator...

Stage 5 Generators:

Fitting DropDuplicatesFeatureGenerator...

Training model for location B...

Train data sample weight sum: 30792.0

Train data number of rows: 30792

Test data sample weight sum: 2051

Test data number of rows: 2051

Useless Original Features (Count: 1): ['location']

These features carry no predictive signal and should be manually investigated.

This is typically a feature which has the same value for all rows.

These features do not need to be present at inference time.

Types of features in original data (raw dtype, special dtypes):

```
('float', []) : 38 | ['absolute_humidity_2m:gm3',  
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',  
'clear_sky_rad:W', ...]
```

```
('int', []) : 1 | ['is_estimated']
```

Types of features in processed data (raw dtype, special dtypes):

```
('float', []) : 37 | ['absolute_humidity_2m:gm3',  
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',  
'clear_sky_rad:W', ...]
```

```
('int', ['bool']) : 2 | ['elevation:m', 'is_estimated']
```

0.2s = Fit runtime

39 features in original data used to generate 39 features in processed data.

Train Data (Processed) Memory Usage: 9.18 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.18s ...

AutoGluon will gauge predictive performance using evaluation metric:

'mean\_absolute\_error'

This metric's sign has been flipped to adhere to being higher\_is\_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval\_metric parameter of Predictor()

User-specified model hyperparameters to be fit:

```
{  
    'NN_TORCH': {},  
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],  
'GBMLarge'],  
    'CAT': {},  
    'XGB': {},  
    'FASTAI': {},  
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',  
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':  
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},  
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',  
'problem_types': ['regression', 'quantile']}}],  
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',  
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':  
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},  
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',  
'problem_types': ['regression', 'quantile']}}],  
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},  
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],  
}
```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 1199.58s of the 1799.81s of remaining time.

```

-45.0801          = Validation score    (-mean_absolute_error)
0.03s            = Training   runtime
0.43s            = Validation runtime
Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 1199.06s of
the 1799.29s of remaining time.
-44.9883          = Validation score    (-mean_absolute_error)
0.03s            = Training   runtime
0.43s            = Validation runtime
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1198.55s of the
1798.78s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-26.235           = Validation score    (-mean_absolute_error)
38.79s           = Training   runtime
20.81s           = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1155.96s of the
1756.2s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-27.6872          = Validation score    (-mean_absolute_error)
44.86s           = Training   runtime
16.33s           = Validation runtime
Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1107.7s of
the 1707.94s of remaining time.
-31.8894          = Validation score    (-mean_absolute_error)
11.38s           = Training   runtime
1.17s            = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1093.5s of the
1693.74s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-29.4372          = Validation score    (-mean_absolute_error)
206.59s          = Training   runtime
0.11s            = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 885.71s of the
1485.94s of remaining time.
-32.6592          = Validation score    (-mean_absolute_error)
1.93s            = Training   runtime
1.16s            = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 880.85s of
the 1481.08s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-35.7575          = Validation score    (-mean_absolute_error)
38.29s           = Training   runtime
0.6s             = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 841.26s of the
1441.5s of remaining time.

```

```

    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -30.4775      = Validation score    (-mean_absolute_error)
    105.3s       = Training    runtime
    24.76s       = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 731.93s of the
1332.17s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -28.9112      = Validation score    (-mean_absolute_error)
    198.41s      = Training    runtime
    0.46s        = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 532.27s of the
1132.5s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -26.5707      = Validation score    (-mean_absolute_error)
    139.44s      = Training    runtime
    26.82s       = Validation runtime
Completed 1/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
985.81s of remaining time.
    -25.5137      = Validation score    (-mean_absolute_error)
    0.82s         = Training    runtime
    0.0s          = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 984.97s of the
984.95s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -23.1597      = Validation score    (-mean_absolute_error)
    10.07s        = Training    runtime
    0.69s         = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 973.33s of the
973.31s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.8809      = Validation score    (-mean_absolute_error)
    2.81s         = Training    runtime
    0.12s         = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ... Training model for up to 969.32s of
the 969.3s of remaining time.
    -21.959       = Validation score    (-mean_absolute_error)
    15.03s        = Training    runtime
    1.19s         = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 952.57s of the
952.55s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with

```

```

ParallelLocalFoldFittingStrategy
    -22.9341          = Validation score    (-mean_absolute_error)
    37.75s           = Training   runtime
    0.06s            = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L2 ... Training model for up to 913.65s of the
913.64s of remaining time.
    -21.9244          = Validation score    (-mean_absolute_error)
    2.39s            = Training   runtime
    1.22s            = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 909.48s of
the 909.47s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.1783          = Validation score    (-mean_absolute_error)
    38.24s           = Training   runtime
    0.63s            = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 869.92s of the
869.91s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.6477          = Validation score    (-mean_absolute_error)
    3.04s            = Training   runtime
    0.13s            = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 865.47s of the
865.46s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.2668          = Validation score    (-mean_absolute_error)
    89.3s            = Training   runtime
    0.6s             = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 774.85s of the
774.84s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.4653          = Validation score    (-mean_absolute_error)
    98.41s           = Training   runtime
    1.73s            = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 671.26s of the
671.25s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -23.0488          = Validation score    (-mean_absolute_error)
    17.41s           = Training   runtime
    1.12s            = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 662.56s of the
662.55s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with

```

```

ParallelLocalFoldFittingStrategy
    -22.8141      = Validation score    (-mean_absolute_error)
    5.29s        = Training    runtime
    0.23s        = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 658.7s of the
658.69s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.8401      = Validation score    (-mean_absolute_error)
    48.76s        = Training    runtime
    0.12s        = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 646.5s of the
646.48s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.0293      = Validation score    (-mean_absolute_error)
    76.47s        = Training    runtime
    1.31s        = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 606.94s of the
606.92s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.5148      = Validation score    (-mean_absolute_error)
    6.21s         = Training    runtime
    0.26s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 602.47s of the
602.45s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.0278      = Validation score    (-mean_absolute_error)
    189.07s       = Training    runtime
    1.11s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 501.31s of the
501.3s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.2095      = Validation score    (-mean_absolute_error)
    204.76s       = Training    runtime
    4.79s         = Validation runtime
Repeating k-fold bagging: 3/20
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 389.89s of the
389.88s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -23.0258      = Validation score    (-mean_absolute_error)
    24.06s        = Training    runtime
    1.58s         = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 381.85s of the

```

```

381.83s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.7665      = Validation score    (-mean_absolute_error)
    8.3s         = Training   runtime
    0.35s        = Validation runtime
Fitting model: CatBoost_BAG_L2 ... Training model for up to 377.65s of the
377.63s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.7782      = Validation score    (-mean_absolute_error)
    62.72s        = Training   runtime
    0.19s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ... Training model for up to 362.45s of
the 362.43s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.9458      = Validation score    (-mean_absolute_error)
    115.45s       = Training   runtime
    1.97s         = Validation runtime
Fitting model: XGBoost_BAG_L2 ... Training model for up to 322.19s of the
322.17s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.4749      = Validation score    (-mean_absolute_error)
    9.33s         = Training   runtime
    0.39s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ... Training model for up to 317.71s of the
317.69s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.9414      = Validation score    (-mean_absolute_error)
    270.03s       = Training   runtime
    1.62s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ... Training model for up to 235.42s of the
235.4s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.1519      = Validation score    (-mean_absolute_error)
    296.67s       = Training   runtime
    6.16s         = Validation runtime
Completed 3/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L3 ... Training model for up to 360.0s of the
139.58s of remaining time.
    -21.4202      = Validation score    (-mean_absolute_error)
    0.7s          = Training   runtime
    0.0s          = Validation runtime
AutoGluon training complete, total runtime = 1661.17s ... Best model:

```

```
"WeightedEnsemble_L3"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_90_B/")
```

```
[ ]: loc = "C"
predictors[2] = fit_predictor_for_location(loc)
```

### 3 Submit

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data
```

```
[ ]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
↪ "location"], left_on=["ds", "location"])

#test_data_merged
```

```
[ ]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in test_data.groupby('location'):
    i = location_map[loc]
    subset = test_data_merged[test_data_merged["location"] == loc].
↪reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)

    # get past predictions
    past_pred = predictors[i].
↪predict(train_data_with_dates[train_data_with_dates["location"] == loc])
```



```

train_data_with_dates.loc[train_data_with_dates["location"] == loc,
↪ "prediction"] = past_pred

```

```

[ ]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
↪ y='y', ax=ax, label="train data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
↪ y='prediction', ax=ax, label="past predictions")

    # title
    ax.set_title(f"Predictions for location {loc}")

```

```

[ ]: # concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df

```

```

[ ]: # Save the submission DataFrame to submissions folder, create new name based on
↪ last submission, format is submission_<last_submission_number + 1>.csv

# Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
↪ index=False)
print("jall1a")

```

```

[ ]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)

```

```

[ ]: # save this notebook to submissions folder
import subprocess

```

```
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↳join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
    ↳ipynb"])
```

```
[ ]: # feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
    ↳time_limit=60*10)
```

```
[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
    ↳time_limit=60*10)
```

```
[ ]: display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)

subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↳join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↳"autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
#             ↳stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
#             ↳strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
#     ↳'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', hello if hello is
#     ↳not None else 'Henrik eller Jørgen'])

# branch_name = new_filename
```

```

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push', '↪
↪ 'origin', branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream', '↪
↪ 'origin', branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])

```