

autogluon_each_location

October 18, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = None
presets = 'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = [] # ["hour", "year"]
use_estimated_diff_attr = False
use_is_estimated_attr = True

use_groups = False
n_groups = 8

auto_stack = False
num_stack_levels = 0
num_bag_folds = 8
num_bag_sets = 20

use_tune_data = True
use_test_data = True
tune_and_test_length = 24*30*6 # 3 months from end
holdout_frac = None
use_bag_holdout = True # Enable this if there is a large gap between score_val_
    ↪ and score_test in stack models.

sample_weight = None # 'sample_weight' # None
weight_evaluation = False
sample_weight_estimated = 1

run_analysis = False

[2]: import pandas as pd
import numpy as np
```

```

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift all columns with "1h" in them by 1 hour, so that for index 16:00,
    # we have the values from 17:00
    # but only for the columns with "1h" in the name
    #X_shifted = X.filter(regex="\dh").shift(-1, axis=1)
    #print(f"Number of columns with 1h in name: {X_shifted.columns}")

    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

    X_shifted = X[X.index.minute==0][columns].copy()
    # loop through all rows and check if index + 1 hour is in the index, if so
    # get that value, else nan
    count1 = 0
    count2 = 0
    for i in range(len(X_shifted)):
        if X_shifted.index[i] + pd.Timedelta('1 hour') in X.index:
            count1 += 1
            X_shifted.iloc[i] = X.loc[X_shifted.index[i] + pd.Timedelta('1
            hour')][columns]
        else:
            count2 += 1
            X_shifted.iloc[i] = np.nan

    print("COUNT1", count1)
    print("COUNT2", count2)

    X_old_unshifted = X[X.index.minute==0][columns]
    # rename X_old_unshifted columns to have _not_shifted at the end
    X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
    columns]

    # put the shifted columns back into the original dataframe
    #X[columns] = X_shifted[columns]

    date_calc = None
    if "date_calc" in X.columns:

```

```

        date_calc = X[X.index.minute == 0]['date_calc']

    # resample to hourly
    print("index: ", X.index[0])
    X = X.resample('H').mean()
    print("index AFTER: ", X.index[0])

    X[columns] = X_shifted[columns]
    #X[X_old_unshifted.columns] = X_old_unshifted

    if date_calc is not None:
        X['date_calc'] = date_calc

    return X

def fix_X(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    # with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    X = feature_engineering(X)

    return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    if weight_evaluation:
        # add sample weights, which are 1 for observed and 3 for estimated
        X_train_observed["sample_weight"] = 1
        X_train_estimated["sample_weight"] = sample_weight_estimated
        X_test["sample_weight"] = sample_weight_estimated

    y_train['ds'] = pd.to_datetime(y_train['time'])

```

```

y_train.drop(columns=['time'], inplace=True)
y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
        X_test["estimated_diff_hours"] = (X_test.index - pd.
to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

        X_train_estimated["estimated_diff_hours"] =
X_train_estimated["estimated_diff_hours"].astype('int64')
        # the filled once will get dropped later anyways, when we drop y nans
        X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
fillna(-50).astype('int64')

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0
        X_train_estimated["is_estimated"] = 1
        X_test["is_estimated"] = 1

    # drop date_calc
    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

```

```

X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

# print number of nans in y
print(f"Number of nans in y: {X_train['y'].isna().sum()}")

X_train["location"] = location
X_test["location"] = location

return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↳X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

Processing location A...

COUNT1 29667

COUNT2 1

index: 2019-06-02 22:00:00

index AFTER: 2019-06-02 22:00:00

```

COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 0
Processing location B...
COUNT1 29232
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 4
Processing location C...
COUNT1 29206
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 6059

```

1 Feature engineering

```

[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
               ↪inplace=True)

```

```

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

print(X_train.head())

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()

        # Calculate the size of each group for this location
        group_size = len(loc_df) // n_groups

        # Create a new 'group' column for this location
        loc_df['group'] = np.repeat(range(n_groups),
    ↪ repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

        # Append to list of grouped DataFrames
        grouped_dfs.append(loc_df)

    # Concatenate all the grouped DataFrames back together
    X_train = pd.concat(grouped_dfs)
    X_train.sort_index(inplace=True)
    print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
    ↪ "dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm",
    ↪ "wind_speed_u_10m:ms", "wind_speed_v_10m:ms"]

X_train.drop(columns=to_drop, inplace=True)

```

```
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)
```

```
absolute_humidity_2m:gm3  air_density_2m:kgm3  \
ds
2019-06-02 22:00:00      7.700      1.22825
2019-06-02 23:00:00      7.700      1.22350
2019-06-03 00:00:00      7.875      1.21975
2019-06-03 01:00:00      8.425      1.21800
2019-06-03 02:00:00      8.950      1.21800
```

```
ceiling_height_agl:m  clear_sky_energy_1h:J  \
ds
2019-06-02 22:00:00      1728.949951      0.000000
2019-06-02 23:00:00      1689.824951      0.000000
2019-06-03 00:00:00      1563.224976      0.000000
2019-06-03 01:00:00      1283.425049      6546.899902
2019-06-03 02:00:00      1003.500000      102225.898438
```

```
clear_sky_rad:W  cloud_base_agl:m  dew_or_rime:idx  \
ds
2019-06-02 22:00:00      0.00      1728.949951      0.0
2019-06-02 23:00:00      0.00      1689.824951      0.0
2019-06-03 00:00:00      0.00      1563.224976      0.0
2019-06-03 01:00:00      0.75      1283.425049      0.0
2019-06-03 02:00:00      23.10      1003.500000      0.0
```

```
dew_point_2m:K  diffuse_rad:W  diffuse_rad_1h:J  ...  \
ds
2019-06-02 22:00:00      280.299988      0.000      0.000000  ...
2019-06-02 23:00:00      280.299988      0.000      0.000000  ...
2019-06-03 00:00:00      280.649994      0.000      0.000000  ...
2019-06-03 01:00:00      281.674988      0.300      7743.299805  ...
2019-06-03 02:00:00      282.500000      11.975      60137.601562  ...
```

```
t_1000hPa:K  total_cloud_cover:p  visibility:m  \
ds
2019-06-02 22:00:00      286.225006      100.000000  40386.476562
2019-06-02 23:00:00      286.899994      100.000000  33770.648438
2019-06-03 00:00:00      286.950012      100.000000  13595.500000
2019-06-03 01:00:00      286.750000      100.000000  2321.850098
2019-06-03 02:00:00      286.450012      99.224998  11634.799805
```

```
wind_speed_10m:ms  wind_speed_u_10m:ms  \
ds
2019-06-02 22:00:00      3.600      -3.575
```


2019-06-02 23:00:00	3.350	-3.350
2019-06-03 00:00:00	3.050	-2.950
2019-06-03 01:00:00	2.725	-2.600
2019-06-03 02:00:00	2.550	-2.350

	wind_speed_v_10m:ms	wind_speed_w_1000hPa:ms	\
ds			
2019-06-02 22:00:00	-0.500	0.0	
2019-06-02 23:00:00	0.275	0.0	
2019-06-03 00:00:00	0.750	0.0	
2019-06-03 01:00:00	0.875	0.0	
2019-06-03 02:00:00	0.925	0.0	

	is_estimated	y	location
ds			
2019-06-02 22:00:00	0	0.00	A
2019-06-02 23:00:00	0	0.00	A
2019-06-03 00:00:00	0	0.00	A
2019-06-03 01:00:00	0	0.00	A
2019-06-03 02:00:00	0	19.36	A

[5 rows x 48 columns]

```
[4]: from autogluon.tabular import TabularDataset, TabularPredictor
from autogluon.timeseries import TimeSeriesDataFrame
import numpy as np
train_data = TabularDataset('X_train_raw.csv')
# set group column of train_data be increasing from 0 to 7 based on time, the
# first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
train_data['ds'] = pd.to_datetime(train_data['ds'])
train_data = train_data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
split_time = pd.to_datetime(train_data["ds"]).max() - pd.
    timedelta(hours=tune_and_test_length)
# 2022-10-28 22:00:00
split_time = pd.to_datetime("2022-10-28 22:00:00")
train_set = TabularDataset(train_data[train_data["ds"] < split_time])
test_set = TabularDataset(train_data[train_data["ds"] >= split_time])
if use_groups:
    test_set = test_set.drop(columns=['group'])
```

```

if do_drop_ds:
    train_set = train_set.drop(columns=['ds'])
    test_set = test_set.drop(columns=['ds'])
    train_data = train_data.drop(columns=['ds'])

def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
        loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

tuning_data = None
if use_tune_data:
    train_data = train_set
    if use_test_data:
        # split test_set in half, use first half for tuning
        tuning_data, test_data = [], []
        for loc in locations:
            loc_test_set = test_set[test_set["location"] == loc]
            # randomly shuffle the loc_test_set
            loc_tuning_data, loc_test_data = pd.DataFrame(), pd.DataFrame()
            for i in range(200):
                # get a part of the test set corresponding to i/100th part of
                the test set and shuffle
                num_bins = len(loc_test_set) // 200
                # set seed to i so that we get the same shuffle every time
                np.random.seed(i)
                current_bin = loc_test_set.iloc[i*num_bins:min((i+1)*num_bins,
                len(loc_test_set))].sample(frac=1)
                loc_tuning_data = pd.concat([loc_tuning_data, current_bin.iloc[:
                len(current_bin)//2]])
                loc_test_data = pd.concat([loc_test_data, current_bin.
                iloc[len(current_bin)//2:]]

                tuning_data.append(loc_tuning_data)
                test_data.append(loc_test_data)
            tuning_data = pd.concat(tuning_data)
            test_data = pd.concat(test_data)
            print("Shapes of tuning and test", tuning_data.shape[0], test_data.
            shape[0], tuning_data.shape[0] + test_data.shape[0])

        else:

```

```

    tuning_data = test_set
    print("Shape of tuning", tuning_data.shape[0])

    # ensure sample weights for your tuning data sum to the number of rows in
    ↪ the tuning data.
    if weight_evaluation:
        tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        train_data = train_set
        test_data = test_set
        print("Shape of test", test_data.shape[0])

    # ensure sample weights for your training (or tuning) data sum to the number of
    ↪ rows in the training (or tuning) data.
    if weight_evaluation:
        train_data = normalize_sample_weights_per_location(train_data)
    if use_test_data:
        test_data = normalize_sample_weights_per_location(test_data)

train_data = TabularDataset(train_data)
if use_tune_data:
    tuning_data = TabularDataset(tuning_data)
if use_test_data:
    test_data = TabularDataset(test_data)

```

Shapes of tuning and test 5000 5400 10400

```

[5]: if run_analysis:
    import autogluon.eda.auto as auto
    auto.dataset_overview(train_data=train_data, test_data=test_data,
    ↪ label="y", sample=None)

```

```

[6]: if run_analysis:
    auto.target_analysis(train_data=train_data, label="y", sample=None)

```

2 Starting

```

[7]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    ↪ filename in os.listdir('submissions') if "submission" in filename]))

```

```

print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)

```

Last submission number: 93
 Now creating submission number: 94
 New filename: submission_94

```
[8]: predictors = [None, None, None]
```

```

[9]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪ train and tune data and test data
    if weight_evaluation:
        print("Train data sample weight sum:",
        ↪ train_data[train_data["location"] == loc]["sample_weight"].sum())
        print("Train data number of rows:", train_data[train_data["location"]
        ↪ == loc].shape[0])
        if use_tune_data:
            print("Tune data sample weight sum:",
            ↪ tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
            print("Tune data number of rows:",
            ↪ tuning_data[tuning_data["location"] == loc].shape[0])
        if use_test_data:
            print("Test data sample weight sum:",
            ↪ test_data[test_data["location"] == loc]["sample_weight"].sum())
            print("Test data number of rows:", test_data[test_data["location"]
            ↪ == loc].shape[0])
    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        # sample_weight=sample_weight,
        # weight_evaluation=weight_evaluation,
        # groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc],
        time_limit=time_limit,

```

```

        presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
        ↪somethin, will be overwritten anyways
        num_bag_sets=num_bag_sets,
        tuning_data=tuning_data[tuning_data["location"] == loc].
        ↪reset_index(drop=True) if use_tune_data else None,
        use_bag_holdout=use_bag_holdout,
        # holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
        ↪drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

```

Presets specified: ['best_quality']
Stack configuration (auto_stack=True): num_stack_levels=1, num_bag_folds=8,
num_bag_sets=1
Beginning AutoGluon training ...
AutoGluon will save models to "AutogluonModels/submission_94_A/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 236.41 GB / 315.93 GB (74.8%)
Train Data Rows: 29667
Train Data Columns: 38
Tuning Data Rows: 2000
Tuning Data Columns: 38
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
Label info (max, min, mean, stddev): (5733.42, 0.0, 674.14552,
1195.53172)
If 'regression' is not the correct problem_type, please manually specify

```

```

the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 132284.33 MB
    Train Data (Original) Memory Usage: 11.21 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
        Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...

Training model for location A...

    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
    This is typically a feature which has the same value for all
rows.
    These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 35 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 35 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', ['bool']) : 1 | ['is_estimated']
    0.1s = Fit runtime
    36 features in original data used to generate 36 features in processed
data.

    Train Data (Processed) Memory Usage: 8.9 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.16s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'

    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()

```

use_bag_holdout=True, will use tuning_data as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge': {},
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
        'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
        {'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
        {'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
        'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
        'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
        {'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
        {'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
        'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
        {'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ...

```
-133.6185      = Validation score    (-mean_absolute_error)
0.03s         = Training runtime
0.39s         = Validation runtime
```

Fitting model: KNeighborsDist_BAG_L1 ...

```
-133.2862      = Validation score    (-mean_absolute_error)
0.03s         = Training runtime
0.36s         = Validation runtime
```

Fitting model: LightGBMXT_BAG_L1 ...

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

```
-109.6322      = Validation score    (-mean_absolute_error)
30.9s         = Training runtime
12.85s        = Validation runtime
```

Fitting model: LightGBM_BAG_L1 ...

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

```
-112.566       = Validation score    (-mean_absolute_error)
30.13s        = Training runtime
14.93s        = Validation runtime
```

Fitting model: RandomForestMSE_BAG_L1 ...

```
-115.8985      = Validation score    (-mean_absolute_error)
8.42s         = Training runtime
```

```

    1.14s      = Validation runtime
Fitting model: CatBoost_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -119.1588      = Validation score      (-mean_absolute_error)
    200.7s      = Training runtime
    0.12s      = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ...
    -116.9235      = Validation score      (-mean_absolute_error)
    1.69s      = Training runtime
    1.13s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -123.2376      = Validation score      (-mean_absolute_error)
    37.05s      = Training runtime
    0.49s      = Validation runtime
Fitting model: XGBoost_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -113.5542      = Validation score      (-mean_absolute_error)
    51.37s      = Training runtime
    2.68s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -105.7007      = Validation score      (-mean_absolute_error)
    94.72s      = Training runtime
    0.37s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -109.2278      = Validation score      (-mean_absolute_error)
    100.24s      = Training runtime
    21.28s      = Validation runtime
Fitting model: WeightedEnsemble_L2 ...
    -104.6917      = Validation score      (-mean_absolute_error)
    0.43s      = Training runtime
    0.0s      = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -105.6881      = Validation score      (-mean_absolute_error)
    3.04s      = Training runtime
    0.17s      = Validation runtime
Fitting model: LightGBM_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with

```



```

ParallelLocalFoldFittingStrategy
    -107.9129      = Validation score    (-mean_absolute_error)
    1.89s         = Training runtime
    0.08s         = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ...
    -109.2965      = Validation score    (-mean_absolute_error)
    13.28s        = Training runtime
    1.22s         = Validation runtime
Fitting model: CatBoost_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -107.5346      = Validation score    (-mean_absolute_error)
    5.29s         = Training runtime
    0.05s         = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L2 ...
    -107.6613      = Validation score    (-mean_absolute_error)
    2.19s         = Training runtime
    1.19s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -110.219       = Validation score    (-mean_absolute_error)
    37.09s        = Training runtime
    0.49s         = Validation runtime
Fitting model: XGBoost_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -107.5087      = Validation score    (-mean_absolute_error)
    2.85s         = Training runtime
    0.11s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -104.6545      = Validation score    (-mean_absolute_error)
    53.87s        = Training runtime
    0.5s          = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -107.8037      = Validation score    (-mean_absolute_error)
    5.25s         = Training runtime
    0.18s         = Validation runtime
Fitting model: WeightedEnsemble_L3 ...
    -104.567       = Validation score    (-mean_absolute_error)
    0.35s         = Training runtime
    0.0s          = Validation runtime
AutoGluon training complete, total runtime = 735.71s ... Best model:
"WeightedEnsemble_L3"

```

```

TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_94_A/")
Evaluation: mean_absolute_error on test data: -108.60384533057054
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -108.60384533057054,
    "root_mean_squared_error": -305.6816918188323,
    "mean_squared_error": -93441.29671322356,
    "r2": 0.8805259503688295,
    "pearsonr": 0.9385599911845793,
    "median_absolute_error": -0.41422155499458313
}

Evaluation on test data:
-108.60384533057054

```


```

[10]: leaderboards = []
if use_test_data:
    lb = predictors[0].leaderboard(test_data[test_data["location"] == loc])
    lb["location"] = loc
    leaderboards.append(lb)
    test_data[test_data["location"] == loc]["y"].plot()
if use_tune_data:
    tuning_data[tuning_data["location"] == loc]["y"].plot()

```

	model	score_test	score_val	pred_time_test
0	WeightedEnsemble_L2	-106.464945	-104.691735	8.610645
34.861226	226.312503		0.004399	0.000636
0.432885	2	True	12	
1	NeuralNetTorch_BAG_L1	-106.483024	-105.700704	0.246203
0.366351	94.717332		0.246203	0.366351
94.717332	1	True	10	
2	WeightedEnsemble_L3	-108.603845	-104.566951	14.854486
56.399232	612.512369		0.002249	0.000729
0.347646	3	True	22	
3	NeuralNetTorch_BAG_L2	-109.282383	-104.654502	14.784469
56.230217	609.129228		0.372648	0.495219
53.866686	2	True	20	
4	XGBoost_BAG_L2	-111.336104	-107.508695	14.499334
55.844983	558.110395		0.087513	0.109986
2.847853	2	True	19	
5	LightGBMXT_BAG_L2	-111.378388	-105.688058	14.479589
55.903284	558.298036		0.067768	0.168286
3.035494	2	True	13	
6	LightGBMLarge_BAG_L2	-111.447478	-107.803747	14.547423

55.910179	560.507911		0.135602	0.175181
5.245369	2	True	21	
7	RandomForestMSE_BAG_L2	-111.668660	-109.296451	15.014299
56.959545	568.539886		0.602477	1.224547
13.277343	2	True	15	
8	CatBoost_BAG_L2	-111.691459	-107.534587	14.454121
55.787292	560.553231		0.042300	0.052295
5.290688	2	True	16	
9	ExtraTreesMSE_BAG_L2	-111.778220	-107.661283	15.008604
56.926856	557.452750		0.596782	1.191859
2.190207	2	True	17	
10	LightGBM_BAG_L2	-112.451150	-107.912871	14.449175
55.818054	557.155247		0.037354	0.083057
1.892704	2	True	14	
11	LightGBMLarge_BAG_L1	-113.161480	-109.227811	4.858676
21.281557	100.236939		4.858676	21.281557
100.236939	1	True	11	
12	NeuralNetFastAI_BAG_L2	-113.675624	-110.219003	15.063398
56.223438	592.356961		0.651576	0.488440
37.094419	2	True	18	
13	LightGBM_BAG_L1	-116.417802	-112.565988	2.965765
14.931696	30.125217		2.965765	14.931696
30.125217	1	True	4	
14	LightGBMXT_BAG_L1	-117.168060	-109.632222	3.465628
12.850172	30.898239		3.465628	12.850172
30.898239	1	True	3	
15	XGBoost_BAG_L1	-117.911577	-113.554171	0.935745
2.682815	51.367466		0.935745	2.682815
51.367466	1	True	9	
16	ExtraTreesMSE_BAG_L1	-118.563553	-116.923479	0.556412
1.130832	1.693071		0.556412	1.130832
1.693071	1	True	7	
17	RandomForestMSE_BAG_L1	-119.504864	-115.898548	0.551370
1.135428	8.415134		0.551370	1.135428
8.415134	1	True	5	
18	CatBoost_BAG_L1	-125.072803	-119.158813	0.103171
0.117834	200.702585		0.103171	0.117834
200.702585	1	True	6	
19	NeuralNetFastAI_BAG_L1	-129.000628	-123.237616	0.644522
0.485936	37.052310		0.644522	0.485936
37.052310	1	True	8	
20	KNeighborsDist_BAG_L1	-134.896380	-133.286181	0.035738
0.362511	0.027107		0.035738	0.362511
0.027107	1	True	2	
21	KNeighborsUnif_BAG_L1	-135.049170	-133.618518	0.048592
0.389867	0.027141		0.048592	0.389867
0.027141	1	True	1	



```
[11]: loc = "B"
      predictors[1] = fit_predictor_for_location(loc)
```

Training model for location B...

Presets specified: ['best_quality']

Stack configuration (auto_stack=True): num_stack_levels=1, num_bag_folds=8, num_bag_sets=1

Beginning AutoGluon training ...

AutoGluon will save models to "AutogluonModels/submission_94_B/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 233.15 GB / 315.93 GB (73.8%)

Train Data Rows: 29218

Train Data Columns: 38

Tuning Data Rows: 1600

Tuning Data Columns: 38

Label Column: y

Preprocessing data ...

AutoGluon infers your prediction problem is: 'regression' (because dtype of

```

label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (1152.3, 0.0, 102.58516, 198.99359)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                130212.9 MB
    Train Data (Original) Memory Usage: 10.91 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
        These features carry no predictive signal and should be manually
investigated.
            This is typically a feature which has the same value for all
rows.
                These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 35 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 35 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', ['bool']) : 1 | ['is_estimated']
    0.1s = Fit runtime
    36 features in original data used to generate 36 features in processed
data.
    Train Data (Processed) Memory Usage: 8.66 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.17s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.

```

To change this, specify the `eval_metric` parameter of `Predictor()` use `bag_holdout=True`, will use `tuning_data` as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ...

-25.9296 = Validation score (-mean_absolute_error)

0.03s = Training runtime

0.37s = Validation runtime

Fitting model: KNeighborsDist_BAG_L1 ...

-26.2267 = Validation score (-mean_absolute_error)

0.03s = Training runtime

0.38s = Validation runtime

Fitting model: LightGBMXT_BAG_L1 ...

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

-22.8699 = Validation score (-mean_absolute_error)

30.66s = Training runtime

13.9s = Validation runtime

Fitting model: LightGBM_BAG_L1 ...

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

-21.5638 = Validation score (-mean_absolute_error)

32.67s = Training runtime

19.25s = Validation runtime

Fitting model: RandomForestMSE_BAG_L1 ...

-19.0379 = Validation score (-mean_absolute_error)

```

    9.39s    = Training    runtime
    1.09s    = Validation runtime
Fitting model: CatBoost_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.5788      = Validation score    (-mean_absolute_error)
    197.16s      = Training    runtime
    0.1s         = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ...
    -19.1835      = Validation score    (-mean_absolute_error)
    1.62s        = Training    runtime
    1.11s        = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -22.355      = Validation score    (-mean_absolute_error)
    35.5s        = Training    runtime
    0.44s        = Validation runtime
Fitting model: XGBoost_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.4734      = Validation score    (-mean_absolute_error)
    92.57s       = Training    runtime
    20.57s       = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -19.4201      = Validation score    (-mean_absolute_error)
    160.61s      = Training    runtime
    0.37s        = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -19.942      = Validation score    (-mean_absolute_error)
    102.46s      = Training    runtime
    22.49s       = Validation runtime
Fitting model: WeightedEnsemble_L2 ...
    -18.2596      = Validation score    (-mean_absolute_error)
    0.41s        = Training    runtime
    0.0s         = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.5565      = Validation score    (-mean_absolute_error)
    4.09s        = Training    runtime
    0.3s         = Validation runtime
Fitting model: LightGBM_BAG_L2 ...

```

```

    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.6628      = Validation score    (-mean_absolute_error)
    2.23s        = Training runtime
    0.12s        = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ...
    -20.7892      = Validation score    (-mean_absolute_error)
    13.43s       = Training runtime
    1.12s       = Validation runtime
Fitting model: CatBoost_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.3994      = Validation score    (-mean_absolute_error)
    9.0s         = Training runtime
    0.06s        = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L2 ...
    -20.5017      = Validation score    (-mean_absolute_error)
    2.02s        = Training runtime
    1.14s        = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.9894      = Validation score    (-mean_absolute_error)
    36.19s       = Training runtime
    0.49s        = Validation runtime
Fitting model: XGBoost_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.7225      = Validation score    (-mean_absolute_error)
    2.93s        = Training runtime
    0.11s        = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.5158      = Validation score    (-mean_absolute_error)
    100.1s       = Training runtime
    0.54s        = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.5379      = Validation score    (-mean_absolute_error)
    67.01s       = Training runtime
    1.66s        = Validation runtime
Fitting model: WeightedEnsemble_L3 ...
    -20.2692      = Validation score    (-mean_absolute_error)
    0.34s        = Training runtime
    0.0s         = Validation runtime
AutoGluon training complete, total runtime = 957.88s ... Best model:

```




```
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_94_B/")
Evaluation: mean_absolute_error on test data: -19.784612134597683
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -19.784612134597683,
    "root_mean_squared_error": -59.48722932746857,
    "mean_squared_error": -3538.7304530588362,
    "r2": 0.8091036613883708,
    "pearsonr": 0.9210702081720571,
    "median_absolute_error": -1.9378460291816468
}

Evaluation on test data:
-19.784612134597683
```

```
[12]: if use_test_data:
        lb = predictors[1].leaderboard(test_data[test_data["location"] == loc])
        test_data[test_data["location"] == loc]["y"].plot()
        lb["location"] = loc
        leaderboards.append(lb)
    if use_tune_data:
        tuning_data[tuning_data["location"] == loc]["y"].plot()
```

	model	score_test	score_val	pred_time_test	pred_time_val
fit_time	pred_time_test_marginal	pred_time_val_marginal	fit_time_marginal		
stack_level	can_infer	fit_order			
0	WeightedEnsemble_L2	-19.784612	-18.259556	1.190219	2.577594
172.023397		0.003609		0.000605	0.406425
2	True	12			
1	RandomForestMSE_BAG_L1	-20.506324	-19.037859	0.465188	1.088299
9.389158		0.465188		1.088299	9.389158
1	True	5			
2	NeuralNetTorch_BAG_L1	-20.636468	-19.420133	0.239440	0.374354
160.612212		0.239440		0.374354	160.612212
1	True	10			
3	ExtraTreesMSE_BAG_L1	-20.913315	-19.183511	0.481981	1.114335
1.615602		0.481981		1.114335	1.615602
1	True	7			
4	CatBoost_BAG_L2	-21.861565	-20.399400	13.679230	80.149340
671.681277		0.042715		0.061000	9.000148
2	True	16			
5	WeightedEnsemble_L3	-21.974210	-20.269160	14.537227	81.824679
774.139362		0.003396		0.000671	0.340890
3	True	22			
6	LightGBMLarge_BAG_L1	-22.002466	-19.942043	4.882097	22.486938

102.458817		4.882097		22.486938	102.458817
1	True	11			
7	ExtraTreesMSE_BAG_L2	-22.133535	-20.501681	14.126836	81.223771
664.699806		0.490320		1.135432	2.018677
2	True	17			
8	LightGBMXT_BAG_L2	-22.159803	-20.556477	13.735883	80.388342
666.766527		0.099368		0.300002	4.085398
2	True	13			
9	LightGBM_BAG_L2	-22.295762	-20.662814	13.682784	80.209733
664.911675		0.046268		0.121393	2.230546
2	True	14			
10	NeuralNetTorch_BAG_L2	-22.335899	-20.515756	14.000797	80.627576
762.779648		0.364281		0.539237	100.098518
2	True	20			
11	XGBoost_BAG_L2	-22.358626	-20.722484	13.722353	80.200207
665.610823		0.085837		0.111868	2.929693
2	True	19			
12	LightGBMLarge_BAG_L2	-22.394402	-20.537913	14.433328	81.745611
729.687910		0.796812		1.657271	67.006781
2	True	21			
13	RandomForestMSE_BAG_L2	-22.512068	-20.789223	14.114488	81.205569
676.115350		0.477973		1.117229	13.434220
2	True	15			
14	NeuralNetFastAI_BAG_L2	-22.805315	-20.989446	14.292718	80.574287
698.871838		0.656202		0.485947	36.190709
2	True	18			
15	XGBoost_BAG_L1	-23.200676	-21.473426	2.839509	20.574657
92.570059		2.839509		20.574657	92.570059
1	True	9			
16	LightGBM_BAG_L1	-23.361430	-21.563789	1.922492	19.250667
32.670249		1.922492		19.250667	32.670249
1	True	4			
17	NeuralNetFastAI_BAG_L1	-23.810966	-22.354976	0.641081	0.443902
35.495719		0.641081		0.443902	35.495719
1	True	8			
18	CatBoost_BAG_L1	-24.726350	-22.578804	0.107755	0.102869
197.158424		0.107755		0.102869	197.158424
1	True	6			
19	LightGBMXT_BAG_L1	-24.872298	-22.869948	1.988891	13.900755
30.658652		1.988891		13.900755	30.658652
1	True	3			
20	KNeighborsDist_BAG_L1	-27.173093	-26.226721	0.030431	0.376924
0.026137		0.030431		0.376924	0.026137
1	True	2			
21	KNeighborsUnif_BAG_L1	-27.275008	-25.929619	0.037650	0.374639
0.026100		0.037650		0.374639	0.026100
1	True	1			



```
[13]: loc = "C"
      predictors[2] = fit_predictor_for_location(loc)
```

Presets specified: ['best_quality']

Stack configuration (auto_stack=True): num_stack_levels=1, num_bag_folds=8,
num_bag_sets=1

Training model for location C...

Beginning AutoGluon training ...

AutoGluon will save models to "AutogluonModels/submission_94_C/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

Disk Space Avail: 229.73 GB / 315.93 GB (72.7%)

Train Data Rows: 23141

Train Data Columns: 38

Tuning Data Rows: 1400

Tuning Data Columns: 38

Label Column: y

Preprocessing data ...

```

AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and label-values can't be converted to int).
    Label info (max, min, mean, stddev): (999.6, 0.0, 82.28807, 171.35018)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                128643.73 MB
    Train Data (Original) Memory Usage: 8.69 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
        These features carry no predictive signal and should be manually
investigated.
            This is typically a feature which has the same value for all
rows.
                These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 35 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 35 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', ['bool']) : 1 | ['is_estimated']
    0.1s = Fit runtime
    36 features in original data used to generate 36 features in processed
data.
    Train Data (Processed) Memory Usage: 6.9 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.15s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.

```

The metric score can be multiplied by -1 to get the metric value.

To change this, specify the `eval_metric` parameter of `Predictor()` `use_bag_holdout=True`, will use `tuning_data` as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ...

-23.8716 = Validation score (-mean_absolute_error)

0.02s = Training runtime

0.83s = Validation runtime

Fitting model: KNeighborsDist_BAG_L1 ...

-24.1578 = Validation score (-mean_absolute_error)

0.02s = Training runtime

0.25s = Validation runtime

Fitting model: LightGBMXT_BAG_L1 ...

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

-21.6402 = Validation score (-mean_absolute_error)

28.93s = Training runtime

14.97s = Validation runtime

Fitting model: LightGBM_BAG_L1 ...

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

-21.8184 = Validation score (-mean_absolute_error)

29.96s = Training runtime

8.21s = Validation runtime

Fitting model: RandomForestMSE_BAG_L1 ...

```

-21.1265          = Validation score    (-mean_absolute_error)
5.06s            = Training   runtime
0.74s           = Validation runtime
Fitting model: CatBoost_BAG_L1 ...
  Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-21.4958          = Validation score    (-mean_absolute_error)
191.58s          = Training   runtime
0.1s            = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ...
-20.8786          = Validation score    (-mean_absolute_error)
1.03s           = Training   runtime
0.76s           = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ...
  Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-19.7355          = Validation score    (-mean_absolute_error)
29.32s          = Training   runtime
0.4s            = Validation runtime
Fitting model: XGBoost_BAG_L1 ...
  Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-21.5175          = Validation score    (-mean_absolute_error)
51.17s          = Training   runtime
3.23s           = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ...
  Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-21.1018          = Validation score    (-mean_absolute_error)
64.76s          = Training   runtime
0.33s           = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ...
  Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-21.2795          = Validation score    (-mean_absolute_error)
88.73s          = Training   runtime
8.6s            = Validation runtime
Fitting model: WeightedEnsemble_L2 ...
-19.2408          = Validation score    (-mean_absolute_error)
0.44s           = Training   runtime
0.0s            = Validation runtime
Fitting 9 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ...
  Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-22.2002          = Validation score    (-mean_absolute_error)
2.27s           = Training   runtime
0.14s           = Validation runtime

```

```

Fitting model: LightGBM_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.4717      = Validation score    (-mean_absolute_error)
    1.78s        = Training runtime
    0.07s        = Validation runtime
Fitting model: RandomForestMSE_BAG_L2 ...
    -21.5425      = Validation score    (-mean_absolute_error)
    8.23s        = Training runtime
    0.79s        = Validation runtime
Fitting model: CatBoost_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.399      = Validation score    (-mean_absolute_error)
    6.9s         = Training runtime
    0.05s        = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L2 ...
    -21.3696      = Validation score    (-mean_absolute_error)
    1.31s        = Training runtime
    0.8s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.3348      = Validation score    (-mean_absolute_error)
    29.79s       = Training runtime
    0.39s        = Validation runtime
Fitting model: XGBoost_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -20.7431      = Validation score    (-mean_absolute_error)
    2.95s        = Training runtime
    0.1s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.4385      = Validation score    (-mean_absolute_error)
    51.7s        = Training runtime
    0.46s        = Validation runtime
Fitting model: LightGBMLarge_BAG_L2 ...
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -21.5175      = Validation score    (-mean_absolute_error)
    5.92s        = Training runtime
    0.19s        = Validation runtime
Fitting model: WeightedEnsemble_L3 ...
    -20.7431      = Validation score    (-mean_absolute_error)
    0.34s        = Training runtime
    0.0s         = Validation runtime

```

AutoGluon training complete, total runtime = 645.08s ... Best model:

"WeightedEnsemble_L2"

TabularPredictor saved. To load, use: predictor =

TabularPredictor.load("AutogluonModels/submission_94_C/")

Evaluation: mean_absolute_error on test data: -19.739339913797004

Note: Scores are always higher_is_better. This metric score can be multiplied by -1 to get the metric value.

Evaluations on test data:

```
{
  "mean_absolute_error": -19.739339913797004,
  "root_mean_squared_error": -48.6566638550214,
  "mean_squared_error": -2367.470937500546,
  "r2": 0.7684410045169413,
  "pearsonr": 0.9031989428540793,
  "median_absolute_error": -0.988118712902069
}
```


Evaluation on test data:

-19.739339913797004

```
[14]: if use_test_data:
      lb = predictors[2].leaderboard(test_data[test_data["location"] == loc])
      test_data[test_data["location"] == loc]["y"].plot()
      lb["location"] = loc
      leaderboards.append(lb)
      if use_tune_data:
          tuning_data[tuning_data["location"] == loc]["y"].plot()
```

	model	score_test	score_val	pred_time_test	pred_time_val
fit_time	pred_time_test_marginal	pred_time_val_marginal	fit_time_marginal		
stack_level	can_infer	fit_order			
0	WeightedEnsemble_L2	-19.739340	-19.240758	1.042274	2.310131
95.566443		0.003901		0.000656	0.435457
2	True	12			
1	NeuralNetFastAI_BAG_L1	-19.785365	-19.735470	0.500163	0.399011
29.319467		0.500163		0.399011	29.319467
1	True	8			
2	NeuralNetTorch_BAG_L1	-19.836243	-21.101790	0.214143	0.325194
64.757850		0.214143		0.325194	64.757850
1	True	10			
3	NeuralNetTorch_BAG_L2	-20.786542	-21.438521	11.972856	38.874038
542.276094		0.328881		0.455749	51.700819
2	True	20			
4	ExtraTreesMSE_BAG_L1	-21.032200	-20.878620	0.300408	0.756305
1.029495		0.300408		0.756305	1.029495
1	True	7			
5	RandomForestMSE_BAG_L1	-21.035793	-21.126537	0.285563	0.743042
5.057846		0.285563		0.743042	5.057846
1	True	5			

6	XGBoost_BAG_L2	-21.092990	-20.743136	11.723059	38.515971
493.527910		0.079084		0.097682	2.952635
2	True	19			
7	WeightedEnsemble_L3	-21.092990	-20.743136	11.724888	38.516615
493.869780		0.001828		0.000645	0.341869
3	True	22			
8	XGBoost_BAG_L1	-21.172547	-21.517455	1.273292	3.230571
51.172070		1.273292		3.230571	51.172070
1	True	9			
9	ExtraTreesMSE_BAG_L2	-21.247823	-21.369577	11.932498	39.216379
491.883744		0.288523		0.798090	1.308468
2	True	17			
10	CatBoost_BAG_L2	-21.290812	-21.399022	11.685564	38.470409
497.472901		0.041589		0.052120	6.897625
2	True	16			
11	CatBoost_BAG_L1	-21.322932	-21.495798	0.100098	0.103473
191.576589		0.100098		0.103473	191.576589
1	True	6			
12	LightGBMLarge_BAG_L1	-21.417781	-21.279517	4.108705	8.595093
88.725821		4.108705		8.595093	88.725821
1	True	11			
13	NeuralNetFastAI_BAG_L2	-21.440529	-21.334762	12.154256	38.805320
520.370076		0.510281		0.387030	29.794801
2	True	18			
14	LightGBM_BAG_L2	-21.474505	-21.471699	11.679493	38.490852
492.359499		0.035518		0.072562	1.784223
2	True	14			
15	RandomForestMSE_BAG_L2	-21.557714	-21.542477	11.923722	39.208714
498.800313		0.279747		0.790424	8.225037
2	True	15			
16	LightGBMXT_BAG_L1	-21.703644	-21.640242	2.594915	14.972013
28.927319		2.594915		14.972013	28.927319
1	True	3			
17	LightGBM_BAG_L1	-21.708364	-21.818396	2.221903	8.213923
29.960548		2.221903		8.213923	29.960548
1	True	4			
18	LightGBMLarge_BAG_L2	-21.710290	-21.517539	11.845297	38.604977
496.496011		0.201322		0.186688	5.920736
2	True	21			
19	LightGBMXT_BAG_L2	-21.886947	-22.200213	11.707611	38.556420
492.845138		0.063636		0.138131	2.269862
2	True	13			
20	KNeighborsUnif_BAG_L1	-23.769434	-23.871633	0.023660	0.828964
0.024173		0.023660		0.828964	0.024173
1	True	1			
21	KNeighborsDist_BAG_L1	-23.799712	-24.157808	0.021125	0.250699
0.024097		0.021125		0.250699	0.024097
1	True	2			



```
[15]: # save leaderboards to csv
pd.concat(leaderboards).to_csv(f"leaderboards/{new_filename}.csv")
```

3 Submit

```
[16]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data
```

Loaded data from: X_train_raw.csv | Columns = 40 / 40 | Rows = 92945 -> 92945

Loaded data from: X_test_raw.csv | Columns = 39 / 39 | Rows = 4608 -> 4608

```
[17]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
```

```
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
↪ "location"], left_on=["ds", "location"])

#test_data_merged
```

Loaded data from: test.csv | Columns = 4 / 4 | Rows = 2160 -> 2160

```
[18]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in test_data.groupby('location'):
    i = location_map[loc]
    subset = test_data_merged[test_data_merged["location"] == loc].
↪ reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)



    # get past predictions
    past_pred = predictors[i].
↪ predict(train_data_with_dates[train_data_with_dates["location"] == loc])
    train_data_with_dates.loc[train_data_with_dates["location"] == loc,
↪ "prediction"] = past_pred
```

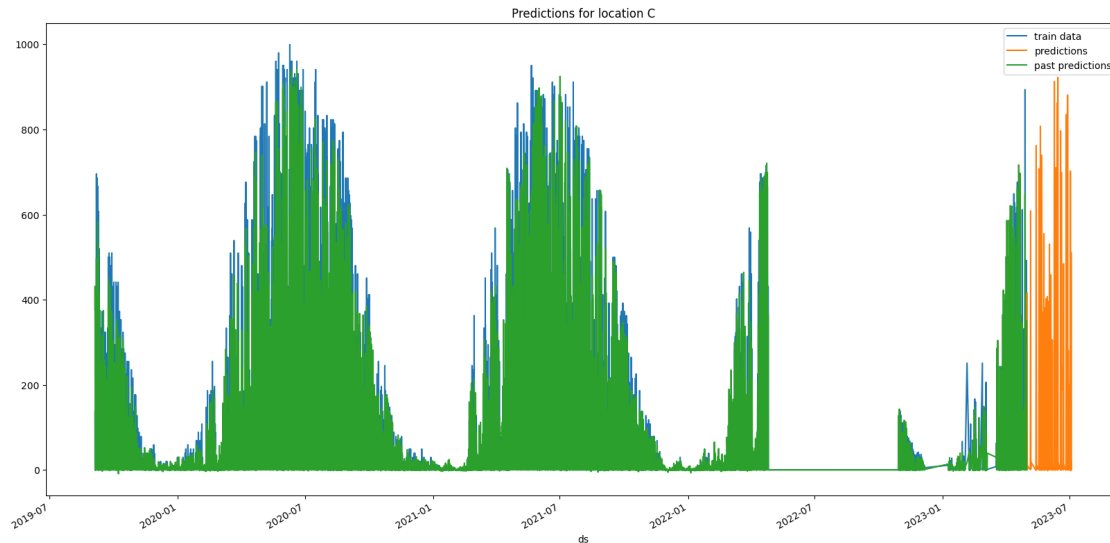
```
[19]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
↪ y='y', ax=ax, label="train data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
↪ y='prediction', ax=ax, label="past predictions")

    # title
    ax.set_title(f"Predictions for location {loc}")
```





```
[20]: # concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df
```

```
[20]:
```

	id	prediction
0	0	0.388048
1	1	0.393996
2	2	0.383476
3	3	56.945789
4	4	347.566467
..
715	2155	46.367020
716	2156	14.130499
717	2157	1.687788
718	2158	0.084798
719	2159	18.746828

[2160 rows x 2 columns]

```
[21]: # Save the submission DataFrame to submissions folder, create new name based on
↳ last submission, format is submission_<last_submission_number + 1>.csv

# Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
↳ index=False)
print("jallia")
```

Saving submission to submissions/submission_94.csv
jall1a

```
[22]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)
```

<IPython.core.display.Javascript object>

```
[23]: # save this notebook to submissions folder
import subprocess
import os
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f'{new_filename}.pdf'), "autogluon_each_location.
    ↪ipynb"])
```

[NbConvertApp] Converting notebook autogluon_each_location.ipynb to pdf
/opt/conda/lib/python3.10/site-packages/nbconvert/utils/pandoc.py:51:
RuntimeWarning: You are using an unsupported version of pandoc (2.9.2.1).
Your version must be at least (2.14.2) but less than (4.0.0).
Refer to <https://pandoc.org/installing.html>.
Continuing with doubts...
check_pandoc_version()
[NbConvertApp] Support files will be in notebook_pdfs/submission_94_files/
[NbConvertApp] Making directory
./notebook_pdfs/submission_94_files/notebook_pdfs
[NbConvertApp] Writing 144295 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 210326 bytes to notebook_pdfs/submission_94.pdf

```
[23]: CompletedProcess(args=['jupyter', 'nbconvert', '--to', 'pdf', '--output',
    'notebook_pdfs/submission_94.pdf', 'autogluon_each_location.ipynb'],
    returncode=0)
```

```
[24]: # feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
```

```

estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
↳time_limit=60*10)

```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'location', 'prediction']

Computing feature importance via permutation shuffling for 36 features using 4392 rows with 10 shuffle sets... Time limit: 600s...

6164.29s = Expected runtime (616.43s per shuffle set)

349.2s = Actual runtime (Completed 1 of 10 shuffle sets) (Early stopping due to lack of time...)

```

[24]:

```

	importance	stddev	p_value	n	p99_high	\
direct_rad_1h:J	1.661192e+02	NaN	NaN	1	NaN	
diffuse_rad_1h:J	3.306991e+01	NaN	NaN	1	NaN	
sun_azimuth:d	3.023113e+01	NaN	NaN	1	NaN	
clear_sky_energy_1h:J	2.874515e+01	NaN	NaN	1	NaN	
clear_sky_rad:W	2.338161e+01	NaN	NaN	1	NaN	
direct_rad:W	1.696939e+01	NaN	NaN	1	NaN	
diffuse_rad:W	1.157282e+01	NaN	NaN	1	NaN	
total_cloud_cover:p	9.132186e+00	NaN	NaN	1	NaN	
cloud_base_agl:m	5.896452e+00	NaN	NaN	1	NaN	
relative_humidity_1000hPa:p	4.958079e+00	NaN	NaN	1	NaN	
snow_water:kgm2	4.532055e+00	NaN	NaN	1	NaN	
effective_cloud_cover:p	4.454154e+00	NaN	NaN	1	NaN	
fresh_snow_6h:cm	3.940948e+00	NaN	NaN	1	NaN	
sun_elevation:d	3.692114e+00	NaN	NaN	1	NaN	
wind_speed_10m:ms	3.464202e+00	NaN	NaN	1	NaN	
visibility:m	3.348638e+00	NaN	NaN	1	NaN	
t_1000hPa:K	3.316352e+00	NaN	NaN	1	NaN	
is_in_shadow:idx	2.676490e+00	NaN	NaN	1	NaN	
precip_5min:mm	1.691140e+00	NaN	NaN	1	NaN	
sfc_pressure:hPa	1.440330e+00	NaN	NaN	1	NaN	
super_cooled_liquid_water:kgm2	1.249810e+00	NaN	NaN	1	NaN	
is_day:idx	1.040715e+00	NaN	NaN	1	NaN	
snow_melt_10min:mm	1.026113e+00	NaN	NaN	1	NaN	
precip_type_5min:idx	1.003837e+00	NaN	NaN	1	NaN	
fresh_snow_3h:cm	8.393493e-01	NaN	NaN	1	NaN	
ceiling_height_agl:m	5.453121e-01	NaN	NaN	1	NaN	
pressure_100m:hPa	2.547247e-01	NaN	NaN	1	NaN	
air_density_2m:kgm3	1.916241e-01	NaN	NaN	1	NaN	
snow_depth:cm	1.884514e-01	NaN	NaN	1	NaN	
is_estimated	5.189622e-08	NaN	NaN	1	NaN	
rain_water:kgm2	-2.983781e-02	NaN	NaN	1	NaN	
fresh_snow_1h:cm	-9.135414e-02	NaN	NaN	1	NaN	
pressure_50m:hPa	-3.242515e-01	NaN	NaN	1	NaN	
msl_pressure:hPa	-3.661329e-01	NaN	NaN	1	NaN	

dew_point_2m:K	-1.109279e+00	NaN	NaN	1	NaN
absolute_humidity_2m:gm3	-1.179244e+00	NaN	NaN	1	NaN

	p99_low
direct_rad_1h:J	NaN
diffuse_rad_1h:J	NaN
sun_azimuth:d	NaN
clear_sky_energy_1h:J	NaN
clear_sky_rad:W	NaN
direct_rad:W	NaN
diffuse_rad:W	NaN
total_cloud_cover:p	NaN
cloud_base_agl:m	NaN
relative_humidity_1000hPa:p	NaN
snow_water:kgm2	NaN
effective_cloud_cover:p	NaN
fresh_snow_6h:cm	NaN
sun_elevation:d	NaN
wind_speed_10m:ms	NaN
visibility:m	NaN
t_1000hPa:K	NaN
is_in_shadow:idx	NaN
precip_5min:mm	NaN
sfc_pressure:hPa	NaN
super_cooled_liquid_water:kgm2	NaN
is_day:idx	NaN
snow_melt_10min:mm	NaN
precip_type_5min:idx	NaN
fresh_snow_3h:cm	NaN
ceiling_height_agl:m	NaN
pressure_100m:hPa	NaN
air_density_2m:kgm3	NaN
snow_depth:cm	NaN
is_estimated	NaN
rain_water:kgm2	NaN
fresh_snow_1h:cm	NaN
pressure_50m:hPa	NaN
msl_pressure:hPa	NaN
dew_point_2m:K	NaN
absolute_humidity_2m:gm3	NaN

```
[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
↪time_limit=60*10)
```

These features in provided data are not utilized by the predictor and will be


```

ignored: ['ds', 'elevation:m', 'location', 'prediction']
Computing feature importance via permutation shuffling for 36 features using
5000 rows with 10 shuffle sets... Time limit: 600s...
      8118.65s          = Expected runtime (811.87s per shuffle set)

```

```

[ ]: display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)

subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
↳ join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
↳ "autoglun_each_location.ipynb"])

[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
↳ stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
↳ strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
↳ 'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', 'hello if hello is
↳ not None else 'Henrik eller Jørgen'])

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote

```

```

# output, success = execute_git_command(git_repo_path, ['push',
↳ 'origin', branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream',
↳ 'origin', branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])

```