

autogluon_each_location

October 18, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*30
presets = 'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = [] # ["hour", "year"]
use_estimated_diff_attr = False
use_is_estimated_attr = True

use_groups = False
n_groups = 8

auto_stack = False
num_stack_levels = 0
num_bag_folds = 8
num_bag_sets = 20

use_tune_data = True
use_test_data = False
tune_and_test_length = 0.25 # 3 months from end
holdout_frac = None
use_bag_holdout = True # Enable this if there is a large gap between score_val_
    ↪ and score_test in stack models.

sample_weight = None # 'sample_weight' # None
weight_evaluation = False
sample_weight_estimated = 1

run_analysis = False

[2]: import pandas as pd
import numpy as np
```

```

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift all columns with "1h" in them by 1 hour, so that for index 16:00,
    # we have the values from 17:00
    # but only for the columns with "1h" in the name
    #X_shifted = X.filter(regex="\dh").shift(-1, axis=1)
    #print(f"Number of columns with 1h in name: {X_shifted.columns}")

    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

    X_shifted = X[X.index.minute==0][columns].copy()
    # loop through all rows and check if index + 1 hour is in the index, if so
    # get that value, else nan
    count1 = 0
    count2 = 0
    for i in range(len(X_shifted)):
        if X_shifted.index[i] + pd.Timedelta('1 hour') in X.index:
            count1 += 1
            X_shifted.iloc[i] = X.loc[X_shifted.index[i] + pd.Timedelta('1
            hour')][columns]
        else:
            count2 += 1
            X_shifted.iloc[i] = np.nan

    print("COUNT1", count1)
    print("COUNT2", count2)

    X_old_unshifted = X[X.index.minute==0][columns]
    # rename X_old_unshifted columns to have _not_shifted at the end
    X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
    columns]

    # put the shifted columns back into the original dataframe
    #X[columns] = X_shifted[columns]

    date_calc = None
    if "date_calc" in X.columns:

```

```

        date_calc = X[X.index.minute == 0]['date_calc']

    # resample to hourly
    print("index: ", X.index[0])
    X = X.resample('H').mean()
    print("index AFTER: ", X.index[0])

    X[columns] = X_shifted[columns]
    #X[X_old_unshifted.columns] = X_old_unshifted

    if date_calc is not None:
        X['date_calc'] = date_calc

    return X

def fix_X(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    # with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    X = feature_engineering(X)

    return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    if weight_evaluation:
        # add sample weights, which are 1 for observed and 3 for estimated
        X_train_observed["sample_weight"] = 1
        X_train_estimated["sample_weight"] = sample_weight_estimated
        X_test["sample_weight"] = sample_weight_estimated

    y_train['ds'] = pd.to_datetime(y_train['time'])

```

```

y_train.drop(columns=['time'], inplace=True)
y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
↳location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
↳handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
↳pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
        X_test["estimated_diff_hours"] = (X_test.index - pd.
↳to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

        X_train_estimated["estimated_diff_hours"] =
↳X_train_estimated["estimated_diff_hours"].astype('int64')
        # the filled once will get dropped later anyways, when we drop y nans
        X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
↳fillna(-50).astype('int64')

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0
        X_train_estimated["is_estimated"] = 1
        X_test["is_estimated"] = 1

    # drop date_calc
    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

```

```

X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

# print number of nans in y
print(f"Number of nans in y: {X_train['y'].isna().sum()}")

X_train["location"] = location
X_test["location"] = location

return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↳X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

```

Processing location A...
COUNT1 29667
COUNT2 1
index: 2019-06-02 22:00:00
index AFTER: 2019-06-02 22:00:00

```

```

COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 0
Processing location B...
COUNT1 29232
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 4
Processing location C...
COUNT1 29206
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 6059

```

1 Feature engineering

```

[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
               inplace=True)

```

```

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

print(X_train.head())

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()

        # Calculate the size of each group for this location
        group_size = len(loc_df) // n_groups

        # Create a new 'group' column for this location
        loc_df['group'] = np.repeat(range(n_groups),
    ↪ repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

        # Append to list of grouped DataFrames
        grouped_dfs.append(loc_df)

    # Concatenate all the grouped DataFrames back together
    X_train = pd.concat(grouped_dfs)
    X_train.sort_index(inplace=True)
    print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
    ↪ "dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm",
    ↪ "wind_speed_u_10m:ms", "wind_speed_v_10m:ms", "snow_melt_10min:mm",
    ↪ "rain_water:kgm2", "dew_point_2m:K", "precip_5min:mm", "absolute_humidity_2m:
    ↪ gm3", "air_density_2m:kgm3"]

```

```

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

```

absolute_humidity_2m:gm3  air_density_2m:kgm3  \
ds
2019-06-02 22:00:00          7.700          1.22825
2019-06-02 23:00:00          7.700          1.22350
2019-06-03 00:00:00          7.875          1.21975
2019-06-03 01:00:00          8.425          1.21800
2019-06-03 02:00:00          8.950          1.21800

```

```

ceiling_height_agl:m  clear_sky_energy_1h:J  \
ds
2019-06-02 22:00:00      1728.949951          0.000000
2019-06-02 23:00:00      1689.824951          0.000000
2019-06-03 00:00:00      1563.224976          0.000000
2019-06-03 01:00:00      1283.425049      6546.899902
2019-06-03 02:00:00      1003.500000     102225.898438

```

```

clear_sky_rad:W  cloud_base_agl:m  dew_or_rime:idx  \
ds
2019-06-02 22:00:00          0.00      1728.949951          0.0
2019-06-02 23:00:00          0.00      1689.824951          0.0
2019-06-03 00:00:00          0.00      1563.224976          0.0
2019-06-03 01:00:00          0.75      1283.425049          0.0
2019-06-03 02:00:00         23.10      1003.500000          0.0

```

```

dew_point_2m:K  diffuse_rad:W  diffuse_rad_1h:J  ...  \
ds
2019-06-02 22:00:00      280.299988          0.000          0.000000  ...
2019-06-02 23:00:00      280.299988          0.000          0.000000  ...
2019-06-03 00:00:00      280.649994          0.000          0.000000  ...
2019-06-03 01:00:00      281.674988          0.300      7743.299805  ...
2019-06-03 02:00:00      282.500000         11.975     60137.601562  ...

```

```

t_1000hPa:K  total_cloud_cover:p  visibility:m  \
ds
2019-06-02 22:00:00      286.225006          100.000000  40386.476562
2019-06-02 23:00:00      286.899994          100.000000  33770.648438
2019-06-03 00:00:00      286.950012          100.000000  13595.500000
2019-06-03 01:00:00      286.750000          100.000000   2321.850098
2019-06-03 02:00:00      286.450012          99.224998  11634.799805

```

```

wind_speed_10m:ms  wind_speed_u_10m:ms  \

```



```

ds
2019-06-02 22:00:00      3.600      -3.575
2019-06-02 23:00:00      3.350      -3.350
2019-06-03 00:00:00      3.050      -2.950
2019-06-03 01:00:00      2.725      -2.600
2019-06-03 02:00:00      2.550      -2.350

      wind_speed_v_10m:ms  wind_speed_w_1000hPa:ms  \
ds
2019-06-02 22:00:00      -0.500      0.0
2019-06-02 23:00:00      0.275      0.0
2019-06-03 00:00:00      0.750      0.0
2019-06-03 01:00:00      0.875      0.0
2019-06-03 02:00:00      0.925      0.0

      is_estimated      y  location
ds
2019-06-02 22:00:00      0  0.00      A
2019-06-02 23:00:00      0  0.00      A
2019-06-03 00:00:00      0  0.00      A
2019-06-03 01:00:00      0  0.00      A
2019-06-03 02:00:00      0  19.36      A

```

[5 rows x 48 columns]

```

[4]: def normalize_sample_weights_per_location(df):
    for loc in locations:
        loc_df = df[df["location"] == loc]
        loc_df["sample_weight"] = loc_df["sample_weight"] /
        ↪ loc_df["sample_weight"].sum() * loc_df.shape[0]
        df[df["location"] == loc] = loc_df
    return df

import pandas as pd
import numpy as np

def split_and_shuffle_data(input_data, num_bins, frac1):
    """
    Splits the input_data into num_bins and shuffles them, then divides the
    ↪ bins into two datasets based on the given fraction for the first set.

    Args:
        input_data (pd.DataFrame): The data to be split and shuffled.
        num_bins (int): The number of bins to split the data into.
        frac1 (float): The fraction of each bin to go into the first output
        ↪ dataset.

```

```

Returns:
    pd.DataFrame, pd.DataFrame: The two output datasets.
    """
    # Validate the input fraction
    if frac1 < 0 or frac1 > 1:
        raise ValueError("frac1 must be between 0 and 1.")

    if frac1==1:
        return input_data, pd.DataFrame()

    # Calculate the fraction for the second output set
    frac2 = 1 - frac1

    # Calculate bin size
    bin_size = len(input_data) // num_bins

    # Initialize empty DataFrames for output
    output_data1 = pd.DataFrame()
    output_data2 = pd.DataFrame()

    for i in range(num_bins):
        # Shuffle the data in the current bin
        np.random.seed(i)
        current_bin = input_data.iloc[i * bin_size: (i + 1) * bin_size].
        ↪sample(frac=1)

        # Calculate the sizes for each output set
        size1 = int(len(current_bin) * frac1)

        # Split and append to output DataFrames
        output_data1 = pd.concat([output_data1, current_bin.iloc[:size1]])
        output_data2 = pd.concat([output_data2, current_bin.iloc[size1:]]

    # Shuffle and split the remaining data
    remaining_data = input_data.iloc[num_bins * bin_size:].sample(frac=1)
    remaining_size1 = int(len(remaining_data) * frac1)

    output_data1 = pd.concat([output_data1, remaining_data.iloc[:
    ↪remaining_size1]])
    output_data2 = pd.concat([output_data2, remaining_data.iloc[remaining_size1:
    ↪]])

    return output_data1, output_data2

```

```

[5]: from autogluon.tabular import TabularDataset, TabularPredictor
    from autogluon.timeseries import TimeSeriesDataFrame
    import numpy as np

```

```

data = TabularDataset('X_train_raw.csv')
# set group column of train_data be increasing from 0 to 7 based on time, the
↳ first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
data['ds'] = pd.to_datetime(data['ds'])
data = data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
# split_time = pd.to_datetime(train_data["ds"]).max() - pd.
↳ Timedelta(hours=tune_and_test_length)
# 2022-10-28 22:00:00
split_time = pd.to_datetime("2022-10-28 22:00:00")
train_set = TabularDataset(data[data["ds"] < split_time])
test_set = TabularDataset(data[data["ds"] >= split_time])

# shuffle test_set and only grab tune_and_test_length percent of it, rest goes
↳ to train_set
test_set, new_train_set = split_and_shuffle_data(test_set, 40,
↳ tune_and_test_length)

print("Length of train set before adding test set", len(train_set))
# add rest to train_set
train_set = pd.concat([train_set, new_train_set])
print("Length of train set after adding test set", len(train_set))
print("Length of test set", len(test_set))

if use_groups:
    test_set = test_set.drop(columns=['group'])

tuning_data = None
if use_tune_data:
    if use_test_data:
        # split test_set in half, use first half for tuning
        tuning_data, test_data = [], []
        for loc in locations:
            loc_test_set = test_set[test_set["location"] == loc]
            # randomly shuffle the loc_test_set

```

```

        loc_tuning_data, loc_test_data =   

↪split_and_shuffle_data(loc_test_set, 40, 0.5)
        tuning_data.append(loc_tuning_data)
        test_data.append(loc_test_data)
        tuning_data = pd.concat(tuning_data)
        test_data = pd.concat(test_data)
        print("Shapes of tuning and test", tuning_data.shape[0], test_data.  

↪shape[0], tuning_data.shape[0] + test_data.shape[0])

    else:
        tuning_data = test_set
        print("Shape of tuning", tuning_data.shape[0])

    # ensure sample weights for your tuning data sum to the number of rows in   

↪the tuning data.
    if weight_evaluation:
        tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        test_data = test_set
        print("Shape of test", test_data.shape[0])

train_data = train_set

# ensure sample weights for your training (or tuning) data sum to the number of   

↪rows in the training (or tuning) data.
if weight_evaluation:
    train_data = normalize_sample_weights_per_location(train_data)
    if use_test_data:
        test_data = normalize_sample_weights_per_location(test_data)

train_data = TabularDataset(train_data)
if use_tune_data:
    tuning_data = TabularDataset(tuning_data)
if use_test_data:
    test_data = TabularDataset(test_data)

```

Length of train set before adding test set 82026

Length of train set after adding test set 87486

Length of test set 5459

Shape of tuning 5459

```
[6]: if run_analysis:
    import autogluon.eda.auto as auto
    auto.dataset_overview(train_data=train_data, test_data=test_data,
        label="y", sample=None)
```

```
[7]: if run_analysis:
    auto.target_analysis(train_data=train_data, label="y", sample=None)
```

2 Starting

```
[8]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

Last submission number: 96
 Now creating submission number: 97
 New filename: submission_97

```
[9]: predictors = [None, None, None]
```

```
[10]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    train and tune data and test data
    if weight_evaluation:
        print("Train data sample weight sum:",
            train_data[train_data["location"] == loc]["sample_weight"].sum())
        print("Train data number of rows:", train_data[train_data["location"]
            == loc].shape[0])
        if use_tune_data:
            print("Tune data sample weight sum:",
                tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
```

```

        print("Tune data number of rows:",
↳tuning_data[tuning_data["location"] == loc].shape[0])
        if use_test_data:
            print("Test data sample weight sum:",
↳test_data[test_data["location"] == loc]["sample_weight"].sum())
            print("Test data number of rows:", test_data[test_data["location"]
↳== loc].shape[0])
        predictor = TabularPredictor(
            label=label,
            eval_metric=metric,
            path=f"AutogluonModels/{new_filename}_{loc}",
            # sample_weight=sample_weight,
            # weight_evaluation=weight_evaluation,
            # groups="group" if use_groups else None,
        ).fit(
            train_data=train_data[train_data["location"] == loc].
↳drop(columns=["ds"]),
            time_limit=time_limit,
            # presets=presets,
            num_stack_levels=num_stack_levels,
            num_bag_folds=num_bag_folds if not use_groups else 2, # just put
↳somethin, will be overwritten anyways
            num_bag_sets=num_bag_sets,
            tuning_data=tuning_data[tuning_data["location"] == loc].
↳reset_index(drop=True).drop(columns=["ds"]) if use_tune_data else None,
            use_bag_holdout=use_bag_holdout,
            # holdout_frac=holdout_frac,
        )

        # evaluate on test data
        if use_test_data:
            # drop sample_weight column
            t = test_data[test_data["location"] == loc]#.
↳drop(columns=["sample_weight"])
            perf = predictor.evaluate(t)
            print("Evaluation on test data:")
            print(perf[predictor.eval_metric.name])

        return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Beginning AutoGluon training ... Time limit = 1800s
AutoGluon will save models to "AutogluonModels/submission_97_A/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12

```

Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 200.38 GB / 315.93 GB (63.4%)
Train Data Rows: 31872
Train Data Columns: 32
Tuning Data Rows: 2187
Tuning Data Columns: 32
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (5733.42, 0.0, 649.68162,
1178.37671)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 132000.72 MB
    Train Data (Original) Memory Usage: 10.42 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...

Training model for location A...

    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 29 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'clear_sky_rad:W', 'cloud_base_agl:m', 'diffuse_rad:W',
...]
        ('int', []) : 1 | ['is_estimated']

```

```

Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 29 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'clear_sky_rad:W', 'cloud_base_agl:m', 'diffuse_rad:W',
...]
    ('int', ['bool']) : 1 | ['is_estimated']
0.1s = Fit runtime
30 features in original data used to generate 30 features in processed
data.

Train Data (Processed) Memory Usage: 7.94 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.16s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'

This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor()
use_bag_holdout=True, will use tuning_data as holdout (will not be used for
early stopping).
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Fitting 11 L1 models ...
Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 1799.84s of
the 1799.84s of remaining time.
-132.836          = Validation score    (-mean_absolute_error)
0.03s            = Training    runtime
0.38s            = Validation runtime
Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 1799.33s of
the 1799.33s of remaining time.
-132.5631         = Validation score    (-mean_absolute_error)
0.03s             = Training    runtime

```


0.39s = Validation runtime

Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1798.84s of the 1798.84s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-93.681 = Validation score (-mean_absolute_error)

28.46s = Training runtime

18.91s = Validation runtime

Fitting model: LightGBM_BAG_L1 ... Training model for up to 1759.66s of the 1759.66s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-97.2854 = Validation score (-mean_absolute_error)

23.71s = Training runtime

5.64s = Validation runtime

Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1731.84s of the 1731.84s of remaining time.

-105.3243 = Validation score (-mean_absolute_error)

7.75s = Training runtime

1.13s = Validation runtime

Fitting model: CatBoost_BAG_L1 ... Training model for up to 1721.7s of the 1721.7s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-103.3471 = Validation score (-mean_absolute_error)

192.01s = Training runtime

0.09s = Validation runtime

Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1528.51s of the 1528.51s of remaining time.

-107.644 = Validation score (-mean_absolute_error)

1.6s = Training runtime

1.14s = Validation runtime

Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1524.54s of the 1524.54s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-108.0147 = Validation score (-mean_absolute_error)

38.33s = Training runtime

0.48s = Validation runtime

Fitting model: XGBoost_BAG_L1 ... Training model for up to 1484.4s of the 1484.4s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-104.4468 = Validation score (-mean_absolute_error)

5.82s = Training runtime

0.3s = Validation runtime

Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1476.48s of the 1476.47s of remaining time.

```

    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -95.8917          = Validation score    (-mean_absolute_error)
    125.34s          = Training    runtime
    0.31s            = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1349.69s of the
1349.69s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -96.3103          = Validation score    (-mean_absolute_error)
    92.93s           = Training    runtime
    16.21s           = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1247.25s of the
1247.25s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -93.7227          = Validation score    (-mean_absolute_error)
    58.55s           = Training    runtime
    38.01s           = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1209.85s of the
1209.85s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -97.5177          = Validation score    (-mean_absolute_error)
    48.72s           = Training    runtime
    10.54s           = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1179.89s of the
1179.89s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -103.1179         = Validation score    (-mean_absolute_error)
    387.72s          = Training    runtime
    0.2s             = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 982.81s of
the 982.8s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -108.2172         = Validation score    (-mean_absolute_error)
    76.99s           = Training    runtime
    0.94s            = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 941.47s of the
941.46s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -103.2888         = Validation score    (-mean_absolute_error)
    13.9s            = Training    runtime
    0.66s            = Validation runtime

```

Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 931.77s of the 931.77s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-94.7112 = Validation score (-mean_absolute_error)
240.82s = Training runtime
0.61s = Validation runtime

Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 814.66s of the 814.66s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-95.9885 = Validation score (-mean_absolute_error)
183.8s = Training runtime
36.7s = Validation runtime

Completed 2/20 k-fold bagging repeats ...

Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the 708.54s of remaining time.

-89.9806 = Validation score (-mean_absolute_error)
0.45s = Training runtime
0.0s = Validation runtime

AutoGluon training complete, total runtime = 1091.93s ... Best model:

"WeightedEnsemble_L2"

TabularPredictor saved. To load, use: predictor =

TabularPredictor.load("AutogluonModels/submission_97_A/")

```
[11]: import matplotlib.pyplot as plt

leaderboards = [None, None, None]
def leaderboard_for_location(i, loc):
    if use_test_data:
        lb = predictors[i].leaderboard(test_data[test_data["location"] == loc])
        lb["location"] = loc
        plt.scatter(test_data[test_data["location"] == loc]["y"].index,
        ↪test_data[test_data["location"] == loc]["y"])
        if use_tune_data:
            plt.scatter(tuning_data[tuning_data["location"] == loc]["y"].index,
            ↪tuning_data[tuning_data["location"] == loc]["y"])
        plt.show()

    return lb
    else:
        return pd.DataFrame()

leaderboards[0] = leaderboard_for_location(0, loc)
```

```
[12]: loc = "B"
predictors[1] = fit_predictor_for_location(loc)
```

```
leaderboards[1] = leaderboard_for_location(1, loc)
```

```
Beginning AutoGluon training ... Time limit = 1800s
AutoGluon will save models to "AutogluonModels/submission_97_B/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 196.72 GB / 315.93 GB (62.3%)
Train Data Rows: 31020
Train Data Columns: 32
Tuning Data Rows: 1797
Tuning Data Columns: 32
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (1152.3, -0.0, 99.56591, 196.469)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 130075.61 MB
    Train Data (Original) Memory Usage: 10.04 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 29 | ['ceiling_height_agl:m',
```

```
'clear_sky_energy_1h:J', 'clear_sky_rad:W', 'cloud_base_agl:m', 'diffuse_rad:W',
...]
```

```
('int', []) : 1 | ['is_estimated']
```

Types of features in processed data (raw dtype, special dtypes):

```
('float', []) : 29 | ['ceiling_height_agl:m',
```

```
'clear_sky_energy_1h:J', 'clear_sky_rad:W', 'cloud_base_agl:m', 'diffuse_rad:W',
...]
```

```
('int', ['bool']) : 1 | ['is_estimated']
```

0.1s = Fit runtime

30 features in original data used to generate 30 features in processed data.

Train Data (Processed) Memory Usage: 7.65 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.16s ...

AutoGluon will gauge predictive performance using evaluation metric:

'mean_absolute_error'

This metric's sign has been flipped to adhere to being higher_is_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor() use_bag_holdout=True, will use tuning_data as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge': {},
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

Training model for location B..

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 1799.84s of the 1799.84s of remaining time.

-26.8714 = Validation score (-mean_absolute_error)

0.03s = Training runtime

```

    0.38s      = Validation runtime
Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 1799.19s of
the 1799.19s of remaining time.
    -26.8453      = Validation score      (-mean_absolute_error)
    0.03s      = Training runtime
    0.39s      = Validation runtime
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1798.71s of the
1798.7s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -17.086      = Validation score      (-mean_absolute_error)
    29.05s      = Training runtime
    18.94s      = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1763.74s of the
1763.73s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -16.5481      = Validation score      (-mean_absolute_error)
    32.3s      = Training runtime
    13.98s      = Validation runtime
Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1726.94s of
the 1726.94s of remaining time.
    -17.8185      = Validation score      (-mean_absolute_error)
    8.83s      = Training runtime
    1.12s      = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1715.93s of the
1715.93s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -17.4238      = Validation score      (-mean_absolute_error)
    194.87s      = Training runtime
    0.1s      = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1519.75s of the
1519.75s of remaining time.
    -17.1018      = Validation score      (-mean_absolute_error)
    1.54s      = Training runtime
    1.13s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1515.93s of
the 1515.93s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -15.344      = Validation score      (-mean_absolute_error)
    39.37s      = Training runtime
    0.47s      = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 1474.64s of the
1474.64s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

-16.6706          = Validation score    (-mean_absolute_error)
82.64s    = Training    runtime
19.6s     = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1386.19s of
the 1386.19s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.0656          = Validation score    (-mean_absolute_error)
185.86s    = Training    runtime
0.31s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1198.95s of the
1198.95s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-15.4634          = Validation score    (-mean_absolute_error)
95.34s     = Training    runtime
19.27s     = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1093.48s of the
1093.48s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-17.0788          = Validation score    (-mean_absolute_error)
58.56s     = Training    runtime
34.22s     = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1057.11s of the
1057.11s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-16.5591          = Validation score    (-mean_absolute_error)
63.31s     = Training    runtime
29.79s     = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1019.24s of the
1019.24s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-17.3786          = Validation score    (-mean_absolute_error)
390.21s    = Training    runtime
0.2s       = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 822.6s of the
822.6s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-15.142 = Validation score    (-mean_absolute_error)
77.13s   = Training    runtime
0.98s    = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 782.26s of the
782.26s of remaining time.

```

```

    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -16.5324      = Validation score    (-mean_absolute_error)
    164.4s       = Training    runtime
    41.45s       = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 691.33s of the
691.33s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.7927      = Validation score    (-mean_absolute_error)
    349.19s       = Training    runtime
    0.65s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 526.39s of the
526.39s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -15.4482      = Validation score    (-mean_absolute_error)
    189.08s       = Training    runtime
    43.7s         = Validation runtime
Completed 2/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
415.81s of remaining time.
    -12.7634      = Validation score    (-mean_absolute_error)
    0.47s         = Training    runtime
    0.0s          = Validation runtime
AutoGluon training complete, total runtime = 1384.68s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_97_B/")

```

```

[13]: loc = "C"
      predictors[2] = fit_predictor_for_location(loc)
      leaderboards[2] = leaderboard_for_location(2, loc)

```

```

Beginning AutoGluon training ... Time limit = 1800s
AutoGluon will save models to "AutogluonModels/submission_97_C/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 192.23 GB / 315.93 GB (60.8%)
Train Data Rows: 24594
Train Data Columns: 32
Tuning Data Rows: 1475
Tuning Data Columns: 32
Label Column: y
Preprocessing data ...

```



```

AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and label-values can't be converted to int).
    Label info (max, min, mean, stddev): (999.6, -0.0, 79.8926, 168.407)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                129625.83 MB
    Train Data (Original) Memory Usage: 7.98 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
        These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 29 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'clear_sky_rad:W', 'cloud_base_agl:m', 'diffuse_rad:W',
...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 29 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'clear_sky_rad:W', 'cloud_base_agl:m', 'diffuse_rad:W',
...]
        ('int', ['bool']) : 1 | ['is_estimated']
    0.1s = Fit runtime
    30 features in original data used to generate 30 features in processed
data.
    Train Data (Processed) Memory Usage: 6.07 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.13s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.

```

The metric score can be multiplied by -1 to get the metric value.

To change this, specify the `eval_metric` parameter of `Predictor()` `use_bag_holdout=True`, will use `tuning_data` as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
        'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
        {'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
        {'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
        'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
        'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
        {'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
        {'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
        'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
        {'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

Fitting 11 L1 models ...

Fitting model: `KNeighborsUnif_BAG_L1` ... Training model for up to 1799.87s of the 1799.86s of remaining time.

Training model for location C...

```
-23.875 = Validation score (-mean_absolute_error)
0.02s   = Training runtime
0.26s   = Validation runtime
```

Fitting model: `KNeighborsDist_BAG_L1` ... Training model for up to 1799.5s of the 1799.5s of remaining time.

```
-23.8095 = Validation score (-mean_absolute_error)
0.02s    = Training runtime
0.27s    = Validation runtime
```

Fitting model: `LightGBMXT_BAG_L1` ... Training model for up to 1799.15s of the 1799.15s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with `ParallelLocalFoldFittingStrategy`

```
-12.2146 = Validation score (-mean_absolute_error)
25.99s   = Training runtime
14.98s   = Validation runtime
```

Fitting model: `LightGBM_BAG_L1` ... Training model for up to 1767.79s of the 1767.78s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with

```

ParallelLocalFoldFittingStrategy
    -13.3767          = Validation score    (-mean_absolute_error)
    24.19s           = Training   runtime
    5.12s            = Validation runtime
Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1739.72s of
the 1739.71s of remaining time.
    -16.5325          = Validation score    (-mean_absolute_error)
    4.66s             = Training   runtime
    0.75s             = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1733.73s of the
1733.72s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.3629          = Validation score    (-mean_absolute_error)
    185.33s           = Training   runtime
    0.09s             = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1547.16s of the
1547.15s of remaining time.
    -16.3493          = Validation score    (-mean_absolute_error)
    0.98s             = Training   runtime
    0.76s             = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1544.76s of
the 1544.76s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.5328          = Validation score    (-mean_absolute_error)
    30.52s            = Training   runtime
    0.44s             = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 1512.48s of the
1512.48s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.5911          = Validation score    (-mean_absolute_error)
    58.44s            = Training   runtime
    4.92s             = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1450.02s of
the 1450.02s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.7827          = Validation score    (-mean_absolute_error)
    86.64s            = Training   runtime
    0.31s             = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1362.02s of the
1362.01s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.8691          = Validation score    (-mean_absolute_error)
    85.79s            = Training   runtime

```

```

10.79s    = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1265.29s of the
1265.29s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.2878    = Validation score    (-mean_absolute_error)
    52.18s     = Training    runtime
    28.82s     = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1231.3s of the
1231.3s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.2583    = Validation score    (-mean_absolute_error)
    48.46s     = Training    runtime
    9.68s      = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1201.56s of the
1201.56s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.326     = Validation score    (-mean_absolute_error)
    371.29s    = Training    runtime
    0.17s      = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1014.3s of
the 1014.3s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.5944    = Validation score    (-mean_absolute_error)
    61.15s     = Training    runtime
    0.84s      = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 981.43s of the
981.43s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.6625    = Validation score    (-mean_absolute_error)
    101.13s    = Training    runtime
    7.07s      = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 933.99s of the
933.99s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.8264    = Validation score    (-mean_absolute_error)
    168.16s    = Training    runtime
    0.66s      = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 850.89s of the
850.89s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

-12.685 = Validation score (-mean_absolute_error)
170.66s = Training runtime
22.33s = Validation runtime
Repeating k-fold bagging: 3/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 751.89s of the
751.89s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.2931 = Validation score (-mean_absolute_error)
78.29s = Training runtime
44.19s = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 716.6s of the 716.6s
of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.2496 = Validation score (-mean_absolute_error)
71.61s = Training runtime
14.21s = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 687.35s of the
687.35s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.3702 = Validation score (-mean_absolute_error)
557.27s = Training runtime
0.25s = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 500.02s of
the 500.02s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-14.5832 = Validation score (-mean_absolute_error)
91.9s = Training runtime
1.25s = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 466.39s of the
466.39s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.7012 = Validation score (-mean_absolute_error)
142.29s = Training runtime
8.29s = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 420.08s of the
420.08s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.7879 = Validation score (-mean_absolute_error)
252.17s = Training runtime
0.96s = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 334.23s of the
334.23s of remaining time.

```

```

Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.6813      = Validation score    (-mean_absolute_error)
259.25s      = Training    runtime
35.99s       = Validation runtime
Completed 3/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
226.04s of remaining time.
-11.8138      = Validation score    (-mean_absolute_error)
0.42s        = Training    runtime
0.0s         = Validation runtime
AutoGluon training complete, total runtime = 1574.41s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_97_C/")

```

```

[14]: # save leaderboards to csv
pd.concat(leaderboards).to_csv(f"leaderboards/{new_filename}.csv")

```

3 Submit

```

[15]: import pandas as pd
import matplotlib.pyplot as plt

train_data_with_dates = TabularDataset('X_train_raw.csv')
train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

test_data = TabularDataset('X_test_raw.csv')
test_data["ds"] = pd.to_datetime(test_data["ds"])
#test_data

```

Loaded data from: X_train_raw.csv | Columns = 34 / 34 | Rows = 92945 -> 92945
Loaded data from: X_test_raw.csv | Columns = 33 / 33 | Rows = 4608 -> 4608

```

[16]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
test_data_merged = pd.merge(test_data, test_ids, how="inner", right_on=["time",
↪ "location"], left_on=["ds", "location"])

#test_data_merged

```

Loaded data from: test.csv | Columns = 4 / 4 | Rows = 2160 -> 2160

```

[17]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,

```

```

        "B": 1,
        "C": 2
    }
    for loc, group in test_data.groupby('location'):
        i = location_map[loc]
        subset = test_data_merged[test_data_merged["location"] == loc].
        ↪reset_index(drop=True)
        #print(subset)
        pred = predictors[i].predict(subset)
        subset["prediction"] = pred
        predictions.append(subset)

        # get past predictions
        past_pred = predictors[i].
        ↪predict(train_data_with_dates[train_data_with_dates["location"] == loc])
        train_data_with_dates.loc[train_data_with_dates["location"] == loc,
        ↪"prediction"] = past_pred

```

```



[18]: # plot predictions for location A, in addition to train data for A
    for loc, idx in location_map.items():
        fig, ax = plt.subplots(figsize=(20, 10))
        # plot train data
        train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
        ↪y='y', ax=ax, label="train data")


        # plot predictions
        predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

        # plot past predictions
        train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
        ↪y='prediction', ax=ax, label="past predictions")

        # title
        ax.set_title(f"Predictions for location {loc}")

```





```
[19]: # clip predictions smaller than 0 to 0
for pred in predictions:
    # print smallest prediction
    print("Smallest prediction:", pred["prediction"].min())
    pred.loc[pred["prediction"] < 0, "prediction"] = 0
    print("Smallest prediction after clipping:", pred["prediction"].min())

# concatenate predictions
submissions_df = pd.concat(predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df
```

```
[19]:      id  prediction
0      0   -1.060897
1      1   -0.598755
2      2   -1.056423
3      3    59.710995
4      4   311.332031
..    ...         ...
715   2155   68.854149
716   2156   39.890915
717   2157    9.410706
718   2158    4.226479
719   2159    1.412523
```

```
[2160 rows x 2 columns]
```

```
[20]: # Save the submission DataFrame to submissions folder, create new name based on
      ↪ last submission, format is submission_<last_submission_number + 1>.csv

      # Save the submission
      print(f"Saving submission to submissions/{new_filename}.csv")
      submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
      ↪ index=False)
      print("jall1a")
```

Saving submission to submissions/submission_97.csv
jall1a

```
[21]: # save this running notebook
      from IPython.display import display, Javascript
      import time

      # hei123

      display(Javascript("IPython.notebook.save_checkpoint();"))

      time.sleep(3)
```

<IPython.core.display.Javascript object>

```
[22]: # save this notebook to submissions folder
      import subprocess
      import os
      subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
      ↪ join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
      ↪ ipynb"])
```

[NbConvertApp] Converting notebook autogluon_each_location.ipynb to pdf
/opt/conda/lib/python3.10/site-packages/nbconvert/utils/pandoc.py:51:
RuntimeWarning: You are using an unsupported version of pandoc (2.9.2.1).
Your version must be at least (2.14.2) but less than (4.0.0).
Refer to <https://pandoc.org/installing.html>.
Continuing with doubts...
check_pandoc_version()
[NbConvertApp] Writing 135589 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 103636 bytes to notebook_pdfs/submission_97.pdf

```
[22]: CompletedProcess(args=['jupyter', 'nbconvert', '--to', 'pdf', '--output',
'notebook_pdfs/submission_97.pdf', 'autogluon_each_location.ipynb'],
returncode=0)
```

```
[23]: # feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
↳time_limit=60*10)
```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'location', 'prediction']

Computing feature importance via permutation shuffling for 30 features using 4392 rows with 10 shuffle sets... Time limit: 600s...

5441.74s = Expected runtime (544.17s per shuffle set)

322.9s = Actual runtime (Completed 1 of 10 shuffle sets) (Early stopping due to lack of time...)

```
[23]:
```

	importance	stddev	p_value	n	p99_high \
direct_rad_1h:J	1.821785e+02	NaN	NaN	1	NaN
clear_sky_energy_1h:J	9.408880e+01	NaN	NaN	1	NaN
clear_sky_rad:W	8.960991e+01	NaN	NaN	1	NaN
diffuse_rad_1h:J	8.077760e+01	NaN	NaN	1	NaN
direct_rad:W	6.765221e+01	NaN	NaN	1	NaN
diffuse_rad:W	6.723790e+01	NaN	NaN	1	NaN
sun_azimuth:d	5.189177e+01	NaN	NaN	1	NaN
sun_elevation:d	3.894327e+01	NaN	NaN	1	NaN
effective_cloud_cover:p	2.830885e+01	NaN	NaN	1	NaN
total_cloud_cover:p	1.854761e+01	NaN	NaN	1	NaN
is_in_shadow:idx	1.564697e+01	NaN	NaN	1	NaN
cloud_base_agl:m	1.305668e+01	NaN	NaN	1	NaN
t_1000hPa:K	1.262343e+01	NaN	NaN	1	NaN
ceiling_height_agl:m	1.259696e+01	NaN	NaN	1	NaN
relative_humidity_1000hPa:p	1.215808e+01	NaN	NaN	1	NaN
snow_water:kgm2	1.171797e+01	NaN	NaN	1	NaN
visibility:m	1.116307e+01	NaN	NaN	1	NaN
wind_speed_10m:ms	1.031712e+01	NaN	NaN	1	NaN
pressure_100m:hPa	8.429827e+00	NaN	NaN	1	NaN
sfc_pressure:hPa	7.851625e+00	NaN	NaN	1	NaN
msl_pressure:hPa	7.671138e+00	NaN	NaN	1	NaN
fresh_snow_6h:cm	6.886095e+00	NaN	NaN	1	NaN
is_day:idx	6.714718e+00	NaN	NaN	1	NaN
pressure_50m:hPa	6.671387e+00	NaN	NaN	1	NaN
precip_type_5min:idx	5.456353e+00	NaN	NaN	1	NaN
super_cooled_liquid_water:kgm2	4.694124e+00	NaN	NaN	1	NaN
fresh_snow_3h:cm	3.501178e+00	NaN	NaN	1	NaN

snow_depth:cm	2.850314e+00	NaN	NaN	1	NaN
fresh_snow_1h:cm	2.693312e+00	NaN	NaN	1	NaN
is_estimated	-1.552544e-08	NaN	NaN	1	NaN

	p99_low
direct_rad_1h:J	NaN
clear_sky_energy_1h:J	NaN
clear_sky_rad:W	NaN
diffuse_rad_1h:J	NaN
direct_rad:W	NaN
diffuse_rad:W	NaN
sun_azimuth:d	NaN
sun_elevation:d	NaN
effective_cloud_cover:p	NaN
total_cloud_cover:p	NaN
is_in_shadow:idx	NaN
cloud_base_agl:m	NaN
t_1000hPa:K	NaN
ceiling_height_agl:m	NaN
relative_humidity_1000hPa:p	NaN
snow_water:kgm2	NaN
visibility:m	NaN
wind_speed_10m:ms	NaN
pressure_100m:hPa	NaN
sfc_pressure:hPa	NaN
msl_pressure:hPa	NaN
fresh_snow_6h:cm	NaN
is_day:idx	NaN
pressure_50m:hPa	NaN
precip_type_5min:idx	NaN
super_cooled_liquid_water:kgm2	NaN
fresh_snow_3h:cm	NaN
snow_depth:cm	NaN
fresh_snow_1h:cm	NaN
is_estimated	NaN

```
[24]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
↳time_limit=60*10)
```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'location', 'prediction']

Computing feature importance via permutation shuffling for 30 features using 5000 rows with 10 shuffle sets... Time limit: 600s...

6395.23s = Expected runtime (639.52s per shuffle set)

KeyboardInterrupt

```
[ ]: display(Javascript("IPython.notebook.save_checkpoint();"))
time.sleep(3)

# subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↳ join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↳ "autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
    ↳ stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8')}.
    ↳ strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
    ↳ 'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', 'hello if hello is
    ↳ not None else 'Henrik eller Jørgen'])

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push',
    ↳ 'origin', branch_name])
```

```
# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream', ↵
↵ 'origin', branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])
# execute_git_command(git_repo_path, ['checkout', 'main'])
```