

autogluon_each_location

October 19, 2023

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*30
presets = 'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = [] # ["hour", "year"]
use_estimated_diff_attr = False
use_is_estimated_attr = True

use_groups = False
n_groups = 8

auto_stack = False
num_stack_levels = 0
num_bag_folds = 8
num_bag_sets = 20

use_tune_data = True
use_test_data = True
tune_and_test_length = 0.5 # 3 months from end
holdout_frac = None
use_bag_holdout = True # Enable this if there is a large gap between score_val_
    ↪ and score_test in stack models.

sample_weight = None # 'sample_weight' # None
weight_evaluation = False #
sample_weight_estimated = 1
sample_weight_may_july = 1

run_analysis = True

shift_predictions_by_average_of_negatives_then_clip = False
```

```
clip_predictions = True
shift_predictions = False
```

```
[2]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift all columns with "1h" in them by 1 hour, so that for index 16:00,
    ↳ we have the values from 17:00
    # but only for the columns with "1h" in the name
    # X_shifted = X.filter(regex="\dh").shift(-1, axis=1)
    # print(f"Number of columns with 1h in name: {X_shifted.columns}")

    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

    X_shifted = X[X.index.minute==0][columns].copy()
    # loop through all rows and check if index + 1 hour is in the index, if so
    ↳ get that value, else nan
    count1 = 0
    count2 = 0
    for i in range(len(X_shifted)):
        if X_shifted.index[i] + pd.Timedelta('1 hour') in X.index:
            count1 += 1
            X_shifted.iloc[i] = X.loc[X_shifted.index[i] + pd.Timedelta('1
    ↳ hour')][columns]
        else:
            count2 += 1
            X_shifted.iloc[i] = np.nan

    print("COUNT1", count1)
    print("COUNT2", count2)

    X_old_unshifted = X[X.index.minute==0][columns]
    # rename X_old_unshifted columns to have _not_shifted at the end
    X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
    ↳ columns]

    # put the shifted columns back into the original dataframe
```

```

X[columns] = X_shifted[columns]

date_calc = None
if "date_calc" in X.columns:
    date_calc = X[X.index.minute == 0]['date_calc']

# resample to hourly
print("index: ", X.index[0])
X = X.resample('H').mean()
print("index AFTER: ", X.index[0])

X[columns] = X_shifted[columns]
#X[X_old_unshifted.columns] = X_old_unshifted

if date_calc is not None:
    X['date_calc'] = date_calc

return X

def fix_X(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    ↪ with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    X = feature_engineering(X)

    return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    if weight_evaluation:
        # add sample weights, which are 1 for observed and 3 for estimated

```

```

X_train_observed["sample_weight"] = 1
X_train_estimated["sample_weight"] = sample_weight_estimated
X_test["sample_weight"] = sample_weight_estimated

y_train['ds'] = pd.to_datetime(y_train['time'])
y_train.drop(columns=['time'], inplace=True)
y_train.sort_values(by='ds', inplace=True)
y_train.set_index('ds', inplace=True)

return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600
        X_test["estimated_diff_hours"] = (X_test.index - pd.
to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

        X_train_estimated["estimated_diff_hours"] =
X_train_estimated["estimated_diff_hours"].astype('int64')
        # the filled once will get dropped later anyways, when we drop y nans
        X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
fillna(-50).astype('int64')

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0
        X_train_estimated["is_estimated"] = 1
        X_test["is_estimated"] = 1

    # drop date_calc
    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)

```

```

X_train = pd.concat([X_train_observed, X_train_estimated])

# clip all y values to 0 if negative
y_train["y"] = y_train["y"].clip(lower=0)

X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

# print number of nans in y
print(f"Number of nans in y: {X_train['y'].isna().sum()}")

X_train["location"] = location
X_test["location"] = location

    return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↳X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

Processing location A...
COUNT1 29667
COUNT2 1
index: 2019-06-02 22:00:00
index AFTER: 2019-06-02 22:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 0
Processing location B...
COUNT1 29232
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 4
Processing location C...
COUNT1 29206
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 6059

1 Feature engineering

```
[3]: import numpy as np
import pandas as pd

X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
               inplace=True)

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

#print(X_train.head())

# If the "sample_weight" column is present and weight_evaluation is True,
# multiply sample_weight with sample_weight_may_july if the ds is between
# 05-01 00:00:00 and 07-03 23:00:00, else add sample_weight as a column to
# X_train
if weight_evaluation:
    if "sample_weight" not in X_train.columns:
        X_train["sample_weight"] = 1

    X_train.loc[((X_train.index.month >= 5) & (X_train.index.month <= 6)) |
               ((X_train.index.month == 7) & (X_train.index.day <= 3)), "sample_weight"] *=
    sample_weight_may_july

print(X_train.iloc[200])
print(X_train[((X_train.index.month >= 5) & (X_train.index.month <= 6)) |
              ((X_train.index.month == 7) & (X_train.index.day <= 3))].head(1))

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    # of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()
```

```

# Calculate the size of each group for this location
group_size = len(loc_df) // n_groups

# Create a new 'group' column for this location
loc_df['group'] = np.repeat(range(n_groups),
↪repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

# Append to list of grouped DataFrames
grouped_dfs.append(loc_df)

# Concatenate all the grouped DataFrames back together
X_train = pd.concat(grouped_dfs)
X_train.sort_index(inplace=True)
print(X_train["group"].head())

to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
↪"dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm",
↪"wind_speed_u_10m:ms", "wind_speed_v_10m:ms", "snow_melt_10min:mm",
↪"rain_water:kgm2", "dew_point_2m:K", "precip_5min:mm", "absolute_humidity_2m:
↪gm3", "air_density_2m:kgm3", "msl_pressure:hPa", "pressure_100m:hPa",
↪"pressure_50m:hPa", "clear_sky_rad:W"]

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

absolute_humidity_2m:gm3	7.825
air_density_2m:kgm3	1.245
ceiling_height_agl:m	2085.774902
clear_sky_energy_1h:J	1685498.875
clear_sky_rad:W	452.100006
cloud_base_agl:m	2085.774902
dew_or_rime:idx	0.0
dew_point_2m:K	280.549988
diffuse_rad:W	140.800003
diffuse_rad_1h:J	538581.625
direct_rad:W	102.599998
direct_rad_1h:J	439453.8125
effective_cloud_cover:p	71.849998
elevation:m	6.0

fresh_snow_12h:cm	0.0
fresh_snow_1h:cm	0.0
fresh_snow_24h:cm	0.0
fresh_snow_3h:cm	0.0
fresh_snow_6h:cm	0.0
is_day:idx	1.0
is_in_shadow:idx	0.0
msl_pressure:hPa	1026.349976
precip_5min:mm	0.0
precip_type_5min:idx	0.0
pressure_100m:hPa	1013.325012
pressure_50m:hPa	1019.450012
prob_rime:p	0.0
rain_water:kgm2	0.0
relative_humidity_1000hPa:p	77.099998
sfc_pressure:hPa	1025.550049
snow_density:kgm3	NaN
snow_depth:cm	0.0
snow_drift:idx	0.0
snow_melt_10min:mm	0.0
snow_water:kgm2	0.0
sun_azimuth:d	93.415253
sun_elevation:d	27.633499
super_cooled_liquid_water:kgm2	0.025
t_1000hPa:K	282.625
total_cloud_cover:p	71.849998
visibility:m	44177.875
wind_speed_10m:ms	2.675
wind_speed_u_10m:ms	-2.3
wind_speed_v_10m:ms	-1.4
wind_speed_w_1000hPa:ms	0.0
is_estimated	0
y	2991.12
location	A

Name: 2019-06-11 06:00:00, dtype: object

	absolute_humidity_2m:gm3	air_density_2m:kgm3	\
ds			
2019-06-02 22:00:00	7.7	1.22825	

	ceiling_height_agl:m	clear_sky_energy_1h:J	\
ds			
2019-06-02 22:00:00	1728.949951	0.0	

	clear_sky_rad:W	cloud_base_agl:m	dew_or_rime:idx	\
ds				
2019-06-02 22:00:00	0.0	1728.949951	0.0	

	dew_point_2m:K	diffuse_rad:W	diffuse_rad_1h:J	...	\
ds					

```

ds
2019-06-02 22:00:00      280.299988      0.0      0.0 ...

      t_1000hPa:K  total_cloud_cover:p  visibility:m  \
ds
2019-06-02 22:00:00      286.225006      100.0  40386.476562

      wind_speed_10m:ms  wind_speed_u_10m:ms  \
ds
2019-06-02 22:00:00      3.6      -3.575

      wind_speed_v_10m:ms  wind_speed_w_1000hPa:ms  \
ds
2019-06-02 22:00:00      -0.5      0.0

      is_estimated    y  location
ds
2019-06-02 22:00:00      0  0.0      A

[1 rows x 48 columns]

```

```

[4]: # Create a plot of X_train showing its "y" and color it based on the value of
      ↪ the sample_weight column.
      #import matplotlib.pyplot as plt
      #import seaborn as sns
      #sns.scatterplot(data=X_train, x=X_train.index, y="y", hue="sample_weight",
      ↪ palette="deep", size=3)
      #plt.show()

```

```

[5]: def normalize_sample_weights_per_location(df):
      for loc in locations:
          loc_df = df[df["location"] == loc]
          loc_df["sample_weight"] = loc_df["sample_weight"] /
          ↪ loc_df["sample_weight"].sum() * loc_df.shape[0]
          df[df["location"] == loc] = loc_df
      return df

import pandas as pd
import numpy as np

def split_and_shuffle_data(input_data, num_bins, frac1):
    """
    Splits the input_data into num_bins and shuffles them, then divides the
    ↪ bins into two datasets based on the given fraction for the first set.

    Args:
        input_data (pd.DataFrame): The data to be split and shuffled.

```

```

    num_bins (int): The number of bins to split the data into.
    frac1 (float): The fraction of each bin to go into the first output_
↳ dataset.

Returns:
    pd.DataFrame, pd.DataFrame: The two output datasets.
"""
# Validate the input fraction
if frac1 < 0 or frac1 > 1:
    raise ValueError("frac1 must be between 0 and 1.")

if frac1==1:
    return input_data, pd.DataFrame()

# Calculate the fraction for the second output set
frac2 = 1 - frac1

# Calculate bin size
bin_size = len(input_data) // num_bins

# Initialize empty DataFrames for output
output_data1 = pd.DataFrame()
output_data2 = pd.DataFrame()

for i in range(num_bins):
    # Shuffle the data in the current bin
    np.random.seed(i)
    current_bin = input_data.iloc[i * bin_size: (i + 1) * bin_size].
↳ sample(frac=1)

    # Calculate the sizes for each output set
    size1 = int(len(current_bin) * frac1)

    # Split and append to output DataFrames
    output_data1 = pd.concat([output_data1, current_bin.iloc[:size1]])
    output_data2 = pd.concat([output_data2, current_bin.iloc[size1:]])

# Shuffle and split the remaining data
remaining_data = input_data.iloc[num_bins * bin_size:].sample(frac=1)
remaining_size1 = int(len(remaining_data) * frac1)

    output_data1 = pd.concat([output_data1, remaining_data.iloc[:
↳ remaining_size1]])
    output_data2 = pd.concat([output_data2, remaining_data.iloc[remaining_size1:
↳ ]])

return output_data1, output_data2

```

```

[6]: from autogluon.tabular import TabularDataset, TabularPredictor
from autogluon.timeseries import TimeSeriesDataFrame
import numpy as np
data = TabularDataset('X_train_raw.csv')
# set group column of train_data be increasing from 0 to 7 based on time, the
  ↳ first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
data['ds'] = pd.to_datetime(data['ds'])
data = data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
# split_time = pd.to_datetime(train_data["ds"]).max() - pd.
  ↳ Timedelta(hours=tune_and_test_length)
# 2022-10-28 22:00:00
split_time = pd.to_datetime("2022-10-28 22:00:00")
train_set = TabularDataset(data[data["ds"] < split_time])
test_set = TabularDataset(data[data["ds"] >= split_time])

# shuffle test_set and only grab tune_and_test_length percent of it, rest goes
  ↳ to train_set
test_set, new_train_set = split_and_shuffle_data(test_set, 40,
  ↳ tune_and_test_length)

print("Length of train set before adding test set", len(train_set))
# add rest to train_set
train_set = pd.concat([train_set, new_train_set])
print("Length of train set after adding test set", len(train_set))
print("Length of test set", len(test_set))

if use_groups:
    test_set = test_set.drop(columns=['group'])

tuning_data = None
if use_tune_data:
    if use_test_data:
        # split test_set in half, use first half for tuning
        tuning_data, test_data = [], []

```

```

    for loc in locations:
        loc_test_set = test_set[test_set["location"] == loc]
        # randomly shuffle the loc_test_set
        loc_tuning_data, loc_test_data = \
↪split_and_shuffle_data(loc_test_set, 40, 0.5)
        tuning_data.append(loc_tuning_data)
        test_data.append(loc_test_data)
        tuning_data = pd.concat(tuning_data)
        test_data = pd.concat(test_data)
        print("Shapes of tuning and test", tuning_data.shape[0], test_data.
↪shape[0], tuning_data.shape[0] + test_data.shape[0])

    else:
        tuning_data = test_set
        print("Shape of tuning", tuning_data.shape[0])

        # ensure sample weights for your tuning data sum to the number of rows in
↪the tuning data.
        if weight_evaluation:
            tuning_data = normalize_sample_weights_per_location(tuning_data)

else:
    if use_test_data:
        test_data = test_set
        print("Shape of test", test_data.shape[0])

train_data = train_set

# ensure sample weights for your training (or tuning) data sum to the number of
↪rows in the training (or tuning) data.
if weight_evaluation:
    train_data = normalize_sample_weights_per_location(train_data)
    if use_test_data:
        test_data = normalize_sample_weights_per_location(test_data)

train_data = TabularDataset(train_data)
if use_tune_data:
    tuning_data = TabularDataset(tuning_data)
if use_test_data:
    test_data = TabularDataset(test_data)

```

Length of train set before adding test set 82026

Length of train set after adding test set 87486

Length of test set 5459

Shapes of tuning and test 2728 2731 5459

```
[7]: if run_analysis:
      import autogluon.eda.auto as auto
      auto.dataset_overview(train_data=train_data, test_data=test_data,
                             label="y", sample=None)
```

train_data dataset summary

	count	unique	top	freq	\
ceiling_height_agl:m	72280	59980			
clear_sky_energy_1h:J	87486	46359			
cloud_base_agl:m	81454	61360			
diffuse_rad:W	87486	11092			
diffuse_rad_1h:J	87486	46319			
direct_rad:W	87486	14181			
direct_rad_1h:J	87486	40118			
ds	87486	36794	2021-02-05 14:00:00	3	
effective_cloud_cover:p	87486	5655			
elevation:m	87486	3			
fresh_snow_1h:cm	87486	39			
fresh_snow_3h:cm	87486	70			
fresh_snow_6h:cm	87486	96			
is_day:idx	87486	5			
is_estimated	87486	2			
is_in_shadow:idx	87486	5			
location	87486	3	A	31872	
precip_type_5min:idx	87486	15			
relative_humidity_1000hPa:p	87486	3799			
sfc_pressure:hPa	87486	3795			
snow_depth:cm	87486	487			
snow_water:kgm2	87486	161			
sun_azimuth:d	87486	83179			
sun_elevation:d	87486	72262			
super_cooled_liquid_water:kgm2	87486	53			
t_1000hPa:K	87486	1989			
total_cloud_cover:p	87486	5556			
visibility:m	87486	85949			
wind_speed_10m:ms	87486	596			
y	87486	11321			

	first	last	mean	\
ceiling_height_agl:m	NaT	NaT	2861.929806	
clear_sky_energy_1h:J	NaT	NaT	530297.395771	
cloud_base_agl:m	NaT	NaT	1740.241802	
diffuse_rad:W	NaT	NaT	40.267497	
diffuse_rad_1h:J	NaT	NaT	145328.6257	
direct_rad:W	NaT	NaT	51.524847	
direct_rad_1h:J	NaT	NaT	185338.05854	

ds	2019-01-01	2023-04-30	22:00:00	
effective_cloud_cover:p	NaT	NaT	67.052836	
elevation:m	NaT	NaT	11.414718	
fresh_snow_1h:cm	NaT	NaT	0.008783	
fresh_snow_3h:cm	NaT	NaT	0.026713	
fresh_snow_6h:cm	NaT	NaT	0.05322	
is_day:idx	NaT	NaT	0.490147	
is_estimated	NaT	NaT	0.06241	
is_in_shadow:idx	NaT	NaT	0.556952	
location	NaT	NaT		
precip_type_5min:idx	NaT	NaT	0.084976	
relative_humidity_1000hPa:p	NaT	NaT	73.779918	
sfc_pressure:hPa	NaT	NaT	1008.035963	
snow_depth:cm	NaT	NaT	0.197574	
snow_water:kgm2	NaT	NaT	0.090839	
sun_azimuth:d	NaT	NaT	179.584247	
sun_elevation:d	NaT	NaT	-0.705998	
super_cooled_liquid_water:kgm2	NaT	NaT	0.058256	
t_1000hPa:K	NaT	NaT	279.675551	
total_cloud_cover:p	NaT	NaT	73.72398	
visibility:m	NaT	NaT	32944.238197	
wind_speed_10m:ms	NaT	NaT	3.032943	
y	NaT	NaT	294.447861	

	std	min	25%	50%	\
ceiling_height_agl:m	2532.377528	27.8	1082.3125	1856.075	
clear_sky_energy_1h:J	831839.646633	0.0	0.0	9084.9	
cloud_base_agl:m	1808.519208	27.5	598.15625	1174.775	
diffuse_rad:W	61.119566	0.0	0.0	1.35	
diffuse_rad_1h:J	218036.296903	0.0	0.0	12531.2	
direct_rad:W	114.236728	0.0	0.0	0.0	
direct_rad_1h:J	406379.39471	0.0	0.0	0.0	
ds					
effective_cloud_cover:p	34.132847	0.0	42.125	79.7	
elevation:m	7.881545	6.0	6.0	7.0	
fresh_snow_1h:cm	0.110515	0.0	0.0	0.0	
fresh_snow_3h:cm	0.277575	0.0	0.0	0.0	
fresh_snow_6h:cm	0.474579	0.0	0.0	0.0	
is_day:idx	0.486133	0.0	0.0	0.5	
is_estimated	0.2419	0.0	0.0	0.0	
is_in_shadow:idx	0.484138	0.0	0.0	1.0	
location					
precip_type_5min:idx	0.32995	0.0	0.0	0.0	
relative_humidity_1000hPa:p	14.200631	19.575	64.4	76.15	
sfc_pressure:hPa	13.038723	941.55	1000.05	1009.0	
snow_depth:cm	1.28439	0.0	0.0	0.0	
snow_water:kgm2	0.240248	0.0	0.0	0.0	
sun_azimuth:d	97.419022	6.983	94.4135	180.00663	

sun_elevation:d	24.006117	-49.932	-17.969875	-0.453875
super_cooled_liquid_water:kgm2	0.106959	0.0	0.0	0.0
t_1000hPa:K	6.551665	258.025	275.1	278.975
total_cloud_cover:p	33.851299	0.0	53.475	92.825
visibility:m	17949.64428	132.375	16564.5125	36910.75
wind_speed_10m:ms	1.758713	0.025	1.675	2.7
y	774.531815	-0.0	0.0	0.0

	75%	max	dtypes \
ceiling_height_agl:m	3916.78125	12285.775	float64
clear_sky_energy_1h:J	831169.675	3006697.2	float64
cloud_base_agl:m	2080.39375	11673.725	float64
diffuse_rad:W	67.15	334.75	float64
diffuse_rad_1h:J	243365.275	1182265.4	float64
direct_rad:W	32.225	683.4	float64
direct_rad_1h:J	122454.125	2445897.0	float64
ds			datetime64[ns]
effective_cloud_cover:p	98.5	100.0	float64
elevation:m	24.0	24.0	float64
fresh_snow_1h:cm	0.0	7.1	float64
fresh_snow_3h:cm	0.0	20.6	float64
fresh_snow_6h:cm	0.0	34.0	float64
is_day:idx	1.0	1.0	float64
is_estimated	0.0	1.0	int64
is_in_shadow:idx	1.0	1.0	float64
location			object
precip_type_5min:idx	0.0	5.0	float64
relative_humidity_1000hPa:p	85.175	100.0	float64
sfc_pressure:hPa	1017.1	1043.725	float64
snow_depth:cm	0.0	18.2	float64
snow_water:kgm2	0.1	5.65	float64
sun_azimuth:d	264.601138	348.48752	float64
sun_elevation:d	16.004499	49.94375	float64
super_cooled_liquid_water:kgm2	0.1	1.375	float64
t_1000hPa:K	284.225	303.25	float64
total_cloud_cover:p	99.9	100.0	float64
visibility:m	48289.05	75326.58	float64
wind_speed_10m:ms	4.05	13.275	float64
y	183.7125	5733.42	float64

	missing_count	missing_ratio	raw_type \
ceiling_height_agl:m	15206	0.173811	float
clear_sky_energy_1h:J			float
cloud_base_agl:m	6032	0.068948	float
diffuse_rad:W			float
diffuse_rad_1h:J			float
direct_rad:W			float
direct_rad_1h:J			float

ds	datetime
effective_cloud_cover:p	float
elevation:m	float
fresh_snow_1h:cm	float
fresh_snow_3h:cm	float
fresh_snow_6h:cm	float
is_day:idx	float
is_estimated	int
is_in_shadow:idx	float
location	object
precip_type_5min:idx	float
relative_humidity_1000hPa:p	float
sfc_pressure:hPa	float
snow_depth:cm	float
snow_water:kgm2	float
sun_azimuth:d	float
sun_elevation:d	float
super_cooled_liquid_water:kgm2	float
t_1000hPa:K	float
total_cloud_cover:p	float
visibility:m	float
wind_speed_10m:ms	float
y	float

	variable_type	special_types
ceiling_height_agl:m	numeric	
clear_sky_energy_1h:J	numeric	
cloud_base_agl:m	numeric	
diffuse_rad:W	numeric	
diffuse_rad_1h:J	numeric	
direct_rad:W	numeric	
direct_rad_1h:J	numeric	
ds		
effective_cloud_cover:p	numeric	
elevation:m	category	
fresh_snow_1h:cm	numeric	
fresh_snow_3h:cm	numeric	
fresh_snow_6h:cm	numeric	
is_day:idx	category	
is_estimated	category	
is_in_shadow:idx	category	
location	category	
precip_type_5min:idx	category	
relative_humidity_1000hPa:p	numeric	
sfc_pressure:hPa	numeric	
snow_depth:cm	numeric	
snow_water:kgm2	numeric	
sun_azimuth:d	numeric	

sun_elevation:d	numeric
super_cooled_liquid_water:kgm2	numeric
t_1000hPa:K	numeric
total_cloud_cover:p	numeric
visibility:m	numeric
wind_speed_10m:ms	numeric
y	numeric

test_data dataset summary

	count	unique	top	freq	\
ceiling_height_agl:m	2100	2065			
clear_sky_energy_1h:J	2731	1138			
cloud_base_agl:m	2374	2332			
diffuse_rad:W	2731	983			
diffuse_rad_1h:J	2731	1138			
direct_rad:W	2731	750			
direct_rad_1h:J	2731	932			
ds	2731	2200	2023-04-10 19:00:00	3	
effective_cloud_cover:p	2731	1348			
elevation:m	2731	3			
fresh_snow_1h:cm	2731	19			
fresh_snow_3h:cm	2731	36			
fresh_snow_6h:cm	2731	50			
is_day:idx	2731	5			
is_estimated	2731	1			
is_in_shadow:idx	2731	5			
location	2731	3	A	1094	
precip_type_5min:idx	2731	10			
relative_humidity_1000hPa:p	2731	1523			
sfc_pressure:hPa	2731	1575			
snow_depth:cm	2731	61			
snow_water:kgm2	2731	61			
sun_azimuth:d	2731	2716			
sun_elevation:d	2731	2652			
super_cooled_liquid_water:kgm2	2731	29			
t_1000hPa:K	2731	774			
total_cloud_cover:p	2731	1126			
visibility:m	2731	2729			
wind_speed_10m:ms	2731	366			
y	2731	895			

	first	last	\
ceiling_height_agl:m	NaT	NaT	
clear_sky_energy_1h:J	NaT	NaT	
cloud_base_agl:m	NaT	NaT	
diffuse_rad:W	NaT	NaT	
diffuse_rad_1h:J	NaT	NaT	
direct_rad:W	NaT	NaT	

direct_rad_1h:J	NaT	NaT
ds	2022-10-28 22:00:00	2023-04-30 19:00:00
effective_cloud_cover:p	NaT	NaT
elevation:m	NaT	NaT
fresh_snow_1h:cm	NaT	NaT
fresh_snow_3h:cm	NaT	NaT
fresh_snow_6h:cm	NaT	NaT
is_day:idx	NaT	NaT
is_estimated	NaT	NaT
is_in_shadow:idx	NaT	NaT
location	NaT	NaT
precip_type_5min:idx	NaT	NaT
relative_humidity_1000hPa:p	NaT	NaT
sfc_pressure:hPa	NaT	NaT
snow_depth:cm	NaT	NaT
snow_water:kgm2	NaT	NaT
sun_azimuth:d	NaT	NaT
sun_elevation:d	NaT	NaT
super_cooled_liquid_water:kgm2	NaT	NaT
t_1000hPa:K	NaT	NaT
total_cloud_cover:p	NaT	NaT
visibility:m	NaT	NaT
wind_speed_10m:ms	NaT	NaT
y	NaT	NaT

	mean	std	min \
ceiling_height_agl:m	3361.682133	2562.862274	28.0
clear_sky_energy_1h:J	285528.142658	573252.521662	0.0
cloud_base_agl:m	1685.111501	1833.56975	27.5
diffuse_rad:W	26.230813	48.360421	0.0
diffuse_rad_1h:J	94649.301538	172723.786046	0.0
direct_rad:W	32.798068	92.244541	0.0
direct_rad_1h:J	118054.008202	329601.815692	0.0
ds			
effective_cloud_cover:p	66.798654	36.717964	0.0
elevation:m	11.193336	7.806119	6.0
fresh_snow_1h:cm	0.023984	0.147555	0.0
fresh_snow_3h:cm	0.069498	0.352141	0.0
fresh_snow_6h:cm	0.13885	0.57722	0.0
is_day:idx	0.378341	0.472171	0.0
is_estimated	1.0	0.0	1.0
is_in_shadow:idx	0.677133	0.452911	0.0
location			
precip_type_5min:idx	0.076254	0.344931	0.0
relative_humidity_1000hPa:p	71.631472	14.652551	21.7
sfc_pressure:hPa	1009.397775	14.318453	971.15
snow_depth:cm	0.119965	0.56196	0.0
snow_water:kgm2	0.080511	0.19473	0.0


sun_azimuth:d	180.975998	94.222121	14.914	
sun_elevation:d	-8.945472	22.095926	-49.887	
super_cooled_liquid_water:kgm2	0.035088	0.082444	0.0	
t_1000hPa:K	275.52987	4.271781	259.975	
total_cloud_cover:p	72.32056	37.085445	0.0	
visibility:m	34504.948017	17242.154257	270.3	
wind_speed_10m:ms	3.109676	1.782531	0.125	
y	179.379421	641.546947	0.0	
	25%	50%	75%	max \
ceiling_height_agl:m	1245.55625	2784.275	4919.18125	12294.9
clear_sky_energy_1h:J	0.0	0.0	220909.2	2551917.2
cloud_base_agl:m	516.7625	1000.8	2066.9875	10813.7
diffuse_rad:W	0.0	0.0	32.0	280.5
diffuse_rad_1h:J	0.0	0.0	116208.0	986147.0
direct_rad:W	0.0	0.0	4.7625	511.7
direct_rad_1h:J	0.0	0.0	24326.65	1844204.9
ds				
effective_cloud_cover:p	36.3125	83.05	99.825	100.0
elevation:m	6.0	7.0	24.0	24.0
fresh_snow_1h:cm	0.0	0.0	0.0	2.3
fresh_snow_3h:cm	0.0	0.0	0.0	4.8
fresh_snow_6h:cm	0.0	0.0	0.0	6.3
is_day:idx	0.0	0.0	1.0	1.0
is_estimated	1.0	1.0	1.0	1.0
is_in_shadow:idx	0.0	1.0	1.0	1.0
location				
precip_type_5min:idx	0.0	0.0	0.0	3.0
relative_humidity_1000hPa:p	61.775	73.55	82.9625	99.775
sfc_pressure:hPa	999.1	1009.5	1020.275	1040.6
snow_depth:cm	0.0	0.0	0.0	4.9
snow_water:kgm2	0.0	0.0	0.1	2.15
sun_azimuth:d	101.047372	179.89075	260.65412	347.81226
sun_elevation:d	-26.72725	-8.178	6.393625	41.09175
super_cooled_liquid_water:kgm2	0.0	0.0	0.0	0.75
t_1000hPa:K	272.6625	275.45	278.5	285.825
total_cloud_cover:p	44.5875	96.775	100.0	100.0
visibility:m	20902.55	36141.85	48867.949	73937.67
wind_speed_10m:ms	1.6	2.8	4.2375	9.9
y	0.0	0.0	40.397706	5043.72
	dtypes	missing_count	missing_ratio	\
ceiling_height_agl:m	float64	631	0.231051	
clear_sky_energy_1h:J	float64			
cloud_base_agl:m	float64	357	0.130721	
diffuse_rad:W	float64			
diffuse_rad_1h:J	float64			
direct_rad:W	float64			

direct_rad_1h:J	float64
ds	datetime64[ns]
effective_cloud_cover:p	float64
elevation:m	float64
fresh_snow_1h:cm	float64
fresh_snow_3h:cm	float64
fresh_snow_6h:cm	float64
is_day:idx	float64
is_estimated	int64
is_in_shadow:idx	float64
location	object
precip_type_5min:idx	float64
relative_humidity_1000hPa:p	float64
sfc_pressure:hPa	float64
snow_depth:cm	float64
snow_water:kgm2	float64
sun_azimuth:d	float64
sun_elevation:d	float64
super_cooled_liquid_water:kgm2	float64
t_1000hPa:K	float64
total_cloud_cover:p	float64
visibility:m	float64
wind_speed_10m:ms	float64
y	float64

	raw_type	variable_type	special_types
ceiling_height_agl:m	float	numeric	
clear_sky_energy_1h:J	float	numeric	
cloud_base_agl:m	float	numeric	
diffuse_rad:W	float	numeric	
diffuse_rad_1h:J	float	numeric	
direct_rad:W	float	numeric	
direct_rad_1h:J	float	numeric	
ds	datetime		
effective_cloud_cover:p	float	numeric	
elevation:m	float	category	
fresh_snow_1h:cm	float	category	
fresh_snow_3h:cm	float	numeric	
fresh_snow_6h:cm	float	numeric	
is_day:idx	float	category	
is_estimated	int	category	
is_in_shadow:idx	float	category	
location	object	category	
precip_type_5min:idx	float	category	
relative_humidity_1000hPa:p	float	numeric	
sfc_pressure:hPa	float	numeric	
snow_depth:cm	float	numeric	
snow_water:kgm2	float	numeric	

sun_azimuth:d	float	numeric
sun_elevation:d	float	numeric
super_cooled_liquid_water:kgm2	float	numeric
t_1000hPa:K	float	numeric
total_cloud_cover:p	float	numeric
visibility:m	float	numeric
wind_speed_10m:ms	float	numeric
y	float	numeric

1.0.1 Feature Distance




The following feature groups are considered as near-duplicates:

Distance threshold: ≤ 0.01 . Consider keeping only some of the columns within each group:

- elevation:m, location - distance 0.00
- diffuse_rad:W, diffuse_rad_1h:J - distance 0.00

Feature interaction between diffuse_rad:W/diffuse_rad_1h:J




```
[8]: if run_analysis:
      auto.target_analysis(train_data=train_data, label="y", sample=None)
```

1.1 Target variable analysis

	count	mean	std	min	25%	50%	75%	max	dtypes	\
y	87486	294.447861	774.531815	-0.0	0.0	0.0	183.7125	5733.42	float64	

	unique	missing_count	missing_ratio	raw_type	special_types
y	11321			float	




1.1.1 Distribution fits for target variable


- none of the [attempted](#) distribution fits satisfy specified minimum p-value threshold: 0.01

1.1.2 Target variable correlations


train_data - spearman correlation matrix; focus: absolute correlation for $y \geq 0.5$




Feature interaction between clear_sky_energy_1h:J/y in train_data




Feature interaction between diffuse_rad:W/y in train_data




Feature interaction between diffuse_rad_1h:J/y in train_data




Feature interaction between `direct_rad_1h:J/y` in `train_data`




Feature interaction between `direct_rad:W/y` in `train_data`




Feature interaction between sun_elevation:d/y in train_data



Feature interaction between is_day:idx/y in train_data



Feature interaction between `is_in_shadow:idx`/`y` in `train_data`



2 Starting

```
[9]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    ↪filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

```
Last submission number: 100
Now creating submission number: 101
New filename: submission_101
```

```
[10]: predictors = [None, None, None]
```

```
[ ]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪train and tune data and test data
    if weight_evaluation:
        print("Train data sample weight sum:",
            ↪train_data[train_data["location"] == loc]["sample_weight"].sum())
        print("Train data number of rows:", train_data[train_data["location"]
            ↪== loc].shape[0])
    if use_tune_data:
        print("Tune data sample weight sum:",
            ↪tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
        print("Tune data number of rows:",
            ↪tuning_data[tuning_data["location"] == loc].shape[0])
    if use_test_data:
        print("Test data sample weight sum:",
            ↪test_data[test_data["location"] == loc]["sample_weight"].sum())
        print("Test data number of rows:", test_data[test_data["location"]
            ↪== loc].shape[0])
    predictor = TabularPredictor(
        label=label,
```

```

        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        # sample_weight=sample_weight,
        # weight_evaluation=weight_evaluation,
        # groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc].
↳drop(columns=["ds"]),
        time_limit=time_limit,
        # presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
↳somethin, will be overwritten anyways
        num_bag_sets=num_bag_sets,
        tuning_data=tuning_data[tuning_data["location"] == loc].
↳reset_index(drop=True).drop(columns=["ds"]) if use_tune_data else None,
        use_bag_holdout=use_bag_holdout,
        # holdout_frac=holdout_frac,
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
↳drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

```

Beginning AutoGluon training ... Time limit = 1800s
AutoGluon will save models to "AutogluonModels/submission_101_A/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 137.83 GB / 315.93 GB (43.6%)
Train Data Rows: 31872
Train Data Columns: 28
Tuning Data Rows: 1093
Tuning Data Columns: 28
Label Column: y

```

```

Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (5733.42, 0.0, 649.68162,
1178.37671)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 132236.17 MB
    Train Data (Original) Memory Usage: 9.03 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 25 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'cloud_base_agl:m', 'diffuse_rad:W',
'diffuse_rad_1h:J', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 25 | ['ceiling_height_agl:m',
'clear_sky_energy_1h:J', 'cloud_base_agl:m', 'diffuse_rad:W',
'diffuse_rad_1h:J', ...]
        ('int', ['bool']) : 1 | ['is_estimated']
    0.1s = Fit runtime
    26 features in original data used to generate 26 features in processed
data.
    Train Data (Processed) Memory Usage: 6.63 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.15s ...
AutoGluon will gauge predictive performance using evaluation metric:

```

'mean_absolute_error'

This metric's sign has been flipped to adhere to being higher_is_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor() use_bag_holdout=True, will use tuning_data as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge'],
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

Training model for location A...

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 1799.85s of the 1799.85s of remaining time.

-140.7607 = Validation score (-mean_absolute_error)

0.03s = Training runtime

0.37s = Validation runtime

Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 1799.34s of the 1799.34s of remaining time.

-140.9568 = Validation score (-mean_absolute_error)

0.03s = Training runtime

0.37s = Validation runtime

Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1798.88s of the 1798.88s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

-95.4279 = Validation score (-mean_absolute_error)

30.28s = Training runtime

12.83s = Validation runtime

Fitting model: LightGBM_BAG_L1 ... Training model for up to 1759.26s of the

1759.26s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

-98.3654 = Validation score (-mean_absolute_error)
23.39s = Training runtime
3.38s = Validation runtime

Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1732.6s of the 1732.6s of remaining time.

-109.0003 = Validation score (-mean_absolute_error)
6.33s = Training runtime
1.19s = Validation runtime

Fitting model: CatBoost_BAG_L1 ... Training model for up to 1723.76s of the 1723.75s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

-105.2989 = Validation score (-mean_absolute_error)
190.92s = Training runtime
0.11s = Validation runtime

Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1531.67s of the 1531.66s of remaining time.

-113.7193 = Validation score (-mean_absolute_error)
1.45s = Training runtime
1.18s = Validation runtime

Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1527.68s of the 1527.68s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

-108.1271 = Validation score (-mean_absolute_error)
38.92s = Training runtime
0.56s = Validation runtime

Fitting model: XGBoost_BAG_L1 ... Training model for up to 1486.97s of the 1486.97s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

-102.6135 = Validation score (-mean_absolute_error)
8.73s = Training runtime
0.34s = Validation runtime

Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1476.11s of the 1476.11s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

-95.0881 = Validation score (-mean_absolute_error)
107.54s = Training runtime
0.37s = Validation runtime

Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1367.18s of the 1367.17s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

-97.7935          = Validation score    (-mean_absolute_error)
90.77s    = Training    runtime
17.6s     = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1266.0s of the
1265.99s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -95.1885          = Validation score    (-mean_absolute_error)
    61.78s    = Training    runtime
    23.97s    = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1228.88s of the
1228.88s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -97.9718          = Validation score    (-mean_absolute_error)
    46.17s    = Training    runtime
    8.24s     = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1201.53s of the
1201.53s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -105.2287         = Validation score    (-mean_absolute_error)
    383.15s = Training    runtime
    0.2s     = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1007.86s of
the 1007.85s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -108.1618         = Validation score    (-mean_absolute_error)
    78.38s    = Training    runtime
    1.1s     = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 965.99s of the
965.98s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -102.2023         = Validation score    (-mean_absolute_error)
    17.15s    = Training    runtime
    0.73s     = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 955.96s of the
955.96s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

[ ]: import matplotlib.pyplot as plt

leaderboards = [None, None, None]

```

```

def leaderboard_for_location(i, loc):
    if use_test_data:
        lb = predictors[i].leaderboard(test_data[test_data["location"] == loc])
        lb["location"] = loc
        plt.scatter(test_data[test_data["location"] == loc]["y"].index,
↳test_data[test_data["location"] == loc]["y"])
        if use_tune_data:
            plt.scatter(tuning_data[tuning_data["location"] == loc]["y"].index,
↳tuning_data[tuning_data["location"] == loc]["y"])
            plt.show()

        return lb
    else:
        return pd.DataFrame()

leaderboards[0] = leaderboard_for_location(0, loc)

```

```

[13]: loc = "B"
predictors[1] = fit_predictor_for_location(loc)
leaderboards[1] = leaderboard_for_location(1, loc)

```

```

-16.3215          = Validation score   (-mean_absolute_error)
385.35s   = Training   runtime
0.21s     = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 854.02s of
the 854.02s of remaining time.
Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-14.193   = Validation score   (-mean_absolute_error)
77.43s    = Training   runtime
1.05s     = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 812.89s of the
812.88s of remaining time.
Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-15.1244      = Validation score   (-mean_absolute_error)
152.45s   = Training   runtime
48.94s    = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 728.57s of the
728.57s of remaining time.
Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-11.751   = Validation score   (-mean_absolute_error)
351.21s   = Training   runtime
0.66s     = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 550.0s of the
550.0s of remaining time.
Fitting 8 child models (S2F1 - S2F8) | Fitting with

```

```


ParallelLocalFoldFittingStrategy
    -14.2513      = Validation score    (-mean_absolute_error)
    186.52s      = Training   runtime
    41.83s       = Validation runtime
Completed 2/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
442.77s of remaining time.
    -11.6597      = Validation score    (-mean_absolute_error)
    0.43s         = Training   runtime
    0.0s          = Validation runtime
AutoGluon training complete, total runtime = 1357.68s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_101_B/")
Evaluation: mean_absolute_error on test data: -14.547869225363064
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -14.547869225363064,
    "root_mean_squared_error": -44.85425668975588,
    "mean_squared_error": -2011.90434319051,
    "r2": 0.909051782885076,
    "pearsonr": 0.9548156385330459,
    "median_absolute_error": -0.37705013155937195
}

Evaluation on test data:
-14.547869225363064

      model  score_test  score_val  pred_time_test  pred_time_val
fit_time  pred_time_test_marginal  pred_time_val_marginal  fit_time_marginal
stack_level  can_infer  fit_order
0  NeuralNetTorch_BAG_L1  -14.489736 -11.751021      0.351410      0.662124
351.210851      0.351410      0.662124      351.210851
1  True      10
1  WeightedEnsemble_L2  -14.547869 -11.659718      1.207992      2.057405
738.400992      0.003429      0.000540      0.426711
2  True      12
2  NeuralNetFastAI_BAG_L1  -16.326178 -14.193033      1.020282      1.049720
77.426473      1.020282      1.049720      77.426473
1  True      8
3  LightGBMLarge_BAG_L1  -16.867807 -14.251304      8.466552      41.832313
186.524971      8.466552      41.832313      186.524971
1  True      11
4  XGBoost_BAG_L1  -17.886484 -15.124446      4.588538      48.941980
152.446251      4.588538      48.941980      152.446251
1  True      9
5  LightGBM_BAG_L1  -18.015132 -15.396378      2.733061      25.483243

```

60.669254		2.733061		25.483243		60.669254
1	True	4				
6	CatBoost_BAG_L1	-18.516337	-16.321547	0.191183		0.214026
385.354122		0.191183		0.214026		385.354122
1	True	6				
7	LightGBMXT_BAG_L1	-19.090083	-15.743320	2.770541		29.846826
58.484428		2.770541		29.846826		58.484428
1	True	3				
8	ExtraTreesMSE_BAG_L1	-19.373523	-15.172927	0.661969		1.180715
1.409308		0.661969		1.180715		1.409308
1	True	7				
9	RandomForestMSE_BAG_L1	-20.019074	-15.782449	0.560164		1.152136
7.220658		0.560164		1.152136		7.220658
1	True	5				
10	KNeighborsDist_BAG_L1	-30.089573	-23.622609	0.017341		1.242610
0.026984		0.017341		1.242610		0.026984
1	True	2				
11	KNeighborsUnif_BAG_L1	-30.135864	-23.684113	0.019941		0.452281
0.028449		0.019941		0.452281		0.028449
1	True	1				



```
[14]: loc = "C"
predictors[2] = fit_predictor_for_location(loc)
leaderboards[2] = leaderboard_for_location(2, loc)
```

```
Beginning AutoGluon training ... Time limit = 1800s
AutoGluon will save models to "AutogluonModels/submission_101_C/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 129.79 GB / 315.93 GB (41.1%)
Train Data Rows: 24594
Train Data Columns: 28
Tuning Data Rows: 737
Tuning Data Columns: 28
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and label-values can't be converted to int).
    Label info (max, min, mean, stddev): (999.6, -0.0, 79.8926, 168.407)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 130059.28 MB
    Train Data (Original) Memory Usage: 6.94 MB (0.0% of available memory)

Training model for location C...

    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 1 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.

    This is typically a feature which has the same value for all
```

rows.

These features do not need to be present at inference time.

Types of features in original data (raw dtype, special dtypes):

```
('float', []) : 25 | ['ceiling_height_agl:m',  
'clear_sky_energy_1h:J', 'cloud_base_agl:m', 'diffuse_rad:W',  
'diffuse_rad_1h:J', ...]  
('int', []) : 1 | ['is_estimated']
```

Types of features in processed data (raw dtype, special dtypes):

```
('float', []) : 25 | ['ceiling_height_agl:m',  
'clear_sky_energy_1h:J', 'cloud_base_agl:m', 'diffuse_rad:W',  
'diffuse_rad_1h:J', ...]  
('int', ['bool']) : 1 | ['is_estimated']
```

0.1s = Fit runtime

26 features in original data used to generate 26 features in processed data.

Train Data (Processed) Memory Usage: 5.09 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.13s ...

AutoGluon will gauge predictive performance using evaluation metric:

'mean_absolute_error'

This metric's sign has been flipped to adhere to being higher_is_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor()
use_bag_holdout=True, will use tuning_data as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{  
    'NN_TORCH': {},  
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],  
'GBMLarge'],  
    'CAT': {},  
    'XGB': {},  
    'FASTAI': {},  
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',  
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':  
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},  
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',  
'problem_types': ['regression', 'quantile']}}],  
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',  
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':  
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},  
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',  
'problem_types': ['regression', 'quantile']}}],  
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},  
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],  
}
```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 1799.87s of the 1799.87s of remaining time.

```

-23.649 = Validation score (-mean_absolute_error)
0.02s   = Training runtime
0.27s   = Validation runtime
Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 1799.52s of
the 1799.52s of remaining time.
-23.7005 = Validation score (-mean_absolute_error)
0.02s    = Training runtime
0.34s    = Validation runtime
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1799.1s of the
1799.1s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.464 = Validation score (-mean_absolute_error)
27.71s  = Training runtime
9.72s   = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1767.42s of the
1767.42s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.7205 = Validation score (-mean_absolute_error)
26.49s   = Training runtime
6.44s    = Validation runtime
Fitting model: RandomForestMSE_BAG_L1 ... Training model for up to 1737.52s of
the 1737.52s of remaining time.
-16.4245 = Validation score (-mean_absolute_error)
3.86s    = Training runtime
0.76s    = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1732.26s of the
1732.25s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.4305 = Validation score (-mean_absolute_error)
184.47s  = Training runtime
0.09s    = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 1546.52s of the
1546.52s of remaining time.
-16.1959 = Validation score (-mean_absolute_error)
0.93s    = Training runtime
0.78s    = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1544.1s of
the 1544.1s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-14.7987 = Validation score (-mean_absolute_error)
31.64s   = Training runtime
0.42s    = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 1510.79s of the
1510.79s of remaining time.

```



```

    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.5135      = Validation score    (-mean_absolute_error)
    43.75s       = Training    runtime
    2.49s        = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1463.86s of
the 1463.86s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.2172      = Validation score    (-mean_absolute_error)
    92.48s        = Training    runtime
    0.26s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1369.98s of the
1369.98s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.8764      = Validation score    (-mean_absolute_error)
    88.53s        = Training    runtime
    12.53s        = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1273.04s of the
1273.03s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.3868      = Validation score    (-mean_absolute_error)
    57.27s        = Training    runtime
    22.75s        = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1238.28s of the
1238.27s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.5816      = Validation score    (-mean_absolute_error)
    52.39s        = Training    runtime
    11.54s        = Validation runtime
Fitting model: CatBoost_BAG_L1 ... Training model for up to 1207.83s of the
1207.83s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.3746      = Validation score    (-mean_absolute_error)
    370.14s       = Training    runtime
    0.18s         = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1020.81s of
the 1020.81s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.7932      = Validation score    (-mean_absolute_error)
    62.91s        = Training    runtime
    0.87s         = Validation runtime

```

Fitting model: XGBoost_BAG_L1 ... Training model for up to 987.16s of the 987.16s of remaining time.
 Fitting 8 child models (S2F1 - S2F8) | Fitting with
 ParallelLocalFoldFittingStrategy
 -13.6574 = Validation score (-mean_absolute_error)
 81.56s = Training runtime
 4.68s = Validation runtime

Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 945.57s of the 945.57s of remaining time.
 Fitting 8 child models (S2F1 - S2F8) | Fitting with
 ParallelLocalFoldFittingStrategy
 -13.3958 = Validation score (-mean_absolute_error)
 164.59s = Training runtime
 0.53s = Validation runtime

Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 871.89s of the 871.89s of remaining time.
 Fitting 8 child models (S2F1 - S2F8) | Fitting with
 ParallelLocalFoldFittingStrategy
 -12.8074 = Validation score (-mean_absolute_error)
 176.26s = Training runtime
 22.76s = Validation runtime

Repeating k-fold bagging: 3/20

Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 772.19s of the 772.19s of remaining time.
 Fitting 8 child models (S3F1 - S3F8) | Fitting with
 ParallelLocalFoldFittingStrategy
 -12.3667 = Validation score (-mean_absolute_error)
 84.13s = Training runtime
 34.2s = Validation runtime

Fitting model: LightGBM_BAG_L1 ... Training model for up to 738.9s of the 738.9s of remaining time.
 Fitting 8 child models (S3F1 - S3F8) | Fitting with
 ParallelLocalFoldFittingStrategy
 -13.555 = Validation score (-mean_absolute_error)
 76.32s = Training runtime
 16.14s = Validation runtime

Fitting model: CatBoost_BAG_L1 ... Training model for up to 709.54s of the 709.54s of remaining time.
 Fitting 8 child models (S3F1 - S3F8) | Fitting with
 ParallelLocalFoldFittingStrategy
 -13.3762 = Validation score (-mean_absolute_error)
 555.1s = Training runtime
 0.27s = Validation runtime

Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 523.16s of the 523.16s of remaining time.
 Fitting 8 child models (S3F1 - S3F8) | Fitting with
 ParallelLocalFoldFittingStrategy
 -14.7958 = Validation score (-mean_absolute_error)


```

    94.98s = Training runtime
    1.28s = Validation runtime
Fitting model: XGBoost_BAG_L1 ... Training model for up to 488.47s of the
488.47s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.612 = Validation score (-mean_absolute_error)
    124.68s = Training runtime
    6.1s = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 440.97s of the
440.97s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.3421 = Validation score (-mean_absolute_error)
    251.81s = Training runtime
    0.79s = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 352.05s of the
352.04s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.8232 = Validation score (-mean_absolute_error)
    262.66s = Training runtime
    33.03s = Validation runtime
Completed 3/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
250.2s of remaining time.
    -11.6491 = Validation score (-mean_absolute_error)
    0.43s = Training runtime
    0.0s = Validation runtime
AutoGluon training complete, total runtime = 1550.25s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_101_C/")
Evaluation: mean_absolute_error on test data: -11.584846974967073
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -11.584846974967073,
    "root_mean_squared_error": -33.28157104597867,
    "mean_squared_error": -1107.662971288526,
    "r2": 0.9119603983629877,
    "pearsonr": 0.9553301989877001,
    "median_absolute_error": -0.5380167663097382
}

Evaluation on test data:
-11.584846974967073

```

	model	score_test	score_val	pred_time_test	pred_time_val
fit_time	pred_time_test_marginal	pred_time_val_marginal	fit_time_marginal		
stack_level	can_infer	fit_order			
0	WeightedEnsemble_L2	-11.584847	-11.649102	15.890272	68.017162
599.024821		0.003410		0.000649	0.426667
2	True	12			
1	LightGBMXT_BAG_L1	-12.414818	-12.366686	5.450568	34.196357
84.132116		5.450568		34.196357	84.132116
1	True	3			
2	LightGBMLarge_BAG_L1	-12.543343	-12.823171	9.914580	33.028136
262.660037		9.914580		33.028136	262.660037
1	True	11			
3	NeuralNetTorch_BAG_L1	-12.844010	-13.342115	0.521714	0.792019
251.806001		0.521714		0.792019	251.806001
1	True	10			
4	LightGBM_BAG_L1	-13.018393	-13.555035	4.026720	16.135706
76.315174		4.026720		16.135706	76.315174
1	True	4			
5	XGBoost_BAG_L1	-13.618762	-13.612037	1.837001	6.100511
124.681797		1.837001		6.100511	124.681797
1	True	9			
6	CatBoost_BAG_L1	-13.908382	-13.376167	0.210222	0.273467
555.101594		0.210222		0.273467	555.101594
1	True	6			
7	NeuralNetFastAI_BAG_L1	-14.484465	-14.795795	1.373520	1.280287
94.977746		1.373520		1.280287	94.977746
1	True	8			
8	ExtraTreesMSE_BAG_L1	-15.827128	-16.195930	0.396986	0.775828
0.931813		0.396986		0.775828	0.931813
1	True	7			
9	RandomForestMSE_BAG_L1	-16.451605	-16.424500	0.334966	0.761000
3.862862		0.334966		0.761000	3.862862
1	True	5			
10	KNeighborsDist_BAG_L1	-23.919280	-23.700527	0.013227	0.338030
0.021528		0.013227		0.338030	0.021528
1	True	2			
11	KNeighborsUnif_BAG_L1	-24.102600	-23.649027	0.028486	0.269659
0.020741		0.028486		0.269659	0.020741
1	True	1			



```
[ ]: # save leaderboards to csv
pd.concat(leaderboards).to_csv(f"leaderboards/{new_filename}.csv")
```

3 Submit

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt

future_test_data = TabularDataset('X_test_raw.csv')
future_test_data["ds"] = pd.to_datetime(future_test_data["ds"])
#test_data
```

Loaded data from: X_test_raw.csv | Columns = 29 / 29 | Rows = 4608 -> 4608

```
[ ]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
future_test_data_merged = pd.merge(future_test_data, test_ids, how="inner",
    ↪right_on=["time", "location"], left_on=["ds", "location"])

#test_data_merged
```

Loaded data from: test.csv | Columns = 4 / 4 | Rows = 2160 -> 2160

```
[ ]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in future_test_data.groupby('location'):
    i = location_map[loc]
    subset = future_test_data_merged[future_test_data_merged["location"] == loc]
    subset.reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)

    # get past predictions
    train_data.loc[train_data["location"] == loc, "prediction"] = pred
    predictors[i].predict(train_data[train_data["location"] == loc])
    if use_tune_data:
        tuning_data.loc[tuning_data["location"] == loc, "prediction"] = pred
        predictors[i].predict(tuning_data[tuning_data["location"] == loc])
    if use_test_data:
        test_data.loc[test_data["location"] == loc, "prediction"] = pred
        predictors[i].predict(test_data[test_data["location"] == loc])

[ ]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data[train_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    label="train data")
    if use_tune_data:
        tuning_data[tuning_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    label="tune data")
    if use_test_data:
        test_data[test_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    label="test data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")



    # plot past predictions
    #train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    y='prediction', ax=ax, label="past predictions")
```

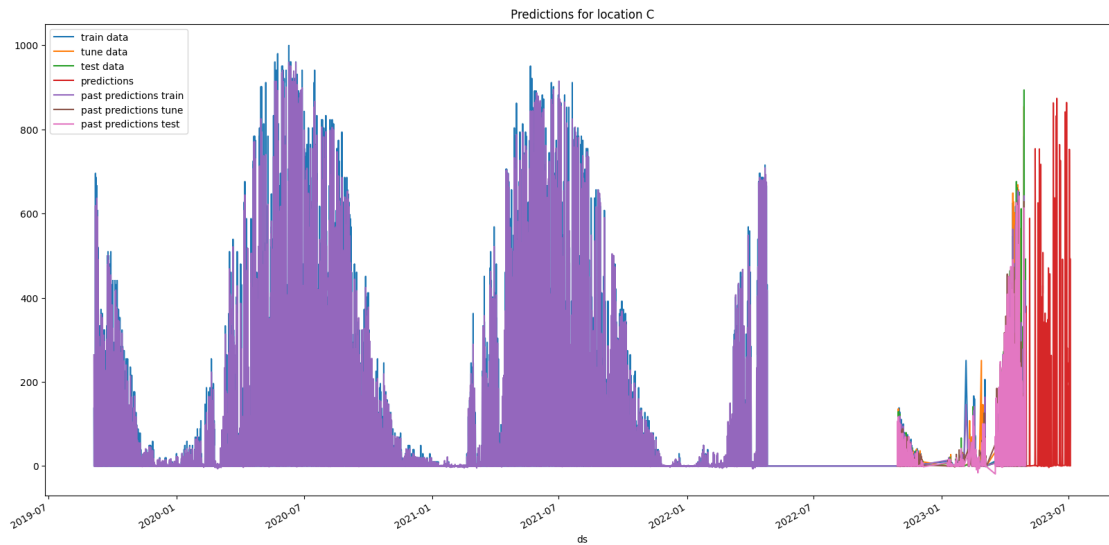
```

train_data[train_data["location"]==loc].plot(x='ds', y='prediction', ax=ax,
↪label="past predictions train")
if use_tune_data:
    tuning_data[tuning_data["location"]==loc].plot(x='ds', y='prediction',
↪ax=ax, label="past predictions tune")
if use_test_data:
    test_data[test_data["location"]==loc].plot(x='ds', y='prediction',
↪ax=ax, label="past predictions test")

# title
ax.set_title(f"Predictions for location {loc}")

```





```
[ ]: temp_predictions = [prediction.copy() for prediction in predictions]
if clip_predictions:
    # clip predictions smaller than 0 to 0
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        pred.loc[pred["prediction"] < 0, "prediction"] = 0
        print("Smallest prediction after clipping:", pred["prediction"].min())

# Instead of clipping, shift all prediction values up by the largest negative
↪ number.
# This way, the smallest prediction will be 0.
elif shift_predictions:
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        pred["prediction"] = pred["prediction"] - pred["prediction"].min()
        print("Smallest prediction after clipping:", pred["prediction"].min())

elif shift_predictions_by_average_of_negatives_then_clip:
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        mean_negative = pred[pred["prediction"] < 0]["prediction"].mean()
        # if not nan
        if mean_negative == mean_negative:
```



```

    pred["prediction"] = pred["prediction"] - mean_negative

    pred.loc[pred["prediction"] < 0, "prediction"] = 0
    print("Smallest prediction after clipping:", pred["prediction"].min())

# concatenate predictions
submissions_df = pd.concat(temp_predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df

```

```

Smallest prediction: -16.878181
Smallest prediction after clipping: 0.0
Smallest prediction: 0.07609784
Smallest prediction after clipping: 0.07609784
Smallest prediction: -3.4177196
Smallest prediction after clipping: 0.0

```

```

[ ]:      id  prediction
0      0      0.000000
1      1      0.000000
2      2      0.000000
3      3     61.788113
4      4    296.051880
..    ...      ...
715   2155    70.078865
716   2156    43.007858
717   2157    10.200218
718   2158     4.401361
719   2159     0.849800

```

```

[2160 rows x 2 columns]

```

```

[ ]: # Save the submission DataFrame to submissions folder, create new name based on
    ↳ last submission, format is submission_<last_submission_number + 1>.csv

    # Save the submission
    print(f"Saving submission to submissions/{new_filename}.csv")
    submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
    ↳ index=False)
    print("jall1a")

```

```

Saving submission to submissions/submission_101.csv
jall1a

```

```

[ ]: train_data_with_dates = TabularDataset('X_train_raw.csv')
    train_data_with_dates["ds"] = pd.to_datetime(train_data_with_dates["ds"])

```

```
# feature importance
location="A"
split_time = pd.Timestamp("2022-10-28 22:00:00")
estimated = train_data_with_dates[train_data_with_dates["ds"] >= split_time]
estimated = estimated[estimated["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=estimated,
↳time_limit=60*10)
```

Loaded data from: X_train_raw.csv | Columns = 30 / 30 | Rows = 92945 -> 92945
 These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'location']

Computing feature importance via permutation shuffling for 26 features using 4392 rows with 10 shuffle sets... Time limit: 600s...

1858.07s = Expected runtime (185.81s per shuffle set)

564.33s = Actual runtime (Completed 5 of 10 shuffle sets) (Early stopping due to lack of time...)

```
[ ]:
```

	importance	stddev	p_value	n \
direct_rad_1h:J	174.758804	1.385776	4.744175e-10	5
clear_sky_energy_1h:J	120.681118	1.767450	5.519356e-09	5
diffuse_rad_1h:J	91.116242	1.237380	4.080410e-09	5
diffuse_rad:W	84.881611	1.184729	4.552924e-09	5
direct_rad:W	80.238447	1.949325	4.176824e-08	5
sun_elevation:d	61.935056	1.761940	7.851108e-08	5
sun_azimuth:d	51.658621	3.061597	1.473570e-06	5
effective_cloud_cover:p	29.954183	1.199781	3.081984e-07	5
sfc_pressure:hPa	20.717220	2.138205	1.342484e-05	5
total_cloud_cover:p	19.316892	0.960869	7.322466e-07	5
is_in_shadow:idx	15.526777	0.518693	1.492284e-07	5
snow_water:kgm2	14.541732	0.706230	6.654835e-07	5
relative_humidity_1000hPa:p	13.896670	0.740659	9.646449e-07	5
t_1000hPa:K	13.122288	1.133759	6.620869e-06	5
visibility:m	12.813578	0.989004	4.225220e-06	5
ceiling_height_agl:m	12.175750	0.728600	1.531388e-06	5
cloud_base_agl:m	12.174280	0.614288	7.752191e-07	5
wind_speed_10m:ms	9.728451	0.535143	1.094294e-06	5
is_day:idx	8.864305	0.477792	1.008965e-06	5
fresh_snow_6h:cm	8.810844	0.434730	7.088907e-07	5
super_cooled_liquid_water:kgm2	5.405043	0.792973	5.403276e-05	5
precip_type_5min:idx	5.397076	0.817897	6.139860e-05	5
fresh_snow_3h:cm	2.969766	0.446444	5.948215e-05	5
snow_depth:cm	2.947989	0.649037	2.646037e-04	5
fresh_snow_1h:cm	2.447736	0.362044	5.579644e-05	5
is_estimated	0.000000	0.000000	5.000000e-01	5

	p99_high	p99_low
direct_rad_1h:J	177.612136	171.905472

clear_sky_energy_1h:J	124.320322	117.041914
diffuse_rad_1h:J	93.664025	88.568460
diffuse_rad:W	87.320985	82.442237
direct_rad:W	84.252133	76.224761
sun_elevation:d	65.562914	58.307197
sun_azimuth:d	57.962491	45.354751
effective_cloud_cover:p	32.424548	27.483818
sfc_pressure:hPa	25.119813	16.314627
total_cloud_cover:p	21.295334	17.338449
is_in_shadow:idx	16.594773	14.458781
snow_water:kgm2	15.995870	13.087594
relative_humidity_1000hPa:p	15.421696	12.371643
t_1000hPa:K	15.456713	10.787864
visibility:m	14.849950	10.777206
ceiling_height_agl:m	13.675948	10.675552
cloud_base_agl:m	13.439108	10.909451
wind_speed_10m:ms	10.830317	8.626585
is_day:idx	9.848084	7.880525
fresh_snow_6h:cm	9.705959	7.915730
super_cooled_liquid_water:kgm2	7.037786	3.772300
precip_type_5min:idx	7.081137	3.713015
fresh_snow_3h:cm	3.889001	2.050532
snow_depth:cm	4.284365	1.611613
fresh_snow_1h:cm	3.193190	1.702281
is_estimated	0.000000	0.000000

```
[ ]: # feature importance
observed = train_data_with_dates[train_data_with_dates["ds"] < split_time]
observed = observed[observed["location"] == location]
predictors[0].feature_importance(feature_stage="original", data=observed,
↳time_limit=60*10)
```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'location']

Computing feature importance via permutation shuffling for 26 features using 5000 rows with 10 shuffle sets... Time limit: 600s...

2130.12s = Expected runtime (213.01s per shuffle set)

569.03s = Actual runtime (Completed 4 of 10 shuffle sets) (Early stopping due to lack of time...)

[]:	importance	stddev	p_value	n	\
direct_rad_1h:J	2.979640e+02	7.742726e+00	2.417003e-06	4	
clear_sky_energy_1h:J	2.209681e+02	9.822432e+00	1.208502e-05	4	
diffuse_rad_1h:J	1.407333e+02	7.063642e+00	1.738850e-05	4	
diffuse_rad:W	1.374164e+02	7.932426e+00	2.643334e-05	4	
sun_elevation:d	1.086858e+02	7.045457e+00	3.740413e-05	4	
direct_rad:W	1.079841e+02	3.567674e+00	4.965927e-06	4	
sun_azimuth:d	1.000221e+02	4.298227e+00	1.091969e-05	4	

effective_cloud_cover:p	5.132260e+01	1.807393e+00	6.013092e-06	4
t_1000hPa:K	4.585019e+01	9.698193e-01	1.303842e-06	4
ceiling_height_agl:m	3.442258e+01	1.177199e+00	5.506973e-06	4
sfc_pressure:hPa	3.296733e+01	1.655336e+00	1.740897e-05	4
visibility:m	3.234283e+01	1.613807e+00	1.708438e-05	4
relative_humidity_1000hPa:p	3.163136e+01	1.637891e+00	1.908995e-05	4
wind_speed_10m:ms	2.991173e+01	1.924733e-01	3.672161e-08	4
cloud_base_agl:m	2.981524e+01	1.391653e+00	1.398871e-05	4
total_cloud_cover:p	2.653299e+01	1.212222e+00	1.311967e-05	4
snow_water:kgm2	2.382866e+01	1.286464e+00	2.163251e-05	4
precip_type_5min:idx	2.226013e+01	1.505111e+00	4.243157e-05	4
super_cooled_liquid_water:kgm2	2.057215e+01	7.752532e-01	7.366942e-06	4
is_in_shadow:idx	1.950323e+01	2.203955e+00	1.966394e-04	4
is_day:idx	1.233222e+01	1.283926e+00	1.540374e-04	4
fresh_snow_6h:cm	2.728599e+00	3.015600e-01	1.840345e-04	4
snow_depth:cm	1.646735e+00	3.652632e-01	1.440103e-03	4
fresh_snow_3h:cm	6.826377e-01	3.024028e-01	1.015541e-02	4
fresh_snow_1h:cm	4.674250e-01	2.723762e-01	2.073566e-02	4
is_estimated	5.015735e-09	1.799561e-08	3.080590e-01	4

	p99_high	p99_low
direct_rad_1h:J	3.205763e+02	2.753517e+02
clear_sky_energy_1h:J	2.496540e+02	1.922821e+02
diffuse_rad_1h:J	1.613624e+02	1.201043e+02
diffuse_rad:W	1.605827e+02	1.142501e+02
sun_elevation:d	1.292618e+02	8.810988e+01
direct_rad:W	1.184033e+02	9.756485e+01
sun_azimuth:d	1.125748e+02	8.746929e+01
effective_cloud_cover:p	5.660100e+01	4.604419e+01
t_1000hPa:K	4.868250e+01	4.301788e+01
ceiling_height_agl:m	3.786053e+01	3.098462e+01
sfc_pressure:hPa	3.780166e+01	2.813300e+01
visibility:m	3.705588e+01	2.762978e+01
relative_humidity_1000hPa:p	3.641475e+01	2.684798e+01
wind_speed_10m:ms	3.047384e+01	2.934962e+01
cloud_base_agl:m	3.387950e+01	2.575098e+01
total_cloud_cover:p	3.007323e+01	2.299276e+01
snow_water:kgm2	2.758572e+01	2.007160e+01
precip_type_5min:idx	2.665574e+01	1.786453e+01
super_cooled_liquid_water:kgm2	2.283625e+01	1.830806e+01
is_in_shadow:idx	2.593978e+01	1.306668e+01
is_day:idx	1.608187e+01	8.582576e+00
fresh_snow_6h:cm	3.609291e+00	1.847907e+00
snow_depth:cm	2.713470e+00	5.800004e-01
fresh_snow_3h:cm	1.565791e+00	-2.005158e-01
fresh_snow_1h:cm	1.262887e+00	-3.280375e-01
is_estimated	5.757109e-08	-4.753962e-08

```
[ ]: # save this running notebook
from IPython.display import display, Javascript
import time

# hei123

display(Javascript("IPython.notebook.save_checkpoint();"))

time.sleep(3)
```

<IPython.core.display.Javascript object>

```
[ ]: # save this notebook to submissions folder
import subprocess
import os
#subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
    ↪ipynb"])
```

```
[NbConvertApp] Converting notebook autogluon_each_location.ipynb to pdf
/opt/conda/lib/python3.10/site-packages/nbconvert/utils/pandoc.py:51:
RuntimeWarning: You are using an unsupported version of pandoc (2.9.2.1).
Your version must be at least (2.14.2) but less than (4.0.0).
Refer to https://pandoc.org/installing.html.
Continuing with doubts...
    check_pandoc_version()
[NbConvertApp] Support files will be in notebook_pdfs/submission_101_files/
[NbConvertApp] Making directory
./notebook_pdfs/submission_101_files/notebook_pdfs
[NbConvertApp] Writing 183909 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 1904761 bytes to notebook_pdfs/submission_101.pdf

[ ]: CompletedProcess(args=['jupyter', 'nbconvert', '--to', 'pdf', '--output',
'notebook_pdfs/submission_101.pdf', 'autogluon_each_location.ipynb'],
returncode=0)
```

```
[ ]: # display(Javascript("IPython.notebook.save_checkpoint();"))
# time.sleep(3)

# subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
    ↪join('notebook_pdfs', f"{new_filename}_with_feature_importance.pdf"),
    ↪"autogluon_each_location.ipynb"])
```

```
[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
#     ↪ stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
#     ↪ strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
#     ↪ 'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', 'hello if hello is
#     ↪ not None else 'Henrik eller Jørgen'])

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"

# execute_git_command(git_repo_path, ['checkout', '-b', branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m', commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push',
#     ↪ 'origin', branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream',
#     ↪ 'origin', branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])
```

[]:

[]:

[]:

[]: