

autogluon_each_location

October 29, 2023

1 Config

```
[1]: # config

label = 'y'
metric = 'mean_absolute_error'
time_limit = 60*60
presets = "experimental_zeroshot_hpo_hybrid"#'best_quality'

do_drop_ds = True
# hour, dayofweek, dayofmonth, month, year
use_dt_attrs = []#["hour", "year"]
use_estimated_diff_attr = False
use_is_estimated_attr = True

drop_night_outliers = True
drop_null_outliers = False

# to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
↳ "dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm",
↳ "wind_speed_u_10m:ms", "wind_speed_v_10m:ms", "snow_melt_10min:mm",
↳ "rain_water:kgm2", "dew_point_2m:K", "precip_5min:mm", "absolute_humidity_2m:
↳ gm3", "air_density_2m:kgm3"]#, "msl_pressure:hPa", "pressure_50m:hPa",
↳ "pressure_100m:hPa"]
to_drop = ["wind_speed_w_1000hPa:ms", "wind_speed_u_10m:ms", "wind_speed_v_10m:
↳ ms"]

excluded_model_types = ['CAT', 'XGB', 'RF']

use_groups = False
n_groups = 8

# auto_stack = True
num_stack_levels = 0
num_bag_folds = None# 8
num_bag_sets = None#20
```

```

use_tune_data = True
use_test_data = True
#tune_and_test_length = 0.5 # 3 months from end
# holdout_frac = None
use_bag_holdout = True # Enable this if there is a large gap between score_val_
↳and score_test in stack models.

sample_weight = None#'sample_weight' #None
weight_evaluation = False#
sample_weight_estimated = 1
sample_weight_may_july = 1

run_analysis = False

shift_predictions_by_average_of_negatives_then_clip = False
clip_predictions = True
shift_predictions = False

```

2 Loading and preprocessing

```

[2]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift all columns with "1h" in them by 1 hour, so that for index 16:00,
    ↳we have the values from 17:00
    # but only for the columns with "1h" in the name
    #X_shifted = X.filter(regex="\dh").shift(-1, axis=1)
    #print(f"Number of columns with 1h in name: {X_shifted.columns}")

    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

    # Filter rows where index.minute == 0
    X_shifted = X[X.index.minute == 0][columns].copy()

    # Create a set for constant-time lookup
    index_set = set(X.index)

```

```

# Vectorized time shifting
one_hour = pd.Timedelta('1 hour')
shifted_indices = X_shifted.index + one_hour
X_shifted.loc[shifted_indices.isin(index_set)] = X.
↪loc[shifted_indices[shifted_indices.isin(index_set)]] [columns]

# set last row to same as second last row
X_shifted.iloc[-1] = X_shifted.iloc[-2]

# Count
count1 = len(shifted_indices[shifted_indices.isin(index_set)])
count2 = len(X_shifted) - count1

print("COUNT1", count1)
print("COUNT2", count2)

# Rename columns
X_old_unshifted = X_shifted.copy()
X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
↪columns]

date_calc = None
# If 'date_calc' is present, handle it
if 'date_calc' in X.columns:
    date_calc = X[X.index.minute == 0]['date_calc']

# resample to hourly
print("index: ", X.index[0])
X = X.resample('H').mean()
print("index AFTER: ", X.index[0])

X[columns] = X_shifted[columns]
#X[X_old_unshifted.columns] = X_old_unshifted

if date_calc is not None:
    X['date_calc'] = date_calc

return X

def fix_X(X, name):

```

```

    # Convert 'date_forecast' to datetime format and replace original column
    ↪with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])
    X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
    X.sort_values(by='ds', inplace=True)
    X.set_index('ds', inplace=True)

    X = feature_engineering(X)

    return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    if weight_evaluation:
        # add sample weights, which are 1 for observed and 3 for estimated
        X_train_observed["sample_weight"] = 1
        X_train_estimated["sample_weight"] = sample_weight_estimated
        X_test["sample_weight"] = sample_weight_estimated

    y_train['ds'] = pd.to_datetime(y_train['time'])
    y_train.drop(columns=['time'], inplace=True)
    y_train.sort_values(by='ds', inplace=True)
    y_train.set_index('ds', inplace=True)

    return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
    ↪location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
    ↪handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_estimated_diff_attr:
        X_train_observed["estimated_diff_hours"] = 0
        X_train_estimated["estimated_diff_hours"] = (X_train_estimated.index -
    ↪pd.to_datetime(X_train_estimated["date_calc"])).dt.total_seconds() / 3600

```

```

X_test["estimated_diff_hours"] = (X_test.index - pd.
↳to_datetime(X_test["date_calc"])).dt.total_seconds() / 3600

X_train_estimated["estimated_diff_hours"] =
↳X_train_estimated["estimated_diff_hours"].astype('int64')
    # the filled once will get dropped later anyways, when we drop y nans
X_test["estimated_diff_hours"] = X_test["estimated_diff_hours"].
↳fillna(-50).astype('int64')

if use_is_estimated_attr:
    X_train_observed["is_estimated"] = 0
    X_train_estimated["is_estimated"] = 1
    X_test["is_estimated"] = 1

# drop date_calc
X_train_estimated.drop(columns=['date_calc'], inplace=True)
X_test.drop(columns=['date_calc'], inplace=True)

y_train["y"] = y_train["pv_measurement"].astype('float64')
y_train.drop(columns=['pv_measurement'], inplace=True)
X_train = pd.concat([X_train_observed, X_train_estimated])

# clip all y values to 0 if negative
y_train["y"] = y_train["y"].clip(lower=0)

X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

# print number of nans in y
print(f"Number of nans in y: {X_train['y'].isna().sum()}")

print(f"Size of estimated after dropping nans:
↳{len(X_train[X_train['is_estimated']==1].dropna(subset=['y']))}")

X_train["location"] = location
X_test["location"] = location

return X_train, X_test
# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []

```

```

# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
    ↪X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

```

Processing location A...
COUNT1 29667
COUNT2 1
index: 2019-06-02 22:00:00
index AFTER: 2019-06-02 22:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 0
Size of estimated after dropping nans: 4418
Processing location B...
COUNT1 29232
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2

```

```

index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 4
Size of estimated after dropping nans: 3625
Processing location C...
COUNT1 29206
COUNT2 1
index: 2019-01-01 00:00:00
index AFTER: 2019-01-01 00:00:00
COUNT1 4392
COUNT2 2
index: 2022-10-28 22:00:00
index AFTER: 2022-10-28 22:00:00
COUNT1 702
COUNT2 18
index: 2023-05-01 00:00:00
index AFTER: 2023-05-01 00:00:00
Number of nans in y: 6059
Size of estimated after dropping nans: 2954

```

2.1 Feature engineering

2.1.1 Remove anomalies

```

[3]: import numpy as np
import pandas as pd

# loop thorough x train[y], keep track of streaks of same values and replace
↳ them with nan if they are too long
# also replace nan with 0

import numpy as np

def replace_streaks_with_nan(df, max_streak_length, column="y"):
    for location in df["location"].unique():
        x = df[df["location"] == location][column].copy()

        last_val = None
        streak_length = 1
        streak_indices = []
        allowed = [0]
        found_streaks = {}

```

```

for idx in x.index:
    value = x[idx]
    # if location == "B":
    #     continue

    if value == last_val and value not in allowed:
        streak_length += 1
        streak_indices.append(idx)
    else:
        streak_length = 1
        last_val = value
        streak_indices.clear()

    if streak_length > max_streak_length:
        found_streaks[value] = streak_length

        for streak_idx in streak_indices:
            x[idx] = np.nan
            streak_indices.clear() # clear after setting to NaN to avoid
↪setting multiple times
        df.loc[df["location"] == location, column] = x

    print(f"Found streaks for location {location}: {found_streaks}")

return df

# deep copy of X_train into x_copy
X_train = replace_streaks_with_nan(X_train.copy(), 3, "y")

```

Found streaks for location A: {}

Found streaks for location B: {3.45: 28, 6.9: 7, 12.9375: 5, 13.8: 8, 276.0: 78, 18.975: 58, 0.8625: 4, 118.1625: 33, 34.5: 11, 183.7125: 1058, 87.1125: 7, 79.35: 34, 7.7625: 12, 27.6: 448, 273.41249999999997: 72, 264.78749999999997: 55, 169.05: 33, 375.1875: 56, 314.8125: 66, 76.7625: 10, 135.4125: 216, 81.9375: 202, 2.5875: 12, 81.075: 210}

Found streaks for location C: {9.8: 4, 29.400000000000002: 4, 19.6: 4}

```

[4]: # print num rows
temprows = len(X_train)
X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
↪inplace=True)
print("Dropped rows: ", temprows - len(X_train))

```

Dropped rows: 9293

```

[5]: import matplotlib.pyplot as plt
import seaborn as sns

```



```

# Filter out rows where y == 0
temp = X_train[X_train["y"] != 0]

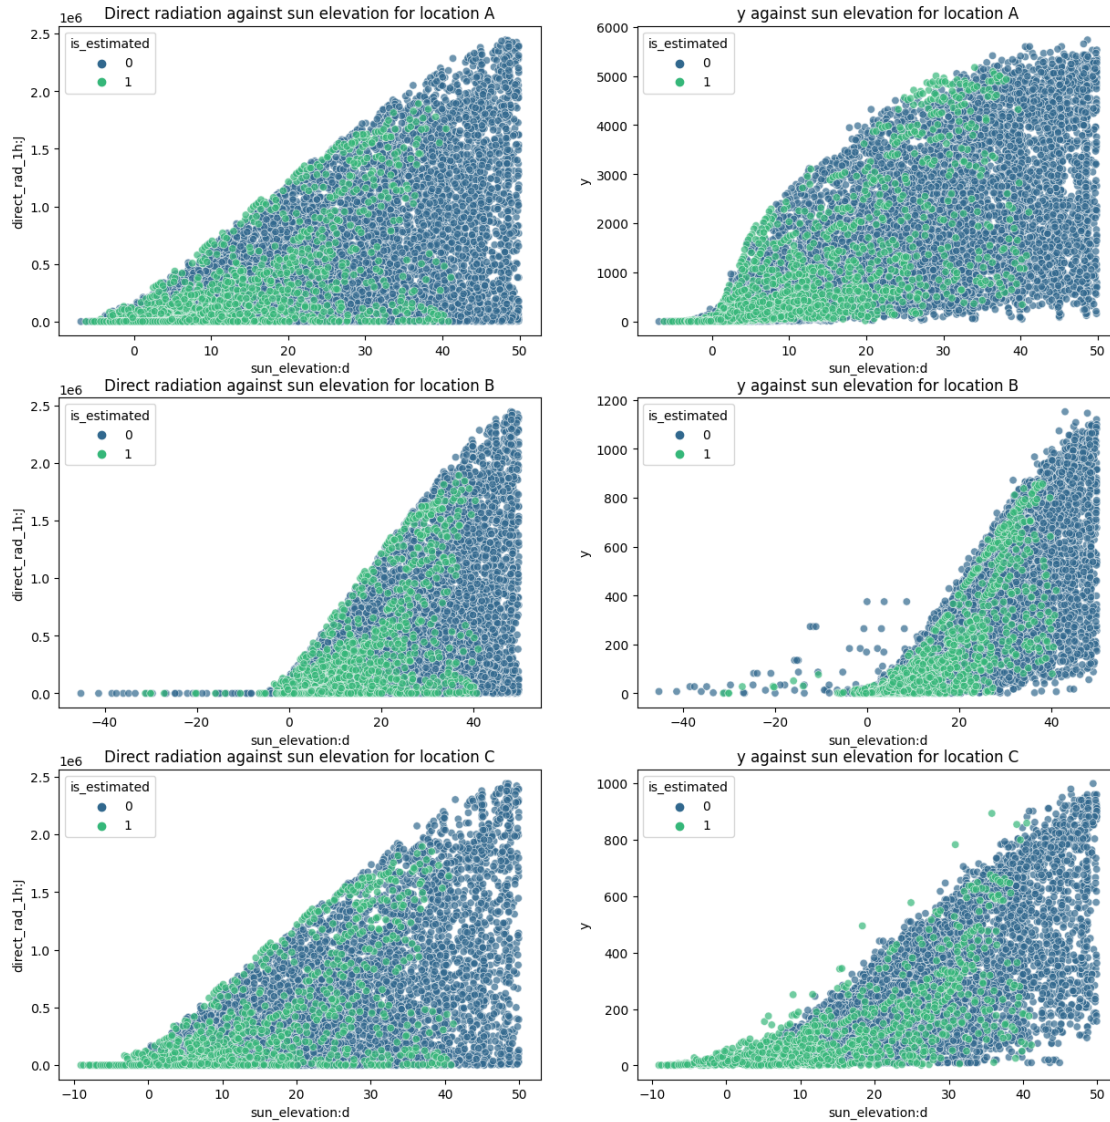
# Plotting
fig, axes = plt.subplots(len(locations), 2, figsize=(15, 5 * len(locations)))

for idx, location in enumerate(locations):
    sns.scatterplot(ax=axes[idx][0], data=temp[temp["location"] == location],
        x="sun_elevation:d", y="direct_rad_1h:J", hue="is_estimated",
        palette="viridis", alpha=0.7)
    axes[idx][0].set_title(f"Direct radiation against sun elevation for
        location {location}")

    sns.scatterplot(ax=axes[idx][1], data=temp[temp["location"] == location],
        x="sun_elevation:d", y="y", hue="is_estimated", palette="viridis", alpha=0.7)
    axes[idx][1].set_title(f"y against sun elevation for location {location}")

# plt.tight_layout()
# plt.show()

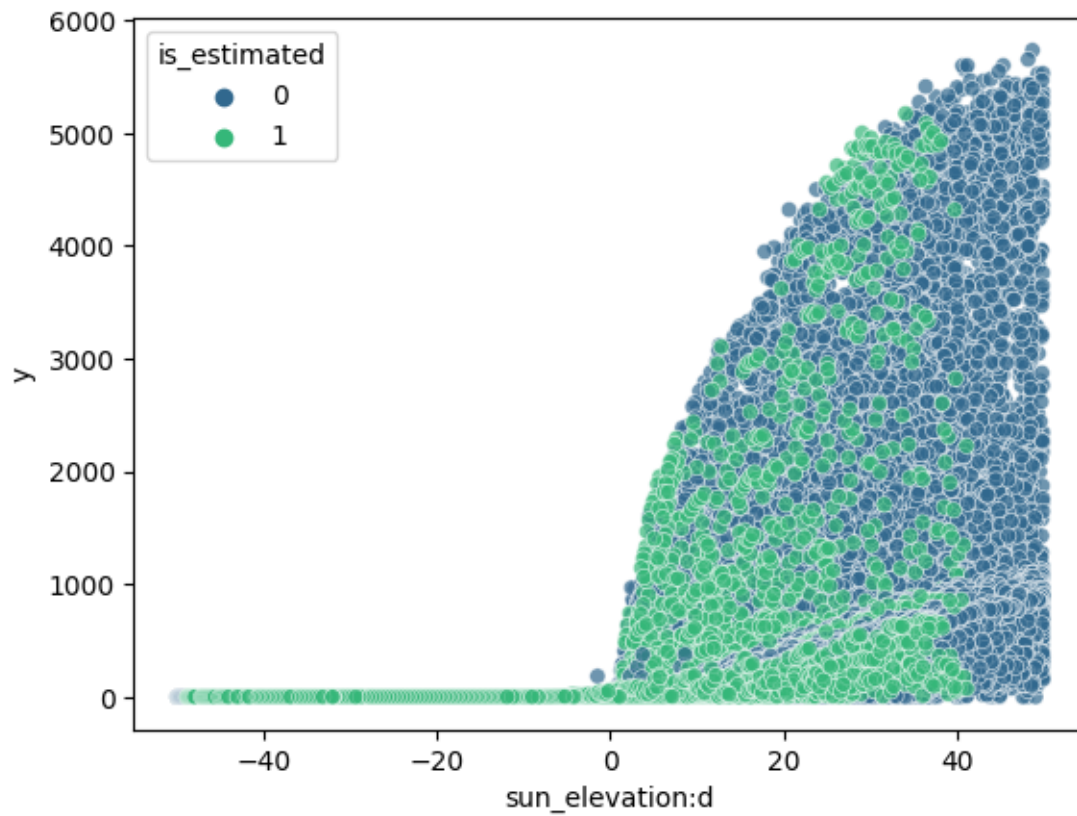
```



```
[6]: thresh = 0.1

# Update "y" values to NaN if they don't meet the criteria
mask = (X_train["direct_rad_1h:J"] <= thresh) & (X_train["diffuse_rad_1h:J"] <=
    ↪ thresh) & (X_train["y"] >= 0.1)
if drop_night_outliers:
    X_train.loc[mask, "y"] = np.nan

# Plot using sns scatterplot
sns.scatterplot(data=X_train, x="sun_elevation:d", y="y", hue="is_estimated",
    ↪ palette="viridis", alpha=0.7)
plt.show()
```




```
[7]: # location B count number of rows with y > 0 and sun_elevation:d < 0
```

```
condition = (X_train["location"] == "B") & (X_train["y"] > 0) &
↳ (X_train["sun_elevation:d"] < 0)
bad = X_train[condition]

bad.plot.scatter(x="sun_elevation:d", y="y")
```

```
[7]: <AxesSubplot: xlabel='sun_elevation:d', ylabel='y'>
```



```
[8]: # set y to nan where y is 0, but direct_rad_1h:J or diffuse_rad_1h:J are > 0
      ↪(or some threshold)
threshold_direct = X_train["direct_rad_1h:J"].max() * 0.01
threshold_diffuse = X_train["diffuse_rad_1h:J"].max() * 0.01
print(f"Threshold direct: {threshold_direct}")
print(f"Threshold diffuse: {threshold_diffuse}")

mask = (X_train["y"] == 0) & ((X_train["direct_rad_1h:J"] > threshold_direct) |
      ↪(X_train["diffuse_rad_1h:J"] > threshold_diffuse)) & (X_train["sun_elevation:
      ↪d"] > 0) & (X_train["fresh_snow_24h:cm"] < 6) & (X_train[['fresh_snow_12h:
      ↪cm', 'fresh_snow_1h:cm', 'fresh_snow_3h:cm', 'fresh_snow_6h:cm']]).
      ↪sum(axis=1) == 0)
print(len(X_train[mask]))

#print(X_train[mask][[x for x in X_train.columns if "snow" in x]])

# show plot where mask is true
#sns.scatterplot(data=X_train[mask], x="sun_elevation:d", y="y",
      ↪hue="is_estimated", palette="viridis", alpha=0.7)
```

```

sns.scatterplot(data=X_train[mask], x="sun_elevation:d", y="fresh_snow_24h:cm",
    hue="is_estimated", palette="viridis", alpha=0.7)
plt.show()

#sns.scatterplot(data=X_train[mask], x="fresh_snow_24h:cm",
    hue="is_estimated", palette="viridis", alpha=0.7)
    y="total_cloud_cover:p",


# set y to nan where mask
if drop_null_outliers:
    X_train.loc[mask, "y"] = np.nan

# show how many rows for each location, and for estimated and not estimated
X_train[mask].groupby(["location", "is_estimated"]).count()["direct_rad_1h:J"]

```

Threshold direct: 24458.97

Threshold diffuse: 11822.505000000001
2599



```
[8]: location  is_estimated
A          0           87
      1          10
B          0        1250
      1          32
C          0        1174
      1          46
Name: direct_rad_1h:J, dtype: int64
```

```
[9]: # print num rows
temprows = len(X_train)
X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
               inplace=True)
print("Dropped rows: ", temprows - len(X_train))
```

Dropped rows: 1876

2.1.2 Other stuff

```
[10]: import numpy as np
import pandas as pd

for attr in use_dt_attrs:
    X_train[attr] = getattr(X_train.index, attr)
    X_test[attr] = getattr(X_test.index, attr)

#print(X_train.head())

# If the "sample_weight" column is present and weight_evaluation is True,
# multiply sample_weight with sample_weight_may_july if the ds is between
# 05-01 00:00:00 and 07-03 23:00:00, else add sample_weight as a column to
# X_train
if weight_evaluation:
    if "sample_weight" not in X_train.columns:
        X_train["sample_weight"] = 1

    X_train.loc[((X_train.index.month >= 5) & (X_train.index.month <= 6)) |
               ((X_train.index.month == 7) & (X_train.index.day <= 3)), "sample_weight"] *=
               sample_weight_may_july

print(X_train.iloc[200])
print(X_train[((X_train.index.month >= 5) & (X_train.index.month <= 6)) |
              ((X_train.index.month == 7) & (X_train.index.day <= 3))].head(1))
```

```

if use_groups:
    # fix groups for cross validation
    locations = X_train['location'].unique() # Assuming 'location' is the name
    ↪ of the column representing locations

    grouped_dfs = [] # To store data frames split by location

    # Loop through each unique location
    for loc in locations:
        loc_df = X_train[X_train['location'] == loc]

        # Sort the DataFrame for this location by the time column
        loc_df = loc_df.sort_index()

        # Calculate the size of each group for this location
        group_size = len(loc_df) // n_groups

        # Create a new 'group' column for this location
        loc_df['group'] = np.repeat(range(n_groups),
    ↪ repeats=[group_size]*(n_groups-1) + [len(loc_df) - group_size*(n_groups-1)])

        # Append to list of grouped DataFrames
        grouped_dfs.append(loc_df)

    # Concatenate all the grouped DataFrames back together
    X_train = pd.concat(grouped_dfs)
    X_train.sort_index(inplace=True)
    print(X_train["group"].head())

X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)

```

```

absolute_humidity_2m:gm3          7.625
air_density_2m:kgm3              1.2215
ceiling_height_agl:m             3644.050049
clear_sky_energy_1h:J            2896336.75
clear_sky_rad:W                  753.849976
cloud_base_agl:m                 3644.050049
dew_or_rime:idx                  0.0

```

dew_point_2m:K	280.475006
diffuse_rad:W	127.475006
diffuse_rad_1h:J	526032.625
direct_rad:W	488.0
direct_rad_1h:J	1718048.625
effective_cloud_cover:p	18.200001
elevation:m	6.0
fresh_snow_12h:cm	0.0
fresh_snow_1h:cm	0.0
fresh_snow_24h:cm	0.0
fresh_snow_3h:cm	0.0
fresh_snow_6h:cm	0.0
is_day:idx	1.0
is_in_shadow:idx	0.0
msl_pressure:hPa	1026.775024
precip_5min:mm	0.0
precip_type_5min:idx	0.0
pressure_100m:hPa	1013.599976
pressure_50m:hPa	1019.599976
prob_rime:p	0.0
rain_water:kgm2	0.0
relative_humidity_1000hPa:p	53.825001
sfc_pressure:hPa	1025.699951
snow_density:kgm3	NaN
snow_depth:cm	0.0
snow_drift:idx	0.0
snow_melt_10min:mm	0.0
snow_water:kgm2	0.0
sun_azimuth:d	222.089005
sun_elevation:d	44.503498
super_cooled_liquid_water:kgm2	0.0
t_1000hPa:K	286.700012
total_cloud_cover:p	18.200001
visibility:m	52329.25
wind_speed_10m:ms	2.6
wind_speed_u_10m:ms	-1.9
wind_speed_v_10m:ms	-1.75
wind_speed_w_1000hPa:ms	0.0
is_estimated	0
y	4367.44
location	A
Name: 2019-06-11 13:00:00, dtype: object	
absolute_humidity_2m:gm3 air_density_2m:kgm3 \	
ds	
2019-06-02 23:00:00	7.7 1.2235
ceiling_height_agl:m clear_sky_energy_1h:J \	
ds	


```

2019-06-02 23:00:00          1689.824951          0.0

          clear_sky_rad:W  cloud_base_agl:m  dew_or_rime:idx  \
ds
2019-06-02 23:00:00          0.0          1689.824951          0.0

          dew_point_2m:K  diffuse_rad:W  diffuse_rad_1h:J  ...  \
ds
2019-06-02 23:00:00          280.299988          0.0          0.0  ...

          t_1000hPa:K  total_cloud_cover:p  visibility:m  \
ds
2019-06-02 23:00:00          286.899994          100.0  33770.648438

          wind_speed_10m:ms  wind_speed_u_10m:ms  \
ds
2019-06-02 23:00:00          3.35          -3.35

          wind_speed_v_10m:ms  wind_speed_w_1000hPa:ms  \
ds
2019-06-02 23:00:00          0.275          0.0

          is_estimated    y  location
ds
2019-06-02 23:00:00          0  0.0          A

[1 rows x 48 columns]

```

```

[11]: # Create a plot of X_train showing its "y" and color it based on the value of
      ↪ the sample_weight column.
      if "sample_weight" in X_train.columns:
          import matplotlib.pyplot as plt
          import seaborn as sns
          sns.scatterplot(data=X_train, x=X_train.index, y="y", hue="sample_weight",
          ↪ palette="deep", size=3)
          plt.show()

```

```

[12]: def normalize_sample_weights_per_location(df):
      for loc in locations:
          loc_df = df[df["location"] == loc]
          loc_df["sample_weight"] = loc_df["sample_weight"] /
          ↪ loc_df["sample_weight"].sum() * loc_df.shape[0]
          df[df["location"] == loc] = loc_df
      return df

import pandas as pd

```

```

def split_and_shuffle_data(input_data, num_bins, frac1):
    """
    Splits the input_data into num_bins and shuffles them, then divides the
    ↪bins into two datasets based on the given fraction for the first set.

    Args:
        input_data (pd.DataFrame): The data to be split and shuffled.
        num_bins (int): The number of bins to split the data into.
        frac1 (float): The fraction of each bin to go into the first output
        ↪dataset.

    Returns:
        pd.DataFrame, pd.DataFrame: The two output datasets.
    """
    # Validate the input fraction
    if frac1 < 0 or frac1 > 1:
        raise ValueError("frac1 must be between 0 and 1.")

    if frac1==1:
        return input_data, pd.DataFrame()

    # Calculate the fraction for the second output set
    frac2 = 1 - frac1

    # Calculate bin size
    bin_size = len(input_data) // num_bins

    # Initialize empty DataFrames for output
    output_data1 = pd.DataFrame()
    output_data2 = pd.DataFrame()

    for i in range(num_bins):
        # Shuffle the data in the current bin
        np.random.seed(i)
        current_bin = input_data.iloc[i * bin_size: (i + 1) * bin_size].
        ↪sample(frac=1)

        # Calculate the sizes for each output set
        size1 = int(len(current_bin) * frac1)

        # Split and append to output DataFrames
        output_data1 = pd.concat([output_data1, current_bin.iloc[:size1]])
        output_data2 = pd.concat([output_data2, current_bin.iloc[size1:]]

    # Shuffle and split the remaining data
    remaining_data = input_data.iloc[num_bins * bin_size:].sample(frac=1)

```

```

    remaining_size1 = int(len(remaining_data) * frac1)

    output_data1 = pd.concat([output_data1, remaining_data.iloc[:
↪remaining_size1]])
    output_data2 = pd.concat([output_data2, remaining_data.iloc[remaining_size1:
↪]])

    return output_data1, output_data2

```

```

[13]: from autogluon.tabular import TabularDataset, TabularPredictor
data = TabularDataset('X_train_raw.csv')
# set group column of train_data be increasing from 0 to 7 based on time, the
↪first 1/8 of the data is group 0, the second 1/8 of the data is group 1, etc.
data['ds'] = pd.to_datetime(data['ds'])
data = data.sort_values(by='ds')

# # print size of the group for each location
# for loc in locations:
#     print(f"Location {loc}:")
#     print(train_data[train_data["location"] == loc].groupby('group').size())

# get end date of train data and subtract 3 months
#split_time = pd.to_datetime(train_data["ds"]).max() - pd.
↪Timedelta(hours=tune_and_test_length)
# 2022-10-28 22:00:00
split_time = pd.to_datetime("2022-10-28 22:00:00")
train_set = TabularDataset(data[data["ds"] < split_time])
estimated_set = TabularDataset(data[data["ds"] >= split_time]) # only estimated

test_set = pd.DataFrame()
tune_set = pd.DataFrame()
new_train_set = pd.DataFrame()

if not use_tune_data:
    raise Exception("Not implemented")

for location in locations:
    loc_data = data[data["location"] == location]
    num_train_rows = len(loc_data)

    tune_rows = 1500.0 # 2500.0
    if use_test_data:
        tune_rows = 1880.0#max(3000.0,
↪len(estimated_set[estimated_set["location"] == location]))

```

```

    holdout_frac = max(0.01, min(0.1, tune_rows / num_train_rows)) *
    ↪ num_train_rows / len(estimated_set[estimated_set["location"] == location])

    print(f"Size of estimated for location {location}:
    ↪ {len(estimated_set[estimated_set['location'] == location])}. Holdout frac
    ↪ should be % of estimated: {holdout_frac}")

    # shuffle and split data
    loc_tune_set, loc_new_train_set =
    ↪ split_and_shuffle_data(estimated_set[estimated_set['location'] == location],
    ↪ 40, holdout_frac)
    print(f"Length of location tune set : {len(loc_tune_set)}")
    new_train_set = pd.concat([new_train_set, loc_new_train_set])

    if use_test_data:
        loc_test_set, loc_tune_set = split_and_shuffle_data(loc_tune_set, 40, 0.
    ↪ 2)
        test_set = pd.concat([test_set, loc_test_set])

    tune_set = pd.concat([tune_set, loc_tune_set])

print("Length of train set before adding test set", len(train_set))
# add rest to train_set
train_set = pd.concat([train_set, new_train_set])
print("Length of train set after adding test set", len(train_set))

if use_groups:
    test_set = test_set.drop(columns=['group'])

tuning_data = tune_set

# number of rows in tuning data for each location
print("Shapes of tuning data", tuning_data.groupby('location').size())

if use_test_data:
    test_data = test_set
    print("Shape of test", test_data.shape[0])

```

```

train_data = train_set

# ensure sample weights for your training (or tuning) data sum to the number of
↳rows in the training (or tuning) data.
if weight_evaluation:
    # ensure sample weights for data sum to the number of rows in the tuning /
    ↳train data.
    tuning_data = normalize_sample_weights_per_location(tuning_data)
    train_data = normalize_sample_weights_per_location(train_data)
    if use_test_data:
        test_data = normalize_sample_weights_per_location(test_data)

train_data = TabularDataset(train_data)
tuning_data = TabularDataset(tuning_data)

if use_test_data:
    test_data = TabularDataset(test_data)

```

```

Size of estimated for location A: 4214. Holdout frac should be % of estimated:
0.4461319411485524
Length of location tune set : 1846
Size of estimated for location B: 3533. Holdout frac should be % of estimated:
0.5321256722332296
Length of location tune set : 1846
Size of estimated for location C: 2923. Holdout frac should be % of estimated:
0.6431748203900103
Length of location tune set : 1841
Length of train set before adding test set 77247
Length of train set after adding test set 82384
Shapes of tuning data location
A    1485
B    1485
C    1481
dtype: int64
Shape of test 1082

```

3 Quick EDA

```

[14]: if run_analysis:
        import autogluon.eda.auto as auto
        auto.dataset_overview(train_data=train_data, test_data=test_data,
        ↳label="y", sample=None)

```

```

[15]: if run_analysis:
        auto.target_analysis(train_data=train_data, label="y", sample=None)

```

4 Modeling

```
[16]: import os

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    ↪filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)
print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

hello = os.environ.get('HELLO')
if hello is not None:
    new_filename += f'_{hello}'

print("New filename:", new_filename)
```

```
Last submission number: 121
Now creating submission number: 122
New filename: submission_122
```

```
[17]: predictors = [None, None, None]
```

```
[18]: def fit_predictor_for_location(loc):
    print(f"Training model for location {loc}...")
    # sum of sample weights for this location, and number of rows, for both
    ↪train and tune data and test data
    if weight_evaluation:
        print("Train data sample weight sum:",
            ↪train_data[train_data["location"] == loc]["sample_weight"].sum())
        print("Train data number of rows:", train_data[train_data["location"]
            ↪== loc].shape[0])
        if use_tune_data:
            print("Tune data sample weight sum:",
                ↪tuning_data[tuning_data["location"] == loc]["sample_weight"].sum())
            print("Tune data number of rows:",
                ↪tuning_data[tuning_data["location"] == loc].shape[0])
        if use_test_data:
            print("Test data sample weight sum:",
                ↪test_data[test_data["location"] == loc]["sample_weight"].sum())
            print("Test data number of rows:", test_data[test_data["location"]
                ↪== loc].shape[0])
        predictor = TabularPredictor(
            label=label,
```

```

        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
        # sample_weight=sample_weight,
        # weight_evaluation=weight_evaluation,
        # groups="group" if use_groups else None,
    ).fit(
        train_data=train_data[train_data["location"] == loc].
↳drop(columns=["ds"]),
        time_limit=time_limit,
        presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds if not use_groups else 2, # just put
↳somethin, will be overwritten anyways
        num_bag_sets=num_bag_sets,
        tuning_data=tuning_data[tuning_data["location"] == loc].
↳reset_index(drop=True).drop(columns=["ds"]) if use_tune_data else None,
        use_bag_holdout=use_bag_holdout,
        # holdout_frac=holdout_frac,
        excluded_model_types=excluded_model_types
    )

    # evaluate on test data
    if use_test_data:
        # drop sample_weight column
        t = test_data[test_data["location"] == loc]#.
↳drop(columns=["sample_weight"])
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Warning: path already exists! This predictor may overwrite an existing predictor! path="AutogluonModels/submission_122_A"

Presets specified: ['experimental_zeroshot_hpo_hybrid']

Stack configuration (auto_stack=True): num_stack_levels=0, num_bag_folds=8, num_bag_sets=20

Beginning AutoGluon training ... Time limit = 3600s

AutoGluon will save models to "AutogluonModels/submission_122_A/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Linux

Platform Machine: x86_64

Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)

```

Disk Space Avail: 155.32 GB / 315.93 GB (49.2%)
Train Data Rows: 30934
Train Data Columns: 44
Tuning Data Rows: 1485
Tuning Data Columns: 44
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (5733.42, 0.0, 673.41535, 1195.24)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 132339.61 MB
    Train Data (Original) Memory Usage: 13.03 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 2 features to boolean dtype as they
only contain 2 unique values.

Training model for location A...

    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 3): ['elevation:m', 'snow_drift:idx',
'location']
        These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',

```



```

'clear_sky_rad:W', ...]
      ('int', ['bool']) : 2 | ['snow_density:kgm3', 'is_estimated']
0.2s = Fit runtime
41 features in original data used to generate 41 features in processed
data.
    Train Data (Processed) Memory Usage: 10.18 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.18s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
use_bag_holdout=True, will use tuning_data as holdout (will not be used for
early stopping).
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}, {'min_samples_leaf': 1,
'max_leaf_nodes': 15000, 'max_features': 0.5, 'ag_args': {'name_suffix': '_r19',
'priority': 20}}],
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}, {'min_samples_leaf': 5,
'max_leaf_nodes': 50000, 'max_features': 0.5, 'ag_args': {'name_suffix': '_r5',
'priority': 19}}],
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {},
'GBMLarge', {'extra_trees': False, 'feature_fraction': 0.7248284762542815,
'learning_rate': 0.07947286942946127, 'min_data_in_leaf': 50, 'num_leaves': 89,
'ag_args': {'name_suffix': '_r158', 'priority': 18}}, {'extra_trees': True,
'feature_fraction': 0.7832570544199176, 'learning_rate': 0.021720607471727896,
'min_data_in_leaf': 3, 'num_leaves': 21, 'ag_args': {'name_suffix': '_r118',
'priority': 17}}, {'extra_trees': True, 'feature_fraction': 0.7113010892989156,
'learning_rate': 0.012535427424259274, 'min_data_in_leaf': 16, 'num_leaves': 48,
'ag_args': {'name_suffix': '_r97', 'priority': 16}}, {'extra_trees': True,
'feature_fraction': 0.45555769907110816, 'learning_rate': 0.009591347321206594,
'min_data_in_leaf': 50, 'num_leaves': 110, 'ag_args': {'name_suffix': '_r71',
'priority': 15}}, {'extra_trees': False, 'feature_fraction':
0.40979710161022476, 'learning_rate': 0.008708890211023034, 'min_data_in_leaf':
3, 'num_leaves': 80, 'ag_args': {'name_suffix': '_r111', 'priority': 14}}],
    'XGB': {},
    'FASTAI': [{}, {'bs': 1024, 'emb_drop': 0.6167722379778131, 'epochs':
44, 'layers': [200, 100, 50], 'lr': 0.053440377855629266, 'ps':

```

```

0.48477211305443607, 'ag_args': {'name_suffix': '_r25', 'priority': 13}}, {'bs':
1024, 'emb_drop': 0.6046989241462619, 'epochs': 48, 'layers': [200, 100, 50],
'lr': 0.00775309042164966, 'ps': 0.09244767444160731, 'ag_args': {'name_suffix':
'_r51', 'priority': 12}}, {'bs': 512, 'emb_drop': 0.6557225316526186, 'epochs':
49, 'layers': [200, 100], 'lr': 0.023627682025564638, 'ps': 0.519566584552178,
'ag_args': {'name_suffix': '_r82', 'priority': 11}}, {'bs': 2048, 'emb_drop':
0.4066210919034579, 'epochs': 43, 'layers': [400, 200], 'lr':
0.0029598312717673434, 'ps': 0.4378695797438974, 'ag_args': {'name_suffix':
'_r121', 'priority': 10}}, {'bs': 128, 'emb_drop': 0.44339037504795686,
'epochs': 31, 'layers': [400, 200, 100], 'lr': 0.008615195908919904, 'ps':
0.19220253419114286, 'ag_args': {'name_suffix': '_r145', 'priority': 9}}, {'bs':
128, 'emb_drop': 0.12106594798980945, 'epochs': 38, 'layers': [200, 100, 50],
'lr': 0.037991970245029975, 'ps': 0.33120008492595093, 'ag_args':
{'name_suffix': '_r173', 'priority': 8}}, {'bs': 128, 'emb_drop':
0.4599138419358, 'epochs': 47, 'layers': [200, 100], 'lr': 0.03888383281136287,
'ps': 0.28193673177122863, 'ag_args': {'name_suffix': '_r128', 'priority': 7}}],
    'CAT': [{}, {'depth': 5, 'l2_leaf_reg': 4.774992314058497,
'learning_rate': 0.038551267822920274, 'ag_args': {'name_suffix': '_r16',
'priority': 6}}, {'depth': 4, 'l2_leaf_reg': 1.9950125740798321,
'learning_rate': 0.028091050379971633, 'ag_args': {'name_suffix': '_r42',
'priority': 5}}, {'depth': 6, 'l2_leaf_reg': 1.8298803017644376,
'learning_rate': 0.017844259810823604, 'ag_args': {'name_suffix': '_r93',
'priority': 4}}, {'depth': 7, 'l2_leaf_reg': 4.81099604606794, 'learning_rate':
0.019085060180573103, 'ag_args': {'name_suffix': '_r44', 'priority': 3}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
Excluded models: ['RF', 'XGB', 'CAT'] (Specified by `excluded_model_types`)
Fitting 21 L1 models ...
Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 3599.82s of
the 3599.82s of remaining time.
    -191.231          = Validation score    (-mean_absolute_error)
    0.04s           = Training    runtime
    1.25s           = Validation runtime
Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 3598.25s of
the 3598.25s of remaining time.
    -192.9182         = Validation score    (-mean_absolute_error)
    0.04s            = Training    runtime
    0.4s             = Validation runtime
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 3597.65s of the
3597.65s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -85.9426          = Validation score    (-mean_absolute_error)
    33.25s           = Training    runtime
    14.42s           = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 3554.93s of the
3554.93s of remaining time.

```

```

    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -90.5139      = Validation score    (-mean_absolute_error)
    30.39s       = Training    runtime
    7.49s        = Validation runtime
Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 3520.37s of the
3520.37s of remaining time.
    -102.5531     = Validation score    (-mean_absolute_error)
    1.81s         = Training    runtime
    1.1s          = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 3515.25s of
the 3515.24s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -103.8453     = Validation score    (-mean_absolute_error)
    38.06s        = Training    runtime
    0.49s         = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 3474.52s of
the 3474.52s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -88.0553      = Validation score    (-mean_absolute_error)
    99.41s        = Training    runtime
    0.36s         = Validation runtime
Fitting model: ExtraTrees_r19_BAG_L1 ... Training model for up to 3373.83s of
the 3373.83s of remaining time.
    -100.9585     = Validation score    (-mean_absolute_error)
    1.13s         = Training    runtime
    1.11s         = Validation runtime
Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 3368.96s of the
3368.96s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -92.787       = Validation score    (-mean_absolute_error)
    49.14s        = Training    runtime
    12.84s        = Validation runtime
Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 3311.47s of the
3311.46s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -85.8912      = Validation score    (-mean_absolute_error)
    21.26s        = Training    runtime
    18.18s        = Validation runtime
Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 3285.14s of the
3285.14s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -86.3739      = Validation score    (-mean_absolute_error)

```

```

34.3s    = Training    runtime
25.28s   = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 3243.85s of the
3243.84s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -92.7495          = Validation score    (-mean_absolute_error)
    82.06s    = Training    runtime
    49.4s     = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 3148.57s of the
3148.57s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -88.6662          = Validation score    (-mean_absolute_error)
    55.43s    = Training    runtime
    26.39s    = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 3082.63s
of the 3082.63s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -104.2684         = Validation score    (-mean_absolute_error)
    23.48s    = Training    runtime
    0.24s     = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 3057.8s
of the 3057.79s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -97.1141          = Validation score    (-mean_absolute_error)
    25.58s    = Training    runtime
    0.24s     = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 3030.67s
of the 3030.66s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -106.6312         = Validation score    (-mean_absolute_error)
    35.99s    = Training    runtime
    0.31s     = Validation runtime
Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 2992.84s
of the 2992.84s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -117.7474         = Validation score    (-mean_absolute_error)
    20.56s    = Training    runtime
    0.22s     = Validation runtime
Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 2970.6s
of the 2970.59s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

-96.1916          = Validation score    (-mean_absolute_error)
88.24s    = Training    runtime
0.89s     = Validation runtime
Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 2880.42s
of the 2880.41s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-100.9066         = Validation score    (-mean_absolute_error)
93.2s    = Training    runtime
0.84s     = Validation runtime
Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 2785.23s
of the 2785.22s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-102.5638         = Validation score    (-mean_absolute_error)
105.23s = Training    runtime
0.8s     = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 2678.0s of the
2678.0s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-87.4068          = Validation score    (-mean_absolute_error)
107.91s = Training    runtime
24.4s    = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 2560.22s of the
2560.21s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-86.1803          = Validation score    (-mean_absolute_error)
67.87s    = Training    runtime
33.59s    = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 2519.62s of the
2519.62s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-90.3028          = Validation score    (-mean_absolute_error)
63.1s    = Training    runtime
14.2s    = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 2482.13s of
the 2482.13s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-103.0736         = Validation score    (-mean_absolute_error)
76.68s    = Training    runtime
1.05s     = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 2441.65s of
the 2441.65s of remaining time.

```

```

    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -86.6093          = Validation score    (-mean_absolute_error)
    208.19s          = Training    runtime
    0.73s            = Validation runtime
Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 2330.96s of the
2330.95s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -91.3734          = Validation score    (-mean_absolute_error)
    92.47s           = Training    runtime
    19.27s           = Validation runtime
Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 2280.33s of the
2280.32s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -85.5189          = Validation score    (-mean_absolute_error)
    45.19s           = Training    runtime
    27.78s           = Validation runtime
Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 2251.74s of the
2251.74s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -86.5744          = Validation score    (-mean_absolute_error)
    73.55s           = Training    runtime
    47.53s           = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 2205.36s of the
2205.36s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -92.7942          = Validation score    (-mean_absolute_error)
    160.78s          = Training    runtime
    98.26s           = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 2108.48s of the
2108.48s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -88.7043          = Validation score    (-mean_absolute_error)
    111.6s           = Training    runtime
    48.38s           = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 2041.96s
of the 2041.96s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -104.3129         = Validation score    (-mean_absolute_error)
    47.03s           = Training    runtime
    0.47s            = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 2016.79s

```

of the 2016.79s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-95.9834 = Validation score (-mean_absolute_error)
50.65s = Training runtime
0.46s = Validation runtime

Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 1989.9s
of the 1989.9s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-106.5279 = Validation score (-mean_absolute_error)
72.81s = Training runtime
0.63s = Validation runtime

Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 1951.35s
of the 1951.34s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-118.3212 = Validation score (-mean_absolute_error)
40.71s = Training runtime
0.46s = Validation runtime

Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 1929.69s
of the 1929.69s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-95.9189 = Validation score (-mean_absolute_error)
177.31s = Training runtime
1.79s = Validation runtime

Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 1838.01s
of the 1838.0s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-99.3103 = Validation score (-mean_absolute_error)
187.28s = Training runtime
1.67s = Validation runtime

Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 1741.46s
of the 1741.46s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-104.2124 = Validation score (-mean_absolute_error)
207.72s = Training runtime
1.6s = Validation runtime

Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1636.48s of the
1636.48s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-87.1874 = Validation score (-mean_absolute_error)
218.28s = Training runtime
47.47s = Validation runtime

Repeating k-fold bagging: 3/20

Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1512.34s of the 1512.34s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-86.1411 = Validation score (-mean_absolute_error)

102.14s = Training runtime

47.86s = Validation runtime

Fitting model: LightGBM_BAG_L1 ... Training model for up to 1471.25s of the 1471.25s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-90.3661 = Validation score (-mean_absolute_error)

93.52s = Training runtime

20.03s = Validation runtime

Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1435.42s of the 1435.41s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-102.877 = Validation score (-mean_absolute_error)

116.2s = Training runtime

1.54s = Validation runtime

Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1393.7s of the 1393.69s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-86.9808 = Validation score (-mean_absolute_error)

319.28s = Training runtime

1.11s = Validation runtime

Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 1280.64s of the 1280.64s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-90.9068 = Validation score (-mean_absolute_error)

134.71s = Training runtime

28.03s = Validation runtime

Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 1228.1s of the 1228.1s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-85.2599 = Validation score (-mean_absolute_error)

68.73s = Training runtime

38.31s = Validation runtime

Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 1199.12s of the 1199.12s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-86.5317 = Validation score (-mean_absolute_error)


```

114.17s = Training runtime
71.39s = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 1150.13s of the
1150.13s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-92.6296 = Validation score (-mean_absolute_error)
238.69s = Training runtime
139.59s = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 1050.86s of the
1050.85s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-88.7444 = Validation score (-mean_absolute_error)
167.53s = Training runtime
71.01s = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 982.17s
of the 982.16s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-103.6838 = Validation score (-mean_absolute_error)
70.2s = Training runtime
0.71s = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 957.32s
of the 957.31s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-95.0394 = Validation score (-mean_absolute_error)
76.12s = Training runtime
0.7s = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 929.97s
of the 929.97s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-107.3447 = Validation score (-mean_absolute_error)
108.83s = Training runtime
0.96s = Validation runtime
Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 891.9s
of the 891.89s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-119.2008 = Validation score (-mean_absolute_error)
60.77s = Training runtime
0.69s = Validation runtime
Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 870.22s
of the 870.21s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

-95.1641          = Validation score    (-mean_absolute_error)
265.62s = Training    runtime
2.7s     = Validation runtime
Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 778.8s
of the 778.8s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -98.4686          = Validation score    (-mean_absolute_error)
    280.87s = Training    runtime
    2.48s     = Validation runtime
Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 682.13s
of the 682.13s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -103.41 = Validation score    (-mean_absolute_error)
    311.95s = Training    runtime
    2.45s     = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 574.81s of the
574.81s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -87.2987          = Validation score    (-mean_absolute_error)
    326.89s = Training    runtime
    70.2s     = Validation runtime
Completed 3/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
448.79s of remaining time.
    -82.2346          = Validation score    (-mean_absolute_error)
    0.53s     = Training    runtime
    0.0s     = Validation runtime
AutoGluon training complete, total runtime = 3151.77s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_122_A/")
Evaluation: mean_absolute_error on test data: -105.21386474786952
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -105.21386474786952,
    "root_mean_squared_error": -334.6204247772324,
    "mean_squared_error": -111970.82867809547,
    "r2": 0.8242712281113946,
    "pearsonr": 0.9122558218574904,
    "median_absolute_error": -2.772814989089966
}
Evaluation on test data:

```

-105.21386474786952


```
[19]: import matplotlib.pyplot as plt
leaderboards = [None, None, None]
def leaderboard_for_location(i, loc):
    if use_tune_data:
        plt.scatter(train_data[(train_data["location"] == loc) &
        ↪(train_data["is_estimated"]==True)]["y"].index,
        ↪train_data[(train_data["location"] == loc) &
        ↪(train_data["is_estimated"]==True)]["y"])
        plt.scatter(tuning_data[tuning_data["location"] == loc]["y"].index,
        ↪tuning_data[tuning_data["location"] == loc]["y"])
        plt.title("Val and Train")
        plt.show()

    if use_test_data:
        lb = predictors[i].leaderboard(test_data[test_data["location"] ==
        ↪loc])
        lb["location"] = loc
        plt.scatter(test_data[test_data["location"] == loc]["y"].index,
        ↪test_data[test_data["location"] == loc]["y"])
        plt.title("Test")

    return lb


return pd.DataFrame()

leaderboards[0] = leaderboard_for_location(0, loc)
```



	model	score_test	score_val	pred_time_test
0	WeightedEnsemble_L2	-105.213865	-82.234560	15.416068
158.172757	893.704767	0.004306		0.000686
0.533880	2	True	22	
1	LightGBMXT_BAG_L1	-105.651709	-86.141116	2.763916
47.861973	102.143539	2.763916		47.861973
102.143539	1	True	3	
2	NeuralNetFastAI_r51_BAG_L1	-106.655413	-95.039358	0.376366
0.700547	76.118525	0.376366		0.700547
76.118525	1	True	15	
3	LightGBM_r118_BAG_L1	-106.870157	-85.259908	2.344336
38.308059	68.731673	2.344336		38.308059
68.731673	1	True	10	
4	LightGBM_r97_BAG_L1	-107.586386	-86.531742	4.333927
71.390422	114.166910	4.333927		71.390422
114.166910	1	True	11	
5	NeuralNetTorch_BAG_L1	-112.454177	-86.980764	0.569817
1.105011	319.282644	0.569817		1.105011
319.282644	1	True	7	

6	LightGBM_r71_BAG_L1	-114.721642	-92.629558	8.428360
139.588237	238.691140		8.428360	139.588237
238.691140	1	True	12	
7	LightGBM_r111_BAG_L1	-115.662826	-88.744429	6.469665
71.013778	167.530970		6.469665	71.013778
167.530970	1	True	13	
8	NeuralNetFastAI_r173_BAG_L1	-116.151949	-98.468557	0.874867
2.481646	280.872100		0.874867	2.481646
280.872100	1	True	19	
9	LightGBMLarge_BAG_L1	-117.513573	-87.298706	9.357328
70.196481	326.894507		9.357328	70.196481
326.894507	1	True	21	
10	NeuralNetFastAI_r145_BAG_L1	-117.535234	-95.164051	0.815387
2.704688	265.617656		0.815387	2.704688
265.617656	1	True	18	
11	LightGBM_BAG_L1	-118.126742	-90.366088	2.076667
20.028206	93.519575		2.076667	20.028206
93.519575	1	True	4	
12	NeuralNetFastAI_BAG_L1	-120.371544	-102.877024	0.527273
1.540404	116.199956		0.527273	1.540404
116.199956	1	True	6	
13	NeuralNetFastAI_r25_BAG_L1	-120.402365	-103.683760	0.377176
0.710119	70.197967		0.377176	0.710119
70.197967	1	True	14	
14	LightGBM_r158_BAG_L1	-120.484917	-90.906798	2.837722
28.030752	134.713080		2.837722	28.030752
134.713080	1	True	9	
15	NeuralNetFastAI_r128_BAG_L1	-122.006312	-103.410039	1.118670
2.450946	311.951262		1.118670	2.450946
311.951262	1	True	20	
16	ExtraTrees_r19_BAG_L1	-126.033456	-100.958547	0.556135
1.114801	1.126558		0.556135	1.114801
1.126558	1	True	8	
17	NeuralNetFastAI_r82_BAG_L1	-127.562204	-107.344693	0.419232
0.961000	108.827765		0.419232	0.961000
108.827765	1	True	16	
18	ExtraTreesMSE_BAG_L1	-130.386831	-102.553116	0.557890
1.103508	1.808400		0.557890	1.103508
1.808400	1	True	5	
19	NeuralNetFastAI_r121_BAG_L1	-139.844186	-119.200840	0.314491
0.688419	60.771383		0.314491	0.688419
60.771383	1	True	17	
20	KNeighborsDist_BAG_L1	-189.567130	-192.918160	0.013305
0.397008	0.035783		0.013305	0.397008
0.035783	1	True	2	
21	KNeighborsUnif_BAG_L1	-191.283846	-191.231007	0.143528
1.246812	0.037004		0.143528	1.246812
0.037004	1	True	1	



```
[20]: loc = "B"
predictors[1] = fit_predictor_for_location(loc)
leaderboards[1] = leaderboard_for_location(1, loc)
```

```
Presets specified: ['experimental_zeroshot_hpo_hybrid']
Stack configuration (auto_stack=True): num_stack_levels=0, num_bag_folds=8,
num_bag_sets=20
Beginning AutoGluon training ... Time limit = 3600s
AutoGluon will save models to "AutogluonModels/submission_122_B/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 147.52 GB / 315.93 GB (46.7%)
Train Data Rows: 27377
Train Data Columns: 44
Tuning Data Rows: 1485
Tuning Data Columns: 44
Label Column: y
Preprocessing data ...
```

```

AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (1152.3, -0.0, 98.11625, 206.48535)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...

Training model for location B...

Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 130392.4 MB
    Train Data (Original) Memory Usage: 11.6 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 2 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 41 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', ['bool']) : 2 | ['snow_density:kgm3', 'is_estimated']
    0.2s = Fit runtime
    42 features in original data used to generate 42 features in processed
data.
    Train Data (Processed) Memory Usage: 9.29 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.22s ...
AutoGluon will gauge predictive performance using evaluation metric:

```

'mean_absolute_error'

This metric's sign has been flipped to adhere to being higher_is_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor() use_bag_holdout=True, will use tuning_data as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}, {'min_samples_leaf': 1,
'max_leaf_nodes': 15000, 'max_features': 0.5, 'ag_args': {'name_suffix': '_r19',
'priority': 20}}],
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}, {'min_samples_leaf': 5,
'max_leaf_nodes': 50000, 'max_features': 0.5, 'ag_args': {'name_suffix': '_r5',
'priority': 19}}],
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {},
'GBMLarge', {'extra_trees': False, 'feature_fraction': 0.7248284762542815,
'learning_rate': 0.07947286942946127, 'min_data_in_leaf': 50, 'num_leaves': 89,
'ag_args': {'name_suffix': '_r158', 'priority': 18}}, {'extra_trees': True,
'feature_fraction': 0.7832570544199176, 'learning_rate': 0.021720607471727896,
'min_data_in_leaf': 3, 'num_leaves': 21, 'ag_args': {'name_suffix': '_r118',
'priority': 17}}, {'extra_trees': True, 'feature_fraction': 0.7113010892989156,
'learning_rate': 0.012535427424259274, 'min_data_in_leaf': 16, 'num_leaves': 48,
'ag_args': {'name_suffix': '_r97', 'priority': 16}}, {'extra_trees': True,
'feature_fraction': 0.45555769907110816, 'learning_rate': 0.009591347321206594,
'min_data_in_leaf': 50, 'num_leaves': 110, 'ag_args': {'name_suffix': '_r71',
'priority': 15}}, {'extra_trees': False, 'feature_fraction':
0.40979710161022476, 'learning_rate': 0.008708890211023034, 'min_data_in_leaf':
3, 'num_leaves': 80, 'ag_args': {'name_suffix': '_r111', 'priority': 14}}],
    'XGB': {},
    'FASTAI': [{}, {'bs': 1024, 'emb_drop': 0.6167722379778131, 'epochs':
44, 'layers': [200, 100, 50], 'lr': 0.053440377855629266, 'ps':
0.48477211305443607, 'ag_args': {'name_suffix': '_r25', 'priority': 13}}, {'bs':
1024, 'emb_drop': 0.6046989241462619, 'epochs': 48, 'layers': [200, 100, 50],
'lr': 0.00775309042164966, 'ps': 0.09244767444160731, 'ag_args': {'name_suffix':
'_r51', 'priority': 12}}, {'bs': 512, 'emb_drop': 0.6557225316526186, 'epochs':
49, 'layers': [200, 100], 'lr': 0.023627682025564638, 'ps': 0.519566584552178,
'ag_args': {'name_suffix': '_r82', 'priority': 11}}, {'bs': 2048, 'emb_drop':
0.4066210919034579, 'epochs': 43, 'layers': [400, 200], 'lr':
0.0029598312717673434, 'ps': 0.4378695797438974, 'ag_args': {'name_suffix':
```



```

'_r121', 'priority': 10}}, {'bs': 128, 'emb_drop': 0.44339037504795686,
'epochs': 31, 'layers': [400, 200, 100], 'lr': 0.008615195908919904, 'ps':
0.19220253419114286, 'ag_args': {'name_suffix': '_r145', 'priority': 9}}, {'bs':
128, 'emb_drop': 0.12106594798980945, 'epochs': 38, 'layers': [200, 100, 50],
'lr': 0.037991970245029975, 'ps': 0.33120008492595093, 'ag_args':
{'name_suffix': '_r173', 'priority': 8}}, {'bs': 128, 'emb_drop':
0.4599138419358, 'epochs': 47, 'layers': [200, 100], 'lr': 0.03888383281136287,
'ps': 0.28193673177122863, 'ag_args': {'name_suffix': '_r128', 'priority': 7}}],
    'CAT': [{}, {'depth': 5, 'l2_leaf_reg': 4.774992314058497,
'learning_rate': 0.038551267822920274, 'ag_args': {'name_suffix': '_r16',
'priority': 6}}, {'depth': 4, 'l2_leaf_reg': 1.9950125740798321,
'learning_rate': 0.028091050379971633, 'ag_args': {'name_suffix': '_r42',
'priority': 5}}, {'depth': 6, 'l2_leaf_reg': 1.8298803017644376,
'learning_rate': 0.017844259810823604, 'ag_args': {'name_suffix': '_r93',
'priority': 4}}, {'depth': 7, 'l2_leaf_reg': 4.81099604606794, 'learning_rate':
0.019085060180573103, 'ag_args': {'name_suffix': '_r44', 'priority': 3}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Excluded models: ['RF', 'XGB', 'CAT'] (Specified by `excluded_model_types`)

Fitting 21 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 3599.78s of the 3599.77s of remaining time.

```

-28.5444      = Validation score    (-mean_absolute_error)
0.03s        = Training   runtime
0.39s        = Validation runtime

```

Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 3599.27s of the 3599.27s of remaining time.

```

-28.798      = Validation score    (-mean_absolute_error)
0.03s        = Training   runtime
0.43s        = Validation runtime

```

Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 3598.74s of the 3598.74s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

```

-13.5683      = Validation score    (-mean_absolute_error)
31.51s       = Training   runtime
15.52s       = Validation runtime

```

Fitting model: LightGBM_BAG_L1 ... Training model for up to 3562.06s of the 3562.05s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy

```

-14.6686      = Validation score    (-mean_absolute_error)
35.32s       = Training   runtime
16.27s       = Validation runtime

```

Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 3522.15s of the 3522.15s of remaining time.

```

-15.393      = Validation score    (-mean_absolute_error)

```

```

1.49s    = Training    runtime
0.91s    = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 3518.87s of
the 3518.87s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.3296          = Validation score    (-mean_absolute_error)
    35.63s    = Training    runtime
    0.44s    = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 3481.59s of
the 3481.59s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.8806          = Validation score    (-mean_absolute_error)
    146.46s = Training    runtime
    0.37s    = Validation runtime
Fitting model: ExtraTrees_r19_BAG_L1 ... Training model for up to 3333.53s of
the 3333.53s of remaining time.
    -15.2571          = Validation score    (-mean_absolute_error)
    0.97s    = Training    runtime
    0.92s    = Validation runtime
Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 3330.74s of the
3330.73s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -15.1218          = Validation score    (-mean_absolute_error)
    45.04s    = Training    runtime
    10.0s     = Validation runtime
Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 3277.52s of the
3277.52s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.8851          = Validation score    (-mean_absolute_error)
    21.08s    = Training    runtime
    10.15s    = Validation runtime
Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 3252.52s of the
3252.51s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.4763          = Validation score    (-mean_absolute_error)
    38.86s    = Training    runtime
    20.27s    = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 3208.16s of the
3208.16s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.8568          = Validation score    (-mean_absolute_error)
    76.57s    = Training    runtime

```

43.67s = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 3118.5s of the 3118.49s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy
-13.4808 = Validation score (-mean_absolute_error)
53.46s = Training runtime
20.25s = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 3057.06s of the 3057.05s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy
-14.6085 = Validation score (-mean_absolute_error)
21.18s = Training runtime
0.22s = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 3034.53s of the 3034.53s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy
-13.4434 = Validation score (-mean_absolute_error)
23.12s = Training runtime
0.21s = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 3009.89s of the 3009.88s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy
-14.5635 = Validation score (-mean_absolute_error)
31.9s = Training runtime
0.27s = Validation runtime
Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 2976.35s of the 2976.34s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy
-16.5206 = Validation score (-mean_absolute_error)
17.71s = Training runtime
0.2s = Validation runtime
Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 2957.09s of the 2957.09s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy
-13.1404 = Validation score (-mean_absolute_error)
78.14s = Training runtime
0.86s = Validation runtime
Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 2876.97s of the 2876.96s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with ParallelLocalFoldFittingStrategy
-14.5364 = Validation score (-mean_absolute_error)

```

80.21s    = Training    runtime
0.82s     = Validation  runtime
Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 2794.86s
of the 2794.86s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.4378          = Validation score    (-mean_absolute_error)
    88.95s    = Training    runtime
    0.83s     = Validation  runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 2704.0s of the
2704.0s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.1201          = Validation score    (-mean_absolute_error)
    107.96s   = Training    runtime
    20.12s    = Validation  runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 2586.48s of the
2586.48s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.4977          = Validation score    (-mean_absolute_error)
    61.58s    = Training    runtime
    28.91s    = Validation  runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 2550.59s of the
2550.59s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.5685          = Validation score    (-mean_absolute_error)
    69.39s    = Training    runtime
    27.32s    = Validation  runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 2510.69s of
the 2510.68s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.2696          = Validation score    (-mean_absolute_error)
    70.5s     = Training    runtime
    0.88s     = Validation  runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 2473.96s of
the 2473.96s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.8787          = Validation score    (-mean_absolute_error)
    261.2s    = Training    runtime
    0.75s     = Validation  runtime
Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 2357.56s of the
2357.56s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with

```

```

ParallelLocalFoldFittingStrategy
    -14.9927          = Validation score    (-mean_absolute_error)
    99.8s            = Training   runtime
    22.26s           = Validation runtime
Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 2292.79s of the
2292.78s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.9312          = Validation score    (-mean_absolute_error)
    42.74s            = Training   runtime
    22.36s           = Validation runtime
Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 2266.71s of the
2266.7s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.4148          = Validation score    (-mean_absolute_error)
    77.57s            = Training   runtime
    40.21s           = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 2220.71s of the
2220.7s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.8498          = Validation score    (-mean_absolute_error)
    152.56s           = Training   runtime
    90.59s           = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 2126.41s of the
2126.4s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.489           = Validation score    (-mean_absolute_error)
    106.84s           = Training   runtime
    40.05s           = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 2061.96s
of the 2061.95s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.2517          = Validation score    (-mean_absolute_error)
    42.33s            = Training   runtime
    0.44s            = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 2039.29s
of the 2039.29s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.3601          = Validation score    (-mean_absolute_error)
    45.84s            = Training   runtime
    0.46s            = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 2014.78s
of the 2014.77s of remaining time.

```

```

    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.5315      = Validation score    (-mean_absolute_error)
    64.49s       = Training    runtime
    0.57s        = Validation runtime
Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 1980.57s
of the 1980.56s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -16.5741      = Validation score    (-mean_absolute_error)
    35.36s        = Training    runtime
    0.41s         = Validation runtime
Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 1961.35s
of the 1961.34s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.2858      = Validation score    (-mean_absolute_error)
    157.01s       = Training    runtime
    1.65s         = Validation runtime
Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 1879.95s
of the 1879.95s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.911       = Validation score    (-mean_absolute_error)
    158.12s       = Training    runtime
    1.56s         = Validation runtime
Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 1799.66s
of the 1799.66s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.2643      = Validation score    (-mean_absolute_error)
    173.77s       = Training    runtime
    1.52s         = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1712.49s of the
1712.48s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.0138      = Validation score    (-mean_absolute_error)
    215.58s       = Training    runtime
    40.7s         = Validation runtime
Repeating k-fold bagging: 3/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1591.2s of the
1591.19s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.4815      = Validation score    (-mean_absolute_error)
    92.05s        = Training    runtime
    43.69s        = Validation runtime

```

Fitting model: LightGBM_BAG_L1 ... Training model for up to 1553.88s of the 1553.88s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-14.4787 = Validation score (-mean_absolute_error)
101.63s = Training runtime
44.43s = Validation runtime

Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1513.18s of the 1513.18s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-13.1896 = Validation score (-mean_absolute_error)
105.4s = Training runtime
1.38s = Validation runtime

Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1476.17s of the 1476.17s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-12.8367 = Validation score (-mean_absolute_error)
399.43s = Training runtime
1.13s = Validation runtime

Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 1335.99s of the 1335.99s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-14.9971 = Validation score (-mean_absolute_error)
146.79s = Training runtime
29.08s = Validation runtime

Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 1278.99s of the 1278.98s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-12.9213 = Validation score (-mean_absolute_error)
63.37s = Training runtime
35.65s = Validation runtime

Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 1251.93s of the 1251.92s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-13.4306 = Validation score (-mean_absolute_error)
115.14s = Training runtime
58.36s = Validation runtime

Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 1205.67s of the 1205.67s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy

-14.8527 = Validation score (-mean_absolute_error)
228.65s = Training runtime

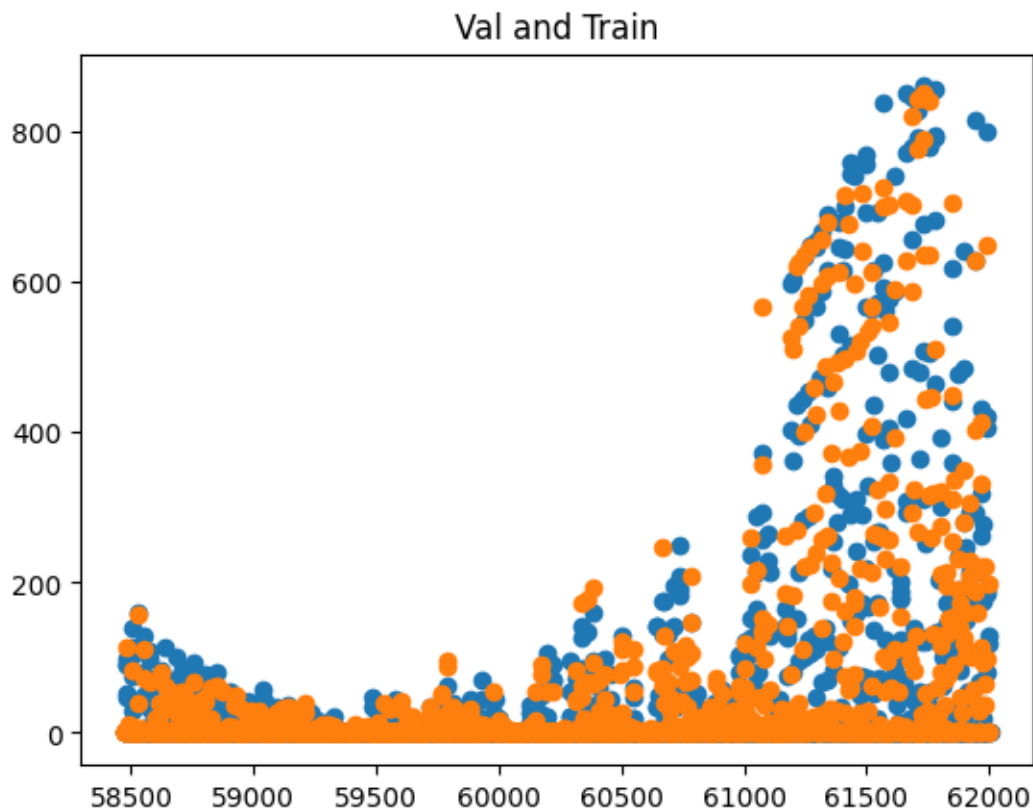
132.81s = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 1109.05s of the 1109.05s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy
-13.482 = Validation score (-mean_absolute_error)
160.09s = Training runtime
55.83s = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 1043.52s of the 1043.52s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy
-14.1615 = Validation score (-mean_absolute_error)
63.32s = Training runtime
0.66s = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 1020.9s of the 1020.9s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy
-13.2851 = Validation score (-mean_absolute_error)
68.48s = Training runtime
0.69s = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 996.38s of the 996.37s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy
-14.5784 = Validation score (-mean_absolute_error)
96.96s = Training runtime
0.88s = Validation runtime
Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 961.84s of the 961.83s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy
-16.5649 = Validation score (-mean_absolute_error)
53.1s = Training runtime
0.62s = Validation runtime
Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 942.45s of the 942.45s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy
-13.4038 = Validation score (-mean_absolute_error)
235.04s = Training runtime
2.71s = Validation runtime
Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 861.35s of the 861.34s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with ParallelLocalFoldFittingStrategy
-13.8874 = Validation score (-mean_absolute_error)


```

    240.45s = Training runtime
    2.3s    = Validation runtime
Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 776.07s
of the 776.07s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.2013 = Validation score (-mean_absolute_error)
    261.75s = Training runtime
    2.21s   = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 685.18s of the
685.17s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.9419 = Validation score (-mean_absolute_error)
    321.07s = Training runtime
    58.51s  = Validation runtime
Completed 3/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
562.81s of remaining time.
    -12.3391 = Validation score (-mean_absolute_error)
    0.53s    = Training runtime
    0.0s     = Validation runtime
AutoGluon training complete, total runtime = 3037.75s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_122_B/")
Evaluation: mean_absolute_error on test data: -10.723597852022728
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -10.723597852022728,
    "root_mean_squared_error": -29.434039222380054,
    "mean_squared_error": -866.3626649446074,
    "r2": 0.9627022467439814,
    "pearsonr": 0.9811782810511718,
    "median_absolute_error": -0.5645688772201538
}


Evaluation on test data:
-10.723597852022728

```



	model	score_test	score_val	pred_time_test
0	WeightedEnsemble_L2	-10.723598	-12.339116	4.378928
40.877352	803.763585	0.003843		0.000623
0.532389	2	True	22	
1	LightGBM_r118_BAG_L1	-11.007801	-12.921285	2.304665
35.652883	63.369035	2.304665		35.652883
63.369035	1	True	10	
2	LightGBMXT_BAG_L1	-11.132298	-13.481468	3.404401
43.693103	92.046542	3.404401		43.693103
92.046542	1	True	3	
3	LightGBM_BAG_L1	-11.142879	-14.478723	2.731642
44.433135	101.633064	2.731642		44.433135
101.633064	1	True	4	
4	NeuralNetTorch_BAG_L1	-11.182017	-12.836698	0.591085
1.127306	399.426073	0.591085		1.127306
399.426073	1	True	7	
5	LightGBMLarge_BAG_L1	-11.261584	-13.941926	9.446423
58.507039	321.065502	9.446423		58.507039
321.065502	1	True	21	

6	LightGBM_r111_BAG_L1	-11.295378	-13.482024	6.365718
55.826619	160.094988		6.365718	55.826619
160.094988	1	True	13	
7	LightGBM_r97_BAG_L1	-11.517776	-13.430644	3.840222
58.361140	115.136348		3.840222	58.361140
115.136348	1	True	11	
8	NeuralNetFastAI_BAG_L1	-12.389999	-13.189636	0.506129
1.384786	105.397018		0.506129	1.384786
105.397018	1	True	6	
9	LightGBM_r71_BAG_L1	-12.534359	-14.852719	8.406466
132.805016	228.645143		8.406466	132.805016
228.645143	1	True	12	
10	LightGBM_r158_BAG_L1	-12.538895	-14.997143	3.252692
29.084249	146.785765		3.252692	29.084249
146.785765	1	True	9	
11	NeuralNetFastAI_r145_BAG_L1	-12.685629	-13.403836	0.973206
2.711753	235.039069		0.973206	2.711753
235.039069	1	True	18	
12	NeuralNetFastAI_r51_BAG_L1	-12.941687	-13.285118	0.357658
0.688112	68.483178		0.357658	0.688112
68.483178	1	True	15	
13	ExtraTreesMSE_BAG_L1	-13.117027	-15.392978	0.447318
0.914828	1.486284		0.447318	0.914828
1.486284	1	True	5	
14	ExtraTrees_r19_BAG_L1	-13.312309	-15.257086	0.448604
0.922027	0.965683		0.448604	0.922027
0.965683	1	True	8	
15	NeuralNetFastAI_r173_BAG_L1	-14.269517	-13.887434	0.817005
2.300970	240.453202		0.817005	2.300970
240.453202	1	True	19	
16	NeuralNetFastAI_r25_BAG_L1	-14.480768	-14.161466	0.359478
0.658887	63.315527		0.359478	0.658887
63.315527	1	True	14	
17	NeuralNetFastAI_r128_BAG_L1	-14.515248	-14.201333	0.804592
2.211224	261.745151		0.804592	2.211224
261.745151	1	True	20	
18	NeuralNetFastAI_r82_BAG_L1	-14.742232	-14.578388	0.399377
0.882292	96.959816		0.399377	0.882292
96.959816	1	True	16	
19	NeuralNetFastAI_r121_BAG_L1	-16.697332	-16.564891	0.297959
0.616057	53.095748		0.297959	0.616057
53.095748	1	True	17	
20	KNeighborsDist_BAG_L1	-23.570593	-28.797984	0.016379
0.429177	0.029572		0.016379	0.429177
0.029572	1	True	2	
21	KNeighborsUnif_BAG_L1	-24.697224	-28.544405	0.014151
0.389905	0.030971		0.014151	0.389905
0.030971	1	True	1	



```
[21]: loc = "C"
      predictors[2] = fit_predictor_for_location(loc)
      leaderboards[2] = leaderboard_for_location(2, loc)
```

```
Presets specified: ['experimental_zeroshot_hpo_hybrid']
Stack configuration (auto_stack=True): num_stack_levels=0, num_bag_folds=8,
num_bag_sets=20
Beginning AutoGluon training ... Time limit = 3600s
AutoGluon will save models to "AutogluonModels/submission_122_C/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP Debian 5.10.197-1 (2023-09-29)
Disk Space Avail: 136.15 GB / 315.93 GB (43.1%)
Train Data Rows: 24073

Training model for location C...

Train Data Columns: 44
Tuning Data Rows: 1481
Tuning Data Columns: 44
```

```

Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and label-values can't be converted to int).
    Label info (max, min, mean, stddev): (999.6, -0.0, 80.87539, 169.67845)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory: 130055.27 MB
    Train Data (Original) Memory Usage: 10.27 MB (0.0% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 2 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 3): ['elevation:m', 'snow_drift:idx',
'location']
        These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 40 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 39 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', ['bool']) : 2 | ['snow_density:kgm3', 'is_estimated']
    0.1s = Fit runtime
    41 features in original data used to generate 41 features in processed
data.
    Train Data (Processed) Memory Usage: 8.02 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.16s ...

```

AutoGluon will gauge predictive performance using evaluation metric:

```
'mean_absolute_error'
```

This metric's sign has been flipped to adhere to being higher_is_better. The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval_metric parameter of Predictor() use_bag_holdout=True, will use tuning_data as holdout (will not be used for early stopping).

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}, {'min_samples_leaf': 1,
'max_leaf_nodes': 15000, 'max_features': 0.5, 'ag_args': {'name_suffix': '_r19',
'priority': 20}}],
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}, {'min_samples_leaf': 5,
'max_leaf_nodes': 50000, 'max_features': 0.5, 'ag_args': {'name_suffix': '_r5',
'priority': 19}}],
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {},
'GBMLarge', {'extra_trees': False, 'feature_fraction': 0.7248284762542815,
'learning_rate': 0.07947286942946127, 'min_data_in_leaf': 50, 'num_leaves': 89,
'ag_args': {'name_suffix': '_r158', 'priority': 18}}, {'extra_trees': True,
'feature_fraction': 0.7832570544199176, 'learning_rate': 0.021720607471727896,
'min_data_in_leaf': 3, 'num_leaves': 21, 'ag_args': {'name_suffix': '_r118',
'priority': 17}}, {'extra_trees': True, 'feature_fraction': 0.7113010892989156,
'learning_rate': 0.012535427424259274, 'min_data_in_leaf': 16, 'num_leaves': 48,
'ag_args': {'name_suffix': '_r97', 'priority': 16}}, {'extra_trees': True,
'feature_fraction': 0.45555769907110816, 'learning_rate': 0.009591347321206594,
'min_data_in_leaf': 50, 'num_leaves': 110, 'ag_args': {'name_suffix': '_r71',
'priority': 15}}, {'extra_trees': False, 'feature_fraction':
0.40979710161022476, 'learning_rate': 0.008708890211023034, 'min_data_in_leaf':
3, 'num_leaves': 80, 'ag_args': {'name_suffix': '_r111', 'priority': 14}}],
    'XGB': {},
    'FASTAI': [{}, {'bs': 1024, 'emb_drop': 0.6167722379778131, 'epochs':
44, 'layers': [200, 100, 50], 'lr': 0.053440377855629266, 'ps':
0.48477211305443607, 'ag_args': {'name_suffix': '_r25', 'priority': 13}}, {'bs':
1024, 'emb_drop': 0.6046989241462619, 'epochs': 48, 'layers': [200, 100, 50],
'lr': 0.00775309042164966, 'ps': 0.09244767444160731, 'ag_args': {'name_suffix':
'_r51', 'priority': 12}}, {'bs': 512, 'emb_drop': 0.6557225316526186, 'epochs':
49, 'layers': [200, 100], 'lr': 0.023627682025564638, 'ps': 0.519566584552178,
'ag_args': {'name_suffix': '_r82', 'priority': 11}}, {'bs': 2048, 'emb_drop':
0.4066210919034579, 'epochs': 43, 'layers': [400, 200], 'lr':
```

```

0.0029598312717673434, 'ps': 0.4378695797438974, 'ag_args': {'name_suffix':
'_r121', 'priority': 10}}, {'bs': 128, 'emb_drop': 0.44339037504795686,
'epochs': 31, 'layers': [400, 200, 100], 'lr': 0.008615195908919904, 'ps':
0.19220253419114286, 'ag_args': {'name_suffix': '_r145', 'priority': 9}}, {'bs':
128, 'emb_drop': 0.12106594798980945, 'epochs': 38, 'layers': [200, 100, 50],
'lr': 0.037991970245029975, 'ps': 0.33120008492595093, 'ag_args':
{'name_suffix': '_r173', 'priority': 8}}, {'bs': 128, 'emb_drop':
0.4599138419358, 'epochs': 47, 'layers': [200, 100], 'lr': 0.03888383281136287,
'ps': 0.28193673177122863, 'ag_args': {'name_suffix': '_r128', 'priority': 7}}],
    'CAT': [{}, {'depth': 5, 'l2_leaf_reg': 4.774992314058497,
'learning_rate': 0.038551267822920274, 'ag_args': {'name_suffix': '_r16',
'priority': 6}}, {'depth': 4, 'l2_leaf_reg': 1.9950125740798321,
'learning_rate': 0.028091050379971633, 'ag_args': {'name_suffix': '_r42',
'priority': 5}}, {'depth': 6, 'l2_leaf_reg': 1.8298803017644376,
'learning_rate': 0.017844259810823604, 'ag_args': {'name_suffix': '_r93',
'priority': 4}}, {'depth': 7, 'l2_leaf_reg': 4.81099604606794, 'learning_rate':
0.019085060180573103, 'ag_args': {'name_suffix': '_r44', 'priority': 3}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}

```

Excluded models: ['RF', 'XGB', 'CAT'] (Specified by `excluded_model_types`)

Fitting 21 L1 models ...

Fitting model: KNeighborsUnif_BAG_L1 ... Training model for up to 3599.84s of the 3599.84s of remaining time.

-19.8149 = Validation score (-mean_absolute_error)

0.03s = Training runtime

0.33s = Validation runtime

Fitting model: KNeighborsDist_BAG_L1 ... Training model for up to 3599.28s of the 3599.28s of remaining time.

-20.1923 = Validation score (-mean_absolute_error)

0.03s = Training runtime

0.25s = Validation runtime

Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 3598.94s of the 3598.94s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

-11.8238 = Validation score (-mean_absolute_error)

29.22s = Training runtime

11.97s = Validation runtime

Fitting model: LightGBM_BAG_L1 ... Training model for up to 3565.0s of the 3565.0s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy

-12.8555 = Validation score (-mean_absolute_error)

32.69s = Training runtime

11.1s = Validation runtime

Fitting model: ExtraTreesMSE_BAG_L1 ... Training model for up to 3528.27s of the 3528.27s of remaining time.

```

-15.4038          = Validation score    (-mean_absolute_error)
1.14s            = Training   runtime
0.77s            = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 3525.69s of
the 3525.68s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.5826          = Validation score    (-mean_absolute_error)
31.1s            = Training   runtime
0.39s            = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 3493.07s of
the 3493.07s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.6436          = Validation score    (-mean_absolute_error)
82.85s           = Training   runtime
0.34s            = Validation runtime
Fitting model: ExtraTrees_r19_BAG_L1 ... Training model for up to 3408.61s of
the 3408.61s of remaining time.
-15.4648          = Validation score    (-mean_absolute_error)
0.84s            = Training   runtime
0.79s            = Validation runtime
Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 3406.26s of the
3406.25s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-14.9423          = Validation score    (-mean_absolute_error)
42.08s           = Training   runtime
4.03s            = Validation runtime
Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 3357.69s of the
3357.69s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-11.2823          = Validation score    (-mean_absolute_error)
20.66s           = Training   runtime
8.37s            = Validation runtime
Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 3333.23s of the
3333.22s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-11.7182          = Validation score    (-mean_absolute_error)
38.59s           = Training   runtime
13.59s           = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 3289.77s of the
3289.77s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.3148          = Validation score    (-mean_absolute_error)

```



```

74.62s = Training runtime
38.76s = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 3202.85s of the
3202.85s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.9079 = Validation score (-mean_absolute_error)
51.5s = Training runtime
19.88s = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 3142.8s
of the 3142.8s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-16.8271 = Validation score (-mean_absolute_error)
18.81s = Training runtime
0.18s = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 3122.6s
of the 3122.59s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.8295 = Validation score (-mean_absolute_error)
20.41s = Training runtime
0.2s = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 3100.76s
of the 3100.76s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-16.8595 = Validation score (-mean_absolute_error)
29.0s = Training runtime
0.24s = Validation runtime
Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 3070.23s
of the 3070.23s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-17.2171 = Validation score (-mean_absolute_error)
16.22s = Training runtime
0.2s = Validation runtime
Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 3052.51s
of the 3052.51s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.8646 = Validation score (-mean_absolute_error)
69.4s = Training runtime
0.76s = Validation runtime
Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 2980.97s
of the 2980.97s of remaining time.
Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy

```

```

-16.2981          = Validation score    (-mean_absolute_error)
72.26s           = Training   runtime
0.73s            = Validation runtime
Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 2906.83s
of the 2906.82s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-16.2728          = Validation score    (-mean_absolute_error)
79.66s           = Training   runtime
0.69s            = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 2825.32s of the
2825.32s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.9072          = Validation score    (-mean_absolute_error)
103.2s           = Training   runtime
11.25s           = Validation runtime
Repeating k-fold bagging: 2/20
Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 2713.82s of the
2713.82s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-11.736           = Validation score    (-mean_absolute_error)
59.15s           = Training   runtime
23.91s           = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 2678.25s of the
2678.25s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.8129          = Validation score    (-mean_absolute_error)
63.47s           = Training   runtime
20.25s           = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 2641.93s of
the 2641.93s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.5089          = Validation score    (-mean_absolute_error)
62.11s           = Training   runtime
0.79s            = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 2609.14s of
the 2609.14s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.5827          = Validation score    (-mean_absolute_error)
161.92s          = Training   runtime
0.67s            = Validation runtime
Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 2528.26s of the
2528.26s of remaining time.

```

```

    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -14.6379      = Validation score    (-mean_absolute_error)
    96.36s       = Training    runtime
    14.77s       = Validation runtime
Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 2465.83s of the
2465.83s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -11.2452      = Validation score    (-mean_absolute_error)
    41.86s       = Training    runtime
    15.46s       = Validation runtime
Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 2440.38s of the
2440.38s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -11.7499      = Validation score    (-mean_absolute_error)
    77.81s       = Training    runtime
    27.73s       = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 2394.07s of the
2394.06s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -13.2916      = Validation score    (-mean_absolute_error)
    147.84s      = Training    runtime
    80.21s       = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 2303.75s of the
2303.75s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.8981      = Validation score    (-mean_absolute_error)
    103.56s      = Training    runtime
    38.56s       = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 2240.44s
of the 2240.43s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -16.7878      = Validation score    (-mean_absolute_error)
    37.49s       = Training    runtime
    0.37s        = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 2220.22s
of the 2220.21s of remaining time.
    Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy
    -12.6092      = Validation score    (-mean_absolute_error)
    40.56s       = Training    runtime
    0.4s         = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 2198.51s

```

of the 2198.51s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-16.9557 = Validation score (-mean_absolute_error)
57.9s = Training runtime
0.5s = Validation runtime

Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 2167.71s
of the 2167.7s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-17.2424 = Validation score (-mean_absolute_error)
32.46s = Training runtime
0.39s = Validation runtime

Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 2149.85s
of the 2149.85s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-13.5793 = Validation score (-mean_absolute_error)
137.72s = Training runtime
1.51s = Validation runtime

Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 2078.85s
of the 2078.84s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-16.6716 = Validation score (-mean_absolute_error)
143.27s = Training runtime
1.42s = Validation runtime

Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 2005.42s
of the 2005.42s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-16.5392 = Validation score (-mean_absolute_error)
154.18s = Training runtime
1.38s = Validation runtime

Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1928.56s of the
1928.56s of remaining time.

Fitting 8 child models (S2F1 - S2F8) | Fitting with
ParallelLocalFoldFittingStrategy

-12.8781 = Validation score (-mean_absolute_error)
206.98s = Training runtime
22.89s = Validation runtime

Repeating k-fold bagging: 3/20

Fitting model: LightGBMXT_BAG_L1 ... Training model for up to 1812.72s of the
1812.72s of remaining time.

Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy

-11.7971 = Validation score (-mean_absolute_error)
90.66s = Training runtime

34.91s = Validation runtime
Fitting model: LightGBM_BAG_L1 ... Training model for up to 1774.63s of the
1774.63s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.7243 = Validation score (-mean_absolute_error)
92.65s = Training runtime
28.48s = Validation runtime
Fitting model: NeuralNetFastAI_BAG_L1 ... Training model for up to 1739.2s of
the 1739.2s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.5271 = Validation score (-mean_absolute_error)
93.22s = Training runtime
1.21s = Validation runtime
Fitting model: NeuralNetTorch_BAG_L1 ... Training model for up to 1705.97s of
the 1705.96s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.3914 = Validation score (-mean_absolute_error)
262.5s = Training runtime
1.03s = Validation runtime
Fitting model: LightGBM_r158_BAG_L1 ... Training model for up to 1603.32s of the
1603.32s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-14.7092 = Validation score (-mean_absolute_error)
143.66s = Training runtime
18.97s = Validation runtime
Fitting model: LightGBM_r118_BAG_L1 ... Training model for up to 1546.97s of the
1546.96s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-11.2132 = Validation score (-mean_absolute_error)
62.96s = Training runtime
22.11s = Validation runtime
Fitting model: LightGBM_r97_BAG_L1 ... Training model for up to 1520.5s of the
1520.5s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-11.7531 = Validation score (-mean_absolute_error)
115.49s = Training runtime
42.1s = Validation runtime
Fitting model: LightGBM_r71_BAG_L1 ... Training model for up to 1474.45s of the
1474.44s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.3005 = Validation score (-mean_absolute_error)

```

221.62s = Training runtime
119.34s = Validation runtime
Fitting model: LightGBM_r111_BAG_L1 ... Training model for up to 1379.96s of the
1379.96s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.8684 = Validation score (-mean_absolute_error)
156.77s = Training runtime
57.32s = Validation runtime
Fitting model: NeuralNetFastAI_r25_BAG_L1 ... Training model for up to 1314.0s
of the 1314.0s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-16.8597 = Validation score (-mean_absolute_error)
56.15s = Training runtime
0.56s = Validation runtime
Fitting model: NeuralNetFastAI_r51_BAG_L1 ... Training model for up to 1293.67s
of the 1293.67s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.5751 = Validation score (-mean_absolute_error)
60.52s = Training runtime
0.6s = Validation runtime
Fitting model: NeuralNetFastAI_r82_BAG_L1 ... Training model for up to 1271.97s
of the 1271.97s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-17.0273 = Validation score (-mean_absolute_error)
87.02s = Training runtime
0.82s = Validation runtime
Fitting model: NeuralNetFastAI_r121_BAG_L1 ... Training model for up to 1240.79s
of the 1240.78s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-17.2 = Validation score (-mean_absolute_error)
48.78s = Training runtime
0.57s = Validation runtime
Fitting model: NeuralNetFastAI_r145_BAG_L1 ... Training model for up to 1222.77s
of the 1222.76s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-13.6868 = Validation score (-mean_absolute_error)
206.19s = Training runtime
2.27s = Validation runtime
Fitting model: NeuralNetFastAI_r173_BAG_L1 ... Training model for up to 1151.3s
of the 1151.29s of remaining time.
Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy


```

```

-16.5738          = Validation score    (-mean_absolute_error)
210.91s = Training    runtime
2.06s   = Validation runtime
Fitting model: NeuralNetFastAI_r128_BAG_L1 ... Training model for up to 1080.78s
of the 1080.77s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-16.8236          = Validation score    (-mean_absolute_error)
216.36s = Training    runtime
1.99s   = Validation runtime
Fitting model: LightGBMLarge_BAG_L1 ... Training model for up to 1015.8s of the
1015.8s of remaining time.
    Fitting 8 child models (S3F1 - S3F8) | Fitting with
ParallelLocalFoldFittingStrategy
-12.833 = Validation score    (-mean_absolute_error)
308.54s = Training    runtime
35.09s  = Validation runtime
Completed 3/20 k-fold bagging repeats ...
Fitting model: WeightedEnsemble_L2 ... Training model for up to 360.0s of the
898.84s of remaining time.
-11.123 = Validation score    (-mean_absolute_error)
0.53s   = Training    runtime
0.0s    = Validation runtime
AutoGluon training complete, total runtime = 2702.23s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_122_C/")
Evaluation: mean_absolute_error on test data: -11.882573633916262
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -11.882573633916262,
    "root_mean_squared_error": -28.984463498596043,
    "mean_squared_error": -840.0991243014464,
    "r2": 0.914844370410137,
    "pearsonr": 0.9592036434639014,
    "median_absolute_error": -0.7651807069778442
}


Evaluation on test data:
-11.882573633916262

```



	model	score_test	score_val	pred_time_test
0	WeightedEnsemble_L2	-11.882574	-11.122967	3.856977
24.068558	386.542620	0.003762		0.000596
0.534351	2	True	22	
1	LightGBM_r118_BAG_L1	-11.977019	-11.213198	2.924395
22.111541	62.962651	2.924395		22.111541
62.962651	1	True	10	
2	LightGBMXT_BAG_L1	-12.137031	-11.797102	2.805304
34.914344	90.661297	2.805304		34.914344
90.661297	1	True	3	
3	LightGBM_r97_BAG_L1	-12.358573	-11.753063	3.886996
42.103179	115.485718	3.886996		42.103179
115.485718	1	True	11	
4	NeuralNetFastAI_BAG_L1	-12.998475	-13.527125	0.475597
1.210410	93.215323	0.475597		1.210410
93.215323	1	True	6	
5	NeuralNetFastAI_r51_BAG_L1	-13.047249	-12.575052	0.355870
0.598878	60.517148	0.355870		0.598878
60.517148	1	True	15	

6	NeuralNetFastAI_r145_BAG_L1	-13.064247	-13.686793	0.743834
2.268717	206.190490		0.743834	2.268717
206.190490	1	True	18	
7	NeuralNetTorch_BAG_L1	-13.236769	-13.391357	0.555999
1.026419	262.499245		0.555999	1.026419
262.499245	1	True	7	
8	LightGBM_BAG_L1	-13.516948	-12.724342	2.450099
28.478944	92.647044		2.450099	28.478944
92.647044	1	True	4	
9	LightGBM_r71_BAG_L1	-13.549919	-13.300531	8.404623
119.343533	221.619780		8.404623	119.343533
221.619780	1	True	12	
10	LightGBM_r111_BAG_L1	-13.830864	-12.868428	6.638808
57.324214	156.768054		6.638808	57.324214
156.768054	1	True	13	
11	LightGBMLarge_BAG_L1	-14.153453	-12.832983	8.209069
35.091078	308.535153		8.209069	35.091078
308.535153	1	True	21	
12	NeuralNetFastAI_r128_BAG_L1	-15.275111	-16.823553	0.742124
1.988479	216.364365		0.742124	1.988479
216.364365	1	True	20	
13	NeuralNetFastAI_r173_BAG_L1	-15.372847	-16.573827	0.768911
2.060926	210.906921		0.768911	2.060926
210.906921	1	True	19	
14	ExtraTreesMSE_BAG_L1	-15.433922	-15.403829	0.302901
0.769768	1.139719		0.302901	0.769768
1.139719	1	True	5	
15	ExtraTrees_r19_BAG_L1	-15.557787	-15.464786	0.339243
0.786726	0.839571		0.339243	0.786726
0.839571	1	True	8	
16	NeuralNetFastAI_r82_BAG_L1	-15.602025	-17.027349	0.392537
0.819656	87.022309		0.392537	0.819656
87.022309	1	True	16	
17	NeuralNetFastAI_r121_BAG_L1	-15.648391	-17.200041	0.307789
0.565804	48.779081		0.307789	0.565804
48.779081	1	True	17	
18	NeuralNetFastAI_r25_BAG_L1	-15.960404	-16.859700	0.375360
0.561592	56.147095		0.375360	0.561592
56.147095	1	True	14	
19	LightGBM_r158_BAG_L1	-16.031353	-14.709153	2.388238
18.966787	143.661812		2.388238	18.966787
143.661812	1	True	9	
20	KNeighborsUnif_BAG_L1	-20.049167	-19.814903	0.016952
0.331125	0.029224		0.016952	0.331125
0.029224	1	True	1	
21	KNeighborsDist_BAG_L1	-20.130194	-20.192291	0.010500
0.250257	0.027984		0.010500	0.250257
0.027984	1	True	2	



```
[22]: # save leaderboards to csv
pd.concat(leaderboards).to_csv(f"leaderboards/{new_filename}.csv")

for i in range(len(predictors)):
    print(f"Predictor {i}:")
    print(predictors[i].
    ↪info()["model_info"]["WeightedEnsemble_L2"]["children_info"]["S1F1"]["model_weights"])
```

Predictor 0:

```
{'LightGBMXT_BAG_L1': 0.29577464788732394, 'NeuralNetTorch_BAG_L1':
0.352112676056338, 'LightGBM_r118_BAG_L1': 0.16901408450704225,
'NeuralNetFastAI_r51_BAG_L1': 0.028169014084507043, 'LightGBMLarge_BAG_L1':
0.15492957746478872}
```

Predictor 1:

```
{'NeuralNetFastAI_BAG_L1': 0.1836734693877551, 'NeuralNetTorch_BAG_L1':
0.3673469387755102, 'LightGBM_r118_BAG_L1': 0.3469387755102041,
'NeuralNetFastAI_r145_BAG_L1': 0.10204081632653061}
```

Predictor 2:

```
{'KNeighborsUnif_BAG_L1': 0.03488372093023256, 'NeuralNetTorch_BAG_L1':
0.011627906976744186, 'LightGBM_r118_BAG_L1': 0.8255813953488372,
```

```
'NeuralNetFastAI_r51_BAG_L1': 0.12790697674418605}
```

5 Submit

```
[23]: import pandas as pd
import matplotlib.pyplot as plt

future_test_data = TabularDataset('X_test_raw.csv')
future_test_data["ds"] = pd.to_datetime(future_test_data["ds"])
#test_data
```

Loaded data from: X_test_raw.csv | Columns = 45 / 45 | Rows = 4608 -> 4608

```
[24]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
future_test_data_merged = pd.merge(future_test_data, test_ids, how="inner",
    ↪right_on=["time", "location"], left_on=["ds", "location"])

#test_data_merged
```

Loaded data from: test.csv | Columns = 4 / 4 | Rows = 2160 -> 2160

```
[25]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in future_test_data.groupby('location'):
    i = location_map[loc]
    subset = future_test_data_merged[future_test_data_merged["location"] == loc]
    ↪loc].reset_index(drop=True)
    #print(subset)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)


# get past predictions
#train_data.loc[train_data["location"] == loc, "prediction"] =
    ↪predictors[i].predict(train_data[train_data["location"] == loc])
    if use_tune_data:
        tuning_data.loc[tuning_data["location"] == loc, "prediction"] =
    ↪predictors[i].predict(tuning_data[tuning_data["location"] == loc])
    if use_test_data:
        test_data.loc[test_data["location"] == loc, "prediction"] =
    ↪predictors[i].predict(test_data[test_data["location"] == loc])
```



```
[26]: # plot predictions for location A, in addition to train data for A
for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data[train_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="train data")
    if use_tune_data:
        tuning_data[tuning_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="tune data")
    if use_test_data:
        test_data[test_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="test data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")

    # plot past predictions
    #train_data_with_dates[train_data_with_dates["location"]==loc].plot(x='ds',
    ↪y='prediction', ax=ax, label="past predictions")
    #train_data[train_data["location"]==loc].plot(x='ds', y='prediction',
    ↪ax=ax, label="past predictions train")
    if use_tune_data:
        tuning_data[tuning_data["location"]==loc].plot(x='ds', y='prediction',
    ↪ax=ax, label="past predictions tune")
    if use_test_data:
        test_data[test_data["location"]==loc].plot(x='ds', y='prediction',
    ↪ax=ax, label="past predictions test")

    # title
    ax.set_title(f"Predictions for location {loc}")
```





```
[27]: temp_predictions = [prediction.copy() for prediction in predictions]
if clip_predictions:
    # clip predictions smaller than 0 to 0
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        pred.loc[pred["prediction"] < 0, "prediction"] = 0
        print("Smallest prediction after clipping:", pred["prediction"].min())
```

```

# Instead of clipping, shift all prediction values up by the largest negative
↳ number.
# This way, the smallest prediction will be 0.
elif shift_predictions:
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        pred["prediction"] = pred["prediction"] - pred["prediction"].min()
        print("Smallest prediction after clipping:", pred["prediction"].min())

elif shift_predictions_by_average_of_negatives_then_clip:
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        mean_negative = pred[pred["prediction"] < 0]["prediction"].mean()
        # if not nan
        if mean_negative == mean_negative:
            pred["prediction"] = pred["prediction"] - mean_negative

        pred.loc[pred["prediction"] < 0, "prediction"] = 0
        print("Smallest prediction after clipping:", pred["prediction"].min())

# concatenate predictions
submissions_df = pd.concat(temp_predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df

```

```

Smallest prediction: -28.481598
Smallest prediction after clipping: 0.0
Smallest prediction: -2.31651
Smallest prediction after clipping: 0.0
Smallest prediction: -3.1237252
Smallest prediction after clipping: 0.0

```

```

[27]:      id  prediction
0      0      0.000000
1      1      0.000000
2      2      0.000000
3      3     30.205181
4      4    311.056702
..    ...      ...
715  2155     62.010193
716  2156     35.529598
717  2157      8.074565

```

```
718 2158    1.756826
719 2159    1.467647
```

```
[2160 rows x 2 columns]
```

```
[28]: # Save the submission DataFrame to submissions folder, create new name based on
      ↪ last submission, format is submission_<last_submission_number + 1>.csv

      # Save the submission
      print(f"Saving submission to submissions/{new_filename}.csv")
      submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
      ↪ index=False)
      print("jall1a")
```

```
Saving submission to submissions/submission_122.csv
jall1a
```

```
[ ]: # feature importance
      # print starting calculating feature importance for location A with big text
      ↪ font
      print("\033[1m" + "Calculating feature importance for location A..." +
      ↪ "\033[0m")
      predictors[0].feature_importance(feature_stage="original",
      ↪ data=test_data[test_data["location"] == "A"], time_limit=60*10)
      print("\033[1m" + "Calculating feature importance for location B..." +
      ↪ "\033[0m")
      predictors[1].feature_importance(feature_stage="original",
      ↪ data=test_data[test_data["location"] == "B"], time_limit=60*10)
      print("\033[1m" + "Calculating feature importance for location C..." +
      ↪ "\033[0m")
      predictors[2].feature_importance(feature_stage="original",
      ↪ data=test_data[test_data["location"] == "C"], time_limit=60*10)
```

```
These features in provided data are not utilized by the predictor and will be
ignored: ['ds', 'elevation:m', 'snow_drift:idx', 'location', 'prediction']
Computing feature importance via permutation shuffling for 41 features using 361
rows with 10 shuffle sets... Time limit: 600s...
```

```
Calculating feature importance for location A...
```

```
6689.13s          = Expected runtime (668.91s per shuffle set)
560.66s = Actual runtime (Completed 8 of 10 shuffle sets) (Early
stopping due to lack of time...)
```

```
These features in provided data are not utilized by the predictor and will be
ignored: ['ds', 'elevation:m', 'location', 'prediction']
Computing feature importance via permutation shuffling for 42 features using 361
rows with 10 shuffle sets... Time limit: 600s...
```

```
Calculating feature importance for location B...
```

1815.09s = Expected runtime (181.51s per shuffle set)

307.41s = Actual runtime (Completed 10 of 10 shuffle sets)

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'snow_drift:idx', 'location', 'prediction']

Computing feature importance via permutation shuffling for 41 features using 360 rows with 10 shuffle sets... Time limit: 600s...

Calculating feature importance for location C...

1336.22s = Expected runtime (133.62s per shuffle set)

```
[ ]: # save this notebook to submissions folder
import subprocess
import os
#subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
↳join('notebook_pdfs', f"{new_filename}_automatic_save.pdf"),
↳"autogluon_each_location.ipynb"])
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.
↳join('notebook_pdfs', f"{new_filename}.pdf"), "autogluon_each_location.
↳ipynb"])

[ ]: # import subprocess

# def execute_git_command(directory, command):
#     """Execute a Git command in the specified directory."""
#     try:
#         result = subprocess.check_output(['git', '-C', directory] + command,
↳stderr=subprocess.STDOUT)
#         return result.decode('utf-8').strip(), True
#     except subprocess.CalledProcessError as e:
#         print(f"Git command failed with message: {e.output.decode('utf-8').
↳strip()}")
#         return e.output.decode('utf-8').strip(), False

# git_repo_path = "."

# execute_git_command(git_repo_path, ['config', 'user.email',
↳'henrikskog01@gmail.com'])
# execute_git_command(git_repo_path, ['config', 'user.name', hello if hello is
↳not None else 'Henrik eller Jørgen'])

# branch_name = new_filename

# # add datetime to branch name
# branch_name += f"_{pd.Timestamp.now().strftime('%Y-%m-%d_%H-%M-%S')}"

# commit_msg = "run result"
```



```

# execute_git_command(git_repo_path, ['checkout', '-b',branch_name])

# # Navigate to your repo and commit changes
# execute_git_command(git_repo_path, ['add', '.'])
# execute_git_command(git_repo_path, ['commit', '-m',commit_msg])

# # Push to remote
# output, success = execute_git_command(git_repo_path, ['push',
↳ 'origin',branch_name])

# # If the push fails, try setting an upstream branch and push again
# if not success and 'upstream' in output:
#     print("Attempting to set upstream and push again...")
#     execute_git_command(git_repo_path, ['push', '--set-upstream',
↳ 'origin',branch_name])
#     execute_git_command(git_repo_path, ['push', 'origin', 'henrik_branch'])

# execute_git_command(git_repo_path, ['checkout', 'main'])

```

[]: