

## short\_1

November 10, 2023

### 1 Install necessary packages

```
[1]: !pip install autogluon matplotlib
```

```
Requirement already satisfied: autogluon in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (0.8.2)
Requirement already satisfied: matplotlib in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (3.6.3)
Collecting autogluon.core[all]==0.8.2 (from autogluon)
  Obtaining dependency information for autogluon.core[all]==0.8.2 from https://files.pythonhosted.org/packages/e0/56/545adb1d388e78591cd7e36de0c8b889c1944de362bdaeec0f31d01890df/autogluon.core-0.8.2-py3-none-any.whl.metadata
  Using cached autogluon.core-0.8.2-py3-none-any.whl.metadata (12 kB)
Collecting autogluon.features==0.8.2 (from autogluon)
  Obtaining dependency information for autogluon.features==0.8.2 from https://files.pythonhosted.org/packages/bb/ea/7892719f78a30aee1bf42c4a0540fbae98bfbd56b85fab79ffc437eb687/autogluon.features-0.8.2-py3-none-any.whl.metadata
  Using cached autogluon.features-0.8.2-py3-none-any.whl.metadata (11 kB)
Collecting autogluon.tabular[all]==0.8.2 (from autogluon)
  Obtaining dependency information for autogluon.tabular[all]==0.8.2 from https://files.pythonhosted.org/packages/f3/dc/0bd8cadd9a5e2f3e5b12caaa6745357d912ffc7b8b75fb4e426a38331028/autogluon.tabular-0.8.2-py3-none-any.whl.metadata
  Using cached autogluon.tabular-0.8.2-py3-none-any.whl.metadata (13 kB)
Collecting autogluon.multimodal==0.8.2 (from autogluon)
  Obtaining dependency information for autogluon.multimodal==0.8.2 from https://files.pythonhosted.org/packages/f4/df/a3921edb866555154d5a53adabcd2268fd3b53071820b6382f2619ad439f3/autogluon.multimodal-0.8.2-py3-none-any.whl.metadata
  Using cached autogluon.multimodal-0.8.2-py3-none-any.whl.metadata (13 kB)
Collecting autogluon.timeseries[all]==0.8.2 (from autogluon)
  Obtaining dependency information for autogluon.timeseries[all]==0.8.2 from https://files.pythonhosted.org/packages/50/38/1768d30684292d064dad4457355d6928581c5b832a3f65a6a023df2ff4ff/autogluon.timeseries-0.8.2-py3-none-any.whl.metadata
  Using cached autogluon.timeseries-0.8.2-py3-none-any.whl.metadata (12 kB)
Requirement already satisfied: numpy<1.27,>=1.21 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (1.24.4)
Requirement already satisfied: scipy<1.12,>=1.5.4 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
```

```

autogluon.core[all]==0.8.2->autogluon) (1.11.3)
Requirement already satisfied: scikit-learn<1.3,>=1.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (1.2.2)
Requirement already satisfied: networkx<4,>=3.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (3.1)
Requirement already satisfied: pandas<1.6,>=1.4.1 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (1.5.3)
Requirement already satisfied: tqdm<5,>=4.38 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (4.66.1)
Requirement already satisfied: requests in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (2.31.0)
Requirement already satisfied: boto3<2,>=1.10 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (1.28.60)
Collecting autogluon.common==0.8.2 (from autogluon.core[all]==0.8.2->autogluon)
  Obtaining dependency information for autogluon.common==0.8.2 from https://file
s.pythonhosted.org/packages/67/c7/aa2bd5708c9329a29245aa0cd955cace0662b92a82eb76
05cf46b6e7e9d5/autogluon.common-0.8.2-py3-none-any.whl.metadata
  Using cached autogluon.common-0.8.2-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: hyperopt<0.2.8,>=0.2.7 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (0.2.7)
Collecting ray[default]<2.4,>=2.3 (from autogluon.core[all]==0.8.2->autogluon)
  Downloading ray-2.3.1-cp310-cp310-macosx_11_0_arm64.whl (28.6 MB)

28.6/28.6 MB 8.8 MB/s eta 0:00:00m eta
0:00:01[36m0:00:01
Requirement already satisfied: pydantic<2.0,>=1.10.4 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.core[all]==0.8.2->autogluon) (1.10.13)
Collecting grpcio<=1.50.0,>=1.42.0 (from autogluon.core[all]==0.8.2->autogluon)
  Downloading grpcio-1.50.0.tar.gz (22.1 MB)

22.1/22.1 MB 9.9 MB/s eta 0:00:00m eta
0:00:010:01:01
  Preparing metadata (setup.py) ... done
Requirement already satisfied: Pillow<9.6,>=9.3 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
autogluon.multimodal==0.8.2->autogluon) (9.5.0)
Requirement already satisfied: jsonschema<4.18,>=4.14 in
/Users/jorgensandhaug/.local/lib/python3.10/site-packages (from
autogluon.multimodal==0.8.2->autogluon) (4.17.3)
Requirement already satisfied: seqeval<1.3.0,>=1.2.2 in

```

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (1.2.2)  
 Requirement already satisfied: evaluate<0.4.0,>=0.2.2 in  
 /Users/jorgensandhaug/.local/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.3.0)  
 Requirement already satisfied: accelerate<0.17,>=0.9 in  
 /Users/jorgensandhaug/.local/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.16.0)  
 Requirement already satisfied: timm<0.10.0,>=0.9.2 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.9.7)  
 Requirement already satisfied: torch<1.14,>=1.9 in  
 /Users/jorgensandhaug/.local/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (1.13.1)  
 Requirement already satisfied: torchvision<0.15.0 in  
 /Users/jorgensandhaug/.local/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.14.1)  
 Requirement already satisfied: scikit-image<0.20.0,>=0.19.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.19.3)  
 Requirement already satisfied: pytorch-lightning<1.10.0,>=1.9.0 in  
 /Users/jorgensandhaug/.local/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (1.9.5)  
 Requirement already satisfied: text-unidecode<1.4,>=1.3 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (1.3)  
 Requirement already satisfied: torchmetrics<0.12.0,>=0.11.0 in  
 /Users/jorgensandhaug/.local/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.11.4)  
 Requirement already satisfied: transformers[sentencepiece]<4.27.0,>=4.23.0 in  
 /Users/jorgensandhaug/.local/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (4.26.1)  
 Requirement already satisfied: nptyping<2.5.0,>=1.4.4 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (2.4.1)  
 Requirement already satisfied: omegaconf<2.3.0,>=2.1.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (2.2.3)  
 Requirement already satisfied: pytorch-metric-learning<2.0,>=1.3.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (1.7.3)  
 Requirement already satisfied: nlpaug<1.2.0,>=1.1.10 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (1.1.11)  
 Requirement already satisfied: nltk<4.0.0,>=3.4.5 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (3.8.1)  
 Requirement already satisfied: openmim<0.4.0,>=0.3.7 in

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.3.7)  
 Requirement already satisfied: defusedxml<0.7.2,>=0.7.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.7.1)  
 Requirement already satisfied: jinja2<3.2,>=3.0.3 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (3.1.2)  
 Requirement already satisfied: tensorboard<3,>=2.9 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (2.14.1)  
 Requirement already satisfied: pytesseract<0.3.11,>=0.3.9 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.multimodal==0.8.2->autogluon) (0.3.10)  
 Requirement already satisfied: lightgbm<3.4,>=3.3 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.tabular[all]==0.8.2->autogluon) (3.3.5)  
 Requirement already satisfied: xgboost<1.8,>=1.6 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.tabular[all]==0.8.2->autogluon) (1.7.6)  
 Requirement already satisfied: fastai<2.8,>=2.3.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.tabular[all]==0.8.2->autogluon) (2.7.12)  
 Requirement already satisfied: catboost<1.3,>=1.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.tabular[all]==0.8.2->autogluon) (1.1.1)  
 Requirement already satisfied: joblib<2,>=1.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.timeseries[all]==0.8.2->autogluon) (1.3.2)  
 Requirement already satisfied: statsmodels<0.15,>=0.13.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.timeseries[all]==0.8.2->autogluon) (0.14.0)  
 Requirement already satisfied: gluonts<0.14,>=0.13.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.timeseries[all]==0.8.2->autogluon) (0.13.5)  
 Requirement already satisfied: statsforecast<1.5,>=1.4.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.timeseries[all]==0.8.2->autogluon) (1.4.0)  
 Requirement already satisfied: mlforecast<0.7.4,>=0.7.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.timeseries[all]==0.8.2->autogluon) (0.7.3)  
 Requirement already satisfied: ujson<6,>=5 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.timeseries[all]==0.8.2->autogluon) (5.8.0)  
 Requirement already satisfied: psutil<6,>=5.7.3 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.common==0.8.2->autogluon.core[all]==0.8.2->autogluon) (5.9.5)  
 Requirement already satisfied: setuptools in

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 autogluon.common==0.8.2->autogluon.core[all]==0.8.2->autogluon) (68.0.0)  
 Requirement already satisfied: contourpy>=1.0.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from matplotlib  
 (1.1.1))  
 Requirement already satisfied: cycler>=0.10 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from matplotlib  
 (0.12.0))  
 Requirement already satisfied: fonttools>=4.22.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from matplotlib  
 (4.43.0))  
 Requirement already satisfied: kiwisolver>=1.0.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from matplotlib  
 (1.4.5))  
 Requirement already satisfied: packaging>=20.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from matplotlib  
 (23.2))  
 Requirement already satisfied: pyparsing>=2.2.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from matplotlib  
 (3.1.1))  
 Requirement already satisfied: python-dateutil>=2.7 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from matplotlib  
 (2.8.2))  
 Requirement already satisfied: six>=1.5 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from python-  
 dateutil>=2.7->matplotlib) (1.16.0)  
 Requirement already satisfied: pyyaml in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 accelerate<0.17,>=0.9->autogluon.multimodal==0.8.2->autogluon) (6.0.1)  
 Requirement already satisfied: botocore<1.32.0,>=1.31.60 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 boto3<2,>=1.10->autogluon.core[all]==0.8.2->autogluon) (1.31.60)  
 Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 boto3<2,>=1.10->autogluon.core[all]==0.8.2->autogluon) (1.0.1)  
 Requirement already satisfied: s3transfer<0.8.0,>=0.7.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 boto3<2,>=1.10->autogluon.core[all]==0.8.2->autogluon) (0.7.0)  
 Requirement already satisfied: graphviz in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 catboost<1.3,>=1.1->autogluon.tabular[all]==0.8.2->autogluon) (0.8.4)  
 Requirement already satisfied: plotly in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 catboost<1.3,>=1.1->autogluon.tabular[all]==0.8.2->autogluon) (5.17.0)  
 Requirement already satisfied: datasets>=2.0.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon) (2.14.5)  
 Requirement already satisfied: dill in

```

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon) (0.3.7)
Requirement already satisfied: xxhash in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon) (3.4.1)
Requirement already satisfied: multiprocessing in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon) (0.70.15)
Requirement already satisfied: fsspec[http]>=2021.05.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon) (2023.6.0)
Requirement already satisfied: huggingface-hub>=0.7.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon) (0.17.3)
Requirement already satisfied: responses<0.19 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon) (0.18.0)
Requirement already satisfied: pip in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (23.2.1)
Requirement already satisfied: fastdownload<2,>=0.0.5 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (0.0.7)
Requirement already satisfied: fastcore<1.6,>=1.5.29 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (1.5.29)
Requirement already satisfied: fastprogress>=0.2.4 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (1.0.3)
Requirement already satisfied: spacy<4 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (3.6.1)
Requirement already satisfied: toolz~=0.10 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
gluonts<0.14,>=0.13.1->autogluon.timeseries[all]==0.8.2->autogluon) (0.12.0)
Requirement already satisfied: typing-extensions~=4.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
gluonts<0.14,>=0.13.1->autogluon.timeseries[all]==0.8.2->autogluon) (4.8.0)
Requirement already satisfied: future in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
hyperopt<0.2.8,>=0.2.7->autogluon.core[all]==0.8.2->autogluon) (0.18.3)
Requirement already satisfied: cloudpickle in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
hyperopt<0.2.8,>=0.2.7->autogluon.core[all]==0.8.2->autogluon) (2.2.1)
Requirement already satisfied: py4j in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
hyperopt<0.2.8,>=0.2.7->autogluon.core[all]==0.8.2->autogluon) (0.10.9.7)
Requirement already satisfied: MarkupSafe>=2.0 in

```

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 jinja2<3.2,>=3.0.3->autogluon.multimodal==0.8.2->autogluon) (2.1.3)  
 Requirement already satisfied: attrs>=17.4.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 jsonschema<4.18,>=4.14->autogluon.multimodal==0.8.2->autogluon) (23.1.0)  
 Requirement already satisfied: pyparsing!=0.17.0,!0.17.1,!0.17.2,>=0.14.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 jsonschema<4.18,>=4.14->autogluon.multimodal==0.8.2->autogluon) (0.19.3)  
 Requirement already satisfied: wheel in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 lightgbm<3.4,>=3.3->autogluon.tabular[all]==0.8.2->autogluon) (0.41.2)  
 Requirement already satisfied: numba in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 mlforecast<0.7.4,>=0.7.0->autogluon.timeseries[all]==0.8.2->autogluon) (0.57.1)  
 Requirement already satisfied: window-ops in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 mlforecast<0.7.4,>=0.7.0->autogluon.timeseries[all]==0.8.2->autogluon) (0.0.14)  
 Requirement already satisfied: gdown>=4.0.0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 nlpaug<1.2.0,>=1.1.10->autogluon.multimodal==0.8.2->autogluon) (4.7.1)  
 Requirement already satisfied: click in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 nltk<4.0.0,>=3.4.5->autogluon.multimodal==0.8.2->autogluon) (8.1.7)  
 Requirement already satisfied: regex>=2021.8.3 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 nltk<4.0.0,>=3.4.5->autogluon.multimodal==0.8.2->autogluon) (2023.10.3)  
 Requirement already satisfied: antlr4-python3-runtime==4.9.\* in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 omegaconf<2.3.0,>=2.1.1->autogluon.multimodal==0.8.2->autogluon) (4.9.3)  
 Requirement already satisfied: colorama in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon) (0.4.6)  
 Requirement already satisfied: model-index in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon) (0.1.11)  
 Requirement already satisfied: rich in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon) (13.6.0)  
 Requirement already satisfied: tabulate in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon) (0.9.0)  
 Requirement already satisfied: pytz>=2020.1 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
 pandas<1.6,>=1.4.1->autogluon.core[all]==0.8.2->autogluon) (2023.3.post1)  
 Requirement already satisfied: lightning-utilities>=0.6.0.post0 in  
 /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from pytorch-  
 lightning<1.10.0,>=1.9.0->autogluon.multimodal==0.8.2->autogluon) (0.9.0)  
 Requirement already satisfied: filelock in

```

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (3.12.4)
Requirement already satisfied: msgpack<2.0.0,>=1.0.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (1.0.7)
Requirement already satisfied: protobuf!=3.19.5,>=3.15.3 in
/Users/jorgensandhaug/.local/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (3.20.2)
Requirement already satisfied: aiosignal in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (1.3.1)
Requirement already satisfied: frozenlist in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (1.4.0)
Requirement already satisfied: virtualenv>=20.0.24 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (20.21.0)
Using cached grpcio-1.49.1-cp310-cp310-macosx_11_0_arm64.whl
Requirement already satisfied: aiohttp>=3.7 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (3.8.5)
Requirement already satisfied: aiohttp-cors in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (0.7.0)
Requirement already satisfied: colorful in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (0.5.5)
Requirement already satisfied: py-spy>=0.2.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (0.3.14)
Requirement already satisfied: gpustat>=1.0.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (1.1.1)
Requirement already satisfied: opencensus in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (0.11.3)
Requirement already satisfied: prometheus-client>=0.7.1 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (0.17.1)
Requirement already satisfied: smart-open in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (6.4.0)
Requirement already satisfied: tensorboardX>=1.9 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (2.6.2.2)
Requirement already satisfied: charset-normalizer<4,>=2 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
requests->autogluon.core[all]==0.8.2->autogluon) (3.3.0)

```



Requirement already satisfied: idna<4,>=2.5 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
requests->autogluon.core[all]==0.8.2->autogluon) (3.4)

Requirement already satisfied: urllib3<3,>=1.21.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
requests->autogluon.core[all]==0.8.2->autogluon) (1.26.17)

Requirement already satisfied: certifi>=2017.4.17 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
requests->autogluon.core[all]==0.8.2->autogluon) (2023.7.22)

Requirement already satisfied: imageio>=2.4.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from scikit-  
image<0.20.0,>=0.19.1->autogluon.multimodal==0.8.2->autogluon) (2.31.1)

Requirement already satisfied: tifffile>=2019.7.26 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from scikit-  
image<0.20.0,>=0.19.1->autogluon.multimodal==0.8.2->autogluon) (2023.9.26)

Requirement already satisfied: PyWavelets>=1.1.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from scikit-  
image<0.20.0,>=0.19.1->autogluon.multimodal==0.8.2->autogluon) (1.4.1)

Requirement already satisfied: threadpoolctl>=2.0.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from scikit-  
learn<1.3,>=1.0->autogluon.core[all]==0.8.2->autogluon) (3.2.0)

Requirement already satisfied: patsy>=0.5.2 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
statsmodels<0.15,>=0.13.0->autogluon.timeseries[all]==0.8.2->autogluon) (0.5.3)

Requirement already satisfied: absl-py>=0.4 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon) (2.0.0)

Requirement already satisfied: google-auth<3,>=1.6.3 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon) (2.23.2)

Requirement already satisfied: google-auth-oauthlib<1.1,>=0.5 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon) (1.0.0)

Requirement already satisfied: markdown>=2.6.8 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon) (3.4.4)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon) (0.7.0)

Requirement already satisfied: werkzeug>=1.0.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon) (3.0.0)

Requirement already satisfied: safetensors in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
timm<0.10.0,>=0.9.2->autogluon.multimodal==0.8.2->autogluon) (0.3.3)

Requirement already satisfied: tokenizers!=0.11.3,<0.14,>=0.11.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from transformers[  
sentencepiece]<4.27.0,>=4.23.0->autogluon.multimodal==0.8.2->autogluon) (0.13.3)

Requirement already satisfied: sentencepiece!=0.1.92,>=0.1.91 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from transformers[  
sentencepiece]<4.27.0,>=4.23.0->autogluon.multimodal==0.8.2->autogluon) (0.1.99)

Requirement already satisfied: multidict<7.0,>=4.5 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
aiohttp>=3.7->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon)  
(6.0.4)

Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
aiohttp>=3.7->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon)  
(4.0.3)

Requirement already satisfied: yarl<2.0,>=1.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
aiohttp>=3.7->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon)  
(1.9.2)

Requirement already satisfied: pyarrow>=8.0.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
datasets>=2.0.0->evaluate<0.4.0,>=0.2.2->autogluon.multimodal==0.8.2->autogluon)  
(13.0.0)

Requirement already satisfied: beautifulsoup4 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
gdown>=4.0.0->nlpaug<1.2.0,>=1.1.10->autogluon.multimodal==0.8.2->autogluon)  
(4.12.2)

Requirement already satisfied: cachetools<6.0,>=2.0.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from google-  
auth<3,>=1.6.3->tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon)  
(5.3.1)

Requirement already satisfied: pyasn1-modules>=0.2.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from google-  
auth<3,>=1.6.3->tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon)  
(0.3.0)

Requirement already satisfied: rsa<5,>=3.1.4 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from google-  
auth<3,>=1.6.3->tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon)  
(4.9)

Requirement already satisfied: requests-oauthlib>=0.7.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from google-auth-  
oauthlib<1.1,>=0.5->tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon)  
(1.3.1)

Requirement already satisfied: nvidia-ml-py>=11.450.129 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
gpustat>=1.0.0->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon)  
(12.535.108)

Requirement already satisfied: blessed>=1.17.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
gpustat>=1.0.0->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon)  
(1.20.0)

Requirement already satisfied: llvmlite<0.41,>=0.40.0dev0 in

```

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
numba>mlforecast<0.7.4,>=0.7.0->autogluon.timeseries[all]==0.8.2->autogluon)
(0.40.1)
Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (3.0.12)
Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (1.0.5)
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (1.0.10)
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (2.0.8)
Requirement already satisfied: preshed<3.1.0,>=3.0.2 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (3.0.9)
Requirement already satisfied: thinc<8.2.0,>=8.1.8 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (8.1.12)
Requirement already satisfied: wasabi<1.2.0,>=0.9.1 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (1.1.2)
Requirement already satisfied: srsly<3.0.0,>=2.4.3 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (2.4.8)
Requirement already satisfied: catalogue<2.1.0,>=2.0.6 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (2.0.10)
Requirement already satisfied: typer<0.10.0,>=0.3.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (0.9.0)
Requirement already satisfied: pathy>=0.10.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (0.10.1)
Requirement already satisfied: langcodes<4.0.0,>=3.2.0 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon) (3.3.0)
Requirement already satisfied: distlib<1,>=0.3.6 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from virtualenv>=2
0.0.24->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (0.3.7)
Requirement already satisfied: platformdirs<4,>=2.4 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from virtualenv>=2
0.0.24->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon) (3.10.0)
Requirement already satisfied: ordered-set in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from model-
index->openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon) (4.1.0)

```

Requirement already satisfied: opencensus-context>=0.1.3 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
opencensus->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon)  
(0.1.3)

Requirement already satisfied: google-api-core<3.0.0,>=1.0.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
opencensus->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->autogluon)  
(2.12.0)

Requirement already satisfied: tenacity>=6.2.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
plotly->catboost<1.3,>=1.1->autogluon.tabular[all]==0.8.2->autogluon) (8.2.3)

Requirement already satisfied: markdown-it-py>=2.2.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
rich->openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
rich->openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon) (2.16.1)

Requirement already satisfied: wcwidth>=0.1.4 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from blessed>=1.17  
.1->gpustat>=1.0.0->ray[default]<2.4,>=2.3->autogluon.core[all]==0.8.2->  
autogluon) (0.2.5)

Requirement already satisfied: googleapis-common-protos<2.0.dev0,>=1.56.2 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from google-api-co  
re<3.0.0,>=1.0.0->opencensus->ray[default]<2.4,>=2.3->  
autogluon.core[all]==0.8.2->autogluon) (1.60.0)

Requirement already satisfied: mdurl~=0.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from markdown-it-  
py>=2.2.0->rich->openmim<0.4.0,>=0.3.7->autogluon.multimodal==0.8.2->autogluon)  
(0.1.0)

Requirement already satisfied: pyasn1<0.6.0,>=0.4.6 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from  
pyasn1-modules>=0.2.1->google-  
auth<3,>=1.6.3->tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon)  
(0.5.0)

Requirement already satisfied: oauthlib>=3.0.0 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from requests-  
oauthlib>=0.7.0->google-auth-  
oauthlib<1.1,>=0.5->tensorboard<3,>=2.9->autogluon.multimodal==0.8.2->autogluon)  
(3.2.2)

Requirement already satisfied: blis<0.8.0,>=0.7.8 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from thinc<8.2.0,>  
=8.1.8->spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon)  
(0.7.10)

Requirement already satisfied: confection<1.0.0,>=0.0.1 in  
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from thinc<8.2.0,>  
=8.1.8->spacy<4->fastai<2.8,>=2.3.1->autogluon.tabular[all]==0.8.2->autogluon)  
(0.1.3)

Requirement already satisfied: soupsieve>1.2 in

```

/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from beautifulsoup
4->gdown>=4.0.0->nlpaug<1.2.0,>=1.1.10->autogluon.multimodal==0.8.2->autogluon)
(2.5)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in
/opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages (from
requests->autogluon.core[all]==0.8.2->autogluon) (1.7.1)
Downloading autogluon.features-0.8.2-py3-none-any.whl (62 kB)

62.1/62.1 kB 5.7 MB/s eta 0:00:00
Downloading autogluon.multimodal-0.8.2-py3-none-any.whl (372 kB)

372.3/372.3 kB 9.4 MB/s eta 0:00:00m eta
0:00:01
Downloading autogluon.common-0.8.2-py3-none-any.whl (61 kB)

61.1/61.1 kB 5.6 MB/s eta 0:00:00
Downloading autogluon.core-0.8.2-py3-none-any.whl (224 kB)

224.0/224.0 kB 2.3 MB/s eta 0:00:00 MB/s eta
0:00:01:01
Downloading autogluon.tabular-0.8.2-py3-none-any.whl (285 kB)

285.7/285.7 kB 9.4 MB/s eta 0:00:00
Downloading autogluon.timeseries-0.8.2-py3-none-any.whl (116 kB)

116.3/116.3 kB 7.7 MB/s eta 0:00:00
Installing collected packages: grpcio, ray, autogluon.common,
autogluon.features, autogluon.core, autogluon.tabular, autogluon.multimodal,
autogluon.timeseries
  Attempting uninstall: ray
    Found existing installation: ray 2.6.3
    Uninstalling ray-2.6.3:
      Successfully uninstalled ray-2.6.3
  Attempting uninstall: autogluon.common
    Found existing installation: autogluon.common 0.8.1
    Uninstalling autogluon.common-0.8.1:
^C
ERROR: Operation cancelled by user

```

## 2 Config

```

[45]: # config

label = 'y'
metric = 'mean_absolute_error'

```

```

time_limit = None
presets = "experimental_zeroshot_hpo_hybrid" #'best_quality'

use_is_estimated_attr = True

drop_night_outliers = True

# to_drop = ["snow_drift:idx", "snow_density:kgm3", "wind_speed_w_1000hPa:ms",
↳ "dew_or_rime:idx", "prob_rime:p", "fresh_snow_12h:cm", "fresh_snow_24h:cm",
↳ "wind_speed_u_10m:ms", "wind_speed_v_10m:ms", "snow_melt_10min:mm",
↳ "rain_water:kgm2", "dew_point_2m:K", "precip_5min:mm", "absolute_humidity_2m:
↳ gm3", "air_density_2m:kgm3"]#, "msl_pressure:hPa", "pressure_50m:hPa",
↳ "pressure_100m:hPa"]
to_drop = []#["wind_speed_w_1000hPa:ms", "wind_speed_u_10m:ms",
↳ "wind_speed_v_10m:ms"]

excluded_model_types = ['CAT', 'XGB', 'RF']

num_stack_levels = 0
num_bag_folds = 0#8# 8
num_bag_sets = 0##20

use_tune_data = True
use_test_data = True
use_bag_holdout = True

clip_predictions = True

```

### 3 Loading and preprocessing

```

[46]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings("ignore")

def feature_engineering(X):
    # shift columns with them by 1 hour, so that for index 16:00, we have the
↳ values from 17:00
    columns = ['clear_sky_energy_1h:J', 'diffuse_rad_1h:J', 'direct_rad_1h:J',
               'fresh_snow_12h:cm', 'fresh_snow_1h:cm', 'fresh_snow_24h:cm',
               'fresh_snow_3h:cm', 'fresh_snow_6h:cm']

```

```

# Filter rows where index.minute == 0
X_shifted = X[X.index.minute == 0][columns].copy()

# Create a set for constant-time lookup
index_set = set(X.index)

# Vectorized time shifting
one_hour = pd.Timedelta('1 hour')
shifted_indices = X_shifted.index + one_hour
X_shifted.loc[shifted_indices.isin(index_set)] = X.
↳loc[shifted_indices[shifted_indices.isin(index_set)]] [columns]

# set last row to same as second last row
X_shifted.iloc[-1] = X_shifted.iloc[-2]

# Rename columns
X_old_unshifted = X_shifted.copy()
X_old_unshifted.columns = [f"{col}_not_shifted" for col in X_old_unshifted.
↳columns]

date_calc = None
# If 'date_calc' is present, handle it
if 'date_calc' in X.columns:
    date_calc = X[X.index.minute == 0]['date_calc']

# resample to hourly
X = X.resample('H').mean()

# overwrite columns with shifted columns
X[columns] = X_shifted[columns]

if date_calc is not None:
    X['date_calc'] = date_calc

return X

def fix_X(X, name):
    # Convert 'date_forecast' to datetime format and replace original column
    ↳with 'ds'
    X['ds'] = pd.to_datetime(X['date_forecast'])

```

```

X.drop(columns=['date_forecast'], inplace=True, errors='ignore')
X.sort_values(by='ds', inplace=True)
X.set_index('ds', inplace=True)

X = feature_engineering(X)

return X

def handle_features(X_train_observed, X_train_estimated, X_test, y_train):
    X_train_observed = fix_X(X_train_observed, "X_train_observed")
    X_train_estimated = fix_X(X_train_estimated, "X_train_estimated")
    X_test = fix_X(X_test, "X_test")

    y_train['ds'] = pd.to_datetime(y_train['time'])
    y_train.drop(columns=['time'], inplace=True)
    y_train.sort_values(by='ds', inplace=True)
    y_train.set_index('ds', inplace=True)

    return X_train_observed, X_train_estimated, X_test, y_train

def preprocess_data(X_train_observed, X_train_estimated, X_test, y_train,
location):
    # convert to datetime
    X_train_observed, X_train_estimated, X_test, y_train =
handle_features(X_train_observed, X_train_estimated, X_test, y_train)

    if use_is_estimated_attr:
        X_train_observed["is_estimated"] = 0
        X_train_estimated["is_estimated"] = 1
        X_test["is_estimated"] = 1

    # drop date_calc
    X_train_estimated.drop(columns=['date_calc'], inplace=True)
    X_test.drop(columns=['date_calc'], inplace=True)

    y_train["y"] = y_train["pv_measurement"].astype('float64')
    y_train.drop(columns=['pv_measurement'], inplace=True)
    X_train = pd.concat([X_train_observed, X_train_estimated])

```



```

    # clip all y values to 0 if negative
    y_train["y"] = y_train["y"].clip(lower=0)

    X_train = pd.merge(X_train, y_train, how="inner", left_index=True,
↳right_index=True)

    X_train["location"] = location
    X_test["location"] = location

    return X_train, X_test

# Define locations
locations = ['A', 'B', 'C']

X_trains = []
X_tests = []
# Loop through locations
for loc in locations:
    print(f"Processing location {loc}...")
    # Read target training data
    y_train = pd.read_parquet(f'{loc}/train_targets.parquet')

    # Read estimated training data and add location feature
    X_train_estimated = pd.read_parquet(f'{loc}/X_train_estimated.parquet')

    # Read observed training data and add location feature
    X_train_observed = pd.read_parquet(f'{loc}/X_train_observed.parquet')

    # Read estimated test data and add location feature
    X_test_estimated = pd.read_parquet(f'{loc}/X_test_estimated.parquet')

    # Preprocess data
    X_train, X_test = preprocess_data(X_train_observed, X_train_estimated,
↳X_test_estimated, y_train, loc)

    X_trains.append(X_train)
    X_tests.append(X_test)

# Concatenate all data and save to csv
X_train = pd.concat(X_trains)
X_test = pd.concat(X_tests)

```

```

Processing location A...
Processing location B...
Processing location C...

```

## 3.1 Feature engineering

### 3.1.1 Remove anomalies

```
[47]: import pandas as pd

def replace_streaks_with_nan(df, max_streak_length, column="y"):
    for location in df["location"].unique():
        x = df[df["location"] == location][column].copy()

        last_val = None
        streak_length = 1
        streak_indices = []
        allowed = [0]
        found_streaks = {}

        for idx in x.index:
            value = x[idx]

            if value == last_val and value not in allowed:
                streak_length += 1
                streak_indices.append(idx)
            else:
                streak_length = 1
                last_val = value
                streak_indices.clear()

            if streak_length > max_streak_length:
                found_streaks[value] = streak_length

                for streak_idx in streak_indices:
                    x[streak_idx] = np.nan
                streak_indices.clear() # clear after setting to NaN to avoid
↪ setting multiple times
                df.loc[df["location"] == location, column] = x

        print(f"Found streaks for location {location}: {found_streaks}")

    return df

X_train = replace_streaks_with_nan(X_train.copy(), 3, "y")
```

Found streaks for location A: {}

Found streaks for location B: {3.45: 28, 6.9: 7, 12.9375: 5, 13.8: 8, 276.0: 78, 18.975: 58, 0.8625: 4, 118.1625: 33, 34.5: 11, 183.7125: 1058, 87.1125: 7, 79.35: 34, 7.7625: 12, 27.6: 448, 273.41249999999997: 72, 264.78749999999997: 55, 169.05: 33, 375.1875: 56, 314.8125: 66, 76.7625: 10, 135.4125: 216, 81.9375: 202, 2.5875: 12, 81.075: 210}

Found streaks for location C: {9.8: 4, 29.400000000000002: 4, 19.6: 4}

```
[48]: # print num rows
temprows = len(X_train)
X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
               inplace=True)
print("Dropped rows: ", temprows - len(X_train))
```

Dropped rows: 9293

```
[49]: thresh = 0.1
mask = (X_train["direct_rad_1h:J"] <= thresh) & (X_train["diffuse_rad_1h:J"] <=
       thresh) & (X_train["y"] >= 0.1)

if drop_night_outliers:
    X_train.loc[mask, "y"] = np.nan
```

```
[50]: temprows = len(X_train)
X_train.dropna(subset=['y', 'direct_rad_1h:J', 'diffuse_rad_1h:J'],
               inplace=True)
print("Dropped rows: ", temprows - len(X_train))
```

Dropped rows: 1876

```
[51]: X_train.drop(columns=to_drop, inplace=True)
X_test.drop(columns=to_drop, inplace=True)

X_train.to_csv('X_train_raw.csv', index=True)
X_test.to_csv('X_test_raw.csv', index=True)
```

```
[52]: def split_and_shuffle_data(input_data, num_bins, frac1):
    # Validate the input fraction
    if frac1 < 0 or frac1 > 1:
        raise ValueError("frac1 must be between 0 and 1.")

    if frac1==1:
        return input_data, pd.DataFrame()

    # Calculate the fraction for the second output set
    frac2 = 1 - frac1

    # Calculate bin size
    bin_size = len(input_data) // num_bins

    # Initialize empty DataFrames for output
    output_data1 = pd.DataFrame()
    output_data2 = pd.DataFrame()
```

```

for i in range(num_bins):
    # Shuffle the data in the current bin
    np.random.seed(i)
    current_bin = input_data.iloc[i * bin_size: (i + 1) * bin_size].
    ↪sample(frac=1)

    # Calculate the sizes for each output set
    size1 = int(len(current_bin) * frac1)

    # Split and append to output DataFrames
    output_data1 = pd.concat([output_data1, current_bin.iloc[:size1]])
    output_data2 = pd.concat([output_data2, current_bin.iloc[size1:]])

    # Shuffle and split the remaining data
    remaining_data = input_data.iloc[num_bins * bin_size:].sample(frac=1)

    remaining_size1 = int(len(remaining_data) * frac1)

    output_data1 = pd.concat([output_data1, remaining_data.iloc[:
    ↪remaining_size1]])
    output_data2 = pd.concat([output_data2, remaining_data.iloc[remaining_size1:
    ↪]])

    return output_data1, output_data2

```

```

[53]: from autogluon.tabular import TabularDataset, TabularPredictor
data = TabularDataset('X_train_raw.csv')
data['ds'] = pd.to_datetime(data['ds'])
data = data.sort_values(by='ds')

split_time = pd.to_datetime("2022-10-28 22:00:00")
train_set = TabularDataset(data[data["ds"] < split_time])
estimated_set = TabularDataset(data[data["ds"] >= split_time]) # only estimated

test_set = pd.DataFrame()
tune_set = pd.DataFrame()
new_train_set = pd.DataFrame()

for location in locations:
    loc_data = data[data["location"] == location]
    num_train_rows = len(loc_data)

    tune_rows = 1500.0 # 2500.0
    if use_test_data:
        tune_rows = 1880.0 # max(3000.0,
    ↪len(estimated_set[estimated_set["location"] == location]))

```

```

    holdout_frac = max(0.01, min(0.1, tune_rows / num_train_rows)) *
    ↪ num_train_rows / len(estimated_set[estimated_set["location"] == location])

    # shuffle and split data
    loc_tune_set, loc_new_train_set =
    ↪ split_and_shuffle_data(estimated_set[estimated_set['location'] == location],
    ↪ 40, holdout_frac)

    new_train_set = pd.concat([new_train_set, loc_new_train_set])

    if use_test_data:
        loc_test_set, loc_tune_set = split_and_shuffle_data(loc_tune_set, 40, 0.
    ↪ 2)
        test_set = pd.concat([test_set, loc_test_set])

    tune_set = pd.concat([tune_set, loc_tune_set])

# add rest to train_set
train_set = pd.concat([train_set, new_train_set])
tuning_data = tune_set
if use_test_data:
    test_data = test_set

train_data = train_set

train_data = TabularDataset(train_data)
tuning_data = TabularDataset(tuning_data)

if use_test_data:
    test_data = TabularDataset(test_data)

```

Loaded data from: X\_train\_raw.csv | Columns = 49 / 49 | Rows = 87917 -> 87917

## 4 Modeling

```

[54]: import os

# if submissions folder does not exist, create it
if not os.path.exists('submissions'):
    os.makedirs('submissions')

# Get the last submission number
last_submission_number = int(max([int(filename.split('_')[1].split('.')[0]) for
    ↪ filename in os.listdir('submissions') if "submission" in filename]))
print("Last submission number:", last_submission_number)

```

```

print("Now creating submission number:", last_submission_number + 1)

# Create the new filename
new_filename = f'submission_{last_submission_number + 1}'

print("New filename:", new_filename)

```

Last submission number: 133  
 Now creating submission number: 134  
 New filename: submission\_134

```
[55]: predictors = [None, None, None]
```

```

[56]: def fit_predictor_for_location(loc):
    # All of these hyperparameters have been found by experimenting with some
    ↪ standard parameters in AutoGluon, and then only using the best ones for each
    ↪ location to make the train time shorter
    r118 = {'extra_trees': True, 'feature_fraction': 0.7832570544199176,
    ↪ 'learning_rate': 0.021720607471727896, 'min_data_in_leaf': 3, 'num_leaves':
    ↪ 21, 'ag_args': {'name_suffix': '_r118', 'priority': 17}}
    r51 = {'bs': 1024, 'emb_drop': 0.6046989241462619, 'epochs': 48, 'layers':
    ↪ [200, 100, 50], 'lr': 0.00775309042164966, 'ps': 0.09244767444160731,
    ↪ 'ag_args': {'name_suffix': '_r51', 'priority': 12}}
    r145 = {'bs': 128, 'emb_drop': 0.44339037504795686, 'epochs': 31, 'layers':
    ↪ [400, 200, 100], 'lr': 0.008615195908919904, 'ps': 0.19220253419114286,
    ↪ 'ag_args': {'name_suffix': '_r145', 'priority': 9}}
    lgbmXT = {'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}
    if loc == "A":
        hyperparameters = {
            'NN_TORCH': {},
            'GBM': [lgbmXT, r118], #, 'GBMLarge'], #, r_118],
            'FASTAI': [r51],
        }
    elif loc == "B":
        hyperparameters = {
            'NN_TORCH': {},
            'XT': [{'criterion': 'squared_error', 'ag_args': {'name_suffix':
    ↪ 'MSE', 'problem_types': ['regression', 'quantile']}}],
            'GBM': [lgbmXT, r118],
            'FASTAI': [{}, r51, r145],
        }
    elif loc == "C":
        hyperparameters = {
            'NN_TORCH': {},
            'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix':
    ↪ 'Unif'}}],

```

```

        'GBM': [lgbmXT, r118],
        'FASTAI': [r51],
    }

    predictor = TabularPredictor(
        label=label,
        eval_metric=metric,
        path=f"AutogluonModels/{new_filename}_{loc}",
    ).fit(
        train_data=train_data[train_data["location"] == loc].
↳drop(columns=["ds"]),
        time_limit=time_limit,
        presets=presets,
        num_stack_levels=num_stack_levels,
        num_bag_folds=num_bag_folds,
        num_bag_sets=num_bag_sets,
        tuning_data=tuning_data[tuning_data["location"] == loc].
↳reset_index(drop=True).drop(columns=["ds"]),
        hyperparameters = hyperparameters,
        use_bag_holdout=use_bag_holdout,
        #excluded_model_types=excluded_model_types
    )

    # evaluate on test data
    if use_test_data:
        t = test_data[test_data["location"] == loc]
        perf = predictor.evaluate(t)
        print("Evaluation on test data:")
        print(perf[predictor.eval_metric.name])

    return predictor

loc = "A"
predictors[0] = fit_predictor_for_location(loc)

```

Warning: path already exists! This predictor may overwrite an existing predictor! path="AutogluonModels/submission\_134\_A"

Presets specified: ['experimental\_zeroshot\_hpo\_hybrid']

Stack configuration (auto\_stack=True): num\_stack\_levels=0, num\_bag\_folds=0, num\_bag\_sets=0

Beginning AutoGluon training ...

AutoGluon will save models to "AutogluonModels/submission\_134\_A/"

AutoGluon Version: 0.8.2

Python Version: 3.10.12

Operating System: Darwin

Platform Machine: arm64

```

Platform Version:   Darwin Kernel Version 22.1.0: Sun Oct  9 20:15:09 PDT 2022;
root:xnu-8792.41.9~2/RELEASE_ARM64_T6000
Disk Space Avail:  101.27 GB / 494.38 GB (20.5%)
Train Data Rows:   30934
Train Data Columns: 47
Tuning Data Rows:  1485
Tuning Data Columns: 47
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (5733.42, 0.0, 673.41535, 1195.24)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                3225.39 MB
    Train Data (Original) Memory Usage: 13.81 MB (0.4% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 2 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 3): ['elevation:m', 'snow_drift:idx',
'location']
        These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 43 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', [])   : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',

```



```

'clear_sky_rad:W', ...]
      ('int', ['bool']) : 2 | ['snow_density:kgm3', 'is_estimated']
0.1s = Fit runtime
44 features in original data used to generate 44 features in processed
data.
    Train Data (Processed) Memory Usage: 10.96 MB (0.3% of available memory)
Data preprocessing and feature engineering runtime = 0.18s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
Warning: use_bag_holdout=True, but bagged mode is not enabled. use_bag_holdout
will be ignored.
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}],
}
Fitting 2 L1 models ...
Fitting model: LightGBMXT ...

[1000] valid_set's l1: 94.1472
[2000] valid_set's l1: 90.6203
[3000] valid_set's l1: 89.7064
[4000] valid_set's l1: 88.5809
[5000] valid_set's l1: 88.1969
[6000] valid_set's l1: 87.7129
[7000] valid_set's l1: 87.4664
[8000] valid_set's l1: 87.4638
[9000] valid_set's l1: 87.3975
[10000] valid_set's l1: 87.2144

-87.2056          = Validation score    (-mean_absolute_error)
59.56s           = Training runtime
0.55s            = Validation runtime
Fitting model: NeuralNetTorch ...
-88.8725          = Validation score    (-mean_absolute_error)
28.64s           = Training runtime
0.01s            = Validation runtime
Fitting model: WeightedEnsemble_L2 ...
-82.5454          = Validation score    (-mean_absolute_error)
0.03s            = Training runtime
0.0s             = Validation runtime
AutoGluon training complete, total runtime = 89.45s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_134_A/")
Evaluation: mean_absolute_error on test data: -108.09115440484666

```

Note: Scores are always higher\_is\_better. This metric score can be multiplied by -1 to get the metric value.

Evaluations on test data:

```
{
    "mean_absolute_error": -108.09115440484666,
    "root_mean_squared_error": -328.46462204134156,
    "mean_squared_error": -107889.00793276135,
    "r2": 0.8306773015067175,
    "pearsonr": 0.9171184125361762,
    "median_absolute_error": -5.514040470123291
}
```

Evaluation on test data:

-108.09115440484666


```
[57]: import matplotlib.pyplot as plt
leaderboards = [None, None, None]
def leaderboard_for_location(i, loc):
    plt.scatter(train_data[(train_data["location"] == loc) &
    ↪(train_data["is_estimated"]==True)]["y"].index,
    ↪train_data[(train_data["location"] == loc) &
    ↪(train_data["is_estimated"]==True)]["y"])
    plt.scatter(tuning_data[tuning_data["location"] == loc]["y"].index,
    ↪tuning_data[tuning_data["location"] == loc]["y"])
    plt.title("Val and Train")
    plt.show()

    if use_test_data:
        lb = predictors[i].leaderboard(test_data[test_data["location"] == loc])
        lb["location"] = loc
        plt.scatter(test_data[test_data["location"] == loc]["y"].index,
        ↪test_data[test_data["location"] == loc]["y"])
        plt.title("Test")


    return lb

return pd.DataFrame()

leaderboards[0] = leaderboard_for_location(0, loc)
```



	model	score_test	score_val	pred_time_test	pred_time_val
fit_time	pred_time_test_marginal	pred_time_val_marginal	fit_time_marginal		
stack_level	can_infer	fit_order			
0	LightGBMXT	-105.076269	-87.205583	0.127806	0.554786
59.564247		0.127806		0.554786	59.564247
1	True	1			
1	WeightedEnsemble_L2	-108.091154	-82.545415	0.144621	0.569622
88.234285		0.001212		0.000299	0.027779
2	True	3			
2	NeuralNetTorch	-121.957260	-88.872490	0.015603	0.014537
28.642259		0.015603		0.014537	28.642259
1	True	2			



```
[58]: loc = "B"
      predictors[1] = fit_predictor_for_location(loc)
      leaderboards[1] = leaderboard_for_location(1, loc)
```

```
Warning: path already exists! This predictor may overwrite an existing
predictor! path="AutogluonModels/submission_134_B"
Presets specified: ['experimental_zeroshot_hpo_hybrid']
Stack configuration (auto_stack=True): num_stack_levels=0, num_bag_folds=0,
num_bag_sets=0
Beginning AutoGluon training ...
AutoGluon will save models to "AutogluonModels/submission_134_B/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Darwin
Platform Machine: arm64
Platform Version: Darwin Kernel Version 22.1.0: Sun Oct 9 20:15:09 PDT 2022;
root:xnu-8792.41.9~2/RELEASE_ARM64_T6000
Disk Space Avail: 101.00 GB / 494.38 GB (20.4%)
Train Data Rows: 27377
Train Data Columns: 47
Tuning Data Rows: 1485
```

```

Tuning Data Columns: 47
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (1152.3, -0.0, 98.11625, 206.48535)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                2794.4 MB
    Train Data (Original) Memory Usage: 12.3 MB (0.4% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 2 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 2): ['elevation:m', 'location']
    These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 44 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', []) : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 43 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', ['bool']) : 2 | ['snow_density:kgm3', 'is_estimated']
    0.1s = Fit runtime
    45 features in original data used to generate 45 features in processed
data.
    Train Data (Processed) Memory Usage: 9.99 MB (0.4% of available memory)
Data preprocessing and feature engineering runtime = 0.14s ...

```

```

AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
Warning: use_bag_holdout=True, but bagged mode is not enabled. use_bag_holdout
will be ignored.
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'XT': [{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}]},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}},
{'extra_trees': True, 'feature_fraction': 0.7832570544199176, 'learning_rate':
0.021720607471727896, 'min_data_in_leaf': 3, 'num_leaves': 21, 'ag_args':
{'name_suffix': '_r118', 'priority': 17}}],
    'FASTAI': [{}, {'bs': 1024, 'emb_drop': 0.6046989241462619, 'epochs':
48, 'layers': [200, 100, 50], 'lr': 0.00775309042164966, 'ps':
0.09244767444160731, 'ag_args': {'name_suffix': '_r51', 'priority': 12}}, {'bs':
128, 'emb_drop': 0.44339037504795686, 'epochs': 31, 'layers': [400, 200, 100],
'lr': 0.008615195908919904, 'ps': 0.19220253419114286, 'ag_args':
{'name_suffix': '_r145', 'priority': 9}}],
}
Fitting 7 L1 models ...
Fitting model: LightGBMXT ...

[1000] valid_set's l1: 14.2773
[2000] valid_set's l1: 14.0181
[3000] valid_set's l1: 13.957
[4000] valid_set's l1: 13.8745
[5000] valid_set's l1: 13.8621
[6000] valid_set's l1: 13.8212
[7000] valid_set's l1: 13.8141
[8000] valid_set's l1: 13.7824
[9000] valid_set's l1: 13.7291
[10000] valid_set's l1: 13.7163

-13.7142          = Validation score    (-mean_absolute_error)
126.65s = Training    runtime
0.26s   = Validation runtime
Fitting model: ExtraTreesMSE ...

-15.4761          = Validation score    (-mean_absolute_error)
3.82s   = Training    runtime
0.05s   = Validation runtime
Fitting model: NeuralNetFastAI ...

-14.9623          = Validation score    (-mean_absolute_error)
17.81s   = Training    runtime
0.02s   = Validation runtime
Fitting model: NeuralNetTorch ...

```

```

-13.3344          = Validation score    (-mean_absolute_error)
20.55s           = Training   runtime
0.01s            = Validation runtime
Fitting model: LightGBM_r118 ...

[1000] valid_set's l1: 13.9149
[2000] valid_set's l1: 13.4226
[3000] valid_set's l1: 13.2471
[4000] valid_set's l1: 13.1382
[5000] valid_set's l1: 13.0259
[6000] valid_set's l1: 12.9035
[7000] valid_set's l1: 12.7973
[8000] valid_set's l1: 12.7556
[9000] valid_set's l1: 12.7006
[10000] valid_set's l1: 12.658

-12.6552          = Validation score    (-mean_absolute_error)
43.66s           = Training   runtime
0.15s            = Validation runtime
Fitting model: NeuralNetFastAI_r51 ...

-14.1868          = Validation score    (-mean_absolute_error)
11.89s           = Training   runtime
0.01s            = Validation runtime
Fitting model: NeuralNetFastAI_r145 ...


-13.7911          = Validation score    (-mean_absolute_error)
38.66s           = Training   runtime
0.02s            = Validation runtime
Fitting model: WeightedEnsemble_L2 ...

-11.9643          = Validation score    (-mean_absolute_error)
0.07s            = Training   runtime
0.0s             = Validation runtime
AutoGluon training complete, total runtime = 264.88s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_134_B/")
Evaluation: mean_absolute_error on test data: -10.771986388368392
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -10.771986388368392,
    "root_mean_squared_error": -29.819237777904075,
    "mean_squared_error": -889.1869416551816,
    "r2": 0.9617196394878705,
    "pearsonr": 0.9806798548041709,
    "median_absolute_error": -0.6749260425567627
}

Evaluation on test data:

```


-10.771986388368392



	model	score_test	score_val	pred_time_test	pred_time_val
fit_time	pred_time_test_marginal	pred_time_val_marginal	fit_time_marginal		
stack_level	can_infer	fit_order			
0	LightGBMXT	-10.416372	-13.714202	0.127264	0.263104
126.649451		0.127264		0.263104	126.649451
1	True	1			
1	WeightedEnsemble_L2	-10.771986	-11.964269	0.327276	0.257611
136.459363		0.001567		0.000245	0.067160
2	True	8			
2	LightGBM_r118	-11.610208	-12.655171	0.093102	0.147992
43.656658		0.093102		0.147992	43.656658
1	True	5			
3	NeuralNetTorch	-12.551883	-13.334396	0.008101	0.010194
20.553850		0.008101		0.010194	20.553850
1	True	4			
4	NeuralNetFastAI_r145	-13.092760	-13.791110	0.017274	0.023354
38.659627		0.017274		0.023354	38.659627
1	True	7			
5	ExtraTreesMSE	-13.103427	-15.476124	0.147653	0.049357
3.823860		0.147653		0.049357	3.823860



1	True	2				
6	NeuralNetFastAI	-14.050613	-14.962268	0.043640	0.019394	
17.809215		0.043640		0.019394	17.809215	
1	True	3				
7	NeuralNetFastAI_r51	-14.650156	-14.186801	0.015939	0.007075	
11.888993		0.015939		0.007075	11.888993	
1	True	6				



```
[59]: loc = "C"
predictors[2] = fit_predictor_for_location(loc)
leaderboards[2] = leaderboard_for_location(2, loc)
```

```
Presets specified: ['experimental_zeroshot_hpo_hybrid']
Stack configuration (auto_stack=True): num_stack_levels=0, num_bag_folds=0,
num_bag_sets=0
Beginning AutoGluon training ...
AutoGluon will save models to "AutogluonModels/submission_134_C/"
AutoGluon Version: 0.8.2
Python Version: 3.10.12
Operating System: Darwin
Platform Machine: arm64
Platform Version: Darwin Kernel Version 22.1.0: Sun Oct 9 20:15:09 PDT 2022;
```

```

root:xnu-8792.41.9~2/RELEASE_ARM64_T6000
Disk Space Avail:   100.63 GB / 494.38 GB (20.4%)
Train Data Rows:    24073
Train Data Columns: 47
Tuning Data Rows:   1481
Tuning Data Columns: 47
Label Column: y
Preprocessing data ...
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and label-values can't be converted to int).
    Label info (max, min, mean, stddev): (999.6, -0.0, 80.87539, 169.67845)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression'])
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                3110.38 MB
    Train Data (Original) Memory Usage: 10.89 MB (0.3% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
            Note: Converting 2 features to boolean dtype as they
only contain 2 unique values.
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Useless Original Features (Count: 3): ['elevation:m', 'snow_drift:idx',
'location']
        These features carry no predictive signal and should be manually
investigated.
        This is typically a feature which has the same value for all
rows.
        These features do not need to be present at inference time.
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 43 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]
        ('int', [])   : 1 | ['is_estimated']
    Types of features in processed data (raw dtype, special dtypes):
        ('float', []) : 42 | ['absolute_humidity_2m:gm3',
'air_density_2m:kgm3', 'ceiling_height_agl:m', 'clear_sky_energy_1h:J',
'clear_sky_rad:W', ...]

```

```

        ('int', ['bool']) : 2 | ['snow_density:kgm3', 'is_estimated']
0.1s = Fit runtime
44 features in original data used to generate 44 features in processed
data.
Train Data (Processed) Memory Usage: 8.64 MB (0.3% of available memory)
Data preprocessing and feature engineering runtime = 0.1s ...
AutoGluon will gauge predictive performance using evaluation metric:
'mean_absolute_error'
This metric's sign has been flipped to adhere to being higher_is_better.
The metric score can be multiplied by -1 to get the metric value.
To change this, specify the eval_metric parameter of Predictor()
Warning: use_bag_holdout=True, but bagged mode is not enabled. use_bag_holdout
will be ignored.
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': {},
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}}],
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}},
{'extra_trees': True, 'feature_fraction': 0.7832570544199176, 'learning_rate':
0.021720607471727896, 'min_data_in_leaf': 3, 'num_leaves': 21, 'ag_args':
{'name_suffix': '_r118', 'priority': 17}}],
    'FASTAI': [{'bs': 1024, 'emb_drop': 0.6046989241462619, 'epochs': 48,
'layers': [200, 100, 50], 'lr': 0.00775309042164966, 'ps': 0.09244767444160731,
'ag_args': {'name_suffix': '_r51', 'priority': 12}}],
}
Fitting 5 L1 models ...
Fitting model: KNeighborsUnif ...
-19.8149          = Validation score    (-mean_absolute_error)
0.03s           = Training   runtime
4.73s           = Validation runtime
Fitting model: LightGBMXT ...
[1000] valid_set's l1: 12.5619
[2000] valid_set's l1: 11.7336
[3000] valid_set's l1: 11.5889
-11.5594          = Validation score    (-mean_absolute_error)
25.45s           = Training   runtime
0.04s           = Validation runtime
Fitting model: NeuralNetTorch ...
-12.2202          = Validation score    (-mean_absolute_error)
43.72s           = Training   runtime
0.01s           = Validation runtime
Fitting model: LightGBM_r118 ...
[1000] valid_set's l1: 13.5054
[2000] valid_set's l1: 12.4438
[3000] valid_set's l1: 11.89
[4000] valid_set's l1: 11.6741

```


```

[5000] valid_set's l1: 11.4779
[6000] valid_set's l1: 11.3679
[7000] valid_set's l1: 11.316
[8000] valid_set's l1: 11.2841
[9000] valid_set's l1: 11.2212
[10000] valid_set's l1: 11.195


-11.1909          = Validation score    (-mean_absolute_error)
82.2s            = Training    runtime
0.28s           = Validation runtime
Fitting model: NeuralNetFastAI_r51 ...
-13.4989          = Validation score    (-mean_absolute_error)
68.52s           = Training    runtime
0.02s           = Validation runtime
Fitting model: WeightedEnsemble_L2 ...
-10.7391          = Validation score    (-mean_absolute_error)
0.08s           = Training    runtime
0.0s            = Validation runtime
AutoGluon training complete, total runtime = 226.87s ... Best model:
"WeightedEnsemble_L2"
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/submission_134_C/")
Evaluation: mean_absolute_error on test data: -11.096687568183471
    Note: Scores are always higher_is_better. This metric score can be
multiplied by -1 to get the metric value.
Evaluations on test data:
{
    "mean_absolute_error": -11.096687568183471,
    "root_mean_squared_error": -26.67684023698341,
    "mean_squared_error": -711.6538050295371,
    "r2": 0.9278640745308446,
    "pearsonr": 0.9652418428251913,
    "median_absolute_error": -0.5633931159973145
}

Evaluation on test data:
-11.096687568183471

```



	model	score_test	score_val	pred_time_test	pred_time_val
fit_time	pred_time_test_marginal	pred_time_val_marginal	fit_time_marginal		
stack_level	can_infer	fit_order			
0	WeightedEnsemble_L2	-11.096688	-10.739141	12.911151	5.084761
220.005957		0.002081		0.000204	0.083005
2	True	6			
1	LightGBM_r118	-11.531620	-11.190891	0.322478	0.276384
82.199734		0.322478		0.276384	82.199734
1	True	4			
2	LightGBMXT	-12.545529	-11.559353	0.133682	0.036625
25.454035		0.133682		0.036625	25.454035
1	True	2			
3	NeuralNetTorch	-13.048661	-12.220150	0.092223	0.014784
43.720339		0.092223		0.014784	43.720339
1	True	3			
4	NeuralNetFastAI_r51	-14.205124	-13.498880	0.032736	0.023557
68.520166		0.032736		0.023557	68.520166
1	True	5			
5	KNeighborsUnif	-20.049167	-19.814903	12.327951	4.733207
0.028678		12.327951		4.733207	0.028678
1	True	1			



```
[60]: # save leaderboards to csv
pd.concat(leaderboards).to_csv(f"leaderboards/{new_filename}.csv")

for i in range(len(predictors)):
    print(f"Predictor {i}:")
    print(predictors[i].
    ↪info()["model_info"]["WeightedEnsemble_L2"]["children_info"]["S1F1"]["model_weights"])
```

Predictor 0:

```
{'LightGBMX': 0.5147058823529411, 'NeuralNetTorch': 0.4852941176470588}
```

Predictor 1:

```
{'ExtraTreesMSE': 0.015151515151515152, 'NeuralNetFastAI': 0.045454545454545456,
'NeuralNetTorch': 0.21212121212121213, 'LightGBM_r118': 0.4393939393939394,
'NeuralNetFastAI_r51': 0.09090909090909091, 'NeuralNetFastAI_r145':
0.19696969696969696}
```

Predictor 2:

```
{'KNeighborsUnif': 0.06756756756756757, 'LightGBMX': 0.13513513513513514,
'NeuralNetTorch': 0.3108108108108108, 'LightGBM_r118': 0.4594594594594595,
'NeuralNetFastAI_r51': 0.02702702702702703}
```

## 5 Submit

```
[61]: import pandas as pd
import matplotlib.pyplot as plt
```

```
future_test_data = TabularDataset('X_test_raw.csv')
future_test_data["ds"] = pd.to_datetime(future_test_data["ds"])
```

Loaded data from: X\_test\_raw.csv | Columns = 48 / 48 | Rows = 4608 -> 4608

```
[62]: test_ids = TabularDataset('test.csv')
test_ids["time"] = pd.to_datetime(test_ids["time"])
# merge test_data with test_ids
future_test_data_merged = pd.merge(future_test_data, test_ids, how="inner",
    ↪right_on=["time", "location"], left_on=["ds", "location"])
```

Loaded data from: test.csv | Columns = 4 / 4 | Rows = 2160 -> 2160

```
[63]: # predict, grouped by location
predictions = []
location_map = {
    "A": 0,
    "B": 1,
    "C": 2
}
for loc, group in future_test_data.groupby('location'):
    i = location_map[loc]
    subset = future_test_data_merged[future_test_data_merged["location"] ==
    ↪loc].reset_index(drop=True)
    pred = predictors[i].predict(subset)
    subset["prediction"] = pred
    predictions.append(subset)

# get past predictions
tuning_data.loc[tuning_data["location"] == loc, "prediction"] =
    ↪predictors[i].predict(tuning_data[tuning_data["location"] == loc])
if use_test_data:
    # get predictions for local test_data
    test_data.loc[test_data["location"] == loc, "prediction"] =
    ↪predictors[i].predict(test_data[test_data["location"] == loc])
```

```
[64]: for loc, idx in location_map.items():
    fig, ax = plt.subplots(figsize=(20, 10))
    # plot train data
    train_data[train_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="train data")
    tuning_data[tuning_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="tune data")
```

```


if use_test_data:
    test_data[test_data["location"]==loc].plot(x='ds', y='y', ax=ax,
    ↪label="test data")

    # plot predictions
    predictions[idx].plot(x='ds', y='prediction', ax=ax, label="predictions")



    # plot past predictions
    tuning_data[tuning_data["location"]==loc].plot(x='ds', y='prediction',
    ↪ax=ax, label="past predictions tune")
    if use_test_data:
        test_data[test_data["location"]==loc].plot(x='ds', y='prediction',
    ↪ax=ax, label="past predictions test")

ax.set_title(f"Predictions for location {loc}")

```







```
[65]: temp_predictions = [prediction.copy() for prediction in predictions]
if clip_predictions:
    # clip predictions smaller than 0 to 0
    for pred in temp_predictions:
        # print smallest prediction
        print("Smallest prediction:", pred["prediction"].min())
        pred.loc[pred["prediction"] < 0, "prediction"] = 0
        print("Smallest prediction after clipping:", pred["prediction"].min())
```

```
# concatenate predictions
submissions_df = pd.concat(temp_predictions)
submissions_df = submissions_df[["id", "prediction"]]
submissions_df
```

```
Smallest prediction: -37.823605
Smallest prediction after clipping: 0.0
Smallest prediction: -2.3058283
Smallest prediction after clipping: 0.0
Smallest prediction: -2.1818242
Smallest prediction after clipping: 0.0
```

```
[65]:      id  prediction
0      0      0.000000
1      1      0.000000
2      2      0.000000
3      3     16.584423
4      4    304.872498
..    ...      ...
715   2155     63.452984
716   2156     34.039520
717   2157      5.225869
718   2158      0.645353
719   2159      0.920492
```

```
[2160 rows x 2 columns]
```

```
[66]: # Save the submission
print(f"Saving submission to submissions/{new_filename}.csv")
submissions_df.to_csv(os.path.join('submissions', f"{new_filename}.csv"),
    ↪index=False)
```

```
Saving submission to submissions/submission_134.csv
```

```
[67]: # feature importance
print("\033[1m" + "Calculating feature importance for location A..." +
    ↪"\033[0m")
print(predictors[0].feature_importance(feature_stage="original",
    ↪data=test_data[test_data["location"] == "A"], time_limit=60*10))
print("\033[1m" + "Calculating feature importance for location B..." +
    ↪"\033[0m")
print(predictors[1].feature_importance(feature_stage="original",
    ↪data=test_data[test_data["location"] == "B"], time_limit=60*10))
print("\033[1m" + "Calculating feature importance for location C..." +
    ↪"\033[0m")
print(predictors[2].feature_importance(feature_stage="original",
    ↪data=test_data[test_data["location"] == "C"], time_limit=60*10))
```

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'snow\_drift:idx', 'location', 'prediction']

Calculating feature importance for location A...

Computing feature importance via permutation shuffling for 44 features using 361 rows with 10 shuffle sets... Time limit: 600s...

88.72s = Expected runtime (8.87s per shuffle set)

53.07s = Actual runtime (Completed 10 of 10 shuffle sets)

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'location', 'prediction']

Computing feature importance via permutation shuffling for 45 features using 361 rows with 10 shuffle sets... Time limit: 600s...

	importance	stddev	p_value	n \
direct_rad:W	7.966799e+01	6.196518	8.202915e-12	10
clear_sky_rad:W	7.633021e+01	12.170163	4.887011e-09	10
diffuse_rad:W	5.799179e+01	10.525372	1.526728e-08	10
clear_sky_energy_1h:J	3.031519e+01	10.381896	3.459234e-06	10
sun_elevation:d	2.939179e+01	7.686730	3.608322e-07	10
sun_azimuth:d	2.394524e+01	12.523500	9.566792e-05	10
diffuse_rad_1h:J	1.427764e+01	4.402150	1.448702e-06	10
direct_rad_1h:J	1.398430e+01	5.747896	1.509840e-05	10
effective_cloud_cover:p	1.397670e+01	4.436740	1.846918e-06	10
total_cloud_cover:p	1.084766e+01	3.430958	1.791842e-06	10
snow_water:kgm2	7.979610e+00	5.356829	5.518101e-04	10
relative_humidity_1000hPa:p	7.824189e+00	2.658841	3.245806e-06	10
ceiling_height_agl:m	5.896515e+00	3.069501	9.241565e-05	10
precip_type_5min:idx	4.407535e+00	4.810146	8.832219e-03	10
cloud_base_agl:m	4.323458e+00	2.144760	6.453748e-05	10
is_day:idx	4.234546e+00	1.755630	1.616799e-05	10
is_in_shadow:idx	4.195547e+00	1.843094	2.545761e-05	10
pressure_100m:hPa	3.815979e+00	2.712117	8.008382e-04	10
wind_speed_10m:ms	3.776024e+00	1.275661	3.092497e-06	10
wind_speed_v_10m:ms	3.342700e+00	2.037866	2.869482e-04	10
visibility:m	3.198419e+00	1.971327	3.094968e-04	10
fresh_snow_6h:cm	3.072502e+00	2.323215	1.184209e-03	10
wind_speed_u_10m:ms	2.988759e+00	4.273209	2.714485e-02	10
pressure_50m:hPa	2.501973e+00	1.564861	3.423738e-04	10
snow_depth:cm	2.467273e+00	1.609154	4.550271e-04	10
msl_pressure:hPa	2.072250e+00	2.394007	1.147504e-02	10
fresh_snow_24h:cm	2.062518e+00	4.520783	9.149065e-02	10
air_density_2m:kgm3	1.980115e+00	1.604045	1.799604e-03	10
sfc_pressure:hPa	1.837239e+00	1.883139	6.515049e-03	10
precip_5min:mm	1.833741e+00	2.442299	2.080664e-02	10
super_cooled_liquid_water:kgm2	1.543328e+00	1.515049	5.233747e-03	10
fresh_snow_1h:cm	1.070496e+00	1.074131	5.853889e-03	10
t_1000hPa:K	9.737481e-01	1.600851	4.328374e-02	10
snow_melt_10min:mm	8.978923e-01	1.390949	3.580554e-02	10

fresh_snow_3h:cm	7.902873e-01	1.406984	5.471397e-02	10
absolute_humidity_2m:gm3	6.890160e-01	0.912385	2.034222e-02	10
rain_water:kgm2	5.316596e-01	0.524287	5.357342e-03	10
fresh_snow_12h:cm	5.251203e-01	3.366204	3.168039e-01	10
dew_point_2m:K	3.666395e-01	1.146486	1.691486e-01	10
prob_rime:p	3.051479e-01	0.548112	5.608503e-02	10
snow_density:kgm3	2.111609e-01	1.576556	3.409184e-01	10
is_estimated	4.490987e-08	0.000000	5.000000e-01	10
wind_speed_w_1000hPa:ms	4.226811e-08	0.000000	5.000000e-01	10
dew_or_rime:idx	-5.206986e-01	1.231628	8.929724e-01	10

	p99_high	p99_low
direct_rad:W	8.603607e+01	7.329990e+01
clear_sky_rad:W	8.883734e+01	6.382308e+01
diffuse_rad:W	6.880859e+01	4.717499e+01
clear_sky_energy_1h:J	4.098454e+01	1.964583e+01
sun_elevation:d	3.729135e+01	2.149223e+01
sun_azimuth:d	3.681550e+01	1.107499e+01
diffuse_rad_1h:J	1.880168e+01	9.753601e+00
direct_rad_1h:J	1.989135e+01	8.077258e+00
effective_cloud_cover:p	1.853628e+01	9.417111e+00
total_cloud_cover:p	1.437362e+01	7.321705e+00
snow_water:kgm2	1.348476e+01	2.474460e+00
relative_humidity_1000hPa:p	1.055665e+01	5.091729e+00
ceiling_height_agl:m	9.051005e+00	2.742025e+00
precip_type_5min:idx	9.350865e+00	-5.357949e-01
cloud_base_agl:m	6.527603e+00	2.119313e+00
is_day:idx	6.038786e+00	2.430306e+00
is_in_shadow:idx	6.089673e+00	2.301421e+00
pressure_100m:hPa	6.603190e+00	1.028768e+00
wind_speed_10m:ms	5.087007e+00	2.465042e+00
wind_speed_v_10m:ms	5.436991e+00	1.248409e+00
visibility:m	5.224328e+00	1.172509e+00
fresh_snow_6h:cm	5.460043e+00	6.849609e-01
wind_speed_u_10m:ms	7.380285e+00	-1.402768e+00
pressure_50m:hPa	4.110162e+00	8.937838e-01
snow_depth:cm	4.120981e+00	8.135643e-01
msl_pressure:hPa	4.532542e+00	-3.880429e-01
fresh_snow_24h:cm	6.708473e+00	-2.583437e+00
air_density_2m:kgm3	3.628574e+00	3.316571e-01
sfc_pressure:hPa	3.772519e+00	-9.804137e-02
precip_5min:mm	4.343663e+00	-6.761803e-01
super_cooled_liquid_water:kgm2	3.100326e+00	-1.366959e-02
fresh_snow_1h:cm	2.174368e+00	-3.337505e-02
t_1000hPa:K	2.618923e+00	-6.714272e-01
snow_melt_10min:mm	2.327354e+00	-5.315699e-01
fresh_snow_3h:cm	2.236228e+00	-6.556536e-01
absolute_humidity_2m:gm3	1.626664e+00	-2.486316e-01

rain_water:kgm2	1.070463e+00	-7.143807e-03
fresh_snow_12h:cm	3.984528e+00	-2.934288e+00
dew_point_2m:K	1.544870e+00	-8.115905e-01
prob_rime:p	8.684358e-01	-2.581400e-01
snow_density:kgm3	1.831368e+00	-1.409047e+00
is_estimated	4.490987e-08	4.490987e-08
wind_speed_w_1000hPa:ms	4.226811e-08	4.226811e-08
dew_or_rime:idx	7.450309e-01	-1.786428e+00

Calculating feature importance for location B...

307.44s = Expected runtime (30.74s per shuffle set)

55.82s = Actual runtime (Completed 10 of 10 shuffle sets)

These features in provided data are not utilized by the predictor and will be ignored: ['ds', 'elevation:m', 'snow\_drift:idx', 'location', 'prediction']

Computing feature importance via permutation shuffling for 44 features using 360 rows with 10 shuffle sets... Time limit: 600s...

	importance	stddev	p_value	n \
clear_sky_rad:W	3.089353e+01	2.862590	3.928157e-11	10
direct_rad:W	1.758631e+01	1.326535	6.232521e-12	10
diffuse_rad:W	1.292878e+01	1.449306	2.149220e-10	10
sun_elevation:d	8.974525e+00	1.129608	6.024637e-10	10
sun_azimuth:d	6.951743e+00	1.180075	8.485674e-09	10
clear_sky_energy_1h:J	6.795127e+00	1.047551	3.631051e-09	10
effective_cloud_cover:p	3.173118e+00	0.654582	4.670364e-08	10
direct_rad_1h:J	3.120853e+00	0.435006	1.489493e-09	10
diffuse_rad_1h:J	3.112522e+00	0.729352	1.407778e-07	10
relative_humidity_1000hPa:p	1.751303e+00	0.483333	5.701084e-07	10
is_in_shadow:idx	1.597443e+00	0.199220	5.550254e-10	10
fresh_snow_24h:cm	1.452799e+00	0.497733	3.470695e-06	10
snow_water:kgm2	1.280799e+00	0.406230	1.833939e-06	10
wind_speed_v_10m:ms	1.147136e+00	0.256486	9.390015e-08	10
sfc_pressure:hPa	1.121787e+00	0.288322	3.110949e-07	10
visibility:m	1.118796e+00	0.305319	5.182709e-07	10
cloud_base_agl:m	1.087815e+00	0.273498	2.576902e-07	10
total_cloud_cover:p	9.360905e-01	0.258885	5.802443e-07	10
pressure_100m:hPa	9.270775e-01	0.325786	4.273716e-06	10
pressure_50m:hPa	8.658717e-01	0.314427	5.585112e-06	10
msl_pressure:hPa	8.484935e-01	0.205324	1.856428e-07	10
ceiling_height_agl:m	6.605458e-01	0.324262	5.964112e-05	10
wind_speed_10m:ms	6.583129e-01	0.189969	8.310688e-07	10
t_1000hPa:K	6.118650e-01	0.236174	9.147006e-06	10
snow_density:kgm3	5.429378e-01	0.505616	3.963566e-03	10
is_day:idx	4.907764e-01	0.092593	2.143246e-08	10
air_density_2m:kgm3	4.703118e-01	0.357817	1.230291e-03	10
wind_speed_u_10m:ms	3.440242e-01	0.276338	1.711148e-03	10
super_cooled_liquid_water:kgm2	3.356328e-01	0.321861	4.632723e-03	10
precip_type_5min:idx	2.899514e-01	0.306324	7.560363e-03	10
dew_point_2m:K	2.881253e-01	0.123478	2.098102e-05	10

fresh_snow_12h:cm	2.772887e-01	0.315221	1.066996e-02	10
fresh_snow_6h:cm	2.697109e-01	0.227036	2.254014e-03	10
precip_5min:mm	2.373693e-01	0.313540	2.014490e-02	10
rain_water:kgm2	2.218524e-01	0.180006	1.816843e-03	10
snow_depth:cm	1.467803e-01	0.140861	4.650676e-03	10
absolute_humidity_2m:gm3	1.399695e-01	0.065372	4.084992e-05	10
prob_rime:p	1.062888e-01	0.199169	6.288190e-02	10
fresh_snow_1h:cm	1.057443e-01	0.099328	4.151120e-03	10
dew_or_rime:idx	7.235418e-02	0.080144	9.468436e-03	10
fresh_snow_3h:cm	7.209251e-02	0.087237	1.405900e-02	10
snow_melt_10min:mm	3.663396e-02	0.228162	3.119229e-01	10
is_estimated	0.000000e+00	0.000000	5.000000e-01	10
snow_drift:idx	-5.283516e-09	0.000000	5.000000e-01	10
wind_speed_w_1000hPa:ms	-5.283516e-09	0.000000	5.000000e-01	10

	p99_high	p99_low
clear_sky_rad:W	3.383538e+01	2.795168e+01
direct_rad:W	1.894957e+01	1.622304e+01
diffuse_rad:W	1.441822e+01	1.143935e+01
sun_elevation:d	1.013541e+01	7.813640e+00
sun_azimuth:d	8.164492e+00	5.738994e+00
clear_sky_energy_1h:J	7.871682e+00	5.718571e+00
effective_cloud_cover:p	3.845824e+00	2.500411e+00
direct_rad_1h:J	3.567903e+00	2.673802e+00
diffuse_rad_1h:J	3.862068e+00	2.362976e+00
relative_humidity_1000hPa:p	2.248019e+00	1.254587e+00
is_in_shadow:idx	1.802180e+00	1.392707e+00
fresh_snow_24h:cm	1.964313e+00	9.412840e-01
snow_water:kgm2	1.698277e+00	8.633209e-01
wind_speed_v_10m:ms	1.410724e+00	8.835485e-01
sfc_pressure:hPa	1.418092e+00	8.254824e-01
visibility:m	1.432569e+00	8.050229e-01
cloud_base_agl:m	1.368886e+00	8.067434e-01
total_cloud_cover:p	1.202143e+00	6.700376e-01
pressure_100m:hPa	1.261883e+00	5.922715e-01
pressure_50m:hPa	1.189005e+00	5.427385e-01
msl_pressure:hPa	1.059502e+00	6.374845e-01
ceiling_height_agl:m	9.937865e-01	3.273051e-01
wind_speed_10m:ms	8.535422e-01	4.630837e-01
t_1000hPa:K	8.545787e-01	3.691512e-01
snow_density:kgm3	1.062553e+00	2.332255e-02
is_day:idx	5.859336e-01	3.956193e-01
air_density_2m:kgm3	8.380361e-01	1.025875e-01
wind_speed_u_10m:ms	6.280135e-01	6.003497e-02
super_cooled_liquid_water:kgm2	6.664051e-01	4.860577e-03
precip_type_5min:idx	6.047566e-01	-2.485390e-02
dew_point_2m:K	4.150219e-01	1.612288e-01
fresh_snow_12h:cm	6.012377e-01	-4.666037e-02

fresh_snow_6h:cm	5.030330e-01	3.638877e-02
precip_5min:mm	5.595911e-01	-8.485250e-02
rain_water:kgm2	4.068427e-01	3.686212e-02
snow_depth:cm	2.915413e-01	2.019220e-03
absolute_humidity_2m:gm3	2.071515e-01	7.278755e-02
prob_rime:p	3.109729e-01	-9.839529e-02
fresh_snow_1h:cm	2.078223e-01	3.666258e-03
dew_or_rime:idx	1.547172e-01	-1.000885e-02
fresh_snow_3h:cm	1.617449e-01	-1.755983e-02
snow_melt_10min:mm	2.711138e-01	-1.978459e-01
is_estimated	0.000000e+00	0.000000e+00
snow_drift:idx	-5.283516e-09	-5.283516e-09
wind_speed_w_1000hPa:ms	-5.283516e-09	-5.283516e-09

Calculating feature importance for location C...

3929.37s = Expected runtime (392.94s per shuffle set)

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
/Users/jorgensandhaug/Desktop/tdt4173/TDT4173/short_1.ipynb Cell 31 line 7

    <a href='vscode-notebook-cell:/Users/jorgensandhaug/Desktop/tdt4173/
↪TDT4173/short_1.ipynb#X43sZmlsZQ%3D%3D?line=4'>5</a> print(predictors[1].
↪feature_importance(feature_stage="original",
↪data=test_data[test_data["location"] == "B"], time_limit=60*10))
    <a href='vscode-notebook-cell:/Users/jorgensandhaug/Desktop/tdt4173/
↪TDT4173/short_1.ipynb#X43sZmlsZQ%3D%3D?line=5'>6</a> print("\033[1m" +
↪"Calculating feature importance for location C..." + "\033[0m")
----> <a href='vscode-notebook-cell:/Users/jorgensandhaug/Desktop/tdt4173/
↪TDT4173/short_1.ipynb#X43sZmlsZQ%3D%3D?line=6'>7</a> print(predictors[2].
↪feature_importance(feature_stage="original",
↪data=test_data[test_data["location"] == "C"], time_limit=60*10))

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/
↪tabular/predictor/predictor.py:2418, in TabularPredictor.
↪feature_importance(self, data, model, features, feature_stage, subsample_size,
↪time_limit, num_shuffle_sets, include_confidence_band, confidence_level,
↪silent)
    2415 if num_shuffle_sets is None:
    2416     num_shuffle_sets = 10 if time_limit else 5
-> 2418 fi_df = self._learner.get_feature_importance(
    2419     model=model,
    2420     X=data,
    2421     features=features,
    2422     feature_stage=feature_stage,
    2423     subsample_size=subsample_size,
    2424     time_limit=time_limit,
    2425     num_shuffle_sets=num_shuffle_sets,
    2426     silent=silent,
    2427 )
    2429 if include_confidence_band:
```

```

2430         if confidence_level <= 0.5 or confidence_level >= 1.0:

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/
↳ tabular/learner/abstract_learner.py:870, in AbstractTabularLearner.
↳ get_feature_importance(self, model, X, y, features, feature_stage,
↳ subsample_size, silent, **kwargs)
    867         X = X.drop(columns=unused_features)
    869         if feature_stage == "original":
--> 870         return trainer._get_feature_importance_raw(
    871             model=model, X=X, y=y, features=features,
↳ subsample_size=subsample_size, transform_func=self.transform_features,
↳ silent=silent, **kwargs
    872         )
    873         X = self.transform_features(X)
    874     else:

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ trainer/abstract_trainer.py:2622, in AbstractTrainer.
↳ _get_feature_importance(self, X, y, model, eval_metric, **kwargs)
    2620 model: AbstractModel = self.load_model(model)
    2621 predict_func_kwargs = dict(model=model)
-> 2622 return compute_permutation_feature_importance(
    2623     X=X,
    2624     y=y,
    2625     predict_func=predict_func,
    2626     predict_func_kwargs=predict_func_kwargs,
    2627     eval_metric=eval_metric,
    2628     quantile_levels=self.quantile_levels,
    2629     **kwargs,
    2630 )

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ utils/utils.py:882, in compute_permutation_feature_importance(X, y,
↳ predict_func, eval_metric, features, subsample_size, num_shuffle_sets,
↳ predict_func_kwargs, transform_func, transform_func_kwargs, time_limit,
↳ silent, log_prefix, importance_as_list, random_state, **kwargs)
    880 else:
    881     X_raw_transformed = X_raw if transform_func is None else
↳ transform_func(X_raw, **transform_func_kwargs)
--> 882 y_pred = predict_func(X_raw_transformed, **predict_func_kwargs)
    884 row_index = 0
    885 for feature in parallel_computed_features:

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ trainer/abstract_trainer.py:737, in AbstractTrainer.predict(self, X, model)
    735     model = self._get_best()
    736     cascade = isinstance(model, list)
--> 737 return self._predict_model(X, model, cascade=cascade)

```



```

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ trainer/abstract_trainer.py:2436, in AbstractTrainer._predict_model(self, X,
↳ model, model_pred_proba_dict, cascade)
    2435 def _predict_model(self, X, model, model_pred_proba_dict=None,
↳ cascade=False):
-> 2436     y_pred_proba = self._predict_proba_model(X=X, model=model,
↳ model_pred_proba_dict=model_pred_proba_dict, cascade=cascade)
    2437     return get_pred_from_proba(y_pred_proba=y_pred_proba,
↳ problem_type=self.problem_type)

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ trainer/abstract_trainer.py:2440, in AbstractTrainer.
↳ _predict_proba_model(self, X, model, model_pred_proba_dict, cascade)
    2439 def _predict_proba_model(self, X, model, model_pred_proba_dict=None,
↳ cascade=False):
-> 2440     return self.get_pred_proba_from_model(model=model, X=X,
↳ model_pred_proba_dict=model_pred_proba_dict, cascade=cascade)

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ trainer/abstract_trainer.py:757, in AbstractTrainer.
↳ get_pred_proba_from_model(self, model, X, model_pred_proba_dict, cascade)
    755 else:
    756     models = [model]
--> 757 model_pred_proba_dict = self.get_model_pred_proba_dict(X=X,
↳ models=models, model_pred_proba_dict=model_pred_proba_dict, cascade=cascade)
    758 if not isinstance(model, str):
    759     model = model.name

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ trainer/abstract_trainer.py:1008, in AbstractTrainer.
↳ get_model_pred_proba_dict(self, X, models, model_pred_proba_dict,
↳ model_pred_time_dict, record_pred_time, use_val_cache, cascade,
↳ cascade_threshold)
    1006     model_pred_proba_dict[model_name] = model.predict_proba(X,
↳ **preprocess_kwargs)
    1007 else:
-> 1008     model_pred_proba_dict[model_name] = model.predict_proba(X)
    1010 if record_pred_time:
    1011     time_end = time.time()

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core /
↳ models/abstract/abstract_model.py:940, in AbstractModel.predict_proba(self, X,
↳ normalize, **kwargs)
    938 if normalize is None:
    939     normalize = self.normalize_pred_probabilities
--> 940 y_pred_proba = self._predict_proba(X=X, **kwargs)
    941 if normalize:
    942     y_pred_proba = normalize_pred_probabilities(y_pred_proba, self.problem_type)

```

```

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/autogluon/core/
↳models/abstract/abstract_model.py:955, in AbstractModel._predict_proba(self, X, **kwargs)
    952 X = self.preprocess(X, **kwargs)
    954 if self.problem_type in [REGRESSION, QUANTILE]:
--> 955     y_pred = self.model.predict(X)
    956     return y_pred
    958 y_pred_proba = self.model.predict_proba(X)

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/sklearn/
↳neighbors/_regression.py:236, in KNeighborsRegressor.predict(self, X)
    220 """Predict the target for the provided data.
    221
    222 Parameters
    (...)
    231     Target values.
    232 """
    233 if self.weights == "uniform":
    234     # In that case, we do not need the distances to perform
    235     # the weighting so we do not compute them.
--> 236     neigh_ind = self.kneighbors(X, return_distance=False)
    237     neigh_dist = None
    238 else:

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/sklearn/
↳neighbors/_base.py:824, in KNeighborsMixin.kneighbors(self, X, n_neighbors, return_distance)
    817 use_pairwise_distances_reductions = (
    818     self._fit_method == "brute"
    819     and ArgKmin.is_usable_for(
    820         X if X is not None else self._fit_X, self._fit_X, self.
↳effective_metric_
    821     )
    822 )
    823 if use_pairwise_distances_reductions:
--> 824     results = ArgKmin.compute(
    825         X=X,
    826         Y=self._fit_X,
    827         k=n_neighbors,
    828         metric=self.effective_metric_,
    829         metric_kwargs=self.effective_metric_params_,
    830         strategy="auto",
    831         return_distance=return_distance,
    832     )
    834 elif (
    835     self._fit_method == "brute" and self.metric == "precomputed" and
↳issparse(X)
    836 ):

```

```

837     results = _kneighbors_from_graph(
838         X, n_neighbors=n_neighbors, return_distance=return_distance
839     )

```

```

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/sklearn/
↳ metrics/_pairwise_distances_reduction/_dispatcher.py:289, in ArgKmin.
↳ compute(cls, X, Y, k, metric, chunk_size, metric_kwargs, strategy,
↳ return_distance)
    277     return ArgKmin64.compute(
    278         X=X,
    279         Y=Y,
    (...)
    285         return_distance=return_distance,
    286     )
    288 if X.dtype == Y.dtype == np.float32:
--> 289     return ArgKmin32.compute(
    290         X=X,
    291         Y=Y,
    292         k=k,
    293         metric=metric,
    294         chunk_size=chunk_size,
    295         metric_kwargs=metric_kwargs,
    296         strategy=strategy,
    297         return_distance=return_distance,
    298     )
    300 raise ValueError(
    301     "Only float64 or float32 datasets pairs are supported at this time, "
    302     f"got: X.dtype={X.dtype} and Y.dtype={Y.dtype}."
    303 )

```

```

File sklearn/metrics/_pairwise_distances_reduction/_argkmin.pyx:584, in sklearn
↳ metrics._pairwise_distances_reduction._argkmin.ArgKmin32.compute()

```

```

File /opt/homebrew/anaconda3/envs/ag/lib/python3.10/site-packages/threadpoolctl
↳ py:440, in _ThreadPoolLimiter.__exit__(self, type, value, traceback)
    437 def __enter__(self):
    438     return self
--> 440 def __exit__(self, type, value, traceback):
    441     self.restore_original_limits()
    443 @classmethod
    444 def wrap(cls, controller, *, limits=None, user_api=None):

```

KeyboardInterrupt:

```

[3]: # save this notebook to submissions folder
import subprocess
import os

```

```
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.  
↳join('notebook_pdfs', "submission_134.pdf"), "short_1.ipynb"])
```

Cell In[3], line 4

```
subprocess.run(["jupyter", "nbconvert", "--to", "pdf", "--output", os.path.  
↳join('notebook_pdfs', f"{submission_134}.pdf"), "short_1.ipynb"])
```

↳  
SyntaxError: f-string: expecting '}'